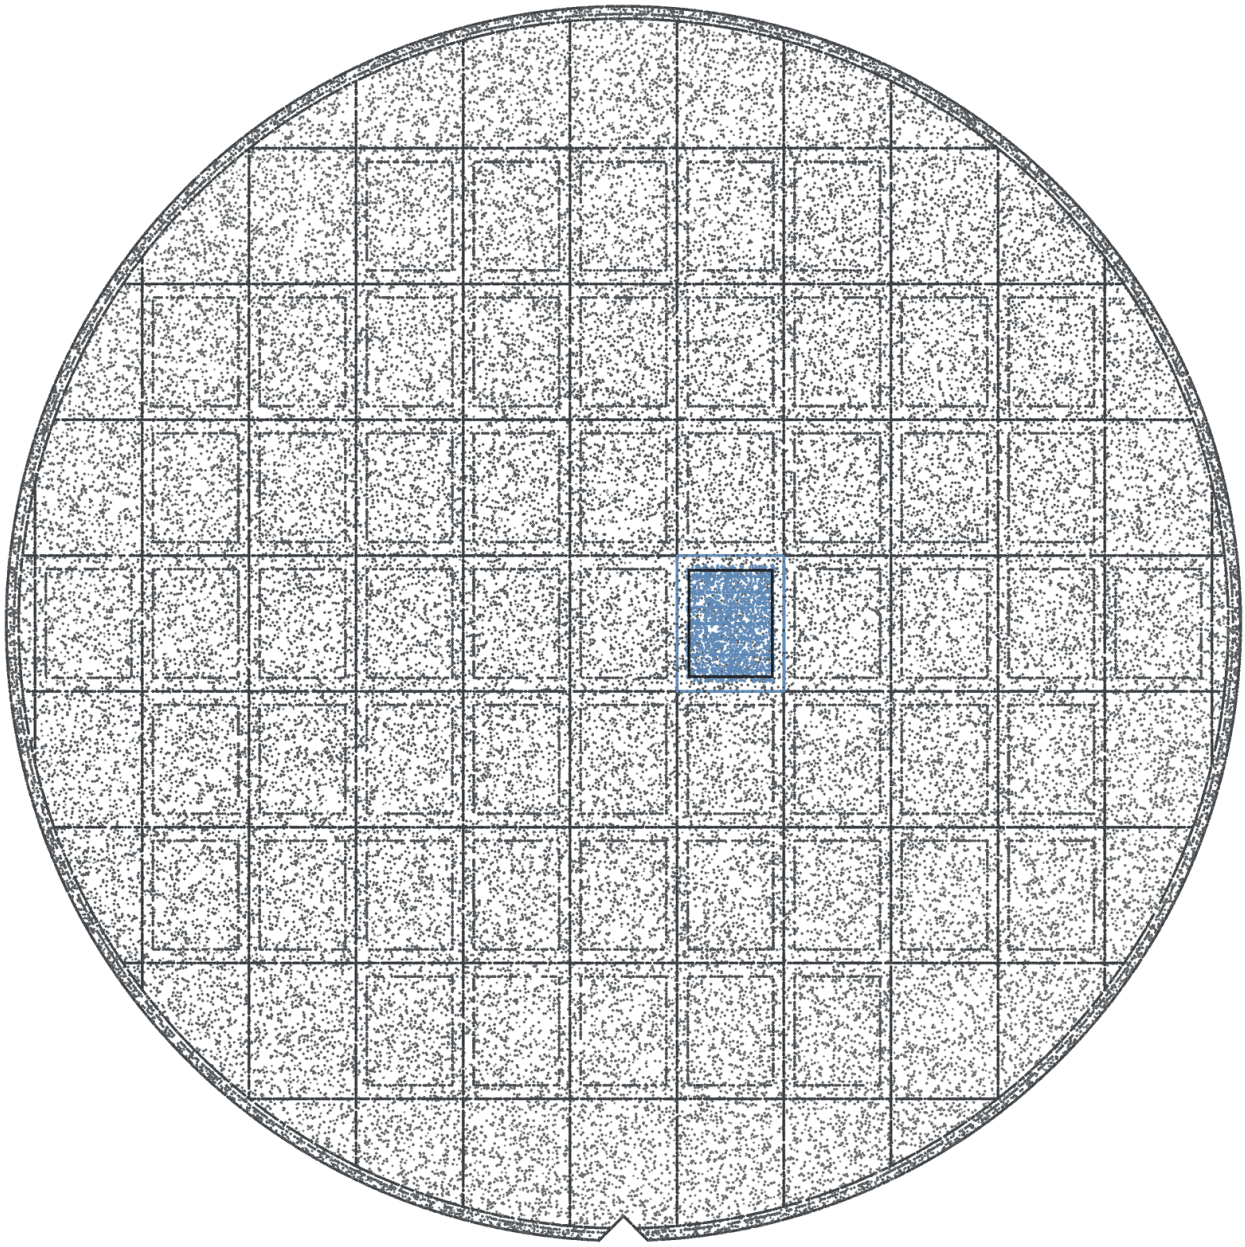


CHIP SUPPLY CHAIN

How AI chips are made

Every stage of the supply chain, and who controls it.

Yashvardhan Sharma with Claude

The Chain: seventeen chapters.

Printed from chipsupplychain.org · Updated 2 October 2026

HOW TO READ A CHAPTER

Each of the fifteen stage chapters opens with a chokepoint card.

Concentration. How many firms credibly supply this stage, rated Extreme, High, Moderate or Low.

Substitutability. How hard it would be for a well-funded state or firm to build an alternative, rated Easy, Moderate, Hard or Very hard. Each rating has its own written test.

Price or market size. What one unit costs, or what the stage earns in a year.

The card also says where China and the United States stand. The chapters follow the order in which a chip is made.

Contents

01 Chip Design, EDA and IP 3

02 Silicon and Wafers 13

03 Chemicals, Gases and Photoresist 23

04 Lithography 33

05 Photomasks and Pellicles 43

06 Deposition and Etch 53

07 Metrology and Inspection 64

08 Transistors and the Front End 74

09 Foundries and Fabs 82

10 Memory and HBM 91

11 Advanced Packaging 101

12 Substrates and PCBs 111

13 Test and Assembly 121

14 Systems and Networking 130

15 Data Centers and Power 139

16 The Economics of AI Chips 148

17 Geopolitics 156

CHAPTER
01
STAGE 1 OF 15

Chip Design, EDA and IP

Synopsys, Cadence and Siemens EDA sell the software used to design every leading-edge AI accelerator. In 2025 they took over 85 percent of an $18 billion market, licensed blocks included.

IN PLAIN TERMS

Chip design turns an idea for a new chip into a complete blueprint. A chip holds billions of transistors, tiny electrical switches. The blueprint fixes where each one sits and how it is wired to the others. The chip runs on a clock, and a signal has to cross those wires within one tick, less than a billionth of a second. Software from three companies, Synopsys, Cadence and Siemens, lays out the blueprint, checks it and signs it off. A chip factory, called a foundry, accepts only designs that have passed that software's checks. The foundry prints each layer of the design onto silicon through a stencil called a mask. Changing the blueprint once the masks are made takes new masks and a few months [2].

IN SHORT

As of 2025, the think tank CSIS found Chinese design tools still far from replacing the full range of advanced capabilities US firms offer, especially for 5 nm and 3 nm chip design [20]. But Chinese design-software companies were winning share at older manufacturing processes, the research firm SemiAnalysis reported [3].

The US government required a license to sell design software to anyone in China from late May to 2 July 2025 [1]. Since October 2022 a narrower requirement has covered software specially made for chips with gate-all-around transistors, a transistor design the rule calls key to processes at 3 nm and below [17].

2 The Dark Side Of The Semiconductor Design Renaissance – Fixed Costs Soaring Due To Photomask Sets, Verification, and Validation. SemiAnalysis · 24 July 2022 [B] newsletter.semianalysis.com/p/the-dark-side-of-the-semiconductor

20 The Double-Edged Sword of Semiconductor Export Controls: Electronic Design Automation. Center for Strategic and International Studies · 6 October 2025 [A] csis.org/analysis/double-edged-sword-semiconductor-export-controls-electronic-design-autom…

3 EDA Market Primer. SemiAnalysis · 21 May 2026 [B] newsletter.semianalysis.com/p/eda-market-primer

1 CADENCE DESIGN SYSTEMS INC, Form 8-K current report for the period ended 2025-07-02 (8-K). U.S. Securities and Exchange Commission (filing by CADENCE DESIGN SYSTEMS INC) · 3 July 2025 [A] sec.gov/Archives/edgar/data/813672/000081367225000093/cdns-20250702.htm

17 Implementation of Certain 2021 Wassenaar Arrangement Decisions on Four Section 1758 Technologies. Federal Register (Commerce Department; Industry and Security Bureau) · 15 August 2022 [A] federalregister.gov/documents/2022/08/15/2022-17125/implementation-of-certain-2021-wassena…

CHOKEPOINT CARD | EXTREME CONCENTRATION | HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
Extreme	**Hard** Only three complete sets of chip design software exist, and a fourth would have to be written from nothing and approved by a foundry.	**SemiAnalysis put a leading-edge photomask set at $40M in May 2026** In 2022 it reported that startups had designed and taped out chips on TSMC 7 nm, then a leading-edge process, for $50-75M

WHO LEADS

- US **Synopsys** $7.054B revenue in the fiscal year ended 31 October 2025, Ansys included from July; 90%+ of the timing-signoff market
- US **Cadence** $5.30B revenue in calendar 2025; 55-60% of the hardware emulator market
- DE **Siemens EDA** An estimated $2.2-2.5B revenue in 2025; 85%+ of physical verification, the final layout check
- GB **Arm** About 50% of processor compute at the top cloud firms, fiscal 2026
- US **Broadcom** $16.7B of AI semiconductor revenue in the third quarter of fiscal 2026, custom accelerators and networking combined

WHERE IT IS MADE

- US **United States** Synopsys, Cadence, Broadcom, Marvell, Nvidia, AMD; Siemens EDA's main sites
- GB **United Kingdom** Arm processor and interconnect designs licensed from Cambridge, and since 2026 its own server chip
- TW **Taiwan** Alchip, Global Unichip and MediaTek: design services that carry a chip to the factory
- IL **Israel** Annapurna Labs, the AWS design house behind Trainium and Graviton
- IN **India** Nearly 20% of the world's chip design engineers, on the Indian government's count

WHY SUBSTITUTION IS SLOW

Three complete sets of chip design software exist. Changing one tool, such as the synthesis tool, means re-running the later steps of the design. The foundry that makes the chip also specifies which signoff tools its customers must use. A newcomer would have to write a fourth set from nothing and win that foundry approval. CSET at Georgetown University found in 2021, from 2019 data, that Chinese design tools largely did not support processes at 14 nm and below. Industry experts interviewed by CSIS in August 2024 judged China unlikely to develop leading design tools for at least another decade. In the short to medium term, they judged, China could replace only specific design steps.

WHERE CHINA STANDS

China designs its own AI accelerators, but Huawei's best, the Ascend 950, delivers about half the computing performance of Nvidia's H100 from 2022, on Epoch AI's

estimate. CSET at Georgetown University found in 2021, from 2019 data, that Chinese design software largely did not support processes at 14 nm and below. In October 2025 CSIS found Chinese design tools still far from replacing the full range of advanced capabilities US firms offer, especially for 5 nm and 3 nm chip design. The US government required a license for design-software sales to China from late May 2025 until it lifted the requirement on 2 July. A narrower license requirement on software made for chips with gate-all-around transistors has applied since October 2022.

WHERE THE US STANDS

The United States has the two largest design-software vendors and the leading custom-chip design houses. Siemens EDA is German-owned, but the American technology inside it still needs US export licenses.

Design is the only stage of the supply chain without a factory, and the US government has already restricted it by letter. In late May 2025 the government required a license to sell chip design software to anyone in China [1]. Cadence cut off its Chinese customers until the government lifted the rule on 2 July [1].

How it works

Design runs in five stages, all in software, and each stage's output is the next one's input.

- **Architecture.** Engineers fix the shape of the chip: the width of the block that multiplies numbers, the amount of fast on-chip memory beside it and the number of high-bandwidth memory stacks along the chip's edges. The output is a plan.
- **RTL.** Engineers write the plan as register-transfer-level code, RTL, in the hardware language SystemVerilog. A register is a small store on the chip that holds one number. On each tick of the chip's clock, every register takes its next value at the same moment. The code says what every register holds on every tick. The output is a text description of the whole chip.
- **Verification.** A large team checks that the code does what the plan says. They run it in simulation and prove parts of it correct with mathematics. They also load it into emulators, cabinets of reprogrammable chips that behave like the design, so the chip's own software can boot before any silicon exists.
- **Physical design.** Software swaps each piece of the code for a logic cell, a small ready-made circuit the foundry has already proved it can make. It then places

1 CADENCE DESIGN SYSTEMS INC, Form 8-K current report for the period ended 2025-07-02 (8-K). U.S. Securities and Exchange Commission (filing by CADENCE DESIGN SYSTEMS INC) · 3 July 2025 A
sec.gov/Archives/edgar/data/813672/000081367225000093/cdns-20250702.htm

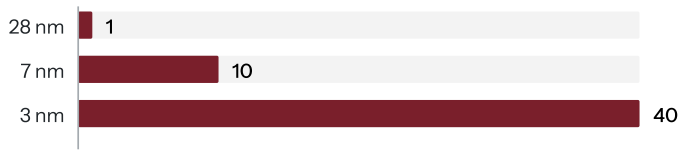
the cells on the chip, wires them together and proves that every signal arrives in time.

- **Tape-out.** Software turns the placed layout into pattern files for the mask shop. First it pre-distorts each shape so that the shape prints as drawn, a step called optical proximity correction. Then it cuts the layout into the files a mask writer reads, a step called mask data preparation. Handing those files to the mask shop is called tape-out.

All that checking exists because each revision after tape-out, called a respin, needs at least some new masks [2]. In 2022 SemiAnalysis put a full mask set at more than $1 million at 28 nm and more than $10 million at 7 nm [2]. The nanometer labels name process generations, called nodes, each finer than the last, and no longer measure anything. That year SemiAnalysis also expected 3 nm sets to push into the $40 million range [2]. Its May 2026 primer still puts a leading-edge mask set at $40 million, and a single leading-edge respin at $50 million to $100 million [3].

The same 2022 SemiAnalysis article reported that startups had taken chips on TSMC 7 nm, then a leading-edge process, from design to tape-out for $50 million to $75 million, everything included [2]. Using a different measure, CSIS put the cost of developing a leading-edge 2 nm chip at about $725 million in September 2026 [4]. It put the cost of developing a 28 nm design in the late 2000s at $48 million [4].

What one photomask set costs, by node, 2022 ($M). The 28 nm and 7 nm figures are minimums $M



Source: SemiAnalysis, The Dark Side of the Semiconductor Design Renaissance, July 2022. The 3 nm figure was a 2022 forecast. SemiAnalysis's May 2026 EDA Market Primer puts a leading-edge mask set at $40M

Variants and trade-offs

Merchant GPU versus custom ASIC

Nvidia designs its Blackwell GPU once and sells it to everyone, so the buyer gets mature software and a resale market. A custom accelerator, an ASIC, is built for one company's own workload. It handles a narrower range of tasks and is tuned to that workload. Google's seventh-generation tensor processing unit, Ironwood,

2 The Dark Side Of The Semiconductor Design Renaissance – Fixed Costs Soaring Due To Photomask Sets, Verification, and Validation. SemiAnalysis · 24 July 2022 [B] newsletter.semianalysis.com/p/the-dark-side-of-the-semiconductor

3 EDA Market Primer. SemiAnalysis · 21 May 2026 [B] newsletter.semianalysis.com/p/eda-market-primer

4 Innovation Lightbulb: The Billion-Dollar Chip: Who Can Afford to Design the Future?. Center for Strategic and International Studies · 11 September 2026 [A] csis.org/analysis/billion-dollar-chip-who-can-afford-design-future

carries 192 GB of memory per chip, six times its predecessor's, and delivers twice the performance per watt [5].

The EDA flow

Chip design software is known as electronic design automation, or EDA. Three complete EDA flows exist, each a set of tools that covers every step from code to signoff [3]. Changing the synthesis tool, which swaps the code for logic cells, means re-running the steps that follow: placing and wiring the cells, signoff and physical verification [3]. Foundries also specify which signoff tools their customers must use for tape-out [3].

Inside some steps of the flow, one vendor holds most of the market. Synopsys alone has more than 90 percent of the market for static timing analysis, the check that proves a chip will run at its rated speed [3].

Each leading vendor's share of its own design-software segment. Lower end of each reported range %



Source: SemiAnalysis, EDA Market Primer, May 2026

Buy the IP or build it

Design houses license ready-made circuit blocks, known as intellectual property or IP, including the processor cores, the interface controllers and the blocks that drive high-bandwidth memory. Arm licenses processor designs. In fiscal 2026 it reported $4.92 billion of revenue, $2.61 billion of it royalties [6]. It said it held about 50 percent of processor compute at the top cloud firms [6].

In 2026 Arm also launched the Arm AGI CPU, its first production silicon product [6]. By May 2026, servers built on it could be ordered from ASRock, Lenovo, Quanta and Supermicro [6].

The memory interface is hard to design, because its signals have to stay on time across thousands of microscopic solder bumps into a stacked memory die.

RISC-V is an open instruction set, the basic list of commands a processor understands, and it needs no license. Small RISC-V processors handle control tasks on the chip, and Nvidia uses them in every GPU [7].

5 Ironwood: The first Google TPU for the age of inference. Google · 9 April 2025 [A] blog.google/innovation-and-ai/infrastructure-and-cloud/google-cloud/ironwood-tpu-age-of-in...

3 EDA Market Primer. SemiAnalysis · 21 May 2026 [B] newsletter.semianalysis.com/p/eda-market-primer

6 Arm delivers record-breaking quarter and full-year results. Arm · 6 May 2026 [A] newsroom.arm.com/news/arm-q4-fye26-results

7 Behind The Scenes of SHD Group's 2026 RISC-V Market Forecast. RISC-V International · 19 June 2026 [A] riscv.org/blog/shd-forecast-2026

Chiplets

The reticle limit caps the size of a single die, one piece of silicon [8]. That limit is the largest area a scanner, the machine that prints the chip's pattern, can expose in one shot. Many large accelerators are built from chiplets, several smaller dies wired together, and UCIe is an open industry standard for the links between them inside a package [9]. Version 3.0 of UCIe adds 48 and 64 GT/s data rates [9]. Broadcom's 3.5D XDSiP platform packs more than 6,000 mm2 of silicon and 12 high-bandwidth memory stacks into one package [8].

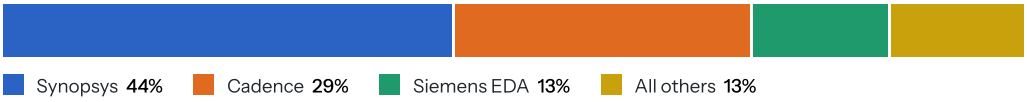
Who makes it

Design software and licensed blocks were an $18 billion market in 2025 [3]. Synopsys, Cadence and Siemens EDA took more than 85 percent of it, counting Ansys as part of Synopsys [3]. No other vendor had more than 5 percent of any core category [3].

Synopsys is the largest. Its own results put revenue at $7.054 billion for the fiscal year to 31 October 2025, with Ansys, bought that July, contributing $756.6 million [10]. Cadence, whose year ends in December, took $5.297 billion in calendar 2025 [11].

The two fiscal years do not line up, so any total for this market depends on which twelve months are counted. It also depends on how much of Ansys it includes, because Ansys sells to other industries too. SemiAnalysis counts Synopsys at $8 billion for calendar 2025 with Ansys included [3]. It puts 31 percent of Ansys's business in semiconductors and high technology, 22 percent in aerospace and 18 percent in automotive [3].

Design software and licensed-block revenue share, 2025. Synopsys includes Ansys %



Source: SemiAnalysis, EDA Market Primer, May 2026. Shares calculated from its calendar 2025 revenue figures, with Synopsys at $8B including Ansys and Siemens EDA at the middle of its $2.2-2.5B estimate

8 3.5D XDSiP Platform Technology. Broadcom · 4 December 2024 [A] docs.broadcom.com/doc/3-5d-xdsip-platform-technology

9 UCIe specifications. UCIe Consortium [A] uciexpress.org/specifications

3 EDA Market Primer. SemiAnalysis · 21 May 2026 [B] newsletter.semianalysis.com/p/eda-market-primer

10 SYNOPSYS INC, Form 8-K current report for the period ended 2025-12-10 (8-K). U.S. Securities and Exchange Commission (filing by SYNOPSYS INC) · 10 December 2025 [A] sec.gov/Archives/edgar/data/883241/000119312525314200/d29055dex991.htm

11 CADENCE DESIGN SYSTEMS INC, Form 10-K annual report for the period ended 2025-12-31 (financial statement R115). U.S. Securities and Exchange Commission (filing by CADENCE DESIGN SYSTEMS INC) · 19 February 2026 [A] sec.gov/Archives/edgar/data/813672/000081367226000016/R115.htm

Some chip firms also build custom accelerators to a customer's design and sell the finished chips. Broadcom booked $16.7 billion of AI semiconductor revenue in the third quarter of fiscal 2026, up 221 percent year over year [12]. That figure covers both its custom AI accelerators and its AI networking chips [12].

In October 2025 OpenAI and Broadcom said they would deploy 10 gigawatts of OpenAI-designed accelerators [13]. Broadcom would turn the design into silicon and supply the Ethernet networking around it [13]. In April 2026 Broadcom extended its Meta partnership to multiple gigawatts of MTIA, Meta's in-house chip, and the first 2 nm AI accelerator [14].

Marvell, another US firm, also builds custom accelerators. Below Broadcom and Marvell are the Taiwanese design-service houses. They carry a customer's circuit list all the way to a finished mask set on TSMC's newest nodes. One of them, Alchip, reported $992 million of revenue in 2025 [15].

Chip design engineers are spread across more countries than these firms are. India alone has nearly 20 percent of the world's chip designers [16].

The chokepoint

Before the 2025 letters, a narrower control already applied. Since 14 October 2022, design software specially made for chips with gate-all-around transistors has

12 Broadcom Inc., Form 8-K current report for the period ended 2026-09-02 (8-K). U.S. Securities and Exchange Commission (filing by Broadcom Inc.) · 2 September 2026 A sec.gov/Archives/edgar/data/0001730168/000173016826000076/avgo-08022026x8kxex99.htm

13 OpenAI and Broadcom announce strategic collaboration to deploy 10 gigawatts of OpenAI-designed AI accelerators. Broadcom A investors.broadcom.com/news-releases/news-release-details/openai-and-broadcom-announce-str…

14 Broadcom Announces Extended Partnership with Meta to Deploy Technology to Support Multi-Gigawatts of Meta's Custom Silicon, MTIA. Broadcom · 14 April 2026 A investors.broadcom.com/news-releases/news-release-details/broadcom-announces-extended-part…

15 Alchip Technologies, financials page. Alchip Technologies A alchip.com/en/Investors/financials

16 India's Semiconductor Vision Gathers Momentum with 3nm Chip Design and Large-Scale Talent Development Initiatives. Press Information Bureau, Government of India A pib.gov.in/PressReleasePage.aspx

needed a US export license for national-security reasons [17]. The rule calls that transistor design key to processes at 3 nm and below [17]. That license requirement covers exports to China [18]. The control carried out a December 2021 decision of the Wassenaar Arrangement, a group of countries that coordinate export controls [17].

In late May 2025 the US Bureau of Industry and Security, the agency that issues export licenses, told the design-software vendors that sales to China now needed a license [1]. Cadence got its letter on 23 May [1]. The letter required a license to export design software classified 3D991 and 3E991 to any party in China, or to a Chinese military end user anywhere [1]. Synopsys got its letter on 29 May and a second letter lifting the requirement on 2 July [19].

Unlike the 2022 control, the 2025 requirement came from the United States alone and took effect immediately. Synopsys's China revenue fell from 16 percent of its total in fiscal 2024 to 12 percent in fiscal 2025 [3]. SemiAnalysis links the decline to tighter export restrictions and to local design-software companies winning share at mature nodes [3]. SemiAnalysis also quoted Synopsys's management saying that companies Synopsys cannot sell to look for alternatives, typically local design-software and licensed-block companies [3].

In October 2025 CSIS found that only a few companies, Cadence, Synopsys and Siemens, make design software that supports the entire chip design flow of more than 40 steps [20]. In a January 2021 report built on 2019 market data, CSET at Georgetown University found that one Chinese firm, Empyrean, could reportedly run a complete design flow only for certain analog and mixed-signal chips [21]. Mixed-signal chips combine analog and digital circuits. For all other chips,

17 Implementation of Certain 2021 Wassenaar Arrangement Decisions on Four Section 1758 Technologies. Federal Register (Commerce Department; Industry and Security Bureau) · 15 August 2022 [A] federalregister.gov/documents/2022/08/15/2022-17125/implementation-of-certain-2021-wassena…

18 Supplement No. 1 to Part 738, Title 15 -- Commerce Country Chart. Electronic Code of Federal Regulations [A] ecfr.gov/current/title-15/subtitle-B/chapter-VII/subchapter-C/part-738/appendix-Supplement…

1 CADENCE DESIGN SYSTEMS INC, Form 8-K current report for the period ended 2025-07-02 (8-K). U.S. Securities and Exchange Commission (filing by CADENCE DESIGN SYSTEMS INC) · 3 July 2025 [A] sec.gov/Archives/edgar/data/813672/000081367225000093/cdns-20250702.htm

19 SYNOPSYS INC, Form 8-K current report for the period ended 2025-07-02 (8-K). U.S. Securities and Exchange Commission (filing by SYNOPSYS INC) · 3 July 2025 [A] sec.gov/Archives/edgar/data/883241/000119312525155294/d80081d8k.htm

3 EDA Market Primer. SemiAnalysis · 21 May 2026 [B] newsletter.semianalysis.com/p/eda-market-primer

20 The Double-Edged Sword of Semiconductor Export Controls: Electronic Design Automation. Center for Strategic and International Studies · 6 October 2025 [A] csis.org/analysis/double-edged-sword-semiconductor-export-controls-electronic-design-autom…

21 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 [A] cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

Chinese designers then relied on US tools [21]. The report found that other Chinese tools covered only narrow parts of the design flow or supplemented US tools, and largely did not support processes at 14 nm and below [21].

In its 2025 annual report, filed in April 2026, Empyrean says it now has complete design flows for analog, memory, radio-frequency and flat-panel display circuits [22]. It says its digital tools cover 80 percent of the main tools for digital design [22]. The same report lists foundry certifications for its circuit simulator at 8, 5 and 4 nm and for its physical verification tool, the final layout check, at 28 nm [22]. But CSIS found that as of 2025 Chinese design tools were still far from replacing the full range of advanced capabilities US firms offer, especially for 5 nm and 3 nm chip design [20].

China can design accelerators, but it cannot finish one at the leading edge without foreign tools, or build what it designs without foreign equipment. SMIC, a Chinese foundry, makes Huawei's chips. In June 2024 CSET found that SMIC's limited 7 nm capacity forced Huawei to choose between making its most advanced phone chips and its AI chips [23].

On Epoch AI's estimates, Huawei's flagship chip for 2026, the Ascend 950, delivers about half the computing performance of Nvidia's H100, which began shipping in 2022 [24]. Epoch puts Nvidia's B300 at 6.75 times the 950's theoretical peak performance [24]. Epoch's median estimate is that Huawei will make about 1.5 million Ascend chips in 2026, against 5.9 million for Nvidia [24]. On those estimates, Huawei's 2026 output would carry less than 4 percent of the computing power Nvidia ships that year [24].

Epoch puts Huawei's chips three to four years behind Nvidia's until at least 2030, for three reasons [24]. SMIC cannot pack transistors as densely as TSMC, which makes Nvidia's. Chinese memory makers are only starting to produce high-bandwidth memory. And Huawei's software for programming its chips is far less mature than Nvidia's [24].

Key evaluation criteria

- **Flow completeness.** How many steps one vendor covers, from the first line of code to final signoff.
- **Foundry certification.** Whether the foundry has validated that exact tool version against its reference flow and process design kit, the file set describing

21 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 A cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

22 北京华大九天科技股份有限公司2025年年度报告全文. CNINFO · 27 April 2026 A static.cninfo.com.cn/finalpage/2026-04-28/1225217315.PDF

20 The Double-Edged Sword of Semiconductor Export Controls: Electronic Design Automation. Center for Strategic and International Studies · 6 October 2025 A csis.org/analysis/double-edged-sword-semiconductor-export-controls-electronic-design-autom…

23 Pushing the Limits: Huawei's AI Chip Tests U.S. Export Controls. Center for Security and Emerging Technology (CSET) · 17 June 2024 A cset.georgetown.edu/publication/pushing-the-limits-huaweis-ai-chip-tests-u-s-export-controls

24 Will Huawei catch up to Nvidia by 2030?. Epoch AI · 24 September 2026 A epoch.ai/publications/huaweis-roadmap-to-2031

what the foundry can print. Without that, the foundry refuses signoff.

- **Verification throughput.** Simulation cycles and emulator capacity per day, which set the schedule more than anything else.
- **IP availability at node.** Whether the licensed blocks have been proven in silicon on that process: the processor cores, the high-speed serial links, PCIe and the memory interface.
- **Design services depth.** Whether a partner can carry a customer's circuit list to a finished mask set on the newest node.
- **Legal exposure.** How much of the flow is US-origin technology, and so subject to export licensing.

FIVE THINGS TO REMEMBER

- Chip design turns an idea into a blueprint, and a factory builds only designs checked in approved software.
- Synopsys, Cadence and Siemens EDA sell the software used to design every leading-edge AI chip.
- To compete with them, a new vendor would have to write a complete set of tools and win a foundry's approval.
- China designs its own AI chips, but in October 2025 CSIS found its design software still far from replacing US tools, especially for 5 nm and 3 nm chips.
- Huawei's best AI chip trails Nvidia's by three to four years, on Epoch AI's estimate.

CHAPTER
02
STAGE 2 OF 15

Silicon and Wafers

Firms based in Japan, Taiwan, Germany and South Korea make most 300 mm wafers used for advanced chips. China makes most polysilicon, almost all of it for solar cells.

IN PLAIN TERMS

Every chip is built on a wafer, a thin round slice of silicon about 30 centimeters across that carries hundreds of chips at once. Making one starts with sand, which is silicon bound to oxygen. The silicon is refined until metal impurities are no more than a few parts per billion [3]. It is then melted, and a small seed crystal is lowered to the surface and drawn slowly back up. The melt freezes onto the seed as one long cylinder, a single crystal with its atoms lined up in one unbroken pattern from end to end. Saws cut the cylinder into discs, and each disc is polished to a mirror finish and made extremely flat across its whole face [3]. In 2025 the top five wafer makers served around 75 percent of the world market, and none of them is American or Chinese [4].

3 SUMCO, production processes page. SUMCO [A] sumcosi.com/english/products/process

4 Investor Presentation: Foundation of Digital Life. Siltronic · April 2026 [A] siltronic.com/fileadmin/investorrelations/2026/Q1/20260429_Siltronic_InvestorPresentation_...

IN SHORT

Japanese companies held 53 percent of the world's silicon wafer sales in 2023, by a Fuji Keizai survey that Japan's trade ministry reproduced [9]. Firms based in Taiwan, South Korea and Europe held most of the rest [9]. China's largest maker of 300 mm wafers, Eswin, says it ranked sixth in the world with about 7.7 percent of global 300 mm shipments in the first half of 2026 [20]. Chips also need electronic-grade polysilicon, the raw material for wafers, and none of the five major suppliers of it that the German wafer maker Siltronic named for 2025 is Chinese [4].

GlobalWafers of Taiwan opened a $3.5 billion plant in Texas in May 2025, the first production line of its kind built in the United States in over twenty years [16]. US Presidential Proclamation 11052 sets a minimum import price of $100 per kilogram for polysilicon ingots and wafers from December 4, 2026 [15].

9 METI, Semiconductor and digital industry strategy: future direction (半導体・デジタル産業戦略の今後の方向性), 23 December 2025. Ministry of Economy, Trade and Industry (Japan) · 23 December 2025 [A] meti.go.jp/policy/mono_info_service/joho/conference/semicon_digital/0014/handeji14-4.pdf

20 西安奕斯伟材料科技股份有限公司2026年半年度报告. CNINFO · 25 August 2026 [A] static.cninfo.com.cn/finalpage/2026-08-26/1225503057.PDF

4 Investor Presentation: Foundation of Digital Life. Siltronic · April 2026 [A] siltronic.com/fileadmin/investorrelations/2026/Q1/20260429_Siltronic_InvestorPresentation_...

16 GlobalWafers America Officially Opens for Business. GlobalWafers · 15 May 2025 [A] sas-globalwafers.com/en/gwc_news_en_20250516

15 Adjusting Imports of Polysilicon and Its Derivatives Into the United States. Federal Register (Executive Office of the President) · 11 August 2026 [A] federalregister.gov/documents/2026/08/11/2026-16400/adjusting-imports-of-polysilicon-and-i...

CHOKEPOINT CARD

HIGH CONCENTRATION

MODERATELY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
High	Moderate Several firms already make 300 mm wafers that fabs have approved, but customers have to test and approve a new wafer plant before they buy from it.	The world wafer market was $11.4bn in 2025, on 12,973 million square inches shipped

WHO LEADS

- JP **Shin-Etsu Handotai** Its parent does not break out wafer sales. The parent's electronics materials segment, which also sells rare earth magnets, encapsulants, photoresists and photomask blanks, had ¥1,015.7bn of sales in the year to March 2026
- JP **SUMCO** ¥409.7bn of 2025 net sales, but an ¥11.8bn loss
- TW **GlobalWafers** NT$60.6bn revenue in 2025
- KR **SK Siltron** One of the top five wafer makers, which together served around 75 percent of the market in 2025
- DE **Siltronic** €1,346.7m revenue in 2025

WHERE IT IS MADE

- JP **Japan** Shin-Etsu and SUMCO crystal growth and polishing
- TW **Taiwan** GlobalWafers headquarters and plants
- DE **Germany** Siltronic plants in Burghausen and Freiberg, and Wacker semiconductor-grade polysilicon
- KR **South Korea** SK Siltron
- US **United States** Hemlock polysilicon in Michigan, and GlobalWafers plants in Sherman, Texas and St Peters, Missouri

WHY SUBSTITUTION IS POSSIBLE

Firms headquartered in Japan, Taiwan, Germany and South Korea make state-of-the-art 300 mm wafers. The top five served around 75 percent of the market in 2025, so losing one leaves four. A new wafer plant can still be built. GlobalWafers of Taiwan spent $3.5 billion on one in Texas, the first wafer production line of its kind in the United States in over twenty years. The slow part is approval. Customers have to test and approve a new plant, a step called qualification, before they buy from it. Siltronic completed key customer qualifications at its new Singapore fab in July 2025. Siltronic also sells under long-term agreements.

WHERE CHINA STANDS

China made about 93 percent of the world's polysilicon in 2023. An estimated 98 percent of China's output was solar grade and 2 percent electronic grade, close to the split for the world as a whole. None of the five major suppliers of electronic-grade silicon that Siltronic names is Chinese. Chinese firms held under 1 percent of the 300 mm wafer market as of 2019. As of 2018, Chinese

producers made only 12 percent of the wafers for Chinese-headquartered 300 mm fabs. Eswin, China's largest 300 mm wafer maker, says it shipped about 7.7 percent of the world's 300 mm wafers in the first half of 2026, up from about 6 percent in 2024.

WHERE THE US STANDS

No major US-headquartered firm makes wafers. Hemlock makes polysilicon in Michigan, and GlobalWafers of Taiwan won CHIPS awards of up to $406 million for wafer plants in Texas and Missouri. Its $3.5 billion Texas plant opened in May 2025. Proclamation 11052, signed August 6, 2026, sets a minimum import price of $100 per kilogram for polysilicon ingots and wafers from December 4, 2026. From the same date it adds an extra 15 percent duty on imported ingots and other listed products made from polysilicon. Firms with approved plans to build US polysilicon, ingot, wafer or solar cell plants can import equipment and some of those products without the new duties.

The world market for silicon wafers was $11.4 billion in 2025, on 12,973 million square inches shipped [1]. That is a small fraction of Nvidia's data center revenue, which reached $193.7 billion in the fiscal year to January 25, 2026 [2].

How it works

Silicon starts as quartzite, a rock that is silicon bound to oxygen. An arc furnace melts it with carbon, and the carbon takes the oxygen away as gas, leaving rough silicon metal. Refiners turn the metal into a gas, trichlorosilane, and distill it, because the impurities boil at other temperatures and stay behind.

In the next step, the Siemens process, the clean gas flows over silicon filaments heated inside a bell jar. On the hot surface the gas breaks apart and its silicon settles onto the filaments. Over days the filaments thicken into gray rods of polysilicon. Everything downstream is made from those rods. Polysilicon comes in two ultra-high-purity grades, one for semiconductors and one for solar cells, and chips use the semiconductor grade [5].

A grower melts the polysilicon at about 1,420 degrees Celsius, hotter than lava, and pulls a single crystal out of it by the Czochralski method [3]. The grower lowers a small seed crystal on the end of a rod to the surface of the melt and draws it back up, slowly turning. Silicon freezes onto the seed in the seed's own atomic pattern, and a cylinder, the ingot, grows beneath it. The pull starts with a thin neck so that flaws from the seed's first contact with the melt stay in the neck and never reach the ingot. Boron or phosphorus stirred into the melt, a step called doping, sets how readily the wafer conducts and whether negative or positive charges carry its current.

1 SEMI Reports 2025 Annual Worldwide Silicon Wafer Shipments and Revenue Results. PR Newswire · 10 February 2026 A prnewswire.com/news-releases/semi-reports-2025-annual-worldwide-silicon-wafer-shipments-an...

2 NVIDIA Announces Financial Results for Fourth Quarter and Fiscal 2026. Nvidia · 25 February 2026 A nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-fourth-quarter-and-fisca...

5 Mineral Commodity Summaries 2026: Silicon. U.S. Geological Survey · 5 February 2026 A pubs.usgs.gov/periodicals/mcs2026/mcs2026-silicon.pdf

3 SUMCO, production processes page. SUMCO A sumcosi.com/english/products/process

A wire saw, a web of fine wire running through abrasive, slices the ingot into discs. The maker then laps each disc, grinding both faces flat between rotating plates. It etches off the damaged surface with chemicals, then polishes the disc to a mirror with colloidal silica, a slurry of silica particles far too small to see. Almost every step throws material away.

Epitaxial wafers get one more layer of near-perfect silicon, grown on top at about 1,200 degrees Celsius [3]. That is still hot enough to glow orange.

Variants and trade-offs

Polished versus epitaxial

A polished wafer is a doped single crystal polished to a mirror finish [3]. An epitaxial wafer is a polished wafer with a thin layer of single-crystal silicon grown on top at about 1,200 degrees Celsius [3]. The transistors are built in that layer, so its quality sets the lowest leakage the chip can reach. It also sets the chip's risk of latch-up, an unintended short circuit between power and ground. SEMI, the industry association, credits the 2025 rise in shipments to AI-driven demand for advanced epitaxial wafers for logic chips and polished wafers for high-bandwidth memory [1].

Silicon on insulator

In a silicon-on-insulator (SOI) wafer, a buried layer of oxide, which is glass, separates a thin working layer of silicon from the bulk beneath. Nearly all SOI wafers are made with Soitec's Smart Cut process [6]. Soitec's SOI product lines include RF-SOI for radio chips and fully depleted SOI (FD-SOI), and its markets run from mobile communications to cloud AI [6]. But SEMI tied the AI-driven growth in logic chips to advanced epitaxial wafers [1].

Who makes it

The Center for Security and Emerging Technology (CSET) at Georgetown University named five key wafer firms in a January 2021 brief [7]. They were Shin-Etsu and SUMCO of Japan, GlobalWafers of Taiwan, Siltronic of Germany and SK Siltron of South Korea [7]. CSET also listed nine Chinese-owned producers, among

3 SUMCO, production processes page. SUMCO A sumcosi.com/english/products/process

1 SEMI Reports 2025 Annual Worldwide Silicon Wafer Shipments and Revenue Results. PR Newswire · 10 February 2026 A prnewswire.com/news-releases/semi-reports-2025-annual-worldwide-silicon-wafer-shipments-an...

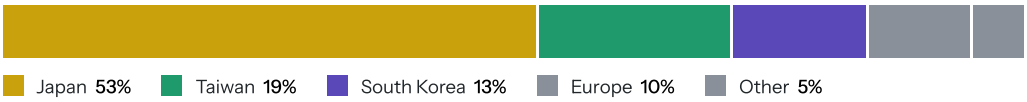
6 Soitec, key figures. Soitec A soitec.com/home/group

7 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 A cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

them Zhonghuan, Simgui, ZingSEMI and Finland's Okmetic, but said they generated little revenue in comparison [7].

As of 2018, five Chinese firms were making or planning 300 mm wafers [7]. In its October 2025 listing prospectus, Eswin, a Chinese wafer maker, counts seven Chinese firms, itself included, making 300 mm wafers at scale [8]. Their combined capacity was 2.94 million wafers a month at the end of 2024 [8].

Silicon wafer sales by supplier headquarters, 2023. Japan held 53%, and the source shows no separate share for China or the United States %



Source: METI, December 2025, from Fuji Keizai data (2023); Japan's share is printed, the others are read off METI's chart

The chart shows shares of 2023 sales from a survey by Fuji Keizai, a Japanese research firm [9]. Japan's Ministry of Economy, Trade and Industry reproduced the figures in a December 2025 strategy deck [9]. The deck prints only Japan's 53 percent, so the other shares are read off its chart [9]. The chart gives no separate share to Chinese or US companies [9]. Siltronic's April 2026 investor presentation says the top five wafer makers serve around 75 percent of the market, based on revenue in the first three quarters of 2025 [4].

The wafer makers' 2025 results show strong AI demand and a slump in demand for chips made on older, mature nodes.

- **Shin-Etsu.** Its parent's electronics materials segment sold ¥1,015.7 billion in the year to March 2026, up 9 percent [10]. AI-related demand stayed strong and demand in other sectors began to rise, lifting sales of silicon wafers, photoresist and photomask blanks [10]. The segment also sells rare earth

7 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 A
cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

8 西安奕材首次公开发行股票并在科创板上市招股说明书. Shanghai Stock Exchange · 22 October 2025 A
static.sse.com.cn/disclosure/listedinfo/announcement/c/new/2025-10-22/688783_20251022_UYMD...

9 METI, Semiconductor and digital industry strategy: future direction (半導体・デジタル産業戦略の今後の方向性), 23 December 2025. Ministry of Economy, Trade and Industry (Japan) · 23 December 2025 A
meti.go.jp/policy/mono_info_service/joho/conference/semicon_digital/0014/handeji14-4.pdf

4 Investor Presentation: Foundation of Digital Life. Siltronic · April 2026 A
siltronic.com/fileadmin/investorrelations/2026/Q1/20260429_Siltronic_InvestorPresentation_...

10 Consolidated Financial Results for the Fiscal Year Ended March 31, 2026. Shin-Etsu Chemical · 27 April 2026 A
shinetsu.co.jp/wp-content/uploads/2025/07/20260428_con_E.pdf

magnets and other materials, and the results release does not break out wafer sales [10]. So the figure cannot be compared with the wafer revenue of the other firms listed.

- **SUMCO.** Sold ¥409.7 billion and still lost ¥11.8 billion, as operating profit fell from ¥36.9 billion to ¥1.3 billion [11]. In the same year it added leading-edge 300 mm capacity and reorganized production of 200 mm and smaller wafers, where demand stayed weak [11].
- **Siltronic.** Lost €77.9 million on €1,346.7 million of sales while investing €369.1 million, with its new Singapore fab as the focus [12].
- **GlobalWafers.** NT$60.6 billion of revenue, down 3.24 percent in local currency [13].

Worldwide, the wafer area shipped rose 5.8 percent in 2025, but revenue fell 1.2 percent [1]. SEMI attributes the softer revenue mostly to traditional chip applications, where demand and pricing had yet to improve [1].

China made 1.50 million metric tons of polysilicon, the raw material, in 2023, about 93 percent of world output [14]. Solar-grade material was an estimated 98 percent of that output and electronic-grade 2 percent [14]. The Commerce Department found in 2026 that semiconductor-grade polysilicon makes up only 2.4 percent of global production, a split close to China's [15]. These figures alone do not show whether Chinese polysilicon goes into chips.

China's polysilicon output by grade, 2023. China made 1.50 million metric tons that year, about 93% of world output %



Solar-grade 98% Electronic-grade 2%

Source: USGS Minerals Yearbook, China, 2023

10 Consolidated Financial Results for the Fiscal Year Ended March 31, 2026. Shin-Etsu Chemical · 27 April 2026 A shinetsu.co.jp/wp-content/uploads/2025/07/20260428_con_E.pdf

11 Notice of Convocation of the 27th Ordinary General Meeting of Shareholders. SUMCO · 6 March 2026 A sumcosi.com/english/pdf/ir/library/shareholders/27/pdf/nc_e_27.pdf

12 Siltronic AG: Robust business performance in 2025 demonstrates resilience despite challenging conditions. Siltronic · 27 August 2026 A siltronic.com/en/press/press-releases/siltronic-ag-robust-business-performance-in-2025-dem…

13 GlobalWafers Reports Full Year 2025 Results. GlobalWafers · 3 March 2026 A sas-globalwafers.com/en/gwc_news_en_20260303

1 SEMI Reports 2025 Annual Worldwide Silicon Wafer Shipments and Revenue Results. PR Newswire · 10 February 2026 A prnewswire.com/news-releases/semi-reports-2025-annual-worldwide-silicon-wafer-shipments-an…

14 2023 Minerals Yearbook: China (advance release). U.S. Geological Survey · February 2026 A pubs.usgs.gov/myb/vol3/2023/myb3-2023-china.pdf

15 Adjusting Imports of Polysilicon and Its Derivatives Into the United States. Federal Register (Executive Office of the President) · 11 August 2026 A federalregister.gov/documents/2026/08/11/2026-16400/adjusting-imports-of-polysilicon-and-i…

Siltronic put the 2025 market for electronic-grade silicon at $1.4 billion, with five major suppliers, none of them Chinese: Wacker, Hemlock, OCI, Tokuyama and Mitsubishi [4].

Siltronic of Germany describes itself as the only Western-based maker among the top five wafer makers, based on revenue in the first three quarters of 2025 [4]. GlobalWafers opened a $3.5 billion plant in Sherman, Texas in May 2025, the first production line of its kind built in the United States in over twenty years [16]. The same day it announced another $4 billion for phases three and four in Sherman [16]. The US government had awarded it up to $406 million under the CHIPS incentives program for wafer plants in Texas and Missouri [17].

Presidential Proclamation 11052, signed August 6, 2026, sets a minimum import price of $100 per kilogram for polysilicon ingots and wafers from December 4, 2026 [15]. From the same date it adds an extra 15 percent duty on imported ingots and other listed products made from polysilicon [15]. A September 2026 Commerce rule meant to stop stockpiling before the duties start lists tariff code 3818.00.0040 among the products it covers [18]. That code covers round, doped, single-crystal silicon wafers [19].

Commerce may approve plans to build US polysilicon, ingot, wafer or solar cell plants whose construction starts by January 20, 2029 [15]. Firms with approved plans can import equipment and some of those products without the new duties [15].

The chokepoint

China makes most of the world's polysilicon, but in 2019 its firms had almost no share of the wafer market. As of 2019 Chinese firms held under 5 percent of the

4 Investor Presentation: Foundation of Digital Life. Siltronic · April 2026 [A] siltronic.com/fileadmin/investorrelations/2026/Q1/20260429_Siltronic_InvestorPresentation_...

16 GlobalWafers America Officially Opens for Business. GlobalWafers · 15 May 2025 [A] sas-globalwafers.com/en/gwc_news_en_20250516

17 Biden-Harris Administration Announces CHIPS Incentives Awards with GlobalWafers to Support Domestic Production of Silicon Wafers. NIST · 17 December 2024 [A] nist.gov/news-events/news/2024/12/biden-harris-administration-announces-chips-incentives-a...

15 Adjusting Imports of Polysilicon and Its Derivatives Into the United States. Federal Register (Executive Office of the President) · 11 August 2026 [A] federalregister.gov/documents/2026/08/11/2026-16400/adjusting-imports-of-polysilicon-and-i...

18 Measures To Restrict Stockpiling of Polysilicon and Polysilicon Derivatives Under Proclamation 11052. Federal Register (Commerce Department; Industry and Security Bureau) · 24 September 2026 [A] federalregister.gov/documents/2026/09/24/2026-19537/measures-to-restrict-stockpiling-of-po...

19 Harmonized Tariff Schedule, heading 3818.00.00. U.S. International Trade Commission [A] hts.usitc.gov/search

wafer market and under 1 percent of the 300 mm market [7]. As of 2018, Chinese producers made 12 percent of the wafers for Chinese-headquartered 300 mm fabs, against 18 percent for 100 to 200 mm fabs [7].

Eswin, mainland China's largest maker of 300 mm wafers, says it ranked sixth in the world with about 7.7 percent of global 300 mm shipments in the first half of 2026 [20]. Its October 2025 Shanghai listing prospectus, drawing on SEMI data, estimates that the five largest makers shipped about 79 percent of 300 mm wafers in 2024 [8]. It puts Eswin at about 6 percent and all other makers at about 15 percent [8].

Making the smoothest, high-purity 300 mm wafers takes considerable know-how that is hard to write down [7]. Customers have to qualify a new wafer plant before they buy from it, and Siltronic completed key customer qualifications at its new Singapore fab in July 2025 [4].

In SEMI figures from November 2020 cited by CSET, 99.7 percent of world fab capacity at 45 nm and below used 300 mm wafers [7]. SEMI estimated that 300 mm wafers made up about 78.8 percent of all wafer area shipped worldwide in 2025, according to Eswin's 2025 annual report summary [21].

The last increase in wafer size, to 300 mm, happened in 2002 [7]. Bigger discs make printing, or photolithography, a bigger share of the processing cost [7]. A 2013 University of Texas lecture cited by CSET put photolithography at half the cost of processing a 300 mm wafer, against a quarter for a 150 mm one [7]. Wafers are still produced in diameters of up to 300 mm [1].

Some GPU dies fill the whole area a lithography tool can print in one exposure. One fatal defect on such a die ruins far more silicon than one on a small chip. Extreme-ultraviolet printing also stays in focus over only a tiny range of heights (see Lithography), so the disc has to be flat at every point.

7 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 A
cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

20 西安奕斯伟材料科技股份有限公司2026年半年度报告. CNINFO · 25 August 2026 A static.cninfo.com.cn/finalpage/2026-08-26/1225503057.PDF

8 西安奕材首次公开发行股票并在科创板上市招股说明书. Shanghai Stock Exchange · 22 October 2025 A
static.sse.com.cn/disclosure/listedinfo/announcement/c/new/2025-10-22/688783_20251022_UYMD...

4 Investor Presentation: Foundation of Digital Life. Siltronic · April 2026 A
siltronic.com/fileadmin/investorrelations/2026/Q1/20260429_Siltronic_InvestorPresentation_...

21 西安奕斯伟材料科技股份有限公司2025年年度报告摘要. Shanghai Stock Exchange · 21 April 2026 A
static.sse.com.cn/disclosure/listedinfo/announcement/c/new/2026-04-21/688783_20260421_6JVC...

1 SEMI Reports 2025 Annual Worldwide Silicon Wafer Shipments and Revenue Results. PR Newswire · 10 February 2026 A
prnewswire.com/news-releases/semi-reports-2025-annual-worldwide-silicon-wafer-shipments-an...

Key evaluation criteria

- **Grade.** Chips need electronic-grade polysilicon, and none of the five major suppliers of it that Siltronic named for 2025 is Chinese [4].
- **Defect density** in the epitaxial layer. A big AI die loses more good silicon per defect than a small chip does. SEMI's wafer group says the move to sub-3nm processes is raising requirements for wafer quality and consistency [1].
- **Flatness and nanotopography**, the fine ripple across the surface. Extreme-ultraviolet exposure stays in focus over only a tiny range of heights, so the surface has to stay within that range across the whole 300 mm wafer.
- **Diameter.** In SEMI figures from November 2020 cited by CSET, 99.7 percent of world fab capacity at 45 nm and below used 300 mm wafers [7]. Of all the wafer area shipped worldwide in 2025, 300 mm wafers made up an estimated 78.8 percent [21].
- **Contract structure**, since Siltronic expects continued price pressure in 2026 on sales outside its existing long-term agreements [12].

FOUR THINGS TO REMEMBER

— Every chip starts on a wafer, a polished disc of pure silicon about 30 centimeters across.

— The top five wafer makers, in Japan, Taiwan, Germany and South Korea, served around 75 percent of the market in 2025.

— Customers have to test and approve a new wafer plant before they buy from it.

— China makes most of the world's polysilicon, but none of the five major suppliers of electronic-grade silicon, the grade chips need, is Chinese.

4 Investor Presentation: Foundation of Digital Life. Siltronic · April 2026 [A]
siltronic.com/fileadmin/investorrelations/2026/Q1/20260429_Siltronic_InvestorPresentation_...

1 SEMI Reports 2025 Annual Worldwide Silicon Wafer Shipments and Revenue Results. PR Newswire · 10 February 2026 [A]
prnewswire.com/news-releases/semi-reports-2025-annual-worldwide-silicon-wafer-shipments-an...

7 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 [A]
cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

21 西安奕斯伟材料科技股份有限公司2025年年度报告摘要. Shanghai Stock Exchange · 21 April 2026 [A]
static.sse.com.cn/disclosure/listedinfo/announcement/c/new/2026-04-21/688783_20260421_6JVC...

12 Siltronic AG: Robust business performance in 2025 demonstrates resilience despite challenging conditions. Siltronic · 27 August 2026 [A] siltronic.com/en/press/press-releases/siltronic-ag-robust-business-performance-in-2025-dem...

CHAPTER 03
STAGE 3 OF 15

Chemicals, Gases and Photoresist

Japanese companies held 78 percent of world photoresist sales in 2023. Supply of chip chemicals, gases and metals has been disrupted three times since 2019.

IN PLAIN TERMS

Photoresist is a coating that changes wherever light hits it. A machine spins a coat of it, thinner than a soap bubble, across the wafer. Light shines through a stencil onto the coat, and a liquid wash then takes away the parts the light reached. What remains is a pattern, and the next machine cuts that pattern into the wafer beneath, a step called etching. A few hundred other gases and liquids lay down layers, strip them off and clean the wafer between steps. All of them have to be pure to about one stray metal atom in every trillion. A chipmaker approves each formula for one factory and one product at a time. A new supplier has to earn that approval from scratch, so adding a second supplier takes years. Japanese companies held 78 percent of world photoresist sales in 2023 [2].

IN SHORT

In July 2019 Japan began requiring an export license for each contract to sell photoresist, hydrogen fluoride or fluorinated polyimide to South Korea [22]. Japan's hydrogen fluoride exports to South Korea then fell 87.9 percent, but its photoresist exports did not [22]. Production of all three chemicals rose inside South Korea [22].

China banned exports of the metals gallium, germanium and antimony to the United States in December 2024 [26]. Germanium metal averaged $4,100 a kilogram in 2025, against $1,392 in 2023 [29]. In November 2025 China suspended the ban until 27 November 2026 [27].

2 METI, Semiconductor and digital industry strategy: future direction (半導体・デジタル産業戦略の今後の方向性), 23 December 2025. Ministry of Economy, Trade and Industry (Japan) · 23 December 2025 [A] meti.go.jp/policy/mono_info_service/joho/conference/semicon_digital/0014/handeji14-4.pdf

22 The impact of export controls on international trade: evidence from the Japan–Korea trade dispute in the semiconductor industry. Research Institute of Economy, Trade and Industry (Japan) · 8 May 2023 [A] rieti.go.jp/en/columns/v01_0201.html

26 商务部公告2024年第46号 关于加强相关两用物项对美国出口管制的公告. Ministry of Commerce of the People's Republic of China [A] mofcom.gov.cn/zwgk/zcfb/art/2024/art_3d5e990b43424e60828030f58a547b60.html

29 Mineral Commodity Summaries 2026: Germanium. U.S. Geological Survey · 5 February 2026 [A] pubs.usgs.gov/periodicals/mcs2026/mcs2026-germanium.pdf

27 商务部公告2025年第72号 公布调整实施商务部公告2024年第46号的决定. Ministry of Commerce of the People's Republic of China [A] mofcom.gov.cn/zwgk/zcfb/art/2025/art_5c68985a6b1a46778e2e8dbff1bb1601.html

CHOKEPOINT CARD

HIGH CONCENTRATION

MODERATELY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
High	Moderate Rated on photoresist, the hardest input to replace at this stage. Japanese companies held 78 percent of world photoresist sales in 2023.	**Wafer fabrication materials were a $45.8bn market in 2025 and packaging materials $27.4bn, $73.2bn in total**

WHO LEADS

- JP **JSR** Owner of Inpria, a maker of metal-oxide resist for extreme-ultraviolet printing. A vehicle of the state-backed Japan Investment Corporation bought 84.36% of JSR in an April 2024 tender offer
- JP **Tokyo Ohka Kogyo** Resist maker with ¥139,668m of sales in the first half of fiscal 2026
- FR **Air Liquide** €2,465m of electronics revenue in 2025, 9.1% of group sales
- US **Entegris** Makes filters, polishing slurries, gas delivery systems and wafer carriers, with $3,196.6m of net sales in 2025
- JP **JX Advanced Metals** Says it holds about 65% of the world market for semiconductor sputtering targets as of fiscal 2026. The targets supply thin metal films such as the barrier and seed layers for copper wiring

WHERE IT IS MADE

- JP **Japan** Photoresist, ultra-pure hydrofluoric acid, sputtering targets, polishing slurries
- US **United States** Entegris filters, wafer carriers and polishing slurries; Qnity Electronics slurries, pads and resist; Air Products gases
- FR **France** Air Liquide electronics gases and precursors
- KR **South Korea** SK Materials nitrogen trifluoride, Dongjin resist, on-site gas for memory fabs
- CN **China** Gallium, germanium and antimony refining; growing wet chemicals and polishing slurries

WHY SUBSTITUTION IS POSSIBLE

Photoresist is the light-sensitive coating that holds the printed pattern on the wafer. Each formula is approved for one layer of one product at one fab. So a new supplier has to develop its version together with the fab, one layer at a time. When Japan required individual export licenses for resist and two other chemicals sold to South Korea in 2019, South Korea imported more from the United States, Taiwan and Belgium. Output of the three chemicals rose at Korean firms and at Japanese companies' affiliates in South Korea. Some of the American supply may have been Japanese material shipped by way of the United States. Japanese resist exports to South Korea did not fall, possibly in part because Japan allowed three-year bulk licenses for some resist sales in December 2019.

WHERE CHINA STANDS

In a January 2021 brief, Georgetown University's CSET found that China made no photoresist for extreme-ultraviolet printing and was still developing resist for 193 nm printing. At the time China held under 5 percent of the photoresist market. In March 2026 the Chinese resist maker Hubei Dinglong said China had yet to reach large-scale production of advanced resist for 248 nm, 193 nm or extreme-ultraviolet printing. In February 2026 the US Geological Survey put China's share of world production of primary low-purity gallium at 99 percent. Since August 2023 China has required licenses to export gallium and germanium. It banned their export to the United States in December 2024, and in November 2025 it suspended that ban until 27 November 2026.

WHERE THE US STANDS

The American firms Entegris, Qnity Electronics and Air Products are strong in filters, containers, polishing slurries and gas delivery. In photoresist, American companies ranked a distant second to Japanese companies by sales in 2023. In 2021 JSR of Japan agreed to buy Inpria, an Oregon maker of resist for extreme ultraviolet.

A chip fab uses a few hundred chemicals and gases, and a shortage of any one of them can stop it. The French gas company Air Liquide alone supplies more than 200 molecules and roughly 50,000 cylinders a year [1].

How it works

The resist has to change only where the light lands. It also has to react fast enough for the scanner, the machine that projects the pattern, to expose each patch of the wafer in milliseconds.

Since the 1990s the standard has been chemically amplified resist. In resists for 248 nm and 193 nm light, a photon absorbed in the film can free an acid molecule, and nothing more happens until the wafer is baked. In the bake, each acid molecule sets off hundreds of reactions around itself that turn the film soluble, so a little light does a lot of chemistry. The wash then takes away the exposed film and leaves the rest. The 248 nm light comes from a krypton fluoride (KrF) laser. The 193 nm light comes from an argon fluoride (ArF) laser and shines through a layer of water, which sharpens the image.

That chemistry works less well with extreme ultraviolet (EUV) light (see Lithography). An absorbed EUV photon instead knocks electrons loose, and a cascade of lower-energy secondary electrons makes the chemical changes in the resist [3]. A 13.5 nm EUV photon carries about fourteen times the energy of a 193 nm one, so the same dose of light arrives as far fewer photons. With fewer photons, random variation in where each one lands matters more. It shows up as ragged line edges and missing contact holes, the small wells that join one layer of wiring to the next.

1 Air Liquide, electronics page. Air Liquide [A] airliquide.com/group/activities/electronics

3 How Scientists Are Accelerating the Development of Next-Gen Microelectronics. Lawrence Berkeley National Laboratory · 21 August 2023 [A] newscenter.lbl.gov/2023/08/21/how-scientists-are-accelerating-next-gen-microelectronics

Resist makers responded with metal-oxide resists, which are built on clusters with a tin-oxide core and absorb more of the EUV light than chemically amplified resists [4]. In 2021 the Japanese resist maker JSR agreed to buy Inpria, an Oregon maker of these resists [5].

The other chemicals and gases put material on the wafer, take it off or keep it clean. Bulk gases arrive by the metric ton, specialty gases in grams. Silane gas carries the silicon for new layers into the chamber that holds the wafer. A plasma tears nitrogen trifluoride into charged fragments, which scrub the chamber walls clean between wafers. Hydrofluoric acid cleans the wafer itself. Polishing slurries grind each new layer flat before the next goes on.

Sputtering targets are slabs of pure metal. A plasma knocks atoms off them, and the atoms settle on the wafer as thin metal films, including the barrier and seed layers for copper wiring [6]. Electroplating then fills in the copper wiring itself [7]. Filters and wafer carriers keep contamination out.

Variants and trade-offs

RESIST CLASS	WAVELENGTH	WHERE IT IS USED	LIMIT
i-line novolak	365 nm	Packaging, mature nodes	Resolution
KrF chemically amplified	248 nm	Implant layers, thick films	Resolution
ArF immersion chemically amplified	193 nm	Most critical layers below 28 nm	Needs several exposures per layer
EUV chemically amplified	13.5 nm	Leading-edge single-exposure layers	Shot noise from too few photons
EUV metal-oxide	13.5 nm	High-NA tools, tightest pitches	Cost, defects, outgassing

Each resist class is made for one kind of lithography, such as EUV, ArF immersion, ArF, KrF or i-line [8].

4 Inpria, company website. Inpria [A] inpria.com

5 JSR Agrees to Acquire EUV Pioneer Inpria Corporation. JSR Corporation · 17 September 2021 [A] jsr.co.jp/jsr_e/news/2021/20210917.html

6 Applied Materials, PVD technology page. Applied Materials [A] appliedmaterials.com/eu/en/semiconductor/semiconductor-technologies/pvd.html

7 SABRE Product Family. Lam Research · 11 August 2017 [A] lamresearch.com/product/sabre

8 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 [A] cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

Who makes it

In a January 2021 brief, Georgetown University's Center for Security and Emerging Technology (CSET) put Japan's share of the semiconductor photoresist market at 90 percent, citing a 2019 estimate [8]. The brief put most of the rest with US and South Korean firms [8]. In a December 2025 strategy presentation, Japan's Ministry of Economy, Trade and Industry (METI) published figures from a separate survey by Fuji Keizai, a Japanese market research firm [2]. That survey counts sales by company nationality and puts Japanese companies at 77 percent in 2021 and 78 percent in 2023, with American and South Korean companies next [2].

Photoresist sales share by company nationality, 2023 %



Japanese companies 78% US, South Korean and other companies 22%

Source: METI, December 2025, from Fuji Keizai data

CSET's January 2021 brief lists five Japanese resist makers: JSR, Tokyo Ohka Kogyo, Shin-Etsu Chemical, Fujifilm and Sumitomo Chemical [8]. It also lists Dongjin of South Korea, the American firms DuPont, EMD Performance Materials and Inpria, and four Chinese firms [8]. Qnity Electronics, the former DuPont electronics business, says in its 2025 annual report that it offers lithography materials for every wavelength, including EUV [9]. In its 2025 annual report summary, the Chinese firm Hubei Dinglong said China had yet to reach large-scale production of advanced resist for 248 nm, 193 nm or EUV printing [10]. The firm also said three of its own resists for 248 nm and 193 nm printing had moved into stable volume supply [10].

8 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 A cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

2 METI, Semiconductor and digital industry strategy: future direction (半導体・デジタル産業戦略の今後の方向性), 23 December 2025. Ministry of Economy, Trade and Industry (Japan) · 23 December 2025 A meti.go.jp/policy/mono_info_service/joho/conference/semicon_digital/0014/handeji14-4.pdf

9 Qnity Electronics, Inc., Form 10-K annual report for the period ended 2025-12-31 (10-K). U.S. Securities and Exchange Commission (filing by Qnity Electronics, Inc.) · 26 February 2026 A sec.gov/Archives/edgar/data/2058873/000205887326000010/q-20251231.htm

10 湖北鼎龙控股股份有限公司2025年年度报告摘要. CNINFO · 26 March 2026 A static.cninfo.com.cn/finalpage/2026-03-27/1225036148.PDF

A vehicle of Japan Investment Corporation, a state-backed fund, bought 84.36 percent of JSR's shares at ¥4,350 each in a tender offer that closed on 16 April 2024 [11]. JSR's shares were delisted from the Tokyo Stock Exchange on 25 June 2024 [12].

Tokyo Ohka Kogyo sold ¥139,668 million in the first half of fiscal 2026, up 25.1 percent, as strong demand for AI servers in data centers raised demand for chips [13]. Fujifilm planned to raise its polishing-slurry capacity at Kumamoto by about 30 percent from January 2025 [14]. It also scheduled new production and quality-evaluation equipment for EUV resist at Shizuoka and Pyeongtaek to start in October 2025 [15].

World sales of wafer fabrication materials reached $45.8 billion in 2025, and sales of packaging materials reached $27.4 billion [16].

In electronic gases, CSET named Merck, Air Products and Air Liquide as the market leaders in January 2021 [8]. Air Liquide's electronics business sold €2,465 million in 2025, 9.1 percent of group revenue [1]. Smaller firms specialize in particular chemicals: SK Materials in nitrogen trifluoride, Kanto Denka in fluorine chemistry, and Stella Chemifa and Morita in ultra-high-purity hydrofluoric acid.

CSET's January 2021 brief, drawing on 2019 market reports, put five firms at more than 60 percent of the wet chemicals market [8]. It put DuPont and Cabot Microelectronics at about 56 percent of a $790 million market for polishing slurries [8]. Counting by company nationality, the same Fuji Keizai survey puts Japanese companies at 51 percent of polishing-slurry sales and 59 percent of high-purity cleaning-liquid sales in 2023 [2]. Both shares were 49 percent in 2021 [2].

11 Announcement Regarding Result of Tender Offer for JSR Corporation (Securities Code: 4185). Japan Investment Corporation · 17 April 2024 [A] jiccapital.co.jp/en/news/.assets/E_20240417_JIC_JICC_PressRelease.pdf

12 Announcement of Delisting of Shares of JSR Corporation. JSR Corporation · 19 June 2024 [A] jsr.co.jp/jsr_e/news/assets/pdf/20240624_01_e.pdf

13 Consolidated Financial Results for the Second Quarter (Interim) of the Fiscal Year Ending December 31, 2026 [J-GAAP]. Tokyo Ohka Kogyo · 5 August 2026 [A] tok.co.jp/application/files/7417/8598/0069/q2_2612_en.pdf

14 Fujifilm to Enhance its Production Capacity of Advanced Semiconductor Material CMP Slurries at the Kumamoto Site. Fujifilm · 5 December 2024 [A] fujifilm.com/jp/en/news/hq/11942

15 Fujifilm Launches EUV Resist and EUV Developer. Fujifilm · 29 October 2024 [A] fujifilm.com/jp/en/news/hq/11842

16 Global Semiconductor Materials Market Revenue Reaches Record $73.2 Billion in 2025, SEMI Reports. PR Newswire · 12 May 2026 [A] prnewswire.com/news-releases/global-semiconductor-materials-market-revenue-reaches-record-...

8 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 [A] cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

1 Air Liquide, electronics page. Air Liquide [A] airliquide.com/group/activities/electronics

2 METI, Semiconductor and digital industry strategy: future direction (半導体・デジタル産業戦略の今後の方向性), 23 December 2025. Ministry of Economy, Trade and Industry (Japan) · 23 December 2025 [A] meti.go.jp/policy/mono_info_service/joho/conference/semicon_digital/0014/handeji14-4.pdf

Cabot Microelectronics renamed itself CMC Materials in October 2020 [17]. The American firm Entegris completed its purchase of CMC Materials on 6 July 2022 [18]. DuPont separated its electronics business into an independent company, Qnity Electronics, on 1 November 2025 [19]. Qnity sells polishing slurries and pads [20].

In sputtering targets for semiconductors, JX Advanced Metals puts its own share of the world market at about 65 percent as of fiscal 2026 [21].

The chokepoint

Supply at this stage has been disrupted three times since 2019, twice by governments and once by a war.

In July 2019 Japan put photoresist, hydrogen fluoride and fluorinated polyimide under individual export licenses for South Korea, so each export contract needed its own license [22]. Japan had supplied more than 90 percent of South Korea's imports of two of the three [22]. Its hydrogen fluoride exports to South Korea then fell 87.9 percent [22]. Japanese photoresist exports to South Korea did not fall, possibly in part because Japan allowed three-year bulk licenses for some resist sales in December 2019 [22].

South Korea bought more hydrogen fluoride from the United States and Taiwan and more photoresist from Belgium [22]. Production of all three chemicals rose inside South Korea [22]. The rise came both from South Korean firms and from Japanese companies' affiliates in South Korea [22].

Researchers at Japan's Research Institute of Economy, Trade and Industry (RIETI) found that Japanese hydrogen fluoride exports to the United States rose [22]. They

17 CMC Materials, Inc., Form 8-K current report for the period ended 2020-09-28 (8-K). U.S. Securities and Exchange Commission (filing by CMC Materials, Inc.) · 1 October 2020 A sec.gov/Archives/edgar/data/1102934/000110293420000048/ccmp-20201001xex991.htm

18 ENTEGRIS INC, Form 8-K current report for the period ended 2022-07-06. U.S. Securities and Exchange Commission (filing by ENTEGRIS INC) · 6 July 2022 A sec.gov/Archives/edgar/data/1101302/000114036122025275/ny20004647x3_ex99-1.htm

19 DuPont Completes Separation of Qnity Electronics. DuPont · 3 November 2025 A investors.dupont.com/news-and-media/press-release-details/2025/DuPont-Completes-Separation...

20 Chemical Mechanical Planarization (CMP) Materials for Semiconductor. Qnity Electronics A qnityelectronics.com/semiconductor-cmp.html

21 Quick Guide to JX Advanced Metals. JX Advanced Metals A jx-nmm.com/english/company/glance

22 The impact of export controls on international trade: evidence from the Japan–Korea trade dispute in the semiconductor industry. Research Institute of Economy, Trade and Industry (Japan) · 8 May 2023 A rieti.go.jp/en/columns/v01_0201.html

say some Japanese supply may have reached South Korea by way of the United States [22]. Japan ended the individual licenses on 23 March 2023, when it put exports of the three chemicals to South Korea under its bulk export license system [23].

In 2022 Russia invaded Ukraine, which supplied about 70 percent of the world's neon and 90 percent of the semiconductor-grade neon used by US industry [24]. Gas mixtures containing neon power the lasers in lithography scanners [24]. Many chipmakers had added to their gas stockpiles after Russia annexed Crimea in 2014 [24]. Samsung, SK Hynix, Renesas and Rohm reported adequate neon supplies from China [24].

At this stage, China's export controls cover metals. On 3 July 2023 China announced that exports of gallium and germanium would need licenses from 1 August [25]. On 3 December 2024 it said it would in principle grant no licenses to export gallium, germanium, antimony or superhard materials to the United States [26]. On 9 November 2025 it suspended that ban until 27 November 2026 [27].

The White House set out a trade deal with China on 1 November 2025 [28]. Under it, China also agreed to issue general licenses for exports of gallium, germanium and antimony for the benefit of US end users and their suppliers [28].

In the United States, germanium goes mainly into fiber optics, infrared optics, semiconductors and solar cells, and radiation detectors [29]. Of the gallium used in the United States, 73 percent went into integrated circuits built on gallium

22 The impact of export controls on international trade: evidence from the Japan–Korea trade dispute in the semiconductor industry. Research Institute of Economy, Trade and Industry (Japan) · 8 May 2023 A rieti.go.jp/en/columns/v01_0201.html

23 Circular Notice on Partial Amendment to the Circular Notice on the Operation of the Cabinet Order on Export Trade Control and other circular notices (Review of Export Control to South Korea). Ministry of Economy, Trade and Industry (Japan) · 23 March 2023 A web.archive.org/web/20250408171837/https://www.meti.go.jp/english/press/2023/0323_001.html

24 Russia's Invasion of Ukraine Impacts Gas Markets Critical to Chip Production. Center for Strategic and International Studies A csis.org/blogs/perspectives-innovation/russias-invasion-ukraine-impacts-gas-markets-critic…

25 商务部 海关总署公告2023年第23号 关于对镓、锗相关物项实施出口管制的公告. Ministry of Commerce of the People's Republic of China A mofcom.gov.cn/zcfb/dwmygl/art/2023/art_52b9a321087f402bb3d310d18b07967e.html

26 商务部公告2024年第46号 关于加强相关两用物项对美国出口管制的公告. Ministry of Commerce of the People's Republic of China A mofcom.gov.cn/zwgk/zcfb/art/2024/art_3d5e990b43424e60828030f58a547b60.html

27 商务部公告2025年第72号 公布调整实施商务部公告2024年第46号的决定. Ministry of Commerce of the People's Republic of China A mofcom.gov.cn/zwgk/zcfb/art/2025/art_5c68985a6b1a46778e2e8dbff1bb1601.html

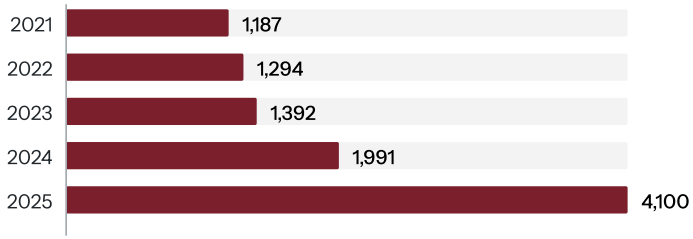
28 Fact Sheet: President Donald J. Trump Strikes Deal on Economic and Trade Relations with China. The White House · 1 November 2025 A whitehouse.gov/fact-sheets/2025/11/fact-sheet-president-donald-j-trump-strikes-deal-on-eco…

29 Mineral Commodity Summaries 2026: Germanium. U.S. Geological Survey · 5 February 2026 A pubs.usgs.gov/periodicals/mcs2026/mcs2026-germanium.pdf

arsenide or gallium nitride [30]. Another 26 percent went into optoelectronic devices such as LEDs and laser diodes [30].

Germanium metal averaged $4,100 a kilogram in 2025, against $1,392 in 2023 [29]. US imports of the metal fell 67 percent in 2025 because of China's ban [29]. The US Geological Survey estimated that a total Chinese ban on gallium and germanium exports could cut US GDP by $3.4 billion [31].

Germanium metal, annual average price ($/kg). China required export licenses from August 2023, banned shipments to the US in December 2024 and suspended the ban from November 2025 to 27 November 2026 $/KG



Source: USGS Mineral Commodity Summaries, February 2026 (Argus Media prices)

A slower risk is the European Union's proposed restriction on PFAS, the long-lived fluorinated chemicals. These chemicals are used throughout a fab. The photoacid generators that free the acid in resist are fluorinated, and so are etch gases, chamber liners, pump seals and filter membranes.

The national authorities of Denmark, Germany, the Netherlands, Norway and Sweden prepared the proposal and updated it after evaluating more than 5,600 comments [32]. ECHA, the EU chemicals agency, published the updated proposal on 20 August 2025 [32]. The updated proposal assessed options that would let electronics and semiconductor makers keep using PFAS where the risks can be controlled [32].

In March 2026 ECHA's two scientific committees backed an EU-wide restriction with targeted derogations, or exemptions [33]. One committee gave its backing in a final opinion and the other in a draft, and their evaluation is due to finish by the

30 Mineral Commodity Summaries 2026: Gallium. U.S. Geological Survey · 5 February 2026 A pubs.usgs.gov/periodicals/mcs2026/mcs2026-gallium.pdf

29 Mineral Commodity Summaries 2026: Germanium. U.S. Geological Survey · 5 February 2026 A pubs.usgs.gov/periodicals/mcs2026/mcs2026-germanium.pdf

31 USGS Critical Minerals Study: Bans on Gallium and Germanium Exports Could Cost the U.S. Billions. U.S. Geological Survey A usgs.gov/news/national-news-release/usgs-critical-minerals-study-bans-gallium-and-germaniu…

32 ECHA publishes updated PFAS restriction proposal. European Chemicals Agency A echa.europa.eu/-/echa-publishes-updated-pfas-restriction-proposal

33 ECHA supports PFAS restriction with targeted derogations. European Chemicals Agency · 26 March 2026 A echa.europa.eu/-/echa-supports-pfas-restriction-with-targeted-derogations

end of 2026 [33]. The European Commission will decide in consultation with the member states [32].

Key evaluation criteria

- **Purity, in parts per trillion** of metallic contamination, a million times finer than parts per million. This requirement keeps new suppliers out.
- **Qualification lock-in**: a resist or slurry is approved for one layer of one product at one fab, so switching suppliers means a new approval for every product that uses it.
- **Photon efficiency** for EUV resists: how much light the resist absorbs sets the dose, and the dose sets how many wafers an hour the scanner prints. Too few photons per dose raise the defect rate, which is the argument for metal-oxide resists.
- **Transportability**: bulk gases are often made on site at the fab, while specialty gases travel in cylinders and can be embargoed.
- **Byproduct exposure**: gallium is recovered mostly as a byproduct of processing bauxite, the ore that aluminum is made from [30]. Germanium comes from certain zinc ores and coal deposits, so supply of both depends on another industry's economics [29].

FOUR THINGS TO REMEMBER

- Photoresist holds the pattern printed by light so the next step can etch it into the wafer.
- Japanese companies held 77 percent of world photoresist sales in 2021 and 78 percent in 2023.
- Each resist is approved for one layer of one product at one factory, so switching to a new supplier takes years.
- China has required licenses to export gallium and germanium since August 2023.

33 ECHA supports PFAS restriction with targeted derogations. European Chemicals Agency · 26 March 2026 [A] echa.europa.eu/-/echa-supports-pfas-restriction-with-targeted-derogations

32 ECHA publishes updated PFAS restriction proposal. European Chemicals Agency [A] echa.europa.eu/-/echa-publishes-updated-pfas-restriction-proposal

30 Mineral Commodity Summaries 2026: Gallium. U.S. Geological Survey · 5 February 2026 [A] pubs.usgs.gov/periodicals/mcs2026/mcs2026-gallium.pdf

29 Mineral Commodity Summaries 2026: Germanium. U.S. Geological Survey · 5 February 2026 [A] pubs.usgs.gov/periodicals/mcs2026/mcs2026-germanium.pdf

CHAPTER 04
STAGE 4 OF 15

Lithography

ASML, a Dutch firm, is the only maker of extreme ultraviolet scanners, the machines that print the finest layers of a chip. It sold 48 of them in 2025.

IN PLAIN TERMS

Lithography prints the pattern of a circuit onto the wafer. A machine holds a stencil of one layer of the circuit up to a lamp and projects the image, four times smaller, onto the silicon. The shorter the wavelength of the light, the finer the lines a machine can print. A machine can still print lines several times narrower than its light's wavelength, and ASML says its 193-nanometer argon fluoride laser light prints features as small as 38 nanometers [3]. Air soaks up extreme ultraviolet light, the shortest in use, so machines that print with it work in a vacuum. Only ASML, a Dutch company, makes them. Export controls are government rules on who a company may sell to. Every advanced chip passes through ASML's machines, so export controls on ASML affect all advanced chipmaking.

IN SHORT

ASML also holds above 80 percent of the market for deep ultraviolet (DUV) lithography, the older kind that uses longer-wavelength light [8]. ASML says the production capacity of Zeiss, its sole supplier of optics, limits how many lithography systems it can build [13].

Export controls bar extreme ultraviolet (EUV) scanners from Chinese chip foundries [21]. In December 2025 Reuters reported a Chinese prototype that had made EUV light but had printed no chip [22]. In July 2026 Reuters reported, citing one unnamed source, that a Shanghai firm with two state shareholders has started producing immersion DUV tools, which shine their light through water [29].

3 All about light and lasers in lithography. ASML [A] asml.com/en/technology/lithography-principles/light-and-lasers

8 Explainer: The $400 million ASML 'printers' key for the AI chip boom. Reuters · 28 January 2026 [B] reuters.com/world/asia-pacific/250-million-asml-printer-behind-nvidias-chips-2026-07-28

13 ASML HOLDING NV, Form 6-K report of foreign private issuer for the period ended 2025-12-31 (6-K). U.S. Securities and Exchange Commission (filing by ASML HOLDING NV) · 25 February 2026 [A] sec.gov/Archives/edgar/data/937966/000162828026011377/asml-2025xannualxreportx.htm

21 Breakthroughs or Boasts? Assessing Recent Chinese Lithography Advancements. Center for Strategic and International Studies [A] csis.org/blogs/strategic-technologies-blog/breakthroughs-or-boasts-assessing-recent-chines…

22 How China built its 'Manhattan Project' to rival the West in AI chips. Reuters · 17 December 2025 [B] reuters.com/world/china/how-china-built-its-manhattan-project-rival-west-ai-chips-2025-12-17

29 China starts production of home-grown immersion DUV chipmaking tools, source says. Reuters · 28 July 2026 [B] reuters.com/world/china/china-starts-production-home-grown-immersion-duv-chipmaking-tools-…

CHOKEPOINT CARD | EXTREME CONCENTRATION | VERY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
Extreme	Very hard	About $200M
	ASML is the only maker of EUV scanners, and Zeiss is its sole supplier of their mirrors. A state-funded Chinese team has made the light but has not printed a chip.	for a standard EUV scanner, $350-400M for the newer High-NA version and about $60M for a DUV tool, according to Reuters. ASML does not publish per-system prices

WHO LEADS

- NL **ASML** 100% of EUV (extreme ultraviolet) and above 80% of DUV (deep ultraviolet). EUV was 48% of its own 2025 system revenue and immersion DUV 42%
- JP **Nikon** 22 new chip lithography systems in the year to March 2026, in a 2025 market Nikon estimates at 570 units
- JP **Canon** i-line and krypton fluoride tools for older, coarser layers, plus a nanoimprint system that Canon says was the first commercialized for chipmaking

WHERE IT IS MADE

- NL **Netherlands** ASML design and final assembly, Veldhoven
- DE **Germany** Zeiss SMT optics, Oberkochen; Trumpf carbon dioxide lasers that drive the light source, Ditzingen
- US **United States** ASML light-source research and manufacturing, San Diego
- JP **Japan** Nikon and Canon deep ultraviolet scanners and steppers

WHY SUBSTITUTION IS SLOW

Nobody else makes EUV scanners. A newcomer would first have to make mirrors a meter wide, polished smooth to within tens of picometers, a picometer being a trillionth of a meter. It would also need a light source that hits tin droplets with a laser 50,000 times a second. An attempt in China, reported in December 2025, had made extreme-ultraviolet light but had printed no chip. In September 2025, before that report, CSIS judged that EUV remained firmly out of reach for China despite massive state investment.

WHERE CHINA STANDS

Export controls block ASML from shipping EUV scanners to Chinese foundries. In September 2025 CSIS judged that EUV remained firmly out of reach for China despite massive state investment. That came before a December 2025 report of a Chinese prototype that made EUV light but had printed no chip. Reuters reported in July 2026 that Shanghai Aishengna Electronic Technology Group, a state-owned firm registered in 2023, has started producing Chinese immersion DUV scanners. SMEE, the older Chinese supplier, has its largest share in i-line scanners, the coarsest kind, and even there holds about 4 percent of the world market.

WHERE THE US STANDS

No American firm makes scanners. But the US government limits who may buy them, through rules on the American-made parts inside each machine and through its entity list of firms that may not be supplied.

ASML sold 327 lithography systems in 2025, 48 of them extreme ultraviolet (EUV) [1]. Nobody else makes EUV systems [2].

How it works

The stencil is called a photomask, and the machine a scanner (see Photomasks and pellicles).

The wafer is coated with photoresist, a film that changes wherever light touches it. The mask holds one layer of the circuit as a pattern of clear and dark areas, drawn at four times its size on the chip. The machine shines light through the mask, shrinks the image four times with a lens, and lands it on the film.

A wash then removes the film where the light hit, leaving the pattern standing on the wafer. The next machine etches that pattern into the wafer or fills it with metal. Then the film is stripped and the next mask goes in. A chip takes dozens of masks, one per layer.

The image from one mask covers a patch about 26 by 33 mm. The wafer moves under the lens one patch at a time, close to a hundred patches per wafer. ASML's own wafers-per-hour rating for an EUV scanner assumes 96 of them [4]. Before the wafer is exposed, the machine measures alignment marks printed in earlier layers so the new layer lines up with them [4]. ASML rates the NXE:3400B's matched-machine overlay, a measure of that alignment, at 2 nm [4].

Finer lines need shorter light. Glass and air both absorb the shortest light in use, at 13.5 nm [5]. So an EUV scanner works in a vacuum and uses mirrors instead of lenses, and its masks are mirrors too [5]. Older machines use light with a wavelength more than fourteen times longer [3]. Immersion scanners, one type of older machine, fill the gap between lens and wafer with water to sharpen the image [6].

When one exposure cannot draw a pattern finely enough, the layer is printed with two, three or four masks instead. Each extra pass costs a mask, machine time and a chance to misalign. ASML says one EUV exposure per layer cuts masks and process steps and lowers cost [2].

1 Financial performance: From ASML's Annual Report 2025. ASML · 24 February 2026 [A] ourbrand.asml.com/m/419103cb23dfeaa4/original/asml-2025-annual-report-financial-performanc…

2 Strategic report: From ASML's Annual Report 2025. ASML · 24 February 2026 [A] ourbrand.asml.com/m/8ab959d4926657b/original/asml-2025-annual-report-strategic-report-sect…

4 TWINSCAN NXE:3400B: EUV lithography systems. ASML [A] asml.com/en/products/euv-lithography-systems/twinscan-nxe3400b

5 ASML EUV lithography systems. ASML [A] asml.com/en/products/euv-lithography-systems

3 All about light and lasers in lithography. ASML [A] asml.com/en/technology/lithography-principles/light-and-lasers

6 Lenses & mirrors - Lithography principles. ASML [A] asml.com/en/technology/lithography-principles/lenses-and-mirrors

Variants and trade-offs

Deep ultraviolet

Deep ultraviolet (DUV) light comes from gas lasers: 248 nanometers from krypton fluoride (KrF), 193 from argon fluoride (ArF), shone through air or water [3]. Most of the machines sold are still DUV tools. Of the 327 systems ASML sold in 2025, 131 were ArF immersion, 78 KrF and 16 ArF dry, all deep ultraviolet, and another 54 were older i-line tools [7]. The i-line tools use 365 nm light from a mercury lamp, and ASML says they print features as small as 220 nm [3].

Reuters puts a DUV scanner at roughly $60 million [8]. In DUV, ASML competes with Nikon and Canon, and it still holds above 80 percent of that market [8]. Among DUV tools, export controls focus on immersion scanners, because without EUV a fab needs NXT:2000i-class scanners and several passes per layer to make leading-edge logic chips. Dutch export licensing covers these scanners.

Extreme ultraviolet

The 13.5 nm light comes from molten tin droplets shot from a generator [3]. A laser hits each droplet twice, first with a low-intensity pulse that flattens it and then with a more powerful one that vaporizes it into plasma [3]. The process repeats 50,000 times a second [3]. Mirrors made by Zeiss, a German optics company, carry the light to the wafer [6]. The largest are a meter across and smooth to within tens of picometers, a picometer being a thousandth of a nanometer [6].

Brighter light means more wafers an hour, and ASML ran the first 1,000-watt EUV source in April 2025 [2]. Its TWINSCAN NXE:3800E ships at its full 220 wafers an hour, 37 percent more than the NXE:3600D [2].

High-NA EUV

High-NA scanners raise the numerical aperture (NA), a measure of how well the optics collect and focus light, from 0.33 to 0.55 [9]. That cuts the smallest printable feature from 13 nm to 8 nm [9]. Their optics shrink the mask image four times in

3 All about light and lasers in lithography. ASML [A] asml.com/en/technology/lithography-principles/light-and-lasers

7 ASML 2025 fourth-quarter and full-year results (investor presentation). ASML · 28 January 2026 [A] ourbrand.asml.com/m/3136300aa4999bc1/original/2026_01_28_Presentation-Investor-Relations-Q...

8 Explainer: The $400 million ASML 'printers' key for the AI chip boom. Reuters · 28 January 2026 [B] reuters.com/world/asia-pacific/250-million-asml-printer-behind-nvidias-chips-2026-07-28

6 Lenses & mirrors - Lithography principles. ASML [A] asml.com/en/technology/lithography-principles/lenses-and-mirrors

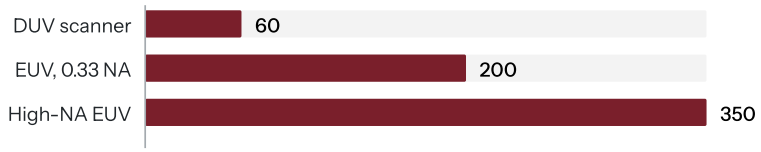
2 Strategic report: From ASML's Annual Report 2025. ASML · 24 February 2026 [A] ourbrand.asml.com/m/8ab959d4926657b/original/asml-2025-annual-report-strategic-report-sect...

9 5 things you should know about High NA in EUV. ASML · 25 January 2024 [A] asml.com/en/company/stories/2024/5-things-high-na-euv

one direction and eight in the other, halving the patch of wafer one exposure covers and doubling the exposures per wafer [9]. The second-generation TWINSCAN EXE:5200B runs at 175 wafers an hour, 60 percent above the EXE:5000 [2].

Reuters puts a standard EUV tool at around $200 million and a High-NA machine at $350 to $400 million, prices ASML does not publish [8].

What one lithography system costs. The High-NA bar shows the low end of a quoted $350-400M range $M



Source: Reuters explainer on ASML's lithography machines, 28 July 2026

Intel and TSMC have chosen differently on High-NA:

- **Intel Foundry** says it uses High-NA in volume production for select layers of some Intel 18A products [10]. It installed the first EXE:5200B [11]. By September 2026 it had run more than one million wafers on High-NA tools, a count that includes tool testing and research as well as volume production [10].
- **TSMC** waited, and in September 2026 said it will use High-NA in high-volume manufacturing from 2030 [12].

ASML sold four High-NA (EXE) systems in 2025, against two in 2024 [1]. In its 2025 annual report, filed in February 2026, ASML said it expected High-NA to start supporting high-volume manufacturing in 2027 [13]. On 8 September 2026 ASML and TSMC announced a move to 12-inch photomasks to remove stitching constraints [12]. Stitching constraints are the design limits that come from printing a large chip in two halves. They aim for a pilot mask line by 2031 and 12-inch High-NA systems in production by 2033 [12].

9 5 things you should know about High NA in EUV. ASML · 25 January 2024 [A] asml.com/en/company/stories/2024/5-things-high-na-euv

2 Strategic report: From ASML's Annual Report 2025. ASML · 24 February 2026 [A] ourbrand.asml.com/m/8ab959d4926657b/original/asml-2025-annual-report-strategic-report-sect…

8 Explainer: The $400 million ASML 'printers' key for the AI chip boom. Reuters · 28 January 2026 [B] reuters.com/world/asia-pacific/250-million-asml-printer-behind-nvidias-chips-2026-07-28

10 Intel Foundry and ASML Accelerate Industry Readiness for High-NA EUV. Intel · 7 September 2026 [A] intel.com/content/www/us/en/newsroom/news/intel-foundry/intel-foundry-asml-accelerate-indu…

11 High NA EUV reaches new readiness milestone with first high-volume Logic product. ASML · 15 July 2026 [A] asml.com/en/news/press-releases/2026/high-na-euv-reaches-new-readiness-milestone

12 ASML and TSMC Announce Initiative to Pioneer Industry Transition to Large-Format Photomasks for High NA EUV. TSMC · 8 September 2026 [A] pr.tsmc.com/english/news/3338

1 Financial performance: From ASML's Annual Report 2025. ASML · 24 February 2026 [A] ourbrand.asml.com/m/419103cb23dfeaa4/original/asml-2025-annual-report-financial-performanc…

13 ASML HOLDING NV, Form 6-K report of foreign private issuer for the period ended 2025-12-31 (6-K). U.S. Securities and Exchange Commission (filing by ASML HOLDING NV) · 25 February 2026 [A] sec.gov/Archives/edgar/data/937966/000162828026011377/asml-2025xannualxreportx.htm

EUV systems ASML sold each year TOOLS



Source: ASML 2025 annual report (2024 and 2025); 2023 from ASML's Q4 2023 investor presentation

Who makes it

ASML's net sales were EUR 32.7 billion in 2025, at a 52.8 percent gross margin [1]. In July 2026 it raised its 2026 forecast to EUR 43 to 45 billion, at a gross margin of 54 to 56 percent [14]. Its 2025 annual report lists 5,100 suppliers [13]. Four of them make the main parts of the EUV machine:

- **Zeiss SMT**, Oberkochen, makes the illumination system and the six-mirror projection optics [15].
- **Trumpf**, Ditzingen, makes the carbon dioxide drive laser, which amplifies a few watts to 40 kilowatts [16].
- **ASML San Diego**, which includes the light-source maker Cymer that ASML bought in 2013, designs the source where the laser turns tin into plasma, and builds the droplet generator [17].
- **VDL ETG** builds the frames that suspend and position the mirrors [18].

ASML says the production capacity of Zeiss, its sole supplier of optics, limits how many lithography systems it can build [13].

Nikon and Canon serve narrower markets. Nikon sold 22 new chip lithography systems in the year to March 2026, in a market it sizes at 570 units for 2025 [19]. Its Precision Equipment business made an operating loss in the year to March 2026 [19]. Canon's alternative is nanoimprint, where the FPA-1200NZ2C stamps the pattern instead of projecting it, down to a 14 nm linewidth [20]. Canon delivered

1 Financial performance: From ASML's Annual Report 2025. ASML · 24 February 2026 [A] ourbrand.asml.com/m/419103cb23dfeaa4/original/asml-2025-annual-report-financial-performanc...

14 Press Release Financial Results Q2 2026. ASML · 14 July 2026 [A] ourbrand.asml.com/asset/c8dbf3fc-4c5e-4406-83f6-27694b138245/Press-Release-Financial-Resul...

13 ASML HOLDING NV, Form 6-K report of foreign private issuer for the period ended 2025-12-31 (6-K). U.S. Securities and Exchange Commission (filing by ASML HOLDING NV) · 25 February 2026 [A] sec.gov/Archives/edgar/data/937966/000162828026011377/asml-2025xannualxreportx.htm

15 EUV lithography and technology. Carl Zeiss [A] zeiss.com/semiconductor-manufacturing-technology/inspiring-technology/euv-lithography.html

16 Good things come in ever-smaller packages. TRUMPF [A] trumpf.com/en_US/products/lasers/euv-drive-laser

17 Explore ASML San Diego. ASML [A] asml.com/en/company/about-asml/locations/san-diego

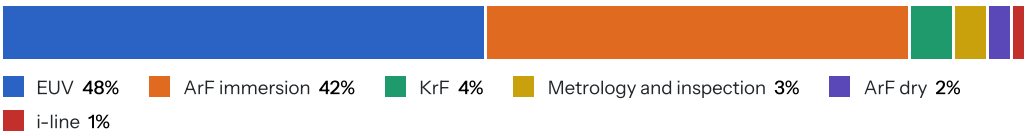
18 VDL ETG builds complex frames for Zeiss. VDL Groep · 22 March 2019 [A] web.archive.org/web/20240726075938/https://www.vdlgroep.com/en/news/vdl-etg-builds-complex...

19 Financial Results: The Year Ended March 31, 2026. Nikon · 8 May 2026 [A] nikon.com/content/dam/web-assets/nikoncom/company/local/global/en/ir/ir_library/result/pdf...

20 Canon delivers FPA -1200NZ2C nanoimprint lithography system for semiconductor manufacturing to the Texas Institute for Electronics. Canon · 26 September 2024 [A] global.canon/en/news/2024/20240926.html

one in 2024 to the Texas Institute for Electronics, to be used for research and prototypes [20].

ASML net system sales by technology, 2025 %



Source: ASML Q4 2025 investor presentation

Controls and China's answer

Export controls bar EUV from Chinese foundries [21]. In September 2025 CSIS judged that EUV remained "firmly out of reach" for China despite massive state investment [21]. That judgment came before Reuters reported a Chinese prototype in December 2025 [22]. For DUV scanners and other tools, the restrictions came in this order:

- **Dutch national licensing** of advanced DUV took effect on 1 September 2023 [23].
- **A US rule** published on 25 October 2023 put certain immersion scanners bound for advanced chipmaking under US export rules if they held any American content at all [24]. The rule did not apply if the country the scanners were first exported from already controlled them [24].
- **A partial license revocation** by the Dutch government, disclosed by ASML on 1 January 2024, stopped shipments of NXT:2050i and NXT:2100i immersion scanners to a few Chinese customers [25].
- **The Netherlands widened its licensing** from 7 September 2024 [23]. On 15 January 2025 it announced that from 1 April 2025 the license requirement would also cover a very limited number of further goods, among them specific measuring and inspection equipment [26].

20 Canon delivers FPA -1200NZ2C nanoimprint lithography system for semiconductor manufacturing to the Texas Institute for Electronics. Canon · 26 September 2024 A global.canon/en/news/2024/20240926.html

21 Breakthroughs or Boasts? Assessing Recent Chinese Lithography Advancements. Center for Strategic and International Studies A csis.org/blogs/strategic-technologies-blog/breakthroughs-or-boasts-assessing-recent-chines…

22 How China built its 'Manhattan Project' to rival the West in AI chips. Reuters · 17 December 2025 B reuters.com/world/china/how-china-built-its-manhattan-project-rival-west-ai-chips-2025-12-17

23 The Netherlands expands export control measure for advanced semiconductor manufacturing equipment. Government of the Netherlands · 6 September 2024 A government.nl/latest/news/2024/09/06/the-netherlands-expands-export-control-measure-advanc…

24 Export Controls on Semiconductor Manufacturing Items. Federal Register (Commerce Department; Industry and Security Bureau) · 25 October 2023 A federalregister.gov/documents/2023/10/25/2023-23049/export-controls-on-semiconductor-manuf…

25 Statement regarding partial revocation export license. ASML · 1 January 2024 A asml.com/en/news/press-releases/2023/statement-regarding-partial-revocation-export-license

26 Klever: export controls on advanced semiconductor manufacturing equipment to be tightened. Government of the Netherlands · 15 January 2025 A government.nl/latest/news/2025/01/15/klever-export-controls-on-advanced-semiconductor-manu…

- **The BIS Affiliates Rule**, issued by the US Bureau of Industry and Security, extended entity-list controls to affiliates owned 50 percent or more by listed entities from 29 September 2025 [27]. BIS then suspended it until 9 November 2026 [28].
- **The EU** added certain Dutch controls to the EU Control List in November 2025 [13]. ASML says this changed nothing for its immersion tools, because they already needed a Dutch license [13].

China's share of ASML's net system sales fell from 41 percent in 2024 to 33 percent in 2025 [7].

On DUV, Reuters reported in July 2026, citing one unnamed source, that Shanghai Aishengna Electronic Technology Group has started producing home-grown immersion DUV tools [29]. The firm was registered in August 2023 with 7 billion yuan of capital and two state shareholders, Shanghai Electric Holding and a Shanghai International Trust subsidiary [29]. It has no website and has taken in teams from SMEE, China's older scanner maker, and the startup Yuliangsheng [29]. SMEE is strongest in i-line, the coarsest kind of scanner, and even there holds about four percent of the world market [21].

On EUV there is less evidence. Claims of Huawei mass production in 2026 remain unverified. Reuters reported in December 2025 that a Shenzhen team of former ASML engineers had built a prototype that made extreme ultraviolet light but had printed no chip [22]. To print a chip, a machine also needs mirrors that carry that light to the wafer, and ASML's sole supplier of those mirrors is Zeiss [13].

27 Expansion of End-User Controls To Cover Affiliates of Certain Listed Entities. Federal Register (Commerce Department; Industry and Security Bureau) · 30 September 2025 [A] federalregister.gov/documents/2025/09/30/2025-19001/expansion-of-end-user-controls-to-cove…

28 One Year Suspension of Expansion of End-User Controls for Affiliates of Certain Listed Entities. Federal Register (Commerce Department; Industry and Security Bureau) · 12 November 2025 [A] federalregister.gov/documents/2025/11/12/2025-19846/one-year-suspension-of-expansion-of-en…

13 ASML HOLDING NV, Form 6-K report of foreign private issuer for the period ended 2025-12-31 (6-K). U.S. Securities and Exchange Commission (filing by ASML HOLDING NV) · 25 February 2026 [A] sec.gov/Archives/edgar/data/937966/000162828026011377/asml-2025xannualxreportx.htm

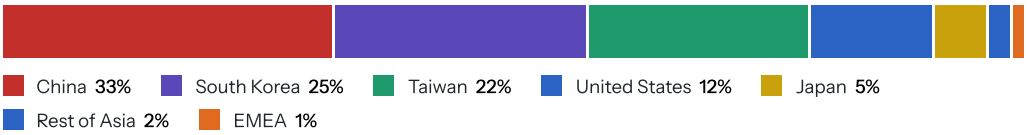
7 ASML 2025 fourth-quarter and full-year results (investor presentation). ASML · 28 January 2026 [A] ourbrand.asml.com/m/3136300aa4999bc1/original/2026_01_28_Presentation-Investor-Relations-Q…

29 China starts production of home-grown immersion DUV chipmaking tools, source says. Reuters · 28 July 2026 [B] reuters.com/world/china/china-starts-production-home-grown-immersion-duv-chipmaking-tools-…

21 Breakthroughs or Boasts? Assessing Recent Chinese Lithography Advancements. Center for Strategic and International Studies [A] csis.org/blogs/strategic-technologies-blog/breakthroughs-or-boasts-assessing-recent-chines…

22 How China built its 'Manhattan Project' to rival the West in AI chips. Reuters · 17 December 2025 [B] reuters.com/world/china/how-china-built-its-manhattan-project-rival-west-ai-chips-2025-12-17

ASML net system sales by the region tools shipped to, 2025 %



Source: ASML Q4 2025 investor presentation

The chokepoint

If ASML stopped shipping, no fab anywhere could add EUV scanners. Tools already in fabs would keep printing for as long as they were serviced. In its 2025 annual report ASML said about 95 percent of the systems it had sold in the previous 30 years were still in use [13]. It guarantees service and spare parts for each platform until at least a committed date, subject to export controls [13]. ASML itself could not quickly rebuild its production, because Zeiss makes its optics and Trumpf its drive laser.

In July 2026 ASML put its 2026 capacity at around 65 standard (low-NA) EUV scanners and said it plans to add 30 percent for 2027 [14].

For AI accelerators, EUV is the point of control, because it prints the finest layers, the ones that set how densely transistors and wires pack [5]. TSMC expects the number of layers needing High-NA to rise as AI designs grow more complex [12].

Key evaluation criteria

- **Resolution.** Set by the Rayleigh criterion, which puts the smallest printable feature at k1 times the wavelength divided by the numerical aperture [30]. The factor k1 depends on the chipmaking process and has a physical limit of 0.25 [30]. EUV reaches 13 nm at 0.33 NA and 8 nm at 0.55 NA [9].
- **Throughput.** Wafers an hour at a stated dose. The NXE:3800E is rated at 220 and the EXE:5200B at 175 [2].
- **Overlay.** How precisely one layer lands on the one below. ASML says the EXE:5200B's improved projection optics, developed with Zeiss, maximize

13 ASML HOLDING NV, Form 6-K report of foreign private issuer for the period ended 2025-12-31 (6-K). U.S. Securities and Exchange Commission (filing by ASML HOLDING NV) · 25 February 2026 [A] sec.gov/Archives/edgar/data/937966/000162828026011377/asml-2025xannualxreportx.htm

14 Press Release Financial Results Q2 2026. ASML · 14 July 2026 [A] ourbrand.asml.com/asset/c8dbf3fc-4c5e-4406-83f6-27694b138245/Press-Release-Financial-Resul…

5 ASML EUV lithography systems. ASML [A] asml.com/en/products/euv-lithography-systems

12 ASML and TSMC Announce Initiative to Pioneer Industry Transition to Large-Format Photomasks for High NA EUV. TSMC · 8 September 2026 [A] pr.tsmc.com/english/news/3338

30 What is the Rayleigh criterion?. ASML [A] asml.com/en/technology/lithography-principles/rayleigh-criterion

9 5 things you should know about High NA in EUV. ASML · 25 January 2024 [A] asml.com/en/company/stories/2024/5-things-high-na-euv

2 Strategic report: From ASML's Annual Report 2025. ASML · 24 February 2026 [A] ourbrand.asml.com/m/8ab959d4926657b/original/asml-2025-annual-report-strategic-report-sect…

imaging and overlay [2].

- **Field size.** The patch of wafer one exposure covers. High-NA halves it [9]. ASML and TSMC plan 12-inch masks to remove stitching constraints [12].
- **Source power.** Brighter light means more wafers an hour, and ASML demonstrated the first 1,000-watt EUV source in April 2025 [2].

FOUR THINGS TO REMEMBER

- Lithography prints each layer of a chip's circuit onto the wafer.
- The finest layers need extreme-ultraviolet (EUV) light, which air and glass absorb, so EUV machines work in a vacuum and use mirrors.
- ASML of the Netherlands is the only maker of EUV machines, and Zeiss alone supplies their mirrors.
- Export controls block EUV sales to Chinese chipmakers, and a Chinese prototype reported in December 2025 had printed no chip.

2 Strategic report: From ASML's Annual Report 2025. ASML · 24 February 2026 [A] ourbrand.asml.com/m/8ab959d4926657b/original/asml-2025-annual-report-strategic-report-sect…

9 5 things you should know about High NA in EUV. ASML · 25 January 2024 [A] asml.com/en/company/stories/2024/5-things-high-na-euv

12 ASML and TSMC Announce Initiative to Pioneer Industry Transition to Large-Format Photomasks for High NA EUV. TSMC · 8 September 2026 [A] pr.tsmc.com/english/news/3338

CHAPTER 05
STAGE 5 OF 15

Photomasks and Pellicles

Only three firms in the world make the blank plates for the masks used to print the finest layers of leading-edge chips. Japan's AGC and Hoya are two of them.

IN PLAIN TERMS

A photomask is the stencil that lithography prints from, one for each layer of a chip. The same mask prints its layer on wafer after wafer, so one mask shapes millions of chips. The most complicated layers are printed with extreme ultraviolet (EUV) light, and the rest can often be printed with deep ultraviolet light shone through a quartz plate. No glass lets EUV light through, so the mask for those layers is a mirror. The mirror is built up from about forty pairs of thin layers, and the light bounces off it. A dark pattern drawn on top of the mirror soaks up the light wherever no line should print. Any speck of dust on the mask repeats on every wafer it prints. Japan's AGC and Hoya are two of the three firms that make those mirrors.

IN SHORT

Hoya says its share of the mask blank market is exceptionally high [10]. On CSET's count, Austria's IMS Nanofabrication held 46.6 percent of the 2025 world market for electron-beam mask writers, which draw the pattern onto a mask [15].

SemiAnalysis put a full set of masks for a leading-edge chip at $40 million in May 2026 [27]. In July 2022 it wrote that more companies will lack the volume to spread the fixed cost of mask sets [26]. On 8 September 2026 ASML and TSMC announced a plan to move from 6-inch to 12-inch masks [8]. They expect the change to raise fab productivity and lower chipmaking costs [8].

10 HOYA Report 2025: Information Technology Business, Review of Operations. HOYA [A] hoya.com/ir/2025/en/review/it.html

15 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

27 EDA Market Primer. SemiAnalysis · 21 May 2026 [B] newsletter.semianalysis.com/p/eda-market-primer

26 The Dark Side Of The Semiconductor Design Renaissance – Fixed Costs Soaring Due To Photomask Sets, Verification, and Validation. SemiAnalysis · 24 July 2022 [B] newsletter.semianalysis.com/p/the-dark-side-of-the-semiconductor

8 ASML and TSMC Announce Initiative to Pioneer Industry Transition to Large-Format Photomasks for High NA EUV. TSMC · 8 September 2026 [A] pr.tsmc.com/english/news/3338

CHOKEPOINT CARD | EXTREME CONCENTRATION | HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
Extreme	Hard AGC says only three firms supply the blank plates for EUV masks, and Lasertec was the first to sell a tool that checks a finished mask in EUV light.	**SemiAnalysis put a leading-edge mask set at $40M in May 2026, and a 7 nm set above $10M in July 2022**

WHO LEADS

- JP **AGC** Ranks itself No.2 worldwide in EUV mask blanks on its own fiscal 2026 estimates, and counts three suppliers in all. In June 2026 it said its EUV blank sales reached ¥40bn in 2024, fell in 2025 and were recovering in 2026
- JP **Hoya** Says its share of mask blanks is exceptionally high, and aims to keep its leading share even if customers gradually start buying EUV blanks from more than one supplier
- JP **Lasertec** Launched the world's first tool that inspects patterned EUV masks at the scanner's own wavelength (2019), and the faster ACTIS A200HiT series (October 2025). JPY 230,485 million of net sales in the year to June 2026
- AT **IMS Nanofabrication** In 2023 Intel called it the established industry leader in multi-beam mask writers, the machines that draw the pattern onto the mask. CSET puts its 2025 share of the $994M market for all electron-beam mask writers at 46.6%, ahead of NuFlare (39.2%) and JEOL (14.2%)
- US **Photronics** Calls itself one of the world's leading photomask makers, with $849.3M of revenue in fiscal 2025

WHERE IT IS MADE

- JP **Japan** Mask blanks (AGC, Hoya), DUV blanks (Shin-Etsu), inspection (Lasertec), writers (NuFlare), pellicles (Mitsui Chemicals)
- AT **Austria** IMS Nanofabrication multi-beam mask writers, Brunn am Gebirge and Vienna
- DE **Germany** Zeiss mask checking in EUV light (AIMS EUV) and mask repair (MeRiT)
- US **United States** Photronics merchant mask shops, KLA inspection
- TW **Taiwan** TSMC's in-house mask shop, and a 10% stake in IMS at the end of 2024

WHY SUBSTITUTION IS SLOW

An EUV mask starts as a blank plate coated with about forty pairs of layers. AGC says it spent more than a decade meeting the blank's requirements, and mass production began only in 2017. Three firms have done it, and the structure of the plate is published, so a newcomer knows what to build. Inspection is harder. In 2019 Lasertec launched the world's first tool that checks a patterned mask in the same EUV light the scanner uses. Germany's Zeiss calls its AIMS the only system that tests how a mask will print under conditions equivalent to the scanner's.

WHERE CHINA STANDS

SMIC runs two mask shops of its own in Shanghai, for g-line, i-line, DUV and ArF steppers and scanners. Photronics names Shenzhen Qingyi Photomask and Shenzhen Newway Photomask among its competitors. Since 2019 Photronics has itself run two mask plants in China, in Hefei and Xiamen. Its Xiamen chip mask business is a joint venture with Dai Nippon Printing. No Chinese firm makes an EUV mask blank, a tool that checks a mask in EUV light, or a multi-beam writer, the machine that draws the finest patterns onto a mask.	WHERE THE US STANDS Photronics calls itself one of the world's leading photomask makers, and KLA makes mask inspection tools. But no American firm makes an EUV mask blank or a tool that checks a mask in EUV light.

Each EUV layer of an AI accelerator is printed from a mask built on a blank plate. Only three companies worldwide supply those plates, and two of them are Japan's AGC and Hoya [1]. AGC says it is the only company in the world that takes an EUV blank all the way from the glass to the finished coating [2].

How it works

ASML says a chip's most complicated layers are made on EUV machines, while the rest can often be printed on older deep ultraviolet machines [3]. A mask starts as a blank, a polished plate with nothing on it yet. For deep ultraviolet the blank is a window of quartz carrying a light-blocking film such as molybdenum silicide, and Shin-Etsu supplies those blanks [4]. For EUV the blank is a mirror of about 40 molybdenum and silicon layer pairs, 7 nm to a pair, laid on glass [5]. A dark tantalum film on top, called the absorber, will carry the pattern [5].

The mask shop coats the blank with a sensitive film called resist. A mask writer then draws the design into the resist with electrons. A single-beam writer flashes one shape at a time, while NuFlare's multi-beam writer for the most advanced masks controls 260,000 beams at once [6].

A wash removes the drawn resist. An etch then removes the blocking film or absorber wherever the resist is gone, and the rest of the resist is stripped. The plate now passes light, or reflects it, in the exact shape of one layer of the chip. On standard scanners the pattern is drawn four times larger than it will print. An inspection tool compares the plate to the design, and a repair tool fixes the stray defects it finds.

Making the plate a mirror creates its own problems:

- **Mask 3D effects.** Light strikes the mirror at a slant, and the absorber has height, so its edges cast small shadows that distort the printed shape. Software has to correct for the shadows in advance.

1 The Japanese glass behind next-generation chips. AGC [A] agc.com/en/hub/pr/the-japanese-glass-behind-next-generation-chips.html

2 AGC to Boost Production Capacity of EUVL Photomask Blanks. AGC · 27 April 2023 [A] agc.com/en/news/detail/1203819_2814.html

3 See ASML's DUV lithography systems. ASML [A] asml.com/en/products/duv-lithography-systems

4 Photomask blanks. Shin-Etsu Chemical · 3 April 2019 [A] shinetsu.co.jp/en/products/electronics-materials/photomask-blanks

5 The refined EUV mask model. arXiv · 19 December 2019 [A] arxiv.org/pdf/1912.09075

6 EB Mask Writer. NuFlare Technology [A] nuflare.co.jp/english/products/beam

- **Write precision.** Every edge of the absorber has to sit within a few nanometers of where the design says.
- **Particles.** Anything on the mirror prints on every wafer until someone finds it.

High-NA scanners, the newest EUV machines, print finer lines with larger mirrors. To keep masks at their usual size, their mirrors shrink the pattern four times in one direction and eight times in the other [7]. So each exposure covers half the area of a standard exposure, and a wafer takes twice as many exposures [7].

On 8 September 2026 ASML and TSMC announced a plan to move from the current 6-inch masks to 12-inch masks [8]. They expect the change to raise fab productivity, lower chipmaking costs and remove stitching constraints, the design limits that come from printing a large chip in two halves [8]. They aim to set up a 12-inch mask pilot line by 2031, so that High-NA scanners using 12-inch masks can enter advanced chip production by 2033 [8].

Variants and trade-offs

Blanks

AGC says it spent more than a decade meeting the blank's requirements, which include near-zero thermal expansion, perfectly smooth and flat surfaces and no particle contamination [1]. It says mass production began only in 2017 [1].

AGC publishes no market share figure for EUV blanks. It says it is one of three companies worldwide that supply them, and the only one that handles every material and process itself [1]. In its data book, AGC ranks itself No.2 worldwide in EUV lithography photomask blanks, based on its own fiscal 2026 estimates [9]. Hoya says its share of the mask blank market is exceptionally high, and that customers may gradually start buying EUV blanks from more than one supplier [10].

In April 2023 AGC said its EUV blank capacity would grow by about 30 percent in 2025 [2]. The growth would come from a line in Fukushima, a project that METI, Japan's industry ministry, selected under a program to strengthen supply chains [2].

7 5 things you should know about High NA in EUV. ASML · 25 January 2024 A asml.com/en/company/stories/2024/5-things-high-na-euv

8 ASML and TSMC Announce Initiative to Pioneer Industry Transition to Large-Format Photomasks for High NA EUV. TSMC · 8 September 2026 A pr.tsmc.com/english/news/3338

1 The Japanese glass behind next-generation chips. AGC A agc.com/en/hub/pr/the-japanese-glass-behind-next-generation-chips.html

9 AGC Data Book. AGC Inc. · 29 May 2026 A agc.com/en/ir/library/outline/pdf/reference.pdf

10 HOYA Report 2025: Information Technology Business, Review of Operations. HOYA A hoya.com/ir/2025/en/review/it.html

2 AGC to Boost Production Capacity of EUVL Photomask Blanks. AGC · 27 April 2023 A agc.com/en/news/detail/1203819_2814.html

AGC also said it was aiming for more than ¥40 billion of sales in the business by 2025 [2]. Its EUV blank sales grew 50 percent in 2023, AGC said in December 2024 [1]. In February 2025 AGC reported that the sales reached ¥40 billion in 2024, a year ahead of the 2025 target [11].

The sales then declined in 2025 and were recovering in 2026, AGC said in June 2026 [12]. In August 2026 AGC set out capacity expansions of 20 to 50 percent across EUV blanks, polishing slurry and lens materials for lithography tools [13]. It gave no separate figure for the blanks, and mass production on the added capacity is to begin in 2028 [13].

Writers

An electron-beam writer draws the pattern into the resist on the blank. Multi-beam writers fire many beams at once, and Japan's NuFlare sells them for the 3 nm node [6].

In September 2023 Intel called Austria's IMS Nanofabrication the established industry leader in the multi-beam mask writing tools that advanced EUV requires [14]. Counting all electron-beam mask writers, CSET, a research center at Georgetown University, put IMS's share of a $994 million world market at 46.6 percent in 2025 [15]. NuFlare had 39.2 percent and JEOL 14.2 percent [15].

Intel agreed in September 2023 to sell roughly 10 percent of IMS to TSMC at a valuation of about $4.3 billion [14]. TSMC's subsidiary TSMC Development held that 10 percent at the end of 2024, at a fair value of $432.8 million in TSMC's accounts [16]. Intel sold a combined 32 percent of IMS in 2023 and still held the other 68 percent on 27 December 2025 [17].

Inspection and repair

A mask defect cuts the number of good chips on every wafer that mask prints, so fabs check masks three times: as blanks, after patterning and again in production.

2 AGC to Boost Production Capacity of EUVL Photomask Blanks. AGC · 27 April 2023 [A] agc.com/en/news/detail/1203819_2814.html

1 The Japanese glass behind next-generation chips. AGC [A] agc.com/en/hub/pr/the-japanese-glass-behind-next-generation-chips.html

11 Financial Results for FY2024. AGC Inc. · 7 February 2025 [A] agc.com/en/ir/library/briefing/pdf/2025_0207e_1.pdf

12 Semiconductor-related Business Briefing: Event Summary. AGC Inc. · 2 June 2026 [A] agc.com/en/ir/library/bizbriefing/pdf/2026_0602esemiconductor_transcript.pdf

13 AGC Company Overview. AGC Inc. · August 2026 [A] agc.com/en/ir/library/outline/pdf/c_overview.pdf

6 EB Mask Writer. NuFlare Technology [A] nuflare.co.jp/english/products/beam

14 Intel to Sell Minority Stake in IMS Nanofabrication Business to TSMC. Intel · 12 September 2023 [A] intc.com/news-events/press-releases/detail/1645/intel-to-sell-minority-stake-in-ims-nanofa…

15 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

16 TAIWAN SEMICONDUCTOR MANUFACTURING CO LTD, Form 6-K report of foreign private issuer for the period ended 2024-12-31 (6-K). U.S. Securities and Exchange Commission (filing by TAIWAN SEMICONDUCTOR MANUFACTURING CO LTD) · 27 February 2025 [A] sec.gov/Archives/edgar/data/1046179/000104617925000019/tsmc2024q4consolidatedfina.htm

17 INTEL CORP, Form 10-K annual report for the period ended 2025-12-27 (10-K). U.S. Securities and Exchange Commission (filing by INTEL CORP) · 23 January 2026 [A] sec.gov/Archives/edgar/data/50863/000005086326000011/intc-20251227.htm

Actinic inspection uses 13.5 nm light, the wavelength of EUV lithography [18]. In 2019 Lasertec launched the world's first actinic inspection system for patterned EUV masks, the ACTIS A150 [18]. In October 2025 it launched the ACTIS A200HiT series, which it says triples the A150's inspection speed [18]. Lasertec reported net sales of JPY 230,485 million in the year to June 2026 [19].

KLA sells optical inspection tools for both EUV masks and optical masks [20]. Zeiss makes AIMS EUV, which checks a mask in EUV light under the scanner's own imaging conditions to judge how a defect will affect printing [21]. Zeiss says the tool can be extended to ASML's High-NA scanner, the EXE:5000 [21]. Zeiss also sells electron-beam repair tools that work on EUV masks [22].

Pellicles

A pellicle is a thin membrane on a frame above the mask. Particles land on it, too far out of focus to print, and the pattern below stays clean.

ASML's EUV pellicle is a membrane 13 nanometers thick that has to hold together at 500 degrees Celsius in vacuum [23]. Every photon it absorbs is light that never reaches the wafer [23]. The light crosses it twice, on the way in and on the way back off the mirror, so absorption counts double. ASML released its first EUV pellicles in 2016 and calls the pellicle one of the breakthroughs that brought EUV lithography to high-volume manufacturing [23].

Carbon nanotube membranes are the next step, aimed at next-generation High-NA, high-output EUV scanners [24]. In 2020 imec, a Belgian chip research center, mounted them on masks and exposed them in an ASML NXE:3300 scanner [25]. The

18 Lasertec Releases ACTIS A200HiT Series Actinic EUV Patterned Mask Inspection System. Lasertec [A] lasertec.co.jp/en/news/2025/20251031_3912.html

19 Lasertec, business report. Lasertec [A] lasertec.co.jp/en/ir/plan/message.html

20 KLA-Tencor Introduces Five Patterning Control Systems for Sub-7nm IC Manufacturing. KLA Corporation · 11 September 2017 [A] ir.kla.com/news-events/press-releases/detail/95/kla-tencor-introduces-five-patterning-cont…

21 Mask Qualification Solutions by ZEISS SMT. Carl Zeiss [A] zeiss.com/semiconductor-manufacturing-technology/products/photomask-solutions/mask-qualifi…

22 Mask Repair Solutions by ZEISS SMT. Carl Zeiss [A] zeiss.com/semiconductor-manufacturing-technology/products/photomask-solutions/mask-repair.…

23 Indistinguishable from magic: the EUV pellicle. ASML · 14 September 2022 [A] asml.com/en/news/stories/2022/the-euv-pellicle-indistinguishable-from-magic

24 Mitsui Chemicals Sets up Production Facilities for CNT Pellicles to Be Used in Next-Gen EUV Lithography. Mitsui Chemicals [A] jp.mitsuichemicals.com/en/release/2024/2024_0528_1/index.htm

25 Imec demonstrates CNT pellicle utilization on EUV scanner. imec · 6 October 2020 [A] imec-int.com/en/press/imec-demonstrates-cnt-pellicle-utilization-euv-scanner

membranes passed up to 97 percent of the light in a single crossing, and imec found their effect on imaging low and correctable [25].

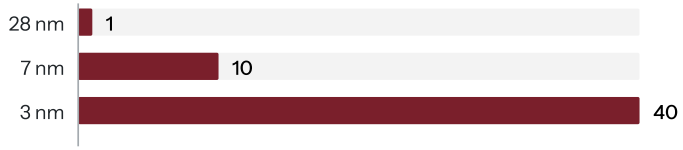
Mitsui Chemicals began commercial production of silicon-based EUV pellicles under an ASML license in 2021 [24]. In 2024 it said it would build a carbon nanotube membrane plant at Iwakuni-Ohtake, tentatively due to finish in December 2025 [24]. The membranes are designed to pass 92 percent or more of the light and to withstand exposure power above 1 kW [24].

Cost and design economics

In July 2022 SemiAnalysis put mask sets above $1 million at 28 nm and above $10 million at 7 nm [26]. It expected 3 nm mask sets to reach the $40 million range [26]. Its May 2026 primer still puts a leading-edge mask set at $40 million [27]. The same primer puts a single respin, remaking a chip after a design flaw, at $50 million to $100 million at leading-edge nodes [27].

Mask costs restrict which firms can use a node. SemiAnalysis wrote that at 5 nm and 3 nm, only firms with large wafer volume see the cost per transistor fall [26]. It wrote that more companies will lack the volume to spread the fixed cost of mask sets [26]. It expected failed chip designs to show up first among AI chip startups, which work on much tighter budgets [26].

What one mask set costs, by node, 2022 ($M). The 28 nm and 7 nm figures are minimums and 3 nm was a projection $M



Source: SemiAnalysis, The Dark Side of the Semiconductor Design Renaissance, July 2022. SemiAnalysis's May 2026 EDA Market Primer puts a leading-edge mask set at $40M

25 Imec demonstrates CNT pellicle utilization on EUV scanner. imec · 6 October 2020 [A] imec-int.com/en/press/imec-demonstrates-cnt-pellicle-utilization-euv-scanner

24 Mitsui Chemicals Sets up Production Facilities for CNT Pellicles to Be Used in Next-Gen EUV Lithography. Mitsui Chemicals [A] jp.mitsuichemicals.com/en/release/2024/2024_0528_1/index.htm

26 The Dark Side Of The Semiconductor Design Renaissance – Fixed Costs Soaring Due To Photomask Sets, Verification, and Validation. SemiAnalysis · 24 July 2022 [B] newsletter.semianalysis.com/p/the-dark-side-of-the-semiconductor

27 EDA Market Primer. SemiAnalysis · 21 May 2026 [B] newsletter.semianalysis.com/p/eda-market-primer

Who makes it

Mask making splits between chipmakers that make masks for their own use and independent merchant firms that sell to them and others [28]. Photronics says that chipmakers' own shops once supplied most of the market, and that production has more recently tended to move toward the independents [28].

Photronics calls itself a leading worldwide maker of masks for chips and flat panel displays [29]. It had $849.3 million of revenue in fiscal 2025, $615.1 million of it from chip masks [29]. The competitors it names in chip and display masks include Dai Nippon Printing, Hoya, Tekscend Photomask and two Chinese firms, Shenzhen Newway Photomask and Shenzhen Qingyi Photomask [28]. It also competes with chipmakers' own mask shops [28].

SMIC, a Chinese foundry, runs two in-house mask shops in Shanghai, for g-line, i-line, DUV and ArF steppers and scanners [30]. Photronics began operating two mask plants in China, in Hefei and Xiamen, in 2019 [28]. Its chip mask business in Xiamen is a joint venture in which Dai Nippon Printing holds 49.99 percent [28]. No Chinese firm makes an EUV blank, an actinic inspection tool or a multi-beam writer.

The chokepoint

Mask blanks are a small, highly concentrated business. Three firms supply EUV blanks worldwide, AGC says, and AGC spent more than a decade meeting the blank's requirements before mass production began in 2017 [1]. AGC said in June

28 PHOTRONICS INC, Form 10-K annual report for the period ended 2025-10-31 (10-K). U.S. Securities and Exchange Commission (filing by PHOTRONICS INC) · 17 December 2025 A sec.gov/Archives/edgar/data/810136/000114036125045801/ef20057458_10k.htm

29 Photronics Reports Full Year and Fourth Quarter Fiscal 2025 Results. GlobeNewswire (release by Photronics) · 10 December 2025 A globenewswire.com/news-release/2025/12/10/3203040/0/en/Photronics-Reports-Full-Year-and-Fo…

30 Mask Service. SMIC A smics.com/en/site/mask

1 The Japanese glass behind next-generation chips. AGC A agc.com/en/hub/pr/the-japanese-glass-behind-next-generation-chips.html

2026 that its EUV blank sales reached ¥40 billion in 2024, declined in 2025 and were recovering in 2026 [12]. In 2019 Lasertec launched the first tool that inspects patterned EUV masks in EUV light [18].

Pellicles already work in high-volume EUV manufacturing [23]. But ASML says no material is completely transparent to EUV light, so any pellicle in the light's path costs customers productivity [23]. In 2024 Mitsui Chemicals said it would set up production of carbon nanotube membranes for next-generation High-NA, high-output EUV scanners [24].

Key evaluation criteria

- **Blank defectivity.** Particles buried in the mirror layers cannot be repaired, so blank yield sets mask yield.
- **Write time and placement.** Intel said in 2023 that advanced EUV requires multi-beam mask writing tools, and called IMS the established industry leader in them [14]. In 2025 IMS held 46.6 percent of the market for all electron-beam mask writers, on CSET's count [15].
- **Actinic inspectability.** Actinic inspection checks a mask in 13.5 nm light, the same wavelength as EUV lithography, and Lasertec launched the first such tool for patterned masks in 2019 [18].
- **Pellicle transmission and lifetime.** Light crosses the pellicle twice, and in 2020 imec measured carbon nanotube membranes passing up to 97 percent of it in a single crossing [25].
- **Mask-set cost per node.** In July 2022 SemiAnalysis put a set above $10M at 7 nm, a cost that limits which designs can use a node [26]. In May 2026 it put a leading-edge set at $40M [27].

FOUR THINGS TO REMEMBER

— A photomask is the stencil for one layer of a chip, and any flaw on it repeats on every wafer.

12 Semiconductor-related Business Briefing: Event Summary. AGC Inc. · 2 June 2026 [A] agc.com/en/ir/library/bizbriefing/pdf/2026_0602esemiconductor_transcript.pdf

18 Lasertec Releases ACTIS A200HiT Series Actinic EUV Patterned Mask Inspection System. Lasertec [A] lasertec.co.jp/en/news/2025/20251031_3912.html

23 Indistinguishable from magic: the EUV pellicle. ASML · 14 September 2022 [A] asml.com/en/news/stories/2022/the-euv-pellicle-indistinguishable-from-magic

24 Mitsui Chemicals Sets up Production Facilities for CNT Pellicles to Be Used in Next-Gen EUV Lithography. Mitsui Chemicals [A] jp.mitsuichemicals.com/en/release/2024/2024_0528_1/index.htm

14 Intel to Sell Minority Stake in IMS Nanofabrication Business to TSMC. Intel · 12 September 2023 [A] intc.com/news-events/press-releases/detail/1645/intel-to-sell-minority-stake-in-ims-nanofa…

15 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

25 Imec demonstrates CNT pellicle utilization on EUV scanner. imec · 6 October 2020 [A] imec-int.com/en/press/imec-demonstrates-cnt-pellicle-utilization-euv-scanner

26 The Dark Side Of The Semiconductor Design Renaissance – Fixed Costs Soaring Due To Photomask Sets, Verification, and Validation. SemiAnalysis · 24 July 2022 [B] newsletter.semianalysis.com/p/the-dark-side-of-the-semiconductor

27 EDA Market Primer. SemiAnalysis · 21 May 2026 [B] newsletter.semianalysis.com/p/eda-market-primer

- Three firms make EUV mask blanks, and Lasertec launched the first tool that checks finished masks in EUV light.
- AGC spent more than a decade learning to make EUV blanks before mass production began in 2017.
- China makes masks only for older machines, in SMIC's own mask shops and at independent firms such as Shenzhen Qingyi and Shenzhen Newway.

CHAPTER
06
STAGE 6 OF 15

Deposition and Etch

Deposition and etch tools build a chip layer by layer. They are among the chipmaking tools where Chinese toolmakers gained the most share from 2019 to 2024.

IN PLAIN TERMS

Deposition and etch build the chip itself, one layer at a time. Deposition lays down a film a few atoms thick. Lithography prints a pattern on it. Etch then eats away everything the pattern does not protect, and the next film goes down on top. Making one chip takes 1,000 to 1,500 steps, the Semiconductor Industry Association estimated in 2025 [3]. Some of the holes the etch cuts are more than 50 times deeper than they are wide, like a finger-wide shaft about a meter deep [4]. The film has to coat them right to the bottom. The etch has to eat one material and leave the one beside it untouched. Applied Materials and Lam Research in America, Tokyo Electron in Japan and ASM International in the Netherlands make most of these machines. Chinese toolmakers have gained share in both kinds of tool since 2019.

IN SHORT

American companies held 60.3 percent of the world market for deposition tools in 2025, Japanese companies 14.0 percent and China-based companies 9.8 percent [8]. In etch and cleaning tools that year, American companies held 52.6 percent, Japanese companies 32.6 percent and China-based companies 10.1 percent [26].

Atomic layer deposition tools coat a hole's walls and floor as evenly as its rim [7]. China-based companies held only 2.6 percent of the world market for them in 2025 [13]. The December 2024 US export rule targets advanced atomic layer deposition tools and the extremely deep, narrow holes that high-bandwidth memory and 3D memory need [18].

3 Testimony of David Isaacs, Semiconductor Industry Association, hearing on “Trade in Critical Supply Chains”. U.S. Senate Committee on Finance · 14 May 2025 [A] finance.senate.gov/download/051425-isaacs-testimony&download=1

4 Lam Research Introduces Lam Cryo™ 3.0 Cryogenic Etch Technology to Accelerate Scaling of 3D NAND for the AI Era. Lam Research · 31 July 2024 [A] newsroom.lamresearch.com/2024-07-31-Lam-Research-Introduces-Lam-Cryo-TM-3-0-Cryogenic-Etch…

8 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

26 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

7 Export Controls on Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 25 October 2023 [A] govinfo.gov/content/pkg/FR-2023-10-25/pdf/2023-23049.pdf

13 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

18 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 5 December 2024 [A] govinfo.gov/content/pkg/FR-2024-12-05/pdf/2024-28270.pdf

CHOKEPOINT CARD

HIGH CONCENTRATION

MODERATELY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
High	Moderate Most deposition and etch steps have three or four established toolmakers, and atomic layer deposition has at least five. China's own toolmakers held 2.6 percent of atomic layer deposition in 2025.	Toolmakers publish no list prices World sales of chipmaking equipment were $135.1B in 2025, with wafer processing tools up 12%.

WHO LEADS

- US **Applied Materials** $28.4B revenue, fiscal 2025; leader in physical and chemical vapor deposition and in epitaxy
- US **Lam Research** $23.2B revenue, fiscal 2026; leader in dry etch and in metal atomic layer deposition
- JP **Tokyo Electron** 23% of dry etch and 38% of tube and low-pressure CVD, 2025
- NL **ASM International** €3.2B revenue, 2025; leader in atomic layer deposition
- CN **Naura** 14.7% of PVD and 5.9% of dry etch, 2025; China's largest toolmaker

WHERE IT IS MADE

- US **United States** Applied Materials (Santa Clara), Lam Research (Fremont)
- JP **Japan** Tokyo Electron, Kokusai Electric, Hitachi High-Tech
- NL **Netherlands** ASM International, Almere; ALD and epitaxy
- KR **South Korea** Large buyer; Jusung, Wonik IPS and Eugene in atomic layer deposition and PSK in dry strip, 2025; Semes in wet processing
- CN **China** Naura, AMEC, Piotech; largest equipment market at $49.3B in 2025

WHY SUBSTITUTION IS POSSIBLE

For most deposition and etch steps, a fab can already buy the tool from three or four makers: Applied Materials, Lam, Tokyo Electron and ASM International. Switching makers means testing the new tool on the fab's own line until it matches the old tool's results. Atomic layer deposition, the most advanced kind of deposition, lays down a film a fraction of a layer of atoms at a time. At least five makers sell it, led by ASM International with 54.5 percent in 2025. China's own toolmakers held 2.6 percent of it in 2025, against 12.4 percent of dry etch.

WHERE CHINA STANDS

Deposition and etch tools are among those where Chinese toolmakers gained the most share between 2019 and 2024. They held 12.4 percent of dry etch, 9.8 percent of deposition and 2.6 percent of atomic layer deposition in 2025, according to CSET's analysis of TechInsights data. CSET's earlier analysis

had them under 1 percent of atomic layer deposition in 2024.

WHERE THE US STANDS

Two of the three largest deposition and etch toolmakers are American. The US government lists some dry etch, atomic layer deposition and deep-hole deposition tools as controlled exports. Its December 2024 foreign direct product rule reaches some tools built abroad with American technology. The rule largely exempts foreign-made tools exported by companies in Japan, the Netherlands and 31 other countries that it lists.

Applied Materials and Lam Research are the two largest deposition and etch suppliers. Applied Materials booked $28.4 billion of revenue in fiscal 2025 [1]. Lam Research booked $23.2 billion in fiscal 2026 [2].

How it works

Most etch steps, and many deposition steps, happen in a vacuum chamber that holds one wafer at a time. A pump empties the chamber of air, and a measured flow of gas comes in. A radio-frequency field, an electric field that reverses millions of times a second, turns the gas into a plasma. The field splits the gas into charged ions and into neutral fragments called radicals, which are more reactive than whole molecules [5]. Batch furnaces such as Tokyo Electron's TELINDY PLUS instead heat as many as 125 wafers at once, usually without plasma, for oxidation, low-pressure chemical vapor deposition and some atomic layer deposition [6].

For deposition, the gas carries the material the film will be made of. The molecules, or the radicals a plasma splits from them, land on the wafer and stick. The film builds from the bottom up, a few atoms at a time.

To coat the floor of a deep hole as thickly as its rim, the machine can work in cycles, a method called atomic layer deposition. First one gas settles in a single layer on every surface, and the pump removes the rest. Then a second gas reacts with that layer and nothing else, and the pump empties the chamber again. Each cycle adds a fraction of a layer of atoms wherever the gas reached, so the film grows as evenly on a hole's walls and floor as on its rim [7].

For etch, the gas attacks one material and leaves the one beside it. Wherever the photoresist pattern covers the wafer, nothing happens. Wherever it does not, the

1 Revenue from Contract with Customer, Excluding Assessed Tax (us-gaap:RevenueFromContractWithCustomerExcludingAssessedTax) reported by APPLIED MATERIALS INC /DE, XBRL company concept data. U.S. Securities and Exchange Commission (XBRL data for APPLIED MATERIALS INC /DE) A data.sec.gov/api/xbrl/companyconcept/CIK0000006951/us-gaap/RevenueFromContractWithCustomer...

2 LAM RESEARCH CORP, Form 8-K current report for the period ended 2026-07-29 (8-K). U.S. Securities and Exchange Commission (filing by LAM RESEARCH CORP) · 29 July 2026 A sec.gov/Archives/edgar/data/707549/000070754926000033/lrcx_exhibitx991xq4x2026.htm

5 Etch Essentials: The Building Blocks of AI Era Microchips (Semi 101). Lam Research A newsroom.lamresearch.com/etch-essentials-semiconductor-manufacturing

6 TELINDY™ Series. Tokyo Electron A tel.com/product/telindy.html

7 Export Controls on Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 25 October 2023 A govinfo.gov/content/pkg/FR-2023-10-25/pdf/2023-23049.pdf

radicals react with the surface and turn it into a gas, which the pump carries away. A second field pulls the ions straight down at the wafer, so the etch cuts vertically and a narrow hole keeps its width as it deepens [5]. When the etch reaches the material underneath, which the gas was chosen not to touch, it slows almost to a stop, and the tool ends the step.

Deposition and etch chambers are separate products built from similar vacuum, gas and radio-frequency parts, and the export controls describe each tool by the process it is designed to run [7].

Accenture and the Global Semiconductor Alliance estimated in 2020 that a part-finished chip can cross international borders 70 or more times on its way to the end customer [3]. The Semiconductor Industry Association repeated the figure in 2025 [3].

Variants and trade-offs

Deposition families

The market sizes and leading shares below are 2025 figures from Georgetown University's Center for Security and Emerging Technology (CSET), based on data from the research firm TechInsights. CSET puts the deposition tool market at $26.2 billion [8]. It puts the dry etch market at $21.9 billion [9].

- **Physical vapor deposition** (PVD), which knocks atoms off a solid target and lands them on the wafer, was a $4.9 billion market in 2025 [10].
- **Chemical vapor deposition** (CVD), which reacts gases at the wafer surface to grow a film, was an $11.2 billion market in 2025 [11]. The plasma-enhanced version (PECVD) runs cooler and made up most of it, $7.5 billion [12].
- **Atomic layer deposition** (ALD) lays down less than one atomic layer at a time, so it coats vertical walls as evenly as horizontal floors. US regulators say that even coating makes it important for building 3D structures, such as 3D DRAM, 3D NAND, and the FinFET and gate-all-around transistors in logic chips [7]. ASM International leads it with 54.5 percent of a $4.1 billion market in 2025 [13].
- **Epitaxy** grows layers that line up with the crystal beneath them, using high-temperature CVD tools [14]. **Electroplating** deposits the copper used for wiring across the wafer [15].

5 Etch Essentials: The Building Blocks of AI Era Microchips (Semi 101). Lam Research [A] newsroom.lamresearch.com/etch-essentials-semiconductor-manufacturing

7 Export Controls on Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 25 October 2023 [A] govinfo.gov/content/pkg/FR-2023-10-25/pdf/2023-23049.pdf

3 Testimony of David Isaacs, Semiconductor Industry Association, hearing on “Trade in Critical Supply Chains”. U.S. Senate Committee on Finance · 14 May 2025 [A] finance.senate.gov/download/051425-isaacs-testimony&download=1

8 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

9 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

10 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

11 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

12 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

13 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

14 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

15 The Semiconductor Supply Chain - Issue Brief. Center for Security and Emerging Technology (CSET) · 21 January 2021 [A] cset.georgetown.edu/wp-content/uploads/The-Semiconductor-Supply-Chain-Issue-Brief-1.pdf

Etch families

- **Conductor etch**, which cuts gates, metals and silicon, was a $12.2 billion market in 2025, led by Lam at 51.2 percent [16].
- **Dielectric etch**, which cuts the insulating oxides and nitrides between them, was an $8.7 billion market in 2025, led by Tokyo Electron at 55.9 percent [17].
- **High-aspect-ratio etch** cuts the deep, narrow channels that let 3D NAND stack layers of memory cells [4]. Lam says slight departures from a channel's intended shape can hurt yield [4]. Lam runs this etch well below 0 degrees Celsius [4]. It says its Cryo 3.0 technology keeps the deviation in a channel's width from top to bottom under 0.1 percent, in channels up to 10 microns deep [4].

AI hardware needs sideways cuts, deeper holes and a second stack of metal wiring on the back of the wafer.

- **Gate-all-around** transistors need highly selective isotropic etch, which cuts sideways [7]. The US Bureau of Industry and Security (BIS) controls isotropic dry etch tools that remove silicon-germanium at least 100 times faster than silicon, a selectivity of 100:1 or more [7].
- **HBM** (high-bandwidth memory) needs holes etched straight through the silicon, called through-silicon vias (TSVs), so chips can be stacked and wired [18]. BIS controls the tools that etch these holes at aspect ratios (depth to width) of 10:1 or more and at rates above 7 microns a minute [18].
- **Backside power delivery** moves the wires that carry power to the back of the wafer [19]. Applied Materials expects it to add another $1 billion per 100,000 wafer starts a month to the wiring market it serves [20].

16 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

17 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

4 Lam Research Introduces Lam Cryo™ 3.0 Cryogenic Etch Technology to Accelerate Scaling of 3D NAND for the AI Era. Lam Research · 31 July 2024 A newsroom.lamresearch.com/2024-07-31-Lam-Research-Introduces-Lam-Cryo-TM-3-0-Cryogenic-Etch…

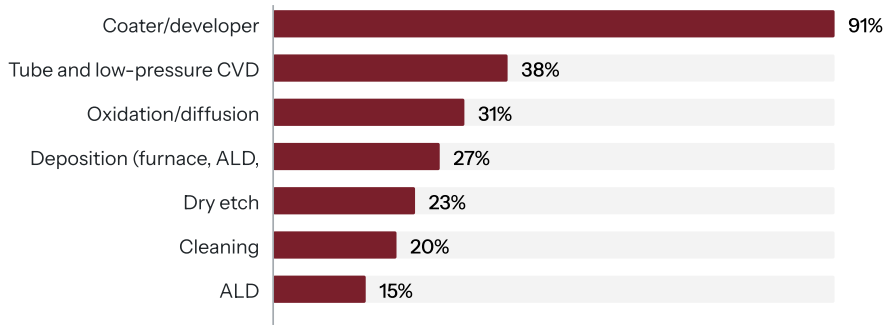
7 Export Controls on Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 25 October 2023 A govinfo.gov/content/pkg/FR-2023-10-25/pdf/2023-23049.pdf

18 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 5 December 2024 A govinfo.gov/content/pkg/FR-2024-12-05/pdf/2024-28270.pdf

19 PowerVia Test Shows Industry-Leading Performance. Intel · 31 January 2025 A download.intel.com/newsroom/archive/2025/en-us-2023-06-05-powervia-test-shows-industrylead…

20 Applied Materials Unveils Chip Wiring Innovations for More Energy-Efficient Computing. GlobeNewswire (release by Applied Materials) · 8 July 2024 A globenewswire.com/news-release/2024/07/08/2909540/0/en/applied-materials-unveils-chip-wiri…

Tokyo Electron world market share by tool type, 2025 %



Source: Tokyo Electron FY2026 Q4 results presentation

Who makes it

The industry association SEMI put world sales of chipmaking equipment at $135.1 billion in 2025, up 15 percent from $117.1 billion [22]. Sales of wafer processing equipment, the category that includes deposition and etch tools, rose 12 percent [22].

No one firm dominates deposition and etch the way ASML dominates lithography. Tokyo Electron's own results presentation gives it 23 percent of dry etch, 27 percent of deposition systems (furnace, ALD and low-pressure CVD tools) and 15 percent of ALD in 2025 [21]. The same presentation gives it 91 percent of coater/developers in 2025 [21]. These machines coat wafers with photoresist and develop the pattern after lithography.

ASM International, the atomic layer deposition specialist, reached a record EUR 3.2 billion of revenue in 2025 as customers built 2 nm gate-all-around capacity [23]. Molybdenum ALD, a new use of its tools that deposits the metal molybdenum, entered volume production that year [23].

22 SEMI Reports Global Semiconductor Equipment Billings Reached $135 Billion in 2025, Up 15% Year-on-Year. SEMI · 7 April 2026 [A] semi.org/en/SEMI-Reports-Global-Semiconductor-Equipment-Billings-Reached-135-Billion-in-2025

21 FY2026 (April 2025-March 2026) Financial Announcement, transcript. Tokyo Electron · 30 April 2026 [A] tel.com/ir/library/report/pjuomj00000000tf-att/fy26q4transcript-e.pdf

23 ASM reports Q4 and full-year 2025 results. ASM International · 3 March 2026 [A] asm.com/media/yvxbavwe/20260303-asm-reports-q4-and-full-year-2025-results.pdf

What the world spent on chipmaking equipment $B



Source: SEMI, Worldwide Semiconductor Equipment Market Statistics

China and the controls

China is the tool vendors' largest market, and Chinese toolmakers have gained share since 2019 [24]. It bought $49.3 billion of equipment in 2025, half a percent less than in 2024, and remained the biggest single market [22]. Taiwan's spending rose 90 percent to $31.5 billion [22]. China still takes a large share of every major vendor's sales.

- **Tokyo Electron**: 34.1 percent of sales in the year to March 2026, down from 41.7 percent [21].
- **Lam Research**: 26 percent in the June 2026 quarter, behind Taiwan at 27 percent [2].
- **Applied Materials**: 28 percent in the third fiscal quarter of 2026, down from 35 percent [25].
- **ASM International**: more than 30 percent of 2025 revenue [23].

Between 2019 and 2024 China-based suppliers went from 2 to 7 percent of deposition and from under 3 to almost 9 percent of etch and clean [24]. In ALD they gained no share and held less than 1 percent in 2024 [24].

CSET's newer analysis of TechInsights data, published in July 2026, puts them at 9.8 percent of deposition in 2025 [8]. They held 10.1 percent of etch and clean that

24 Inside Beijing's Chipmaking Offensive. Center for Security and Emerging Technology (CSET) · 14 July 2025 A cset.georgetown.edu/article/inside-beijings-chipmaking-offensive

22 SEMI Reports Global Semiconductor Equipment Billings Reached $135 Billion in 2025, Up 15% Year-on-Year. SEMI · 7 April 2026 A semi.org/en/SEMI-Reports-Global-Semiconductor-Equipment-Billings-Reached-135-Billion-in-2025

21 FY2026 (April 2025-March 2026) Financial Announcement, transcript. Tokyo Electron · 30 April 2026 A tel.com/ir/library/report/pjuomj00000000tf-att/fy26q4transcript-e.pdf

2 LAM RESEARCH CORP, Form 8-K current report for the period ended 2026-07-29 (8-K). U.S. Securities and Exchange Commission (filing by LAM RESEARCH CORP) · 29 July 2026 A sec.gov/Archives/edgar/data/707549/000070754926000033/lrcx_exhibitx991xq4x2026.htm

25 Applied Materials Announces Third Quarter 2026 Results. GlobeNewswire (release by Applied Materials) · 13 August 2026 A globenewswire.com/news-release/2026/08/13/3344890/0/en/applied-materials-announces-third-q...

23 ASM reports Q4 and full-year 2025 results. ASM International · 3 March 2026 A asm.com/media/yvxbavwe/20260303-asm-reports-q4-and-full-year-2025-results.pdf

8 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

year [26]. In dry etch alone their share was 12.4 percent [9]. In ALD they held 2.6 percent in 2025, most of it from Piotech [13].

The Chinese toolmaker AMEC alone held 12.4 percent of dielectric etch in 2025 [17]. Its revenue was RMB 12.385 billion that year, up 36.6 percent, and it spent RMB 3.744 billion on research, 30.2 percent of sales [27].

The United States tightened its export controls at least five times between 2022 and 2025.

- **October 2022.** The 7 October interim final rule required a license for certain chipmaking equipment, and for some activities by US persons that support the making of certain chips in China [28]. It also widened foreign-produced item rules for 28 Chinese entities already on the Entity List, the US list of foreign firms that exporters need a license to supply [28].
- **October 2023.** The October rule, effective 17 November, rewrote the equipment entries [7]. It added sideways (isotropic) and vertical (anisotropic) dry etch (ECCN 3B001.c.1), spatial ALD (3B001.d.9), low-fluorine tungsten ALD and CVD (3B001.d.10) and carbon hard mask PECVD (3B001.d.5) [7].
- **December 2024.** BIS controlled 24 more types of equipment and three types of software, and added 140 entities to the Entity List [29]. The new equipment entries include TSV etch, and deposition into features with aspect ratios above 200:1 [18]. BIS also wrote a foreign direct product rule that reaches some tools

26 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

9 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

13 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

17 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

27 AMEC holds its 2025 annual results briefing (中微公司成功举办2025年度业绩说明会). Advanced Micro-Fabrication Equipment (AMEC) · 1 April 2026 A amec-inc.com/news/708.html

28 Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items; Supercomputer and Semiconductor End Use; Entity List Modification. U.S. Government Publishing Office (Federal Register) · 13 October 2022 A govinfo.gov/content/pkg/FR-2022-10-13/pdf/2022-21658.pdf

7 Export Controls on Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 25 October 2023 A govinfo.gov/content/pkg/FR-2023-10-25/pdf/2023-23049.pdf

29 Commerce Strengthens Export Controls to Restrict China's Capability to Produce Advanced Semiconductors for Military Applications. U.S. Bureau of Industry and Security · 2 December 2024 A bis.gov/press-release/commerce-strengthens-export-controls-restrict-chinas-capability-prod…

18 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 5 December 2024 A govinfo.gov/content/pkg/FR-2024-12-05/pdf/2024-28270.pdf

built outside the United States [29]. The rule largely exempts foreign-made tools exported by firms located in Japan, the Netherlands and the 31 other countries it lists, whose governments can control these tools themselves [18]. Japan added advanced chipmaking equipment to its control list on 23 July 2023 [18].

- **September 2025.** BIS removed the Chinese units of Samsung and SK hynix, and Intel Semiconductor (Dalian), from the Validated End-User program, effective 31 December 2025 [30]. From that date, US tools bound for them need a license [30]. BIS estimated that the change would add 1,000 license applications a year [30].
- **September 2025.** The Affiliates Rule extended Entity List restrictions to any company at least 50 percent owned by a listed entity [31]. BIS then suspended it until 9 November 2026 [32].

The rules apply to the toolmakers as well as to their customers. Applied Materials took a $253 million charge to settle an export controls compliance matter with BIS [25].

Of the tools in the chart below, China-based suppliers have their largest share in dry stripping, which clears leftover photoresist off the wafer and was a $314.6 million market in 2025 [33]. Mattson Technology held 22.6 percent of it and Naura 8.3 percent [33]. Mattson is a US-based firm that Beijing E-Town bought in 2016 [24].

29 Commerce Strengthens Export Controls to Restrict China's Capability to Produce Advanced Semiconductors for Military Applications. U.S. Bureau of Industry and Security · 2 December 2024 A bis.gov/press-release/commerce-strengthens-export-controls-restrict-chinas-capability-prod…

18 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 5 December 2024 A govinfo.gov/content/pkg/FR-2024-12-05/pdf/2024-28270.pdf

30 Revocation of Validated End-User Authorizations in the People's Republic of China. U.S. Government Publishing Office (Federal Register) · 2 September 2025 A govinfo.gov/content/pkg/FR-2025-09-02/html/2025-16735.htm

31 Expansion of End-User Controls To Cover Affiliates of Certain Listed Entities. U.S. Government Publishing Office (Federal Register) · 30 September 2025 A govinfo.gov/content/pkg/FR-2025-09-30/pdf/2025-19001.pdf

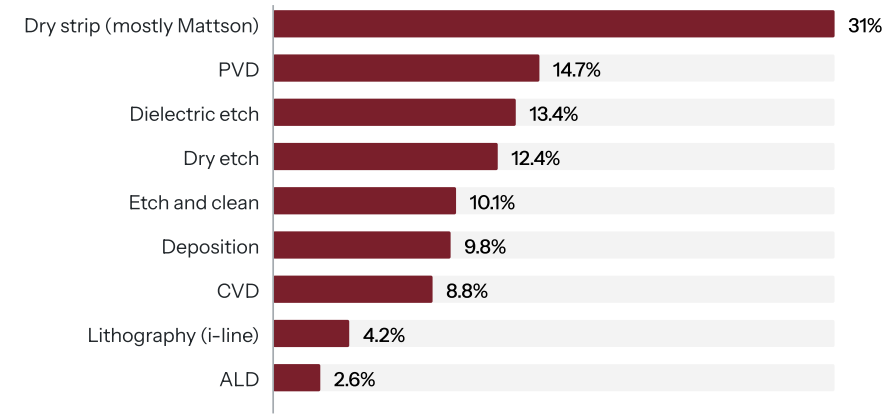
32 One Year Suspension of Expansion of End-User Controls for Affiliates of Certain Listed Entities. U.S. Government Publishing Office (Federal Register) · 12 November 2025 A govinfo.gov/content/pkg/FR-2025-11-12/pdf/2025-19846.pdf

25 Applied Materials Announces Third Quarter 2026 Results. GlobeNewswire (release by Applied Materials) · 13 August 2026 A globenewswire.com/news-release/2026/08/13/3344890/0/en/applied-materials-announces-third-q…

33 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

24 Inside Beijing's Chipmaking Offensive. Center for Security and Emerging Technology (CSET) · 14 July 2025 A cset.georgetown.edu/article/inside-beijings-chipmaking-offensive

China-based suppliers' share of global tool segments, 2025 %



Source: CSET analysis of TechInsights data (published July 2026), ETO Chip Explorer

The chokepoint

Deposition and etch tools are a chokepoint, but a weaker one than lithography. Most deposition and etch steps have three or four credible vendors. Atomic layer deposition has at least five, led by ASM International at 54.5 percent in 2025 [13]. China's own toolmakers held only 2.6 percent of atomic layer deposition in 2025 [13]. The December 2024 rule targets advanced atomic layer deposition tools and the extremely deep, narrow holes that HBM and 3D memory need [18].

Key evaluation criteria

- **Selectivity.** How hard the process attacks one material and spares another. BIS controls wet processing at a silicon-germanium-to-silicon ratio of 100:1 [7].
- **Conformality and aspect ratio.** Whether a film reaches the bottom of a hole, and how deep that hole is against its width. Controls start at 10:1 for TSV etch and 200:1 for 3D DRAM deposition [18].
- **Uniformity.** How much results vary across a 300 mm wafer. The TSV etch control is set at within-wafer etch depth nonuniformity of 2 percent or less [18].
- **Throughput.** Wafers per hour, which sets the cost per layer and how many lots a fab can run.
- **Particles and chamber matching.** How many stray particles these tools leave on a wafer, each a possible defect, and whether one chamber behaves like the next. Metrology and inspection tools measure both (see Metrology and inspection).

13 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

18 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 5 December 2024 A govinfo.gov/content/pkg/FR-2024-12-05/pdf/2024-28270.pdf

7 Export Controls on Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 25 October 2023 A govinfo.gov/content/pkg/FR-2023-10-25/pdf/2023-23049.pdf

FOUR THINGS TO REMEMBER

- Deposition lays down thin films and etch cuts them into patterns. Making one chip takes 1,000 to 1,500 steps, by a 2025 industry estimate.
- Four American, Japanese and Dutch firms make most of these tools, and none dominates.
- Chinese makers held 12.4 percent of dry etch and 9.8 percent of deposition in 2025, on CSET's count.
- Chinese makers held 2.6 percent of atomic layer deposition in 2025, most of it from Piotech.

CHAPTER 07
STAGE 7 OF 15

Metrology and Inspection

KLA, an American firm, held 56.8 percent of the market for tools that measure and inspect wafers in 2025 by CSET's count, about five times its nearest rival.

IN PLAIN TERMS

Metrology (measurement) and inspection are a chip factory's quality control. Between the steps that build a chip, another machine looks the wafer over. It checks whether the lines are the right width, whether this layer landed squarely on the last one, and whether a speck of dirt has killed a circuit. The main chip in an AI accelerator, the one that does the calculating, is one of the largest cut from a wafer. A bigger chip is a bigger target for a speck. On a chip that size, halving the stray specks lifts the share of flawless chips from about half to about seven in ten. In 2025 one American firm, KLA, held about five times the market share of its nearest rival. So a chip factory, or fab, cut off from KLA would find fewer of the faults that make its chips fail.

IN SHORT

By the count of CSET, a Georgetown University research center, defect inspection was the largest kind of measuring and inspection tool in 2025, a $6.7 billion market [1]. KLA held 85.9 percent of it [1]. Chinese makers held 0.7 percent of the whole market that year by the same count [1]. Yet China took 26 percent of KLA's revenue in the June 2026 quarter [2].

Since October 2022 the US has required a license for any US-made tool bound for Chinese fabs that make advanced logic or memory chips [20]. A December 2024 rule added specific tools to the export control list, among them inspection tools that can find defects of 21 nm or smaller [21].

1 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

2 Letter to Shareholders Q4 Fiscal 2026. KLA Corporation · 27 July 2026 [A] d1io3yog0oux5.cloudfront.net/_b9ee755a5f60dd0fb3f9e27967aed6af/klatencor/db/1117/10668/let…

20 Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items; Supercomputer and Semiconductor End Use; Entity List Modification. U.S. Government Publishing Office (Federal Register) · 13 October 2022 [A] govinfo.gov/content/pkg/FR-2022-10-13/pdf/2022-21658.pdf

21 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 5 December 2024 [A] govinfo.gov/content/pkg/FR-2024-12-05/pdf/2024-28270.pdf

CHOKEPOINT CARD

HIGH CONCENTRATION

MODERATELY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
High	Moderate KLA is the largest maker of measuring and inspecting tools, but every kind has a second maker, so losing KLA would slow a fab without stopping it.	**A $15.7B process control tool market in 2025, about $8.9B of it KLA's (56.8%)** KLA's total fiscal 2026 revenue of $13.579B also includes services (22 percent of revenue in the June 2026 quarter), PCB and component inspection, and specialty semiconductor process tools. KLA puts the 2025 wafer equipment market at about $120B.

WHO LEADS

- US **KLA** 56.8% of process control tools in 2025 by CSET's count, about 5 times Applied Materials; Gartner put KLA at 56.5% in 2024; $13.579B total revenue, fiscal 2026
- US **Applied Materials** 11.2% of process control tools in 2025: electron-beam defect review, line-width measurement and electron-beam metrology
- JP **Lasertec** 6.4% of process control tools in 2025, almost all of it mask inspection
- JP **Hitachi High-Tech** 6.2% of process control tools in 2025, most of it CD-SEM line-width measurement, where it holds 71.9%
- NL **ASML** 5.9% of process control tools in 2025: YieldStar layer alignment and HMI electron-beam metrology
- IL **Nova** 4.3% of process control tools in 2025; $880.6M revenue, 2025
- US **Onto Innovation** 2.6% of process control tools in 2025; $1,005.3M revenue, 2025

WHERE IT IS MADE

- US **United States** KLA (Milpitas), Onto Innovation (Wilmington, MA), Applied Materials
- NL **Netherlands** ASML YieldStar and HMI e-beam, Veldhoven and San Jose
- JP **Japan** Hitachi High-Tech CD-SEM; Lasertec mask inspection; Rigaku X-ray
- IL **Israel** Nova (Rehovot) and Camtek (Migdal Haemek)
- CN **China** Skyverse (Shenzhen), RSIC and other domestic makers; 0.7% of the market in 2025 by CSET's count, though Skyverse's own 2025 revenue came to about 1.8%

WHY SUBSTITUTION IS POSSIBLE

Metrology and inspection tools measure and inspect wafers so a fab can find and fix faults in its process. KLA is the biggest maker, but eight other firms each hold more than 1 percent of the market. Applied Materials and Hitachi each lead at least one kind of tool. A fab that loses KLA and buys from the others keeps running. But it finds fewer flaws and fixes them more slowly. A country barred from buying from any of them has to

rely on its own makers. Chinese firms held 0.7 percent of the market for these tools in 2025 by CSET's count.

WHERE CHINA STANDS

Chinese makers are still small in metrology and inspection, but one of them is growing fast. CSET's count, built on TechInsights data, puts all Chinese firms at 0.7 percent of the market for these tools in 2025, against 0.2 percent of lithography tools. Skyverse, a Shenzhen maker of inspection and metrology tools, reported 2.05 billion yuan of revenue in 2025, about $285 million and up 48.75 percent on 2024. Skyverse's revenue alone is about 1.8 percent of CSET's market, so the CSET count may miss some domestic sales. In its 2025 annual report Skyverse says Chinese makers still hold a low share of China's own market for these tools, where KLA holds the largest share. CSET's July 2025 survey of Chinese share gains from 2019 to 2024 covers deposition, etch, polishing, lithography, packaging and test, but leaves out measurement and inspection.

WHERE THE US STANDS

KLA is the largest maker of metrology and inspection tools in the world. CSET's Supply Chain Explorer puts it at 56.8 percent of the 2025 market, about five times Applied Materials. Citing Gartner, KLA puts its 2024 share at 56.5 percent, or 6.5 times its nearest rival. In 2025 American firms held 71.7 percent of this market and Japanese firms 14.1 percent. Since October 2022 the US has required a license for any US tool bound for China's most advanced fabs. A December 2024 US rule added specific metrology and inspection tools to the export control list. It names defect inspection tools that find defects 21 nm across or smaller, high-accuracy overlay and wafer shape metrology, and film and line-width metrology for non-planar transistors.

Process control tools, the machines that measure and inspect wafers, were a $15.7 billion market in 2025 [1]. KLA puts the whole wafer equipment market, including advanced packaging, at about $120 billion that year [2].

CSET's count for 2025, built on TechInsights data, puts KLA at 56.8 percent, about five times Applied Materials' 11.2 percent [1]. At its March 2026 investor day KLA cited Gartner's figures for 2024, which put its share at 56.5 percent, 6.5 times its nearest competitor [3]. In April 2026 KLA said its 2025 share was about 7 times the nearest competitor's [4]. The figures use different data and years, so KLA's lead over the next firm ranges from about 5 to about 7 times.

How it works

An inspection tool sweeps a beam of light across the wafer, photographs it, and compares the image of each chip with the image of its neighbors. The chips on a

1 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

2 Letter to Shareholders Q4 Fiscal 2026. KLA Corporation · 27 July 2026 [A] d1io3yog0oux5.cloudfront.net/_b9ee755a5f60dd0fb3f9e27967aed6af/klatencor/db/1117/10668/let…

3 Compounding Sustainable Outperformance. KLA Corporation · 13 March 2026 [A] d1io3yog0oux5.cloudfront.net/_7791115a123b86b3f10b1a5eb5210224/klatencor/db/1166/10653/fil…

4 Letter to Shareholders Q3 Fiscal 2026. KLA Corporation · 29 April 2026 [A] d1io3yog0oux5.cloudfront.net/_a357bfc9113388e37f3bfcb2ea2f0b64/klatencor/db/1117/10655/let…

wafer are meant to be identical, so anything that shows in one and not the others is a defect. The tool records where each defect is. For faults too small for light to show, a second tool scans a beam of electrons across a small patch and builds a picture from the electrons that bounce back.

A metrology tool measures known shapes. It shines light on a test pattern printed beside the circuits and works out from the scattered light how wide the lines are and how thick the film is. It also reads marks printed in this layer and the last to measure how far the two are out of line, which chipmakers call overlay.

The fab pays for all this to protect yield, the share of chips on a wafer that work. A wafer loses value two ways: whole regions ruined by mishandling, misalignment or over-etching, and single circuits killed by stray material or an airborne particle [5].

A bigger chip is a bigger target for the second kind, and the loss grows faster than the area. Chipmakers call each chip on a wafer a die. In the standard formula, die yield is (1 + D0A/a) to the power minus a. Here D0 is defects per square centimeter, A is die area, and a is a clustering constant of 3 to 4 that says how much defects bunch together. At a defect density of 0.8 per square centimeter, with a set to 3, a 1 square centimeter die yields 49.2 percent and a 2 by 2 centimeter die 11.3 percent [5]. Because the bigger die yields less and fits fewer times on a wafer, it costs 20.1 times as much as the small one [5].

By the same formula, an 8 square centimeter die, about as much as one exposure can print, yields about 49 percent at 0.10 defects per square centimeter, with a set to 3 [5]. Halving the defect density to 0.05 lifts that yield to about 69 percent [5]. Engineers move a fab from the first figure to the second by finding the process step that is losing dies and fixing it, work called yield learning.

Those percentages are the share of dies with no killer defect. A chip built from many identical blocks can carry spares and switch off a faulty block [5]. Nvidia's full H100 design has 144 of its main computing blocks, called streaming multiprocessors, and the SXM5 version it ships has 132 turned on [6]. So the share of dies that can be sold is higher than the formula gives. But the formula shows how fast flawless dies get scarce as chips grow.

Variants and trade-offs

The tools trade sensitivity against speed. All the market sizes and shares below are CSET's, for 2025 [1].

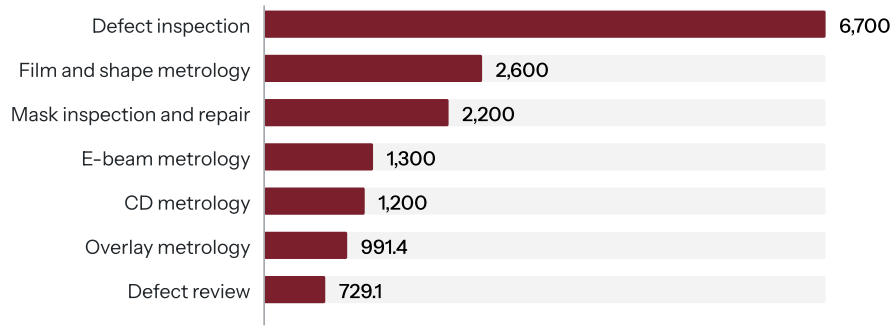
5 Behrooz Parhami, Dependable Computing: A Multilevel Approach, part 2. University of California, Santa Barbara · 11 October 2020 A web.ece.ucsb.edu/~parhami/docs_folder/f33-book-dep-comp-pt2.pdf

6 NVIDIA Hopper Architecture In-Depth. Nvidia · 22 March 2022 A developer.nvidia.com/blog/nvidia-hopper-architecture-in-depth

1 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

- **Defect inspection** scans a whole wafer with light or electrons and compares each die against its neighbors. It is by far the largest family, at $6.7 billion, and KLA held 85.9 percent of it. ASML's HMI eScan 1100 runs 25 electron beams at once, up to 15 times the throughput of single-beam inspection tools, and finds patterning defects down to 7 nm [7].
- **Film and shape metrology** measures how thick each film is and how flat the wafer sits, without cutting it open. In this $2.6 billion family KLA held 45.6 percent, Nova 24.6 and Onto Innovation 10.9. Onto paid about $720 million in August 2026 for 27 percent of the Japanese X-ray tool maker Rigaku [8].
- **Mask inspection and repair** has its own tools, because a defect on a mask, the stencil the circuit pattern is printed from, repeats on every die. This $2.2 billion family is one of two where KLA has a close rival. KLA held 45.9 percent, Lasertec 42 and Zeiss 10.1. Lasertec also makes actinic inspection tools, which check EUV masks with the same 13.5 nm light that prints them [9]. Lasertec calls its ACTIS A150 the world's first EUV patterned mask inspection system [9]. KLA's own actinic mask inspection tool was still at an early demo stage at its March 2026 investor day [3].
- **Overlay metrology** measures whether this layer sits on the one below. Its $991.4 million market was split almost evenly between KLA at 51.5 percent and ASML at 47.9.
- **Critical dimension metrology**, also called line-width measurement, measures how wide a printed feature came out. It is the one family Japan leads. Hitachi held 71.9 percent of this $1.2 billion market and Applied Materials 27.5.
- **Electron-beam metrology** resolves detail too fine for light. Applied Materials led this $1.3 billion family at 45.1 percent, with ASML at 35.6 and KLA at 18.8.

Process control market by tool family, 2025 $M



Source: ETO Supply Chain Explorer, CSET

7 The HMI eScan 1100. ASML [A] asml.com/en/products/metrology-and-inspection-systems/hmi-escan-1100

8 Onto Innovation Completes Strategic Investment in Rigaku Holdings Corporation. Onto Innovation · 10 August 2026 [A] investors.ontoinnovation.com/news/news-details/2026/Onto-Innovation-Completes-Strategic-In...

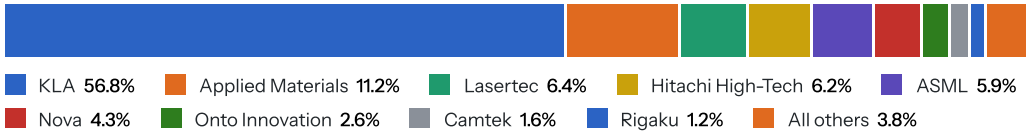
9 ACTIS A150. Lasertec · 12 September 2019 [A] lasertec.co.jp/en/products/semiconductor/actis_a150.html

3 Compounding Sustainable Outperformance. KLA Corporation · 13 March 2026 [A] d1io3yog0oux5.cloudfront.net/_7791115a123b86b3f10b1a5eb5210224/klatencor/db/1166/10653/fil...

Who makes it

KLA held 56.8 percent of a $15.7 billion market in 2025 [1]. KLA says its share grew 3.6 percentage points from 2021 to 2025 [4]. It says it gained share again in 2025 across mask, optical wafer and e-beam inspection [4]. Its revenue for fiscal 2026, the year to June 2026, was $13.579 billion, up from $12.156 billion in fiscal 2025 [10].

Process control tool revenue share, 2025 %



KLA 56.8% Applied Materials 11.2% Lasertec 6.4% Hitachi High-Tech 6.2% ASML 5.9%
Nova 4.3% Onto Innovation 2.6% Camtek 1.6% Rigaku 1.2% All others 3.8%

Source: ETO Supply Chain Explorer, CSET

The rest are specialists, and none holds more than 11.2 percent, by CSET's count for 2025 [1].

- **Applied Materials** is the largest of them at 11.2 percent. It sells e-beam defect review tools, CD-SEMs (electron microscopes that measure line width) and e-beam metrology tools next to the deposition and etch tools that create the defects.
- **Lasertec** holds 6.4 percent, almost all of it mask inspection. **Hitachi High-Tech** holds 6.2 percent, most of it from CD-SEMs. **ASML** holds 5.9 percent through YieldStar overlay and HMI e-beam metrology.
- **Nova** holds 4.3 percent, through film metrology, and took $880.6 million of revenue in 2025 [11].
- **Onto Innovation** holds 2.6 percent on $1,005.3 million of 2025 revenue [12]. **Camtek** holds 1.6 percent on $496.1 million of 2025 revenue [13].

KLA's patterning line, which includes several metrology products and mask inspection, grew 61 percent year on year in the June 2026 quarter [2]. In that quarter Taiwan took 31 percent of KLA's revenue and China 26 percent [2].

1 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

4 Letter to Shareholders Q3 Fiscal 2026. KLA Corporation · 29 April 2026 [A] d1io3yog0oux5.cloudfront.net/_a357bfc9113388e37f3bfcb2ea2f0b64/klatencor/db/1117/10655/let...

10 Revenue from Contract with Customer, Excluding Assessed Tax (us-gaap:RevenueFromContractWithCustomerExcludingAssessedTax) reported by KLA CORP, XBRL company concept data. U.S. Securities and Exchange Commission (XBRL data for KLA CORP) [A] data.sec.gov/api/xbrl/companyconcept/CIK0000319201/us-gaap/RevenueFromContractWithCustomer...

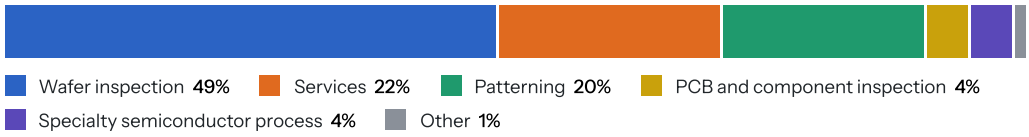
11 Revenues (us-gaap:Revenues) reported by NOVA LTD., XBRL company concept data. U.S. Securities and Exchange Commission (XBRL data for NOVA LTD.) [A] data.sec.gov/api/xbrl/companyconcept/CIK0001109345/us-gaap/Revenues.json

12 Onto Innovation Reports 2025 Fourth Quarter and Full Year Results. Onto Innovation · 19 February 2026 [A] investors.ontoinnovation.com/news/news-details/2026/Onto-Innovation-Reports-2025-Fourth-Qu...

13 CAMTEK LTD, Form 6-K report of foreign private issuer for the period ended 2026-02-18 (6-K). U.S. Securities and Exchange Commission (filing by CAMTEK LTD) · 18 February 2026 [A] sec.gov/Archives/edgar/data/1109138/000117891326000549/zk2634405.htm

2 Letter to Shareholders Q4 Fiscal 2026. KLA Corporation · 27 July 2026 [A] d1io3yog0oux5.cloudfront.net/_b9ee755a5f60dd0fb3f9e27967aed6af/klatencor/db/1117/10668/let...

KLA revenue by product line, June 2026 quarter %



Source: KLA letter to shareholders, Q4 FY2026

KLA says process control's share of fab spending is still rising [4]. In memory chips, KLA credits the rise to higher performance specifications, less redundancy (fewer spare circuits) and more logic content in devices built for high-performance computing [4].

KLA's systems analyze data with AI and also filter out noise [4]. In April 2026 KLA said it had added AI to every product line over the previous five years [4].

Packaging, AI and China

KLA says the AI build-out is making advanced packaging, and the process control it needs, more important [2]. KLA became the largest supplier of process control tools for advanced wafer-level packaging in 2025, adding 14 percentage points of share [4]. It expects about $1.1 billion of packaging process control revenue in 2026, up more than 70 percent [2]. Onto signed a volume purchase agreement estimated at over $240 million with an HBM maker for 2D inspection and 3D bump metrology through 2027 [12].

Chinese makers are still small in process control, but one of them is growing fast. CSET's count, built on TechInsights data, puts all Chinese firms at 0.7 percent of a $15.7 billion market in 2025 [1]. Skyverse (Shenzhen Zhongke Feice), a maker of wafer inspection and metrology tools, reported 2.05 billion yuan of revenue in 2025, about $285 million [14]. That was up 48.75 percent on 2024 and more than double its 2023 revenue [14]. Skyverse's revenue alone is about 1.8 percent of CSET's market, more than twice CSET's figure for all Chinese firms, so the CSET count may miss some domestic sales [1]. In its 2025 annual report Skyverse says Chinese

4 Letter to Shareholders Q3 Fiscal 2026. KLA Corporation · 29 April 2026 A d1io3yog0oux5.cloudfront.net/_a357bfc9113388e37f3bfcb2ea2f0b64/klatencor/db/1117/10655/let...

2 Letter to Shareholders Q4 Fiscal 2026. KLA Corporation · 27 July 2026 A d1io3yog0oux5.cloudfront.net/_b9ee755a5f60dd0fb3f9e27967aed6af/klatencor/db/1117/10668/let...

12 Onto Innovation Reports 2025 Fourth Quarter and Full Year Results. Onto Innovation · 19 February 2026 A investors.ontoinnovation.com/news/news-details/2026/Onto-Innovation-Reports-2025-Fourth-Qu...

1 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

14 深圳中科飞测科技股份有限公司2025年年度报告摘要. CNINFO · 24 April 2026 A static.cninfo.com.cn/finalpage/2026-04-25/1225182755.PDF

makers still hold a low share of China's market for these tools and that KLA holds the largest share of it [14].

By CSET's 2025 count, Chinese firms held 0.6 percent of defect inspection tools [15]. They held 0.4 percent of line-width measurement tools in 2025 [16]. In film and shape metrology they held 2 percent in 2025, all of it from Raintree Scientific Instruments, or RSIC [17]. They held no share of mask inspection and repair tools in 2025 [18]. Nor did they hold any share that year of the tools that inspect wafers during advanced packaging [19].

In October 2022 the US began requiring a license for any US-made item, process control tools included, bound for certain Chinese fabs [20]. The rule covered fabs that made logic chips at 16/14 nm or below, NAND memory with 128 or more layers, or DRAM at 18 nm half-pitch or below [20].

A December 2024 rule added specific process control tools to the export control list. It controls tools for defect inspection and measurement of patterned 300 mm wafers if they can find defects of 21 nm or smaller [21]. They must do this using light with a wavelength below 400 nm, an electron beam resolving 1.65 nm, a cold field emission source, or two or more beams [21]. The same rule controls overlay metrology accurate to 0.5 nm or better if the tool is built into the track (the machine that coats and develops the light-sensitive resist) or can switch wavelengths quickly [21]. It also controls stand-alone wafer shape metrology that feeds corrections to immersion or EUV scanners [21]. A separate new entry, 3B994,

14 深圳中科飞测科技股份有限公司2025年年度报告摘要. CNINFO · 24 April 2026 A static.cninfo.com.cn/finalpage/2026-04-25/1225182755.PDF

15 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

16 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

17 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

18 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

19 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) A chipexplorer.eto.tech

20 Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items; Supercomputer and Semiconductor End Use; Entity List Modification. U.S. Government Publishing Office (Federal Register) · 13 October 2022 A govinfo.gov/content/pkg/FR-2022-10-13/pdf/2022-21658.pdf

21 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 5 December 2024 A govinfo.gov/content/pkg/FR-2024-12-05/pdf/2024-28270.pdf

covers optical film and line-width metrology for 300 mm wafers with software for measuring non-planar transistors[21].

BIS, the US Commerce Department bureau that wrote the rule, placed these tools in new entries for items it says also have legitimate uses in older chipmaking [21]. So they carry narrower license requirements than tools built only for advanced chips [21]. A new foreign direct product rule extends US jurisdiction to specified foreign-made chipmaking tools bound for Macau or for a Country Group D:5 destination such as China [22].

The chokepoint

One firm, KLA, held more than half the process control market in 2025. Several firms can build a good film metrology tool, but KLA leads most kinds of process control tool, and it says its share grew from 2021 to 2025. A fab that can buy scanners and etchers but not sensitive inspection tools finds fewer of the defects that lower its yield, so its yield rises more slowly.

Key evaluation criteria

- **Sensitivity.** The smallest killer defect a tool can find. ASML's multibeam system is specified to 7 nm patterning defects [7].
- **Throughput.** Wafers per hour, which sets how many lots a fab can afford to look at. ASML's multibeam system gives up to 15 times the throughput of single-beam tools [7].
- **Nuisance rate.** The share of flagged events that turn out not to be real defects. AI classification is meant to reduce it.
- **Sampling plan.** How many wafers in each lot the fab checks, and how many spots on each wafer. The plan sets what process control costs the fab, so it is a cost decision as well as a technical one.
- **Time to result.** How fast a measurement reaches the process engineer, which sets the speed of yield learning.

FOUR THINGS TO REMEMBER

— Metrology and inspection tools check each wafer between steps and tell a fab why its chips fail.

21 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. U.S. Government Publishing Office (Federal Register) · 5 December 2024 [A] govinfo.gov/content/pkg/FR-2024-12-05/pdf/2024-28270.pdf

22 Commerce Strengthens Export Controls to Restrict China's Capability to Produce Advanced Semiconductors for Military Applications. U.S. Bureau of Industry and Security · 2 December 2024 [A] bis.gov/press-release/commerce-strengthens-export-controls-restrict-chinas-capability-prod…

7 The HMI eScan 1100. ASML [A] asml.com/en/products/metrology-and-inspection-systems/hmi-escan-1100

- AI chips are among the largest made, so finding defects matters most for them.
- CSET's count puts KLA of the United States at 56.8 percent of this tool market in 2025, about five times the share of its nearest rival.
- Chinese makers held 0.7 percent of this tool market in 2025 on CSET's count, though one Chinese maker, Skyverse, reported revenue equal to about 1.8 percent of it.

CHAPTER
08
STAGE 8 OF 15

Transistors and the Front End

Only TSMC, Intel and Samsung make 2 nm-class transistors in volume. TSMC plans its A14 process for 2028, with more than 20 percent higher logic density than N2.

IN PLAIN TERMS

The front end of a chip factory is the part that builds the transistors, the billions of tiny switches in every chip. Each switch works like a valve. Current runs along a narrow strip of silicon called the channel, and a small voltage on a gate above the strip opens or shuts the flow. Shorter strips switch faster, and more of them fit on a chip. But once the strip gets short enough, the valve stops sealing. Current leaks through even when the gate is off, and the chip burns power doing nothing. The newest kind of transistor, called gate-all-around, fixes this by wrapping the gate around all four sides of the strip. Building them takes a long sequence of steps that have to run in exact order. Only three companies, TSMC, Intel and Samsung, make these switches in volume. In 2024 the United States put a worldwide license requirement on the know-how for building them, with an exception for countries that impose matching controls.

IN SHORT

Intel's 18A process shipped gate-all-around transistors together with backside power, a second power network built underneath them [6]. TSMC says its A14 process will give up to 15 percent more speed or up to 30 percent less power than N2 [23]. Rapidus, a state-backed newcomer in Japan, said in July 2025 that it would begin 2 nm mass production in 2027 [14].

China's SMIC makes 7 nm-class chips without extreme-ultraviolet (EUV) lithography, one way of printing chip patterns [13]. China has announced $43 billion to develop its own EUV [22]. In September 2025 the Center for Strategic and International Studies judged that EUV remains firmly out of reach [22].

6 Intel Foundry Details Process Milestones and Future Innovation at VLSI Symposium. Intel · 16 June 2026 [A] intc.com/news-events/press-releases/detail/1772/intel-foundry-details-process-milestones-a...

23 TSMC Unveils Next-Generation A14 Process at North America Technology Symposium. TSMC · 24 April 2025 [A] pr.tsmc.com/english/news/3228

14 Rapidus Achieves Significant Milestone at its State-of-the-Art Foundry with Prototyping of Leading-Edge 2nm GAA Transistors. Rapidus · 18 July 2025 [A] rapidus.inc/en/news_topics/information/rapidus-achieves-significant-milestone-at-its-state...

13 Is SMIC N+3's Metal Pitch Smaller than Intel 18A's?. SemiAnalysis · 14 June 2026 [B] newsletter.semianalysis.com/p/steel-smic-n3-teardown

22 Breakthroughs or Boasts? Assessing Recent Chinese Lithography Advancements. Center for Strategic and International Studies [A] csis.org/blogs/strategic-technologies-blog/breakthroughs-or-boasts-assessing-recent-chines...

CHOKEPOINT CARD | EXTREME CONCENTRATION | HARD TO SUBSTITUTE

CONCENTRATION

Extreme

SUBSTITUTABILITY

Hard

Three firms make gate-all-around transistors, the newest kind, in volume, and none sells its process recipe.

PRICE OR MARKET SIZE

No foundry publishes prices for its newest wafers, so cost per transistor cannot be worked out from public figures

TSMC does publish density gains. It rates N2P at 1.2 times the logic density of N3E, and A14 at more than 20 percent above N2.

WHO LEADS

- TW **TSMC** N2 in volume production since 4Q25, TSMC's first nanosheet node
- US **Intel** 18A in high-volume manufacturing. 18A-P adds 9% performance at the same power
- KR **Samsung Foundry** First into production with gate-all-around nanosheets, starting initial 3 nm output in June 2022. In the SF2 family, SF2Z adds backside power
- CN **SMIC** N+2 and N+3, printed in several passes on older deep-ultraviolet tools, no EUV
- JP **Rapidus** State-backed 2 nm entrant targeting mass production in 2027

WHERE IT IS MADE

- TW **Taiwan** TSMC N2 at Fab 20 Hsinchu and Fab 22 Kaohsiung
- US **United States** Intel 18A at Fab 52, Arizona. TSMC Arizona on N4, with N3 targeted for the second half of 2027
- KR **South Korea** Samsung Hwaseong and Pyeongtaek
- JP **Japan** Rapidus Chitose, Hokkaido
- CN **China** SMIC Shanghai and Beijing 300 mm lines

WHY SUBSTITUTION IS SLOW

A newcomer has to work out its own process recipe, which sets the steps, their order and the tools for each. Three firms have written one. Rapidus, a state-backed newcomer in Japan, aims to move from prototypes to volume production in 2027. As of June 2026 it had raised 424.95 billion yen in equity and received state research subsidies. China cannot buy extreme-ultraviolet scanners. It made its 7 nm chips on older deep-ultraviolet machines stockpiled before export controls took effect. Relying on those machines keeps it behind the newest chips.

WHERE CHINA STANDS

SMIC ships a process it calls N+3, measured at 113.4 million transistors per square millimeter. It prints its tightest layers in two or four passes because it cannot buy an extreme-ultraviolet scanner. Its narrowest wires sit closer together than Intel 18A's, a comparison SemiAnalysis calls an incomplete, cherry-picked metric.

WHERE THE US STANDS

Intel runs the only American-owned line that makes the newest transistors. Its 18A process brought gate-all-around transistors and backside power to market together. Backside power feeds the transistors from underneath. Intel is also the first to use High-NA EUV scanners, the newest kind, in volume production.

Nvidia's Blackwell GPU holds 208 billion switches on two dies, each die the largest a lithography machine can print in one shot [1]. Every switch has to turn on fully, turn off completely, and keep doing it for years. Making the switches is called the front end of line. Wiring them together is the back end of line.

How it works

The front end builds the transistors in a fixed order. First it marks where each transistor will sit and fires atoms of boron or phosphorus into the silicon there. This step, called ion implantation, makes those patches conduct the way the design needs.

Next it shapes the channel. In a FinFET, the older design, the fab etches the silicon into thin fins that stand up from the wafer. The gate then goes over the top and both sides of each fin, three sides instead of one.

In a gate-all-around transistor, the fab instead grows alternating layers of silicon and silicon-germanium and cuts the stack into a narrow bar. It then dissolves the silicon-germanium away, leaving thin silicon sheets held at their ends like shelves. It coats every sheet with an insulating film a few atoms thick, so the gate can control the channel without touching it. Gate metal then fills the gaps above, below and beside each sheet, covering all four sides. Last come the contacts, small metal plugs down to the two ends of the channel and to the gate, ready for the wiring layers above.

Each of those steps is a separate deposition, etch or implant. They have to run in exact order, because later layers cover earlier ones. Once gate metal surrounds a sheet, the insulator can no longer go between them.

Current leaks two ways, straight through the gate insulator and along the channel from end to end while the gate is off. The second is called short-channel leakage, and wrapping the gate around the channel reduces it. Three fixes deal with the two leaks:

- **Better insulator.** A thicker hafnium-based film, paired with a metal gate, lets fewer electrons slip straight through yet controls the channel as tightly as before. Intel replaced the silicon dioxide insulator with it at 45 nm in 2007 [2]. It said the new film cut leakage through the insulator more than tenfold and, with the metal gate, leakage along the channel more than fivefold [2].
- **More gate.** Fins and nanosheets wrap gate metal around more sides of the channel.
- **Thinner channel.** The fab thins the silicon until no current can flow where the gate cannot reach it.

Variants and trade-offs

Planar and FinFET

A planar transistor lays the gate on one side of a flat channel. As the channel shortens, the gate loses control of it. A fin stands the channel on edge so the gate

1 NVIDIA Blackwell Architecture. Nvidia · 18 December 2025 A nvidia.com/en-us/data-center/technologies/blackwell-architecture

2 Intel's Transistor Technology Breakthrough Represents Biggest Change to Computer Chips in 40 Years. Intel · 27 January 2007 A intel.com/pressroom/archive/releases/2007/20070128comp.htm

covers three sides. But each fin adds a fixed amount of current, so a designer who needs more current has to add a whole fin.

Gate-all-around nanosheets

A nanosheet transistor lays the channel flat again, stacks two to four thin sheets and wraps gate metal around each one. The sheets can be drawn at different widths, so current can be set in finer steps. TSMC sells this as NanoFlex, which lets designers mix short logic cells for small area with tall cells for performance [3].

TSMC's N2, its first nanosheet node, entered volume production in the fourth quarter of 2025 [4]. TSMC rates the follow-on N2P at 18 percent more speed at the same power, 36 percent less power at the same speed, and 1.2 times the logic density of its earlier N3E [5]. Intel's version, called RibbonFET, shipped on its 18A process [6]. An improved version, 18A-P, adds another 9 percent performance at the same power [6]. Samsung began initial production of its 3 nm gate-all-around process in June 2022, more than three years before TSMC's N2 [7].

Those numbers are the foundries' own claims until someone measures the silicon. Each foundry picks its own baseline, and none publishes a density figure that can be compared with a rival's.

TSMC's published chip density gains, each over the baseline TSMC names, with A16 at its stated maximum %



Source: TSMC advanced technologies, HPC platform

Backside power delivery

In a conventional chip, power and signal share one stack of copper wiring above the transistors, where wide power rails crowd out the signal wires. With backside power, the fab grinds the wafer thin and builds a second network for power underneath it.

Intel shipped backside power first, as PowerVia on 18A [6]. It cut routed area by 11 percent and dynamic voltage droop tenfold [6]. Droop is the sag in supply voltage when a block of logic switches on at once. TSMC's Super Power Rail connects the backside network directly to each transistor's source and drain, the two ends of its channel. Against N2P, TSMC rates A16, which uses Super Power Rail, at 8 to 10

3 TSMC Celebrates 30th North America Technology Symposium with Innovations Powering AI with Silicon Leadership. TSMC · 24 April 2024 [A] pr.tsmc.com/english/news/3136

4 2nm Technology. TSMC [A] tsmc.com/english/dedicatedFoundry/technology/logic/l_2nm

5 HPC Platform – Advanced Technologies. TSMC [A] tsmc.com/english/dedicatedFoundry/technology/platform_HPC_tech_advancedTech

6 Intel Foundry Details Process Milestones and Future Innovation at VLSI Symposium. Intel · 16 June 2026 [A] intc.com/news-events/press-releases/detail/1772/intel-foundry-details-process-milestones-a…

7 Samsung Begins Chip Production Using 3nm Process Technology With GAA Architecture. Samsung Global Newsroom · 30 June 2022 [A] news.samsung.com/global/samsung-begins-chip-production-using-3nm-process-technology-with-g…

percent more speed, 15 to 20 percent less power and up to 1.10 times the chip density [5].

TSMC said in April 2026 that two more processes, A13 and A12, are scheduled to enter production in 2029 [8]. A13 is a direct shrink of A14 with 6 percent area savings, and A12 adds Super Power Rail to the A14 platform [8].

TSMC's logic roadmap beyond N2, from its own announcements

- 2028-01-01
  N2U, 3-4% faster or 8-10% lower power than N2P
- 2028-01-01
  A14, up to 15% faster and more than 20% denser than N2
- 2029-01-01
  A13, a direct shrink of A14 for 6% area savings
- 2029-01-01
  A12, the A14 platform with Super Power Rail backside power

Source: TSMC, 2026 North America Technology Symposium

CFET and what comes after

A basic logic circuit needs two transistors that switch on opposite signals, and in production chips the two sit side by side. Complementary field-effect transistors, or CFETs, stack the two on top of each other [9]. The research institute imec sees them as an attractive design for logic nodes beyond 1 nm [9].

The back end of line

Wiring has become as hard to make as the transistor. Copper seeps into the insulator around it, so every wire needs a barrier, a liner and a cap [10]. As the most critical wiring levels drop below a 20 nm pitch, those extra layers take a large share of the wire's width [10]. Pitch is the distance from the center of one wire to the center of the next. Two metals are moving into those narrow levels.

- **Ruthenium** shows promise for use without a barrier, which would keep the wire's resistance low [10]. imec has demonstrated ruthenium lines at a 16 nm pitch [11].

5 HPC Platform – Advanced Technologies. TSMC [A] tsmc.com/english/dedicatedFoundry/technology/platform_HPC_tech_advancedTech

8 TSMC Debuts A13 Technology at 2026 North America Technology Symposium. TSMC · 23 April 2026 [A] pr.tsmc.com/english/news/3302

9 CFET (complementary FET). imec [A] imec-int.com/en/articles/imec-puts-complementary-fet-cfet-logic-technology-roadmap

10 Semi-damascene metallization. imec [A] imec-int.com/en/articles/semi-damascene-metallization-inflection-point-back-end-line-proce…

11 16nm Ru lines using semi-damascene integration approach. imec · 3 June 2025 [A] imec-int.com/en/press/imec-demonstrates-16nm-pitch-ru-lines-record-low-resistance-obtained…

- **Molybdenum** conducts worse than copper in bulk, with a resistivity of 5.3 microhm-cm against copper's 1.68 [12]. But in a thin film it beats tungsten by up to 30 percent and needs no barrier either, so chipmakers are putting it into some contacts and local wiring layers [12].

Who makes it

TSMC, Intel and Samsung build gate-all-around nodes with extreme-ultraviolet (EUV) lithography. China's SMIC builds 7 nm-class nodes without it [13]. In July 2025 Rapidus, a state-backed newcomer in Japan, said it had started making 2 nm prototypes at its plant in Chitose, Hokkaido, and would begin mass production in 2027 [14]. By June 2026 it had raised 424.95 billion yen in equity, the latest 150 billion yen from a state agency [15]. It has also received government research subsidies since fiscal 2022 [15].

A node's name, such as 7 nm, does not measure any feature on the chip. SMIC's N+3, its third-generation 7 nm process, sets its narrowest wires 32.5 nm apart, tighter than Intel 18A's 36 nm [13]. It packs 113.4 million transistors per square millimeter, against 107.7 for TSMC's N6 [13].

AI accelerators have usually trailed the leading edge, because a die as large as theirs comes out working reliably only on a process that has run for years. Blackwell uses a custom TSMC 4NP process [1]. AMD's MI455X, launched on 23 July 2026, lists TSMC 2 nm and 3 nm silicon [16]. It spreads its compute across eight separate dies [17].

SRAM, the fast memory on the logic die, is the part whose failure to shrink hurts AI chips most. TSMC claims 22 percent more SRAM density from N3E to N2 [18]. But the gain comes mostly from the circuits around the memory array, and the cell that holds one bit shrinks little or not at all [18]. See Memory and HBM.

12 Molybdenum's Role in Ultra-Fast Computing: The Metal Behind the Speed. Entegris · 18 May 2026 [A] blog.entegris.com/molybdenums-role-in-ultra-fast-computing-the-metal-behind-the-speed

13 Is SMIC N+3's Metal Pitch Smaller than Intel 18A's?. SemiAnalysis · 14 June 2026 [B] newsletter.semianalysis.com/p/steel-smic-n3-teardown

14 Rapidus Achieves Significant Milestone at its State-of-the-Art Foundry with Prototyping of Leading-Edge 2nm GAA Transistors. Rapidus · 18 July 2025 [A] rapidus.inc/en/news_topics/information/rapidus-achieves-significant-milestone-at-its-state…

15 Rapidus Completes 150 Billion Yen Funding Round from Japan Government. Rapidus · 5 June 2026 [A] rapidus.inc/en/news_topics/information/20260605-2

1 NVIDIA Blackwell Architecture. Nvidia · 18 December 2025 [A] nvidia.com/en-us/data-center/technologies/blackwell-architecture

16 AMD Instinct™ MI455X GPUs. AMD [A] amd.com/en/products/accelerators/instinct/mi400/mi455x.html

17 AMD Instinct MI455X GPU. AMD · 20 July 2026 [A] amd.com/content/dam/amd/en/documents/products/accelerators/instinct/amd-instinct-mi455x_br…

18 Clash of the Foundries: Gate All Around + Backside Power at 2nm. SemiAnalysis · 1 October 2024 [B] newsletter.semianalysis.com/p/clash-of-the-foundries

The chokepoint

TSMC, Intel and Samsung do not sell their process recipes, which set the steps, their order and the tools for each. In September 2024 the United States put a worldwide license requirement on the know-how for building gate-all-around transistors, with an exception for countries that impose matching controls [19].

China launched its third National Integrated Circuit Fund in May 2024 with 344 billion yuan, about $47 billion [20]. But China cannot buy EUV scanners.

A Shanghai state firm began building domestic immersion scanners in 2026 and planned about five that year and roughly 20 in 2027 [21]. Those are deep-ultraviolet machines, a generation behind EUV. SMIC gets fine features out of them by printing its tightest layers in several passes. That many scanners keep those lines running but come nowhere near replacing ASML, the Dutch scanner maker.

Domestic EUV is further off. China has announced $43 billion to develop it [22]. In September 2025 the Center for Strategic and International Studies (CSIS) judged that EUV remains firmly out of reach despite massive state investment [22].

See Lithography and Geopolitics.

Key evaluation criteria

- **Electrostatic control.** How far the gate wraps the channel. It sets how much the switch leaks when off, and the lowest supply voltage it can run on.
- **Density.** Transistors per square millimeter, comparable only inside one foundry's own numbering and for one kind of cell.
- **Interconnect resistance.** Below a 20 nm pitch, copper wiring adds more and more delay from its resistance and capacitance [11].
- **Power delivery.** The voltage lost carrying power across the die. On Intel's 18A, backside power frees area and cuts the droop tenfold [6].
- **SRAM scaling.** The cell that holds one bit has barely shrunk since the 5 nm node, which caps how much SRAM an accelerator can afford [18].
- **Cost per transistor.** No foundry publishes wafer prices, so it cannot be worked out from public figures.

19 Commerce Control List Additions and Revisions; Implementation of Controls on Advanced Technologies Consistent With Controls Implemented by International Partners. Federal Register (Commerce Department; Industry and Security Bureau) · 6 September 2024 [A] federalregister.gov/documents/2024/09/06/2024-19633/commerce-control-list-additions-and-re…

20 China launches $47bn chip fund to counter U.S. restrictions. Nikkei Asia · 27 May 2024 [B] asia.nikkei.com/spotlight/supply-chain/china-launches-47bn-chip-fund-to-counter-u.s.-restr…

21 China starts production of home-grown immersion DUV chipmaking tools, source says. Reuters · 28 July 2026 [B] reuters.com/world/china/china-starts-production-home-grown-immersion-duv-chipmaking-tools-…

22 Breakthroughs or Boasts? Assessing Recent Chinese Lithography Advancements. Center for Strategic and International Studies [A] csis.org/blogs/strategic-technologies-blog/breakthroughs-or-boasts-assessing-recent-chines…

11 16nm Ru lines using semi-damascene integration approach. imec · 3 June 2025 [A] imec-int.com/en/press/imec-demonstrates-16nm-pitch-ru-lines-record-low-resistance-obtained…

6 Intel Foundry Details Process Milestones and Future Innovation at VLSI Symposium. Intel · 16 June 2026 [A] intc.com/news-events/press-releases/detail/1772/intel-foundry-details-process-milestones-a…

18 Clash of the Foundries: Gate All Around + Backside Power at 2nm. SemiAnalysis · 1 October 2024 [B] newsletter.semianalysis.com/p/clash-of-the-foundries

FOUR THINGS TO REMEMBER

- The newest transistors wrap the gate around all four sides of the channel to stop current leaking.
- Only TSMC, Intel and Samsung make them in volume.
- No firm sells its process recipe, so a newcomer has to develop its own.
- China's SMIC makes 7 nm-class chips without EUV by printing its tightest layers in several passes.

CHAPTER
09
STAGE 9 OF 15

Foundries and Fabs

TSMC's six largest plants are all in Taiwan. In 2025 they could make more than 13 million wafers a year, out of more than 17 million for the whole company.

IN PLAIN TERMS

A foundry is a chip factory, or fab, that makes chips for other companies. Those companies send in their designs, and the foundry runs them all through the same building on the same tools. The building costs tens of billions of dollars, and the air inside is cleaner than in an operating room. Even so, each wafer, the silicon disc the chips are built on, moves between hundreds of machines over several months inside sealed boxes and never touches that air. No chip company sells enough of one product to fill such a factory on its own, so chip companies share a handful of them. TSMC runs the biggest, and all six of its largest plants are in Taiwan, so a political crisis over Taiwan would affect the whole chip supply.

IN SHORT

In 2022 Taiwan held 69 percent of the world's capacity for advanced logic chips, those below 10 nanometers, and the United States almost none [24]. In July 2026 TSMC raised its planned investment in Arizona to $265 billion across 12 facilities [18].

The industry association SEMI projected in 2026 that China would keep the world's largest installed fab capacity in 2026 and 2027 [25]. Little of China's capacity is advanced [24]. China's SMIC cannot buy an EUV scanner, the machine that prints the finest layers, yet it makes 7 nm-class chips by printing those layers in several passes on older deep-ultraviolet scanners [10].

24 **Emerging Resilience in the Semiconductor Supply Chain.** Boston Consulting Group and Semiconductor Industry Association · May 2024 [A] semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc…

18 **TSMC Announces Additional $100 Billion Investment In Arizona.** City of Phoenix · 16 July 2026 [A] phoenix.gov/newsroom/ced-news/tsmc-announces-additional--100-billion-investment-in-arizona…

25 **World Fab Forecast.** SEMI [A] semi.org/en/products-services/market-data/world-fab-forecast

10 **Is SMIC N+3's Metal Pitch Smaller than Intel 18A's?.** SemiAnalysis · 14 June 2026 [B] newsletter.semianalysis.com/p/steel-smic-n3-teardown

CHOKEPOINT CARD | EXTREME CONCENTRATION | HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
Extreme	Hard Intel and Samsung also run lines for the newest chips. But TSMC has not built one of its largest fabs outside Taiwan, and the slow part of building one is training the workforce.	**In 2023 BCG estimated that a new 2 nm-class fab completed in 2026 would cost $35-43 billion to own over ten years** In March 2026 SIA put the cost to build and equip a leading-edge fab at as much as $20-25 billion. In July 2026 TSMC's planned investment in Arizona rose to $265 billion

WHO LEADS

- TW **TSMC** 77% of wafer revenue from 7 nm and below, 2Q26; six gigafabs, all in Taiwan
- US **Intel Foundry** $5.8bn segment revenue, only $293m of it from outside customers; $2.089bn operating loss, 2Q26
- KR **Samsung Foundry** Hwaseong and Pyeongtaek; an initial $17bn minimum at Taylor, Texas, now part of over $37bn planned in Texas (CHIPS award, Dec 2024)
- CN **SMIC** $9.327bn revenue in 2025 at 93.5% utilization
- JP **Rapidus** State-backed 2 nm entrant, Chitose

WHERE IT IS MADE

- TW **Taiwan** TSMC's six gigafabs and UMC; TSMC's N2 and N3 lines
- KR **South Korea** Samsung Foundry, Hwaseong and Pyeongtaek
- CN **China** SMIC and Hua Hong. The largest installed capacity of any region, almost all of it on older nodes. SEMI forecast in 2026 that it would stay the largest in 2026 and 2027
- US **United States** Intel Foundry, GlobalFoundries, TSMC Arizona, Samsung Taylor
- JP **Japan** TSMC Kumamoto, Rapidus Chitose

WHY SUBSTITUTION IS SLOW

A newcomer with enough money can pay for the building. In 2023 BCG put the ten-year cost of owning a new 2 nm-class fab finished in 2026 at $35 to $43 billion. In March 2026 SIA put the cost to build and equip a leading-edge fab at as much as $20 to $25 billion. Intel and Samsung already run lines for the newest chips, so a newcomer would be copying something that has been done. Building a workforce that has run one process for years takes longer than putting up the plant. Even TSMC has yet to build one of its largest fabs anywhere but Taiwan.

WHERE CHINA STANDS

SMIC and Hua Hong are China's largest foundries. SMIC had revenue of $9.327 billion in 2025 with its lines 93.5 percent full. It cannot buy an EUV scanner, the machine that prints the finest layers of leading-edge chips. Yet it makes 7 nm-class chips on its N+2 and N+3 processes by printing the tightest layers in several

passes on older deep-ultraviolet scanners. Georgetown University's CSET reported in June 2024 that Huawei's Ascend 910B AI chip is believed to be made on N+2. But SemiAnalysis judged in April 2025 that most 910B and 910C chips were made at TSMC. In the third-quarter 2026 update of its World Fab Forecast, SEMI projected that China would keep the world's largest installed fab capacity in 2026 and 2027. Little of that capacity is advanced.	WHERE THE US STANDS Intel Foundry is the only American-owned foundry that makes the newest chips. In the second quarter of 2026 it lost $2.089 billion on $5.8 billion of segment revenue. Only $293 million of that revenue came from outside customers. The federal government bought a 9.9 percent stake in Intel in August 2025. The stake is passive, with no board seat or governance rights. The government also agreed to vote with Intel's board on shareholder matters, with limited exceptions.

TSMC calls a plant a GIGAFAB once it can run more than 100,000 wafers a month, each a 300 mm disc, twelve inches across [1]. It operates six of them, and every one is in Taiwan [2].

How it works

Inside, a chip is built up in layers, and the same loop of steps repeats for each one. Wafers travel in a lot, a batch of 25 in a sealed pod called a FOUP that only a machine opens. An overhead track carries the pod from tool to tool.

At each tool a robot lifts one wafer out and the tool does its one job: depositing a film, printing a pattern, etching it, cleaning it or measuring it. Then the robot puts the wafer back. The pod moves on to the next tool, often across the building. When a layer is finished, the loop starts again with the next one. A chip takes hundreds of those visits, so a wafer spends months inside.

The pod keeps the wafer clean for all of that time because the wafer never meets the room's air. That air is also filtered until it carries almost no dust.

Every design in the building runs on the same tools because each tool is set by a recipe, the list of gases, temperatures and times for one step. A foundry offers each customer the same list of recipes on shared machines. The tools run day and night, because an idle tool earns nothing. The finished wafer is tested (see Test and assembly), then cut and packaged.

TSMC, founded in 1987, pioneered the pure-play model, in which a foundry makes only its customers' products [3]. In 2023 the Boston Consulting Group (BCG) estimated that a new fab for chips at 2 nm and below, completed in 2026, would cost $35 to $43 billion to own over ten years [4]. That is 33 to 66 percent above 2023 costs [4]. In March 2026 the Semiconductor Industry Association (SIA) told the US

1 GIGAFAB® Facilities. TSMC [A] tsmc.com/english/dedicatedFoundry/manufacturing/gigafab

2 Fab Capacity. TSMC [A] tsmc.com/english/dedicatedFoundry/manufacturing/fab_capacity

3 TSMC, company profile page. TSMC [A] tsmc.com/english/aboutTSMC/company_profile

4 Navigating the Costly Economics of Chip Making. Boston Consulting Group · 28 September 2023 [A] bcg.com/publications/2023/navigating-the-semiconductor-manufacturing-costs

Senate that building and equipping a leading-edge fab takes as much as $20 to $25 billion of capital [5].

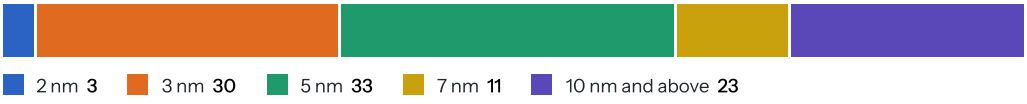
Together, the plants TSMC and its subsidiaries run had capacity for more than 17 million wafers a year in 2025, counted as if each were the standard 12-inch size [2]. The six largest plants, called gigafabs, had a combined capacity of more than 13 million [1]. All six are in Taiwan [2]. The other plants are five smaller fabs in Taiwan, two in China, two in the United States and one in Japan [2].

Variants and trade-offs

Pure-play foundries

TSMC and UMC, both based in Taiwan, make chips only for other companies. Chips at 7 nm and below, meaning the newest processes, were 77 percent of TSMC's wafer revenue in the second quarter of 2026 [6]. That share split into 33 percent at 5 nm, 30 at 3 nm, 11 at 7 nm and 3 at 2 nm [6]. TSMC's revenue for the quarter was $40.20 billion, at a 67.7 percent gross margin [6].

TSMC wafer revenue by node, 2Q26 % OF WAFER REVENUE



2 nm 3 3 nm 30 5 nm 33 7 nm 11 10 nm and above 23

Source: TSMC 2Q26 results, filed with the SEC

Integrated device makers turned foundry

Intel and Samsung design and sell their own chips and rent out capacity too. Renting out capacity is harder for them, because customers dislike handing designs to a competitor and the company's own products take the best capacity. Intel Foundry reported $5.8 billion of segment revenue in the second quarter of 2026 and an operating loss of $2.089 billion [7]. Only $293 million of that revenue came from outside customers [8]. Intel says substantially all of its foundry business still makes Intel's own products [8].

5 Testimony of the Semiconductor Industry Association (SIA) before the Senate Committee on Environment & Public Works. U.S. Senate Committee on Environment and Public Works · 4 March 2026 [A] epw.senate.gov/public/_cache/files/a/c/ac4552ff-551b-40d5-bea6-19dd17a1cfea/1F7BCD91016CBE...

2 Fab Capacity. TSMC [A] tsmc.com/english/dedicatedFoundry/manufacturing/fab_capacity

1 GIGAFAB® Facilities. TSMC [A] tsmc.com/english/dedicatedFoundry/manufacturing/gigafab

6 TAIWAN SEMICONDUCTOR MANUFACTURING CO LTD, Form 6-K report of foreign private issuer for the period ended 2026-06-30 (6-K). U.S. Securities and Exchange Commission (filing by TAIWAN SEMICONDUCTOR MANUFACTURING CO LTD) · 16 July 2026 [A] sec.gov/Archives/edgar/data/0001046179/000104617926000451/a2q26e_withguidancexfinal.htm

7 Intel Reports Second-Quarter 2026 Financial Results. Intel · 23 July 2026 [A] intc.com/news-events/press-releases/detail/1776/intel-reports-second-quarter-2026-financia...

8 INTEL CORP, Form 10-Q quarterly report for the period ended 2026-06-27 (10-Q). U.S. Securities and Exchange Commission (filing by INTEL CORP) · 24 July 2026 [A] sec.gov/Archives/edgar/data/50863/000005086326000157/intc-20260627.htm

National champions

China's SMIC and Hua Hong and Japan's Rapidus have government backing. SMIC's revenue rose 16.2 percent to $9.327 billion in 2025, at a 21 percent gross margin [9]. Its lines ran 93.5 percent full, on monthly capacity above a million wafers counted in the smaller 8-inch size [9]. SMIC cannot buy an EUV scanner, yet it makes 7 nm-class chips on its N+2 and N+3 processes by printing the tightest layers in several passes on older deep-ultraviolet scanners [10].

Georgetown University's Center for Security and Emerging Technology (CSET) reported in June 2024 that Huawei's Ascend 910B AI chip is believed to be made on N+2 [11]. It also reported that SMIC's limited 7 nm capacity forced Huawei to choose between its most advanced phone chips and its AI chips [11]. But the research firm SemiAnalysis judged in April 2025 that most 910B and 910C chips were made on TSMC's 7 nm process, from wafers Huawei bought through another company [12].

In February 2026 Rapidus raised 267.6 billion yen from the Japanese government and 32 companies to reach 2 nm mass production by 2027 [13].

Mature-node capacity

Most wafers come off older lines. These make power devices, analog chips and microcontrollers, and they compete on price. In January 2025 the industry association SEMI forecast that foundry capacity at all nodes would rise from 11.3 million wafers a month in 2024 to a record 12.6 million in 2025, counted in the 8-inch size [14]. Its June 2024 forecast had credited much of that growth to Intel's new foundry business and China's expansion [15]. But in July 2025 Intel scaled back its own fab plans, saying it would no longer go ahead with planned projects in Germany and Poland and would slow construction in Ohio [16].

Capacity and cost

In March 2025 TSMC added $100 billion to the $65 billion it had already committed in Phoenix, for three more fabs, two packaging plants and an R&D center [17]. In July 2026 it added another $100 billion for four more fabs at 2 nm or below, taking

9 SMIC Announces 2025 Annual Results. SMIC [A] smics.com/en/site/news_read/7951

10 Is SMIC N+3's Metal Pitch Smaller than Intel 18A's?. SemiAnalysis · 14 June 2026 [B] newsletter.semianalysis.com/p/steel-smic-n3-teardown

11 Pushing the Limits: Huawei's AI Chip Tests U.S. Export Controls. Center for Security and Emerging Technology (CSET) · 17 June 2024 [A] cset.georgetown.edu/publication/pushing-the-limits-huaweis-ai-chip-tests-u-s-export-controls

12 Huawei AI CloudMatrix 384. SemiAnalysis · 16 April 2025 [B] newsletter.semianalysis.com/p/huawei-ai-cloudmatrix-384-chinas-answer-to-nvidia-gb200-nvl72

13 Rapidus Secures 267.6 Billion Yen in Funding from Japan Government and Private Sector Companies. Rapidus · 27 February 2026 [A] rapidus.inc/en/news_topics/information/rapidus-secures-267-6-billion-yen-in-funding-from-j...

14 Eighteen New Semiconductor Fabs to Start Construction in 2025, SEMI Reports. SEMI [A] semi.org/en/semi-press-release/eighteen-new-semiconductor-fabs-to-start-construction-in-20...

15 Global Semiconductor Fab Capacity Projected to Expand 6% in 2024 and 7% in 2025, SEMI Reports. SEMI · 18 June 2024 [A] semi.org/en/news-media-press-releases/semi-press-releases/global-semiconductor-fab-capacit...

16 Intel Reports Second-Quarter 2025 Financial Results. Intel · 24 July 2025 [A] intc.com/news-events/press-releases/detail/1745/intel-reports-second-quarter-2025-financia...

17 TSMC Intends to Expand Its Investment in the United States to US$165 Billion to Power the Future of AI. TSMC · 4 March 2025 [A] pr.tsmc.com/english/news/3210

its total investment in Arizona to $265 billion across 12 facilities [18]. All of TSMC's gigafabs are still in Taiwan [2].

Some customers pay TSMC in advance to hold capacity at its fabs [19]. These deposits stood at NT$189.9 billion at the end of 2025, down from NT$291.1 billion a year earlier [19]. In 2025 the four largest AI chip designers used around 90 percent of the world's CoWoS packaging capacity and HBM memory supply, against only 12 percent of advanced logic production [20]. See Advanced packaging.

Who makes it

The US government owns part of Intel. In August 2025 the government bought 433.3 million newly issued Intel shares at $20.47 each, or $8.9 billion in all, for a 9.9 percent stake [21]. It paid with $5.7 billion of CHIPS Act grants that Intel had not yet received and $3.2 billion from the Secure Enclave program [21].

The stake is passive, with no board seat [21]. The government agreed to vote with Intel's board on matters that need shareholder approval, with limited exceptions [21]. The deal also gave the government a five-year warrant, a right to buy a further 5 percent of Intel at $20 a share [21]. The government can use that right only if Intel's ownership of its foundry business falls below 51 percent [21].

In 2026 Intel began high-volume manufacturing of some of its Panther Lake processors using ASML's High-NA extreme ultraviolet scanners [7]. It also put its 18A-P process into risk production, the trial runs a process goes through before it is signed off [7].

18 TSMC Announces Additional $100 Billion Investment In Arizona. City of Phoenix · 16 July 2026 [A] phoenix.gov/newsroom/ced-news/tsmc-announces-additional--100-billion-investment-in-arizona...

2 Fab Capacity. TSMC [A] tsmc.com/english/dedicatedFoundry/manufacturing/fab_capacity

19 TAIWAN SEMICONDUCTOR MANUFACTURING CO LTD, Form 20-F annual report for the period ended 2025-12-31 (20-F). U.S. Securities and Exchange Commission (filing by TAIWAN SEMICONDUCTOR MANUFACTURING CO LTD) · 16 April 2026 [A] sec.gov/Archives/edgar/data/1046179/000162828026025362/tsm-20251231.htm

20 Advanced packaging and HBM, not logic dies, were the bottlenecks on AI chip production in 2025. Epoch AI · 12 March 2026 [A] epoch.ai/data-insights/ai-chip-supply-chain-constraints

21 Intel and Trump Administration Reach Historic Agreement to Accelerate American Technology and Manufacturing Leadership. Intel · 22 August 2025 [A] intc.com/news-events/press-releases/detail/1748/intel-and-trump-administration-reach-histo...

7 Intel Reports Second-Quarter 2026 Financial Results. Intel · 23 July 2026 [A] intc.com/news-events/press-releases/detail/1776/intel-reports-second-quarter-2026-financia...

In 2022 Samsung broke ground in Taylor, Texas, on a plant with an initial minimum investment of $17 billion, split into $6 billion for buildings and $11 billion for equipment [22]. In December 2024 the Commerce Department awarded Samsung up to $4.745 billion of CHIPS Act funding, to support more than $37 billion of planned investment in Central Texas [23]. That plan covers two new leading-edge logic fabs and a research fab in Taylor, and an expansion of Samsung's existing plant in Austin [23].

The chokepoint

In 2022 South Korea and Taiwan held almost all of the world's logic capacity below 10 nm [24]. The United States had almost no capacity below 10 nm that year [24]. In May 2024 SIA and BCG projected that the US share of that capacity would reach 28 percent by 2032 [24]. They also projected that new fab investment worldwide would reach around $2.3 trillion from 2024 to 2032 [24].

Across all nodes, old and new, SEMI forecast in June 2024 that China would have installed capacity of 10.1 million wafers a month in 2025, the most of any region [15]. It forecast 5.8 million for Taiwan, out of a world total of 33.7 million, all counted in the 8-inch size [15]. In the third-quarter 2026 update of its World Fab Forecast, SEMI projected that China would keep the largest installed capacity in 2026 and 2027 [25]. Little of China's capacity is advanced [24].

22 Taylor, US Fab. Samsung Electronics · 11 September 2025 A semiconductor.samsung.com/sas/company/taylor

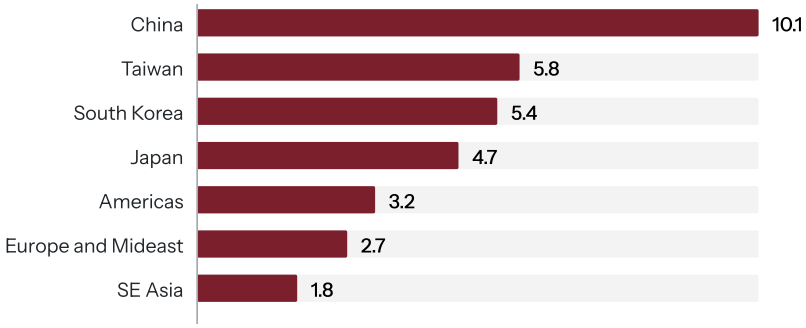
23 Biden-Harris Administration Announces CHIPS Incentives Award with Samsung Electronics to Solidify U.S. Leadership in Leading-Edge Semiconductor Production. NIST · 20 December 2024 A nist.gov/news-events/news/2024/12/biden-harris-administration-announces-chips-incentives-a…

24 Emerging Resilience in the Semiconductor Supply Chain. Boston Consulting Group and Semiconductor Industry Association · May 2024 A semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc…

15 Global Semiconductor Fab Capacity Projected to Expand 6% in 2024 and 7% in 2025, SEMI Reports. SEMI · 18 June 2024 A semi.org/en/news-media-press-releases/semi-press-releases/global-semiconductor-fab-capacit…

25 World Fab Forecast. SEMI A semi.org/en/products-services/market-data/world-fab-forecast

SEMI's June 2024 forecast of installed wafer fab capacity by region in 2025, all nodes MILLION 8-INCH-EQUIVALENT WAFERS PER MONTH



Source: SEMI World Fab Forecast, June 2024

Chip production in Taiwan faces three risks:

- **Seismic.** Taiwan sits on an active plate boundary, and after an earthquake in January 2025 TSMC had to scrap some wafers in process [26].
- **Power.** A gigafab is a large industrial load on a small island grid.
- **Political.** Taiwan held 69 percent of the world's logic capacity below 10 nm in 2022 [24]. In December 2025 SemiAnalysis put Taiwan's share of the world's advanced-node capacity at over 90 percent, though it did not say which node sizes count as advanced [27]. So a blockade or conflict would cut most of that output off from the world market at once. See Geopolitics.

Building capacity outside Taiwan is slow, and TSMC has yet to build a gigafab anywhere else [2].

Key evaluation criteria

- **Leading-edge share.** The share of revenue from the newest processes. At TSMC, 7 nm and below made up 77 percent of wafer revenue in the second quarter of 2026 [6].
- **Yield at ramp.** How many good chips a new process gives in its first year. If too few work, the process loses money and has to be subsidized.
- **Ten-year cost of ownership.** $35 to $43 billion for a new 2 nm-class fab finished in 2026, in BCG's 2023 estimate, and rising [4].
- **Utilization.** How full the lines run. SMIC's lines ran 93.5 percent full in 2025, at a 21 percent gross margin [9].
- **Customer concentration.** A foundry owned by a chipmaker competes with its own customers, so some of them will not trust it with a flagship design.

26 TSMC January 2025 Revenue Report and the Statement on the Impact of Earthquake. TSMC · 10 February 2025 [A] pr.tsmc.com/english/news/3204

24 Emerging Resilience in the Semiconductor Supply Chain. Boston Consulting Group and Semiconductor Industry Association · May 2024 [A] semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc…

27 TSMC Overseas Fabs – A Success?. SemiAnalysis · 1 December 2025 [B] newsletter.semianalysis.com/p/tsmc-overseas-fabs-a-success

2 Fab Capacity. TSMC [A] tsmc.com/english/dedicatedFoundry/manufacturing/fab_capacity

6 TAIWAN SEMICONDUCTOR MANUFACTURING CO LTD, Form 6-K report of foreign private issuer for the period ended 2026-06-30 (6-K). U.S. Securities and Exchange Commission (filing by TAIWAN SEMICONDUCTOR MANUFACTURING CO LTD) · 16 July 2026 [A] sec.gov/Archives/edgar/data/0001046179/000104617926000451/a2q26e_withguidancexfinal.htm

4 Navigating the Costly Economics of Chip Making. Boston Consulting Group · 28 September 2023 [A] bcg.com/publications/2023/navigating-the-semiconductor-manufacturing-costs

9 SMIC Announces 2025 Annual Results. SMIC [A] smics.com/en/site/news_read/7951

- **Geographic exposure.** How much of a buyer's supply comes from one island, with a single power grid and a single fault system.

FOUR THINGS TO REMEMBER

- A foundry makes chips for other companies, because no chip company can fill a fab alone.
- All six of TSMC's largest plants are in Taiwan.
- Outside Taiwan, the slowest part of setting up leading-edge chipmaking is training a workforce to run the process.
- SEMI projected in 2026 that China would keep the world's largest installed fab capacity in 2026 and 2027, but little of it is advanced.

CHAPTER
10
STAGE 10 OF 15

Memory and HBM

SK hynix, Samsung and Micron make nearly all high bandwidth memory (HBM), the stacked memory beside AI chips. SK hynix took 63.2 percent of HBM revenue in 2025.

IN PLAIN TERMS

High bandwidth memory, or HBM, is where an AI chip keeps the numbers it is working on, and it is built to hand them over fast. A memory maker etches tiny holes into each memory wafer, fills them with copper and grinds the wafer until the copper shows through its back. It tests each of the wafer's chips and stacks up to sixteen good ones into a cube, with every layer wired to the ones below. The cube then sits beside the processor, the chip that does the calculating. Stacking is the hard part because each chip is ground thinner than a human hair and the finished cube has to stay within a set height. One chip that fails after stacking can scrap the whole cube. Only SK hynix, Samsung and Micron make HBM in volume, and a single entry in the US export rules covers HBM from all three.

IN SHORT

US-based Micron went from 5.8 percent of HBM revenue in 2024 to 23.1 percent in the first quarter of 2026 [13]. Over the same period, Samsung's share fell from 37.8 percent to 20.5 percent [13].

Since December 2024, a US license has been required, unless an exception applies, to export to China any HBM moving more than 2 GB per second per square millimeter of its area [1]. The Institute for Progress reported in October 2025 that China's CXMT planned to produce HBM3 in 2026 [22]. Even so, the institute estimated that the United States and its partners would make 70 times as much HBM as China in 2026 [22].

13 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 A sec.gov/Archives/edgar/data/2120882/000119312526299963/d32785d424b4.htm

1 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. Federal Register (Commerce Department; Industry and Security Bureau) · 5 December 2024 A federalregister.gov/documents/2024/12/05/2024-28270/foreign-produced-direct-product-rule-a...

22 Should the US Sell Blackwell Chips to China?. Institute for Progress · 25 October 2025 A ifp.org/the-b30a-decision

CHOKEPOINT CARD | EXTREME CONCENTRATION | HARD TO SUBSTITUTE

CONCENTRATION

Extreme

SUBSTITUTABILITY

Hard

Three firms ship nearly all HBM, and no fourth firm has both the newest DRAM and a line that stacks it.

PRICE OR MARKET SIZE

An HBM stack moves data over a path 16 times wider than a DDR5 memory module's

That width costs chip area and lowers the share of stacks that come out working.

WHO LEADS

- KR **SK hynix** 63.2% of HBM revenue in 2025. Says it was first to complete HBM4 development, in September 2025, at over 10 Gbps per pin
- KR **Samsung** 19.3% of HBM revenue in 2025. Says it shipped the industry's first commercial HBM4 in February 2026, at 11.7 Gbps per pin on a 4 nm logic base die
- US **Micron** 17.4% of HBM revenue in 2025. Its HBM4 runs above 11 Gbps per pin

WHERE IT IS MADE

- KR **South Korea** SK hynix in Cheongju and Icheon, Samsung in Pyeongtaek and Hwaseong
- US **United States** Micron headquarters and R&D, and an SK hynix packaging plant under construction in Indiana
- TW **Taiwan** Micron DRAM and HBM production. TSMC builds the HBM4 base die for SK hynix
- JP **Japan** Micron Hiroshima, 1-gamma DRAM
- SG **Singapore** Micron HBM advanced packaging plant, adding capacity from 2027

WHY SUBSTITUTION IS SLOW

Market share moves quickly among the three HBM makers. Micron went from 5.8 percent of HBM revenue in 2024 to 23.1 percent in the first quarter of 2026. A fourth maker would find it harder to break in. It would need the newest DRAM, a line that drills through and thins wafers, and a way to stack the thinned chips without warping them, all inside one company. Its stack would then have to pass a customer's qualification tests. No Chinese firm ships HBM3E or HBM4 in volume yet.

WHERE CHINA STANDS

No Chinese firm ships HBM3E or HBM4 in volume. In December 2024 CSIS called CXMT, China's strongest DRAM maker, one of only two Chinese firms with a credible path to producing advanced HBM. As of October 2025, CXMT planned to produce HBM3 in 2026. Huawei said in September 2025 that its Ascend 950 chips would use two HBM designs of its own, HiBL 1.0 and HiZQ 2.0. It said the 950DT, due in the fourth quarter of 2026, would carry 144 GB of HiZQ 2.0 at 4 TB/s. The announcement did not say who would manufacture the memory.

WHERE THE US STANDS

Micron is the only American HBM maker. BIS, the US export control office, controls HBM with a memory bandwidth density above 2 GB per second per square millimeter under control number ECCN 3A090.c. Shipping it to Macau or a Country Group D:5 destination such as China needs a license unless an exception applies. BIS said in December 2024 that every HBM stack then in production exceeded that threshold. License Exception HBM (15 CFR 740.25) applies to exporters headquartered in the United States or a Country Group A:5 ally. It lets them ship HBM below 3.3 GB per second

per square millimeter without a license. The HBM must be bought directly by the designer of the chip it will be packaged into and must go straight to the packaging site. The control excludes co-packaged chips that combine HBM and logic and whose main function is processing. The accelerator rules may cover those instead.

An AI accelerator moves numbers from memory to its arithmetic units and back. The cost of arithmetic has fallen faster than the cost of moving numbers. The US export control office wrote in December 2024 that a processor cannot reach its full capabilities unless its memory capacity and bandwidth rise with its speed [1].

How it works

Each layer in the cube is a DRAM chip, the ordinary kind of computer memory. Each layer is ground thin and bonded to the one below through tiny metal bumps called microbumps. The whole stack sits on a base die, a chip that manages it. The copper-filled holes are called through-silicon vias, or TSVs. They let a signal reach any layer without a wire running around the outside.

An HBM3E cube has a 1,024-bit interface, meaning it exchanges data with the processor over 1,024 wires at once [2]. That is sixteen times wider than a standard DDR5 memory module, the memory stick in a PC [2]. HBM4 doubles that to 2,048 wires [3]. Ordinary memory cannot be made that wide, because the pins would not fit on a circuit board. The stack can, since its wires run down through the TSVs and across a silicon interposer to the processor beside it, with no board in the way [2].

A chip needs the width most when it answers a prompt. Training reuses each stored number across a huge multiplication, so a chip can fetch little and compute a lot. But to generate one word of an answer, the chip reads the whole model from memory, plus the whole record of the conversation so far. Then it does a little arithmetic and waits for the next read. That record grows with the length of the conversation and the number of users served at once. So memory size and speed usually set how many people one chip can serve.

Variants and trade-offs

GENERATION	INTERFACE	BANDWIDTH PER STACK	CAPACITY PER STACK
HBM3E	1,024-bit, 16 channels (32 pseudo-channels)	over 1.2 TB/s	24 GB (8 chips), 36 GB (12 chips)

1 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. Federal Register (Commerce Department; Industry and Security Bureau) · 5 December 2024 A federalregister.gov/documents/2024/12/05/2024-28270/foreign-produced-direct-product-rule-a...

2 High-bandwidth memory (HBM). Micron Technology A micron.com/products/memory/hbm

3 JEDEC® and Industry Leaders Collaborate to Release JESD270-4 HBM4 Standard: Advancing Bandwidth, Efficiency, and Capacity for AI and HPC. JEDEC Solid State Technology Association · 16 April 2025 A jedec.org/news/pressreleases/jedec%C2%AE-and-industry-leaders-collaborate-release-jesd270-...

HBM4	2,048-bit, 32 channels	over 2.0 TB/s, up to 3.3 TB/s	up to 64 GB (16 chips, 32 Gb dies)

HBM3E

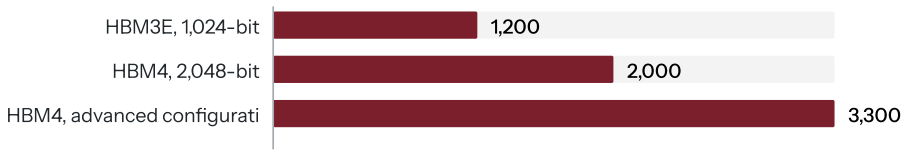
Micron's HBM3E comes in two sizes, 24 GB with eight chips stacked and 36 GB with twelve [4]. SK hynix's January 2026 market outlook says major research and brokerage analysts expect HBM3E to be about two thirds of all HBM shipments in 2026 [5].

HBM4

The standards body JEDEC published JESD270-4 in April 2025, setting the 2,048-bit interface and 2 TB/s per stack at 8 Gb/s per pin [3]. SK hynix says it was first to complete HBM4 development, in September 2025, at over 10 Gb/s per pin [6]. Samsung says it began mass production and shipped the first commercial HBM4 in February 2026, at 11.7 Gb/s per pin, with a 4 nm logic base die [7]. Micron's parts run above 11 Gb/s for more than 2.8 TB/s per stack [8].

Through HBM3E, SK hynix made its own base dies, the chips at the bottom of each stack [9]. For HBM4 it chose TSMC's advanced logic process, so the base die can hold more functions [9]. As a result, TSMC, which does not make HBM, makes the base die for SK hynix's HBM4.

Peak bandwidth per HBM stack GB/S



Source: Siemens EDA, HBM3E and HBM4 IC design guide, April 2026

4 Micron, HBM3E product page. Micron Technology [A] micron.com/products/memory/hbm/hbm3e

5 2026 Market Outlook – “Focus on the HBM-Led Memory Supercycle”. SK hynix [A] news.skhynix.com/en/2026-market-outlook-focus-on-the-hbm-led-memory-supercycle

3 JEDEC® and Industry Leaders Collaborate to Release JESD270-4 HBM4 Standard: Advancing Bandwidth, Efficiency, and Capacity for AI and HPC. JEDEC Solid State Technology Association · 16 April 2025 [A] jedec.org/news/pressreleases/jedec%C2%AE-and-industry-leaders-collaborate-release-jesd270-...

6 SK hynix Completes World’s First HBM4 Development and Readies Mass Production. SK hynix [A] news.skhynix.com/en/sk-hynix-completes-worlds-first-hbm4-development-and-readies-mass-prod...

7 Samsung Ships Industry-First Commercial HBM4 With Ultimate Performance for AI Computing. Samsung Global Newsroom · 12 February 2026 [A] news.samsung.com/global/samsung-ships-industry-first-commercial-hbm4

8 Micron, HBM4 product page. Micron Technology [A] micron.com/products/memory/hbm/hbm4

9 SK hynix Partners with TSMC to Strengthen HBM Technological Leadership. SK hynix [A] news.skhynix.com/sk-hynix-partners-with-tsmc-to-strengthen-hbm-technological-leadership

How a stack is built

- **Drill and thin.** Etch vertical holes, the through-silicon vias, into a memory wafer, fill them with copper, then grind the wafer from behind until the copper comes through.
- **Stack and join.** Thermocompression bonding joins the chips through metal bumps under heat and pressure, and it comes in several forms [11]. SK hynix uses Advanced MR-MUF, which stacks the chips, injects a liquid protective material between them and hardens it [6]. SK hynix says the process controls warping and dissipates heat well [6]. Samsung said in February 2024 that its twelve-high HBM3E uses TC-NCF, thermocompression with a thin non-conductive film between the chips [12]. Samsung expected TC-NCF to help limit warping as chips get thinner [12].
- **Stay under the height limit.** HBM4 raised the package specification height of a finished cube from 720 micrometers to 775, still under a millimeter [11].

Hybrid bonding is the next step, but it is not ready. It would press copper pad straight onto copper pad, with no solder and no filler, and cut the spacing between connections from about 20 micrometers to under 1 [11]. A tech note on SK hynix's newsroom says full-scale adoption in HBM is most likely at HBM4E or HBM5, where stacks are projected to pass twenty layers [11].

Who makes it

Three firms took more than 90 percent of the revenue from ordinary computer memory, or DRAM, in the first quarter of 2026 [13]. HBM comes from the same three, SK hynix, Samsung and Micron [13]. No other company has both leading-edge DRAM and a stacking line of its own.

SK hynix was the first company to mass-produce several generations of HBM, including HBM3 and HBM3E [13]. UBS, as quoted in SK hynix's January 2026 market

11 [Tech Note] Hybrid Bonding: Evolving into a Foundational Technology for Improving Semiconductor Performance. SK hynix [A] news.skhynix.com/en/tech-note-series-ep2

6 SK hynix Completes World's First HBM4 Development and Readies Mass Production. SK hynix [A] news.skhynix.com/en/sk-hynix-completes-worlds-first-hbm4-development-and-readies-mass-prod…

12 Samsung Develops Industry-First 36GB HBM3E 12H DRAM. Samsung Global Newsroom · 27 February 2024 [A] news.samsung.com/global/samsung-develops-industry-first-36gb-hbm3e-12h-dram

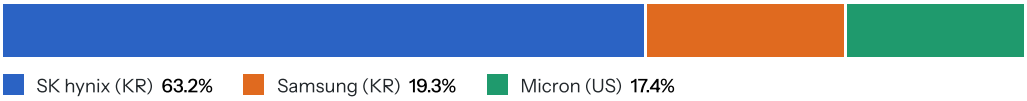
13 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 [A] sec.gov/Archives/edgar/data/2120882/000119312526299963/d32785d424b4.htm

outlook, expected SK hynix to take about 70 percent of the HBM4 market for Nvidia's Rubin platform in 2026 [5]. Samsung says it shipped the first commercial HBM4 in February 2026 [7].

SK hynix broke ground in August 2026 on an advanced packaging plant in West Lafayette, Indiana, an investment of over $4 billion [14]. Micron is modernizing its DRAM and HBM production capacity in Taiwan [15]. It has also broken ground on an HBM advanced packaging plant in Singapore to expand its packaging capacity from 2027 [15].

Of 2025 HBM revenue, SK hynix took 63.2 percent, Samsung 19.3 and Micron 17.4 [13]. The split is public because SK hynix registered shares in the United States [13]. Its US prospectus credits the research firm IDC. Market shares move fast. In 2024 the same three stood at 56.4, 37.8 and 5.8 percent [13]. In the first quarter of 2026 they stood at 56.4, 20.5 and 23.1 percent [13].

Global HBM revenue by supplier, 2025 %



SK hynix (KR) 63.2% Samsung (KR) 19.3% Micron (US) 17.4%

Source: SK hynix Form 424B4 prospectus, 10 July 2026, on IDC data

HBM market forecast for 2026, Bank of America $B



Source: Bank of America estimate cited in SK hynix's 2026 market outlook, January 2026

5 2026 Market Outlook – "Focus on the HBM-Led Memory Supercycle". SK hynix [A] news.skhynix.com/en/2026-market-outlook-focus-on-the-hbm-led-memory-supercycle

7 Samsung Ships Industry-First Commercial HBM4 With Ultimate Performance for AI Computing. Samsung Global Newsroom · 12 February 2026 [A] news.samsung.com/global/samsung-ships-industry-first-commercial-hbm4

14 SK hynix breaks ground on $4 billion advanced packaging production facility in Purdue Research Park. Purdue University · 27 August 2026 [A] purdue.edu/newsroom/2026/Q3/sk-hynix-breaks-ground-on-4-billion-advanced-packaging-product…

15 MICRON TECHNOLOGY INC, Form 10-K annual report for the period ended 2025-08-28 (10-K). U.S. Securities and Exchange Commission (filing by MICRON TECHNOLOGY INC) · 3 October 2025 [A] sec.gov/Archives/edgar/data/723125/000072312525000028/mu-20250828.htm

13 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 [A] sec.gov/Archives/edgar/data/2120882/000119312526299963/d32785d424b4.htm

Bank of America, in an estimate SK hynix quoted in January 2026, put the 2026 HBM market at $54.6 billion, up 58 percent on 2025 [5]. SK hynix says HBM needs more wafers than traditional DRAM, the kind used in PCs and phones [13]. It also says the industry's shift of production capacity to HBM has left traditional DRAM undersupplied [13]. The shortage, it says, helped drive a significant price recovery for that DRAM from the third quarter of 2025 [13].

The chokepoint

Every leading accelerator made outside China depends on Korean or American HBM. Nvidia's GB300 NVL72 rack carries 20 TB of HBM3E at up to 576 TB/s [16]. Each Nvidia Rubin GPU carries up to 288 GB of HBM4 at up to 22 TB/s [17]. AMD's MI355X carries 288 GB of HBM3E at 8 TB/s [18].

On 5 December 2024 the Bureau of Industry and Security (BIS), the US export control office, published a rule at 89 FR 96790. It added an export classification, ECCN 3A090.c, for "High bandwidth memory (HBM) having a ' memory bandwidth density' greater than 2 gigabytes per second per square millimeter". Its technical note defines that density as the memory bandwidth in gigabytes per second divided by the area of the package or stack in square millimeters. The rule also covers stacks made outside the United States, through a foreign direct product rule, which applies US controls to goods made abroad with American technology [1]. BIS stated in the rule that all HBM stacks then in production exceeded that threshold [1].

Under 15 CFR 742.6(b)(10)(ii), BIS reviews license applications to send 3A090.c items to Macau or a Country Group D:5 destination such as China [19]. It reviews them with a presumption of approval when neither the buyer nor its ultimate

5 2026 Market Outlook – “Focus on the HBM-Led Memory Supercycle”. SK hynix [A] news.skhynix.com/en/2026-market-outlook-focus-on-the-hbm-led-memory-supercycle

13 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 [A] sec.gov/Archives/edgar/data/2120882/000119312526299963/d32785d424b4.htm

16 NVIDIA GB300 NVL72. Nvidia · 23 July 2026 [A] nvidia.com/en-us/data-center/gb300-nvl72

17 Inside NVIDIA Rubin GPU Architecture: Powering the Era of Agentic AI. Nvidia · 21 July 2026 [A] developer.nvidia.com/blog/inside-nvidia-rubin-gpu-architecture-powering-the-era-of-agentic…

18 AMD Instinct™ MI355X GPUs. AMD [A] amd.com/en/products/accelerators/instinct/mi350/mi355x.html

1 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. Federal Register (Commerce Department; Industry and Security Bureau) · 5 December 2024 [A] federalregister.gov/documents/2024/12/05/2024-28270/foreign-produced-direct-product-rule-a…

19 15 CFR 742.6, Regional stability. Electronic Code of Federal Regulations [A] ecfr.gov/current/title-15/subtitle-B/chapter-VII/subchapter-C/part-742/section-742.6

parent is headquartered in one of those places, and with a presumption of denial otherwise [19].

Two provisions narrow the control. License Exception HBM, at 15 CFR 740.25, applies to exporters headquartered in the United States or a Country Group A:5 ally [20]. It lets them ship HBM below 3.3 GB per second per square millimeter without a license [20]. The HBM must be bought directly by the designer of the chip it will be packaged into and must go straight to the packaging site [20]. The rule also excludes co-packaged chips that combine HBM and logic and whose main function is processing, which other chip controls may cover instead [1].

The think tank CSIS said in December 2024 that the rules still allowed HBM2 to go to China with end-use and end-user restrictions [21]. It also said Nvidia's H20, whose HBM3 is packaged with the processor, was then still exportable under the accelerator thresholds [21].

In December 2024 CSIS called China's strongest DRAM maker, CXMT, one of only two Chinese firms with a credible path to producing advanced HBM [21]. The Institute for Progress reported in October 2025 that CXMT was planning to produce HBM3 in 2026 [22]. Huawei said in September 2025 that its Ascend 950 chips would use two HBM designs of its own, HiBL 1.0 and HiZQ 2.0 [23]. It said the 950DT, due in the fourth quarter of 2026, would carry 144 GB of HiZQ 2.0 at 4 TB/s [23]. The announcement did not say who would manufacture the memory [23]. SK hynix says HBM needs advanced production processes, such as TSV packaging, and must pass the rigorous testing and approval processes its customers set [13].

19 15 CFR 742.6, Regional stability. Electronic Code of Federal Regulations [A] ecfr.gov/current/title-15/subtitle-B/chapter-VII/subchapter-C/part-742/section-742.6

20 15 CFR 740.25 -- License Exception High Bandwidth Memory (HBM).. Electronic Code of Federal Regulations [A] ecfr.gov/current/title-15/subtitle-B/chapter-VII/subchapter-C/part-740/section-740.25

1 Foreign-Produced Direct Product Rule Additions, and Refinements to Controls for Advanced Computing and Semiconductor Manufacturing Items. Federal Register (Commerce Department; Industry and Security Bureau) · 5 December 2024 [A] federalregister.gov/documents/2024/12/05/2024-28270/foreign-produced-direct-product-rule-a...

21 Understanding the Biden Administration's Updated Export Controls. Center for Strategic and International Studies · 11 December 2024 [A] csis.org/analysis/understanding-biden-administrations-updated-export-controls

22 Should the US Sell Blackwell Chips to China?. Institute for Progress · 25 October 2025 [A] ifp.org/the-b30a-decision

23 Groundbreaking SuperPoD Interconnect: Leading a New Paradigm for AI Infrastructure. Huawei · 18 September 2025 [A] huawei.com/en/news/2025/9/hc-xu-keynote-speech

13 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 [A] sec.gov/Archives/edgar/data/2120882/000119312526299963/d32785d424b4.htm

Key evaluation criteria

- **Bandwidth per stack** sets how fast a chip can serve answers: over 1.2 TB/s on HBM3E, over 2 TB/s on HBM4 [10].
- **Capacity per stack** sets how large a model fits on one accelerator, up to 64 GB with sixteen chips stacked in the JEDEC HBM4 standard [3]. Samsung said in February 2026 that its twelve-high HBM4 comes in 24 to 36 GB and that sixteen-high parts would take it to 48 GB [7].
- **Stack height**, 775 micrometers in the HBM4 package specification, caps how many layers fit in a stack [11].
- **Stack yield** is the share of finished cubes that work, and one chip that fails after stacking can scrap the whole cube. Micron sends only dies that pass testing on to assembly and tests the finished cube again [8].
- **Thermal path** is the route heat takes out of the stack, and every added layer makes cooling harder [11]. Heat builds up in the stack, and removing it becomes a major challenge at 16 to 20 layers [11].
- **Qualification** means passing each customer's tests. SK hynix credits its experience with complex qualification processes for its standing as a trusted HBM supplier [13].

FOUR THINGS TO REMEMBER

— HBM is memory stacked into a cube beside the processor and wired to pass data fast.

— An AI chip reads its whole model from memory for each word, so memory speed limits it.

— SK hynix, Samsung and Micron make nearly all HBM.

— Since December 2024 US controls have covered the HBM stacks then in production, with two exceptions. No Chinese

10 HBM3e and HBM4: IC design guide for next-generation high bandwidth memory. Siemens Digital Industries Software · 24 April 2026 [A] blogs.sw.siemens.com/semiconductor-packaging/2026/04/24/hbm3e-hbm4-ic-design-guide

3 JEDEC® and Industry Leaders Collaborate to Release JESD270-4 HBM4 Standard: Advancing Bandwidth, Efficiency, and Capacity for AI and HPC. JEDEC Solid State Technology Association · 16 April 2025 [A] jedec.org/news/pressreleases/jedec%C2%AE-and-industry-leaders-collaborate-release-jesd270-...

7 Samsung Ships Industry-First Commercial HBM4 With Ultimate Performance for AI Computing. Samsung Global Newsroom · 12 February 2026 [A] news.samsung.com/global/samsung-ships-industry-first-commercial-hbm4

11 [Tech Note] Hybrid Bonding: Evolving into a Foundational Technology for Improving Semiconductor Performance. SK hynix [A] news.skhynix.com/en/tech-note-series-ep2

8 Micron, HBM4 product page. Micron Technology [A] micron.com/products/memory/hbm/hbm4

13 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 [A] sec.gov/Archives/edgar/data/2120882/000119312526299963/d32785d424b4.htm

firm ships HBM3E or HBM4 in volume.

CHAPTER
11
STAGE 11 OF 15

Advanced Packaging

TSMC packages nearly every AI accelerator. Epoch AI found that this packaging and high bandwidth memory were the main bottlenecks on AI chip production in 2025.

IN PLAIN TERMS

Packaging joins finished chips into one component and mounts them on a base, called a substrate, wired almost as finely as the chips themselves. A lithography machine, which prints a chip's circuits, can cover only a rectangle about the size of a postage stamp in one shot. Many leading accelerators, Nvidia's Blackwell among them, need more logic than fits in that rectangle. So they are built from several chips, called dies, set side by side on an interposer, which wires them together and sits on the substrate. Stacks of memory chips sit on the same interposer, a few millimeters away. The interposer is either one piece of silicon or layers of fine copper wiring that often carry small silicon bridges where the links are densest. Interposer, chips and substrate all expand at different rates when heated, so the package can warp and its joints crack. Preventing that in high volume is so hard that TSMC, in Taiwan, packages nearly every AI accelerator.

IN SHORT

Epoch AI estimates that in 2025 Nvidia used 60.3 percent of all capacity for CoWoS, TSMC's interposer packaging [4]. On the July 2026 earnings call, TSMC's chief executive, C.C. Wei, said its packaging capacity was so tight that it limited customers' growth [12]. Asked about Intel's EMIB-T packaging, he said he hoped it would take some of the load off TSMC [12].

Ranked by 2024 revenue, four of the top ten outsourced firms that package and test chips are headquartered in China [22]. But in 2024 Chinese firms had 7 percent of the market for advanced packaging fabrication tools, against 65 percent for US firms [22].

4 Advanced packaging and HBM, not logic dies, were the bottlenecks on AI chip production in 2025. Epoch AI · 12 March 2026 [A] epoch.ai/data-insights/ai-chip-supply-chain-constraints

12 Q2 2026 Taiwan Semiconductor Manufacturing Co Ltd Earnings Call (Chinese, English) — edited transcript. LSEG StreetEvents (transcript of a TSMC earnings call), via TSMC · 16 July 2026 [A] investor.tsmc.com/english/encrypt/files/encrypt_file/reports/2026-08/3e494f0c14dd0890f897a...

22 Inside Beijing's Chipmaking Offensive. Center for Security and Emerging Technology (CSET) · 14 July 2025 [A] cset.georgetown.edu/article/inside-beijings-chipmaking-offensive

CHOKEPOINT CARD

EXTREME CONCENTRATION

MODERATELY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
Extreme	Moderate Intel's EMIB can do the same job as TSMC's CoWoS. But the shortage is in making enough working packages at the volume Nvidia orders.	About 65,000 to 75,000 CoWoS wafers a month at the end of 2025, up from 35,000 to 40,000 a year earlier, in industry estimates from SemiWiki and TrendForce compiled by Epoch AI. TSMC does not publish its CoWoS capacity. At its 2024 symposium it projected 60 percent annual growth for the next few years. Epoch AI estimates Nvidia alone used 60.3 percent of the world's CoWoS capacity in 2025

WHO LEADS

- TW **TSMC** Almost all CoWoS for leading AI accelerators
- US **Intel** Foveros 3D packaging in volume at Rio Rancho, New Mexico; EMIB 2.5D bridges
- TW **ASE Technology** NT$389.2B assembly, test and materials revenue, 2025; the leading provider of outsourced assembly and test
- US **Amkor** $6.71B net sales, 2025; the largest American assembly and test firm
- CN **JCET** RMB 38.87B revenue, 2025; a leading Chinese assembly and test firm
- SG **ASMPT** $532.1M advanced packaging revenue, 2025; thermocompression bonders

WHERE IT IS MADE

- TW **Taiwan** TSMC advanced packaging fabs and a cluster of outsourced assembly and test firms; almost all 2.5D capacity for AI accelerators
- US **United States** Intel Foveros packaging in volume at Rio Rancho, New Mexico; Amkor plant in Peoria, Arizona, from 2028; two TSMC advanced packaging plants planned in Arizona
- CN **China** JCET, Tongfu and HT-Tech; four of the top ten outsourced assembly and test firms by 2024 revenue are Chinese
- NL **Netherlands** Besi hybrid and die bonders
- JP **Japan** Disco grinders and dicers; TEL and Shibaura bonding tools

WHY SUBSTITUTION IS POSSIBLE

None of the packaging steps is unusual. Intel's EMIB is another 2.5D option, and Amazon's Trainium and Google's TPU were adopting it to different degrees as of March 2026. But a newcomer has to buy enough bonding machines and run enough practice lots to make the interposer at 5.5 times the area a scanner prints in one shot. It then has to attach the memory stacks, make sure they work, and do all of that at the volume Nvidia alone orders. Amkor started building its plant

in Peoria, Arizona in October 2025 and expects production to start in early 2028, more than two years later. For a newcomer, the wait is two to five years.

WHERE CHINA STANDS

Four of the top ten outsourced assembly and test firms by 2024 revenue are headquartered in China. JCET alone had revenue of RMB 38.87 billion in 2025. Between 2019 and 2024, Chinese firms gained 5 percentage points of share in test tools and 4 points in fabrication tools built for advanced packaging, especially in etch and clean tools. But they lost share in bonding and other conventional assembly and packaging tools, the category where CSET counts hybrid and thermocompression bonders. In 2024 Chinese firms had 7 percent of the market for advanced packaging fabrication tools, against 65 percent for US firms and 15 percent for Japanese firms.

WHERE THE US STANDS

Intel opened Fab 9 in Rio Rancho, New Mexico, in January 2024 to make its Foveros 3D packages in volume. It called the plant the only US factory producing the most advanced packaging at scale. Intel also offers its EMIB 2.5D packaging to outside customers. SemiAnalysis reported in March 2026 that Amazon's Trainium and Google's TPU were adopting EMIB to different degrees. No US site yet packages Nvidia-class accelerators with HBM on CoWoS. Amkor is building a $7B packaging and test plant in Peoria, Arizona, with up to $407M of CHIPS Act money, and expects production to start in early 2028.

A lithography scanner prints a rectangle 26 mm by 33 mm, or 858 mm2, in one exposure [1]. That area is called the reticle field. Every leading AI accelerator is bigger than that. Packaging is the stage that lets an accelerator be larger than one reticle field, and the packaging method limits how much larger it can be.

How it works

A second reason for the interposer is high bandwidth memory (HBM). An HBM4 stack connects through 2,048 data signals, packed too tightly for a standard resin substrate, so the memory sits beside the logic on the interposer [2]. Even Nvidia's H100, whose single logic die is 814 mm2 and fits inside one reticle field, carries five HBM stacks [3]. Epoch AI counts TSMC's CoWoS packaging among the three inputs the H100 needs, along with its logic die and its HBM [4].

CoWoS stands for chip on wafer on substrate, the order in which the parts are joined. For CoWoS-S, TSMC makes a thin silicon wafer that carries wiring and vertical copper vias but no transistors. Vias are metal-filled holes that join one layer of wiring to the next. TSMC bonds the logic dies and memory stacks face down onto that wafer with microbumps. These are dots of solder 30 to 40 microns apart, about a third of a hair's width, and each dot is one electrical joint. Then

1 TWINSCAN NXE:3400C. ASML [A] asml.com/en/products/euv-lithography-systems/twinscan-nxe3400c

2 New JEDEC® SPHBM4 Standard Enables HBM4-Class Bandwidth on Organic Substrates. JEDEC Solid State Technology Association · 13 July 2026 [A] jedec.org/news/pressreleases/new-jedec%C2%AE-sphbm4-standard-enables-hbm4-class-bandwidth-...

3 NVIDIA Hopper Architecture In-Depth. Nvidia · 22 March 2022 [A] developer.nvidia.com/blog/nvidia-hopper-architecture-in-depth

4 Advanced packaging and HBM, not logic dies, were the bottlenecks on AI chip production in 2025. Epoch AI · 12 March 2026 [A] epoch.ai/data-insights/ai-chip-supply-chain-constraints

TSMC cuts the interposer out of the wafer and attaches it to the resin substrate, which later goes onto a circuit board.

Each of those packaging steps is ordinary on its own. But one cracked joint scraps a package holding eight tested HBM stacks and two large, expensive logic dies.

Variants and trade-offs

FAMILY	CARRIER	SIZE LIMIT	USED FOR
CoWoS-S	Full silicon interposer	3.3x reticle, about 2,700 mm2	H100-class parts
CoWoS-R	Resin interposer with fine redistribution wiring	Above 3.3x reticle	Cost-sensitive, fewer HBM stacks
CoWoS-L	Redistribution wiring with local silicon bridges	5.5x reticle in production, 9.5x in development	Blackwell, Rubin

2.5D: silicon, RDL and bridges

In 2.5D packaging the chips sit side by side and connect through an interposer or silicon bridges beneath them. TSMC has run CoWoS in volume since 2012 [5]. CoWoS-S has the finest wiring, but its interposer is one unbroken piece of silicon whose yield (the share of interposers that work) falls as the area grows. CoWoS-L replaces that piece of silicon with redistribution layers (RDL) of fine copper wiring and embeds small silicon bridges where the wiring is densest [6]. Intel's EMIB embeds similar silicon bridges in the resin substrate, and SemiAnalysis reported in March 2026 that Amazon's Trainium and Google's TPU were adopting it to different degrees [7].

Blackwell was the first high-volume design on CoWoS-L, and packaging problems at TSMC and in Nvidia's design limited its supply [8]. The bridges that carry the 10 TB/s link between the two logic dies must be placed very precisely [8]. The dies, bridges, interposer and substrate also expand at different rates, and that mismatch caused warpage [8].

3D: SoIC and Foveros

TSMC's SoIC and Intel's Foveros bond dies face to face, copper pad to copper pad, with no solder. That lets the connections sit far closer together than microbumps allow. TSMC put 3 nm SoIC into volume production in 2025 [6]. An A14-on-A14 version, stacking dies made on TSMC's 1.4 nm process, is due in 2029 [9]. It will carry 1.8 times the die-to-die connection density of the 2 nm version [9].

5 TSMC, CoWoS technology page. TSMC [A] 3dfabric.tsmc.com/english/dedicatedFoundry/technology/cowos.htm

6 TSMC 2025 Annual Report, chapter 5. TSMC [A] investor.tsmc.com/static/annualReports/2025/english/pdf/2025_tsmc_ar_e_ch5.pdf

7 The Great AI Silicon Shortage. SemiAnalysis · 12 March 2026 [B] newsletter.semianalysis.com/p/the-great-ai-silicon-shortage

8 Nvidia's Blackwell Reworked - Shipment Delays & GB200A Reworked Platforms. SemiAnalysis · 4 August 2024 [B] semianalysis.com/2024/08/04/nvidias-blackwell-reworked-shipment

9 TSMC Debuts A13 Technology at 2026 North America Technology Symposium. TSMC · 23 April 2026 [A] pr.tsmc.com/english/news/3302

In a stack the bottom die overheats and the top die is hard to power. Because of the heat, 3D designs often put cache memory underneath and compute circuits on top.

SoIC bonding also brings optics into the package. TSMC's COUPE uses SoIC to join a silicon photonics chip, which sends data as light, to its electrical control chip [6]. TSMC said in April 2026 that COUPE on substrate would begin production in 2026 [9]. It also said that putting COUPE inside the package gives twice the power efficiency of a pluggable version on the circuit board [9].

Two ways off the silicon interposer

Panel-level packaging replaces the round wafer with a rectangular panel, which wastes far less area. See Test and assembly. JEDEC, the semiconductor standards body, published a standard called SPHBM4 on 13 July 2026. It puts HBM4 memory on a base chip carrying 512 data signals instead of 2,048 [2]. With fewer signals, the connections can be spaced widely enough to mount the stack on a standard resin substrate instead of the silicon interposer [2].

Capacity

TSMC does not publish its CoWoS capacity [4]. Industry estimates from SemiWiki and TrendForce, compiled by Epoch AI, put it at about 65,000 to 75,000 wafers a month at the end of 2025, up from 35,000 to 40,000 a year earlier [4].

At its 2024 symposium TSMC said CoWoS capacity would grow at a 60 percent compound annual rate over the next few years [10]. But when asked about that rate on the July 2024 earnings call, TSMC's chief executive, C.C. Wei, gave no figure [10]. He said only that he would keep adding capacity wherever he could, and that he could not yet balance supply and demand [10].

6 TSMC 2025 Annual Report, chapter 5. TSMC [A] investor.tsmc.com/static/annualReports/2025/english/pdf/2025_tsmc_ar_e_ch5.pdf

9 TSMC Debuts A13 Technology at 2026 North America Technology Symposium. TSMC · 23 April 2026 [A] pr.tsmc.com/english/news/3302

2 New JEDEC® SPHBM4 Standard Enables HBM4-Class Bandwidth on Organic Substrates. JEDEC Solid State Technology Association · 13 July 2026 [A] jedec.org/news/pressreleases/new-jedec%C2%AE-sphbm4-standard-enables-hbm4-class-bandwidth-...

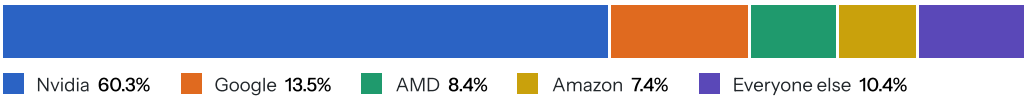
4 Advanced packaging and HBM, not logic dies, were the bottlenecks on AI chip production in 2025. Epoch AI · 12 March 2026 [A] epoch.ai/data-insights/ai-chip-supply-chain-constraints

10 Q2 2024 Taiwan Semiconductor Manufacturing Co Ltd Earnings Call (Chinese, English) — edited transcript. Refinitiv StreetEvents (transcript of a TSMC earnings call), via TSMC · 18 July 2024 [A] investor.tsmc.com/english/encrypt/files/encrypt_file/reports/2024-08/5122725a56670882d777a...

In October 2025 Wei would still say only that TSMC was working to increase capacity in 2026 and would probably give the real number in 2026 [11]. He said front-end capacity (making the chips) and back-end capacity (packaging them) were both, in his words, very tight [11].

Epoch AI compared each chip designer's use of components with world supply [4]. It puts Nvidia at 60.3 percent of all CoWoS capacity in 2025, with a 90 percent confidence interval of 56.5 to 64.3 [4]. Google took 13.5 percent, AMD 8.4 and Amazon 7.4 [4]. Together those four used around 90 percent of global CoWoS capacity and HBM supply, with HBM measured by value, but only 12 percent of advanced logic die production [4].

Estimated share of CoWoS capacity used, 2025 %



Nvidia 60.3% Google 13.5% AMD 8.4% Amazon 7.4% Everyone else 10.4%

Source: Epoch AI, advanced packaging and HBM were the bottlenecks on AI chip production in 2025

In March 2026 the research firm SemiAnalysis judged that front-end capacity had become the dominant bottleneck and that CoWoS constraints were easing [7]. It said TSMC was planning CoWoS capacity around how many wafers N3, its 3 nm process, could produce, because added packaging capacity is of no use without front-end wafers to fill it [7].

TSMC gave a different account four months later. On the July 2026 earnings call, Wei said TSMC's packaging capacity was so tight that it limited customers' growth, and that back-end capacity was still in shortage [12]. Asked about Intel's EMIB-T, he said he welcomed the extra flexibility and hoped it would take some of the load off TSMC [12].

11 Q3 2025 Taiwan Semiconductor Manufacturing Co Ltd Earnings Call (Chinese, English) — edited transcript. Refinitiv StreetEvents (transcript of a TSMC earnings call), via TSMC · 16 October 2025 [A] investor.tsmc.com/english/encrypt/files/encrypt_file/reports/2025-10/6860312f04fd291d0f26b...

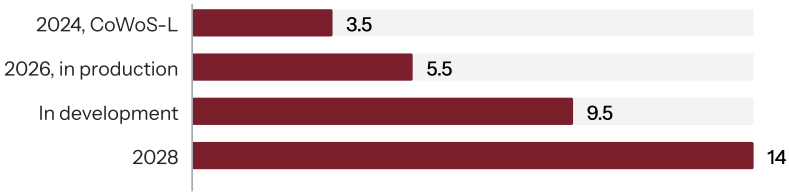
4 Advanced packaging and HBM, not logic dies, were the bottlenecks on AI chip production in 2025. Epoch AI · 12 March 2026 [A] epoch.ai/data-insights/ai-chip-supply-chain-constraints

7 The Great AI Silicon Shortage. SemiAnalysis · 12 March 2026 [B] newsletter.semianalysis.com/p/the-great-ai-silicon-shortage

12 Q2 2026 Taiwan Semiconductor Manufacturing Co Ltd Earnings Call (Chinese, English) — edited transcript. LSEG StreetEvents (transcript of a TSMC earnings call), via TSMC · 16 July 2026 [A] investor.tsmc.com/english/encrypt/files/encrypt_file/reports/2026-08/3e494f0c14dd0890f897a...

Who makes it

CoWoS interposer size on TSMC's roadmap RETICLES



Source: TSMC 2025 annual report and 2026 technology symposium

Three kinds of firm package chips. TSMC, which makes chips for others, and Intel, which makes its own, do it in their own back-end fabs.
Outsourced assembly and test firms do the rest. Intel opened Fab 9 in Rio Rancho, New Mexico, in January 2024 to make its Foveros 3D packages in volume [13]. It called the plant the only US factory producing the most advanced packaging at scale [13].

ASE Technology Holding, the leading provider of assembly and test services, had revenue of NT$645.4 billion in 2025, about 40 percent of it from electronics manufacturing services [14]. Its assembly, test and materials business brought in NT$389.2 billion, up from NT$325.9 billion in 2024 [14]. Amkor, the largest US-headquartered assembly and test firm, had net sales of $6.71 billion in 2025 [15].

The packaging tools are a narrower chokepoint than the assembly firms.

- **Thermocompression bonding.** Pressing dies together under heat and force. ASMPT's advanced packaging revenue reached $532.1 million in 2025, with thermocompression bonder revenue up about 146 percent [16]. ASMPT targets 35 to 40 percent of a market it puts at $1.6 billion in 2028 [16].
- **Hybrid bonding.** Copper pads bond directly, with no solder. Applied Materials co-developed its Kinex hybrid bonding system with Besi, which Applied calls the industry leader in hybrid bonding [17].
- **Grinding and dicing.** Disco held 87 percent of the $1.9 billion world market for dicing tools in 2025, by CSET's count from TechInsights data [18]. The category

13 Intel Opens Fab 9 in New Mexico. Intel [A] intel.com/content/www/us/en/newsroom/news/intel-opens-fab-9-new-mexico.html

14 ASE Technology Holding Co., Ltd. Reports Its Unaudited Consolidated Financial Results for the Fourth Quarter and the Full Year of 2025. PR Newswire · 5 February 2026 [A] prnewswire.com/news-releases/ase-technology-holding-co-ltd-reports-its-unaudited-consolida...

15 AMKOR TECHNOLOGY, INC., Form 8-K current report for the period ended 2026-02-09 (8-K). U.S. Securities and Exchange Commission (filing by AMKOR TECHNOLOGY, INC.) · 9 February 2026 [A] sec.gov/Archives/edgar/data/1047127/000104712726000007/amkr123125erex-991.htm

16 ASMPT Announces 2025 Annual Results AI-Driven Structural Growth Underpins Group Performance. ASMPT · 4 March 2026 [A] asmpt.com/en/investor-relations/news-events/asmpt-announces-2025-annual-results

17 Applied Materials, Kinex integrated die-to-wafer hybrid bonding system product page. Applied Materials [A] appliedmaterials.com/eu/en/product-library/kinex-integrated-die-to-wafer-hybrid-bonding-sy...

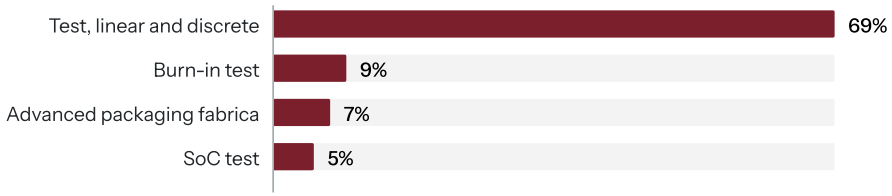
18 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

covers the grinders that thin wafers and the saws that cut them into chips [18]. See Test and assembly.

TSMC's Arizona commitment reached $265 billion in July 2026 and covers two advanced packaging plants alongside 10 fabs [19]. Amkor started building a $7 billion campus in Peoria in October 2025 and expects production from early 2028 [20]. The Commerce Department awarded Amkor up to $407 million in CHIPS Act funding for the site in December 2024 [21]. In October 2025 Wei said a large outsourced packaging partner, already breaking ground in Arizona, was on an earlier schedule than TSMC's own two packaging fabs there [11].

The chokepoint

China-based suppliers' share of selected tool segments, 2024 %



Source: CSET, Inside Beijing's Chipmaking Offensive

Four of the top ten outsourced assembly and test firms by 2024 revenue are headquartered in China [22]. JCET, a leading Chinese assembly and test firm, had record revenue of RMB 38.87 billion in 2025, up 8.1 percent [23]. JCET says its XDFOI platform has entered high-volume manufacturing with redistribution-layer interposer, silicon interposer and silicon bridge options [24].

Between 2019 and 2024, Chinese firms gained 5 percentage points of share in test tools and 4 points in fabrication tools built for advanced packaging [22]. But they lost share in bonding and other conventional assembly and packaging tools, the

18 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

19 TSMC Announces Additional $100 Billion Investment in Arizona. Arizona Commerce Authority · 16 July 2026 [A] azcommerce.com/news-events/news/2026/7/tsmc-announcement

20 Amkor Technology Breaks Ground on New Semiconductor Advanced Packaging and Test Campus in Arizona; Expands Investment to $7 Billion. Amkor Technology [A] ir.amkor.com/news-releases/news-release-details/amkor-technology-breaks-ground-new-semicon...

21 AMKOR TECHNOLOGY, INC., Form 10-K annual report for the period ended 2025-12-31 (10-K). U.S. Securities and Exchange Commission (filing by AMKOR TECHNOLOGY, INC.) · 20 February 2026 [A] sec.gov/Archives/edgar/data/1047127/000104712726000014/amkr-20251231.htm

11 Q3 2025 Taiwan Semiconductor Manufacturing Co Ltd Earnings Call (Chinese, English) — edited transcript. Refinitiv StreetEvents (transcript of a TSMC earnings call), via TSMC · 16 October 2025 [A] investor.tsmc.com/english/encrypt/files/encrypt_file/reports/2025-10/6860312f04fd291d0f26b...

22 Inside Beijing's Chipmaking Offensive. Center for Security and Emerging Technology (CSET) · 14 July 2025 [A] cset.georgetown.edu/article/inside-beijings-chipmaking-offensive

23 JCET Releases 2025 Annual Report, Posts Record-High Full-Year Revenue and Higher Profit Before Tax. JCET Group · 9 April 2026 [A] jcetglobal.com/en/site/news-detail

24 JCET Accelerates Strategic Shift Toward High-End Advanced Packaging, 2025 Advanced Packaging Revenue Hits Record High. JCET Group · 21 April 2026 [A] jcetglobal.com/en/site/news-detail

category where CSET counts hybrid and thermocompression bonders [22]. Within the advanced packaging fabrication tools, Chinese firms gained 10 points in etch and clean tools and 3 points in deposition tools, and lost 2 points each in lithography and polishing (CMP) tools [22]. In 2024 Chinese firms had 7 percent of the market for advanced packaging fabrication tools, against 65 percent for US firms and 15 percent for Japanese firms [22].

Each generation on TSMC's roadmap needs a larger interposer.

- **2028.** A 14-reticle interposer of roughly 12,000 mm2, about the area of a compact disc, carrying about ten large compute dies and 20 HBM stacks [9].
- **2029.** More than 14 reticles, plus a 40-reticle System-on-Wafer [9].

Key evaluation criteria

- **Interposer area** limits how much logic and memory fit in one package, and stood at 5.5x reticle on CoWoS-L in April 2026 [9].
- **Package yield**, the share of finished packages that work, determines whether a generation ships on time. Nobody publishes it.
- **Warpage control** means managing the different expansion rates of die, bridge, interposer and substrate. That mismatch caused warpage in Blackwell, the first high-volume CoWoS-L design [8].
- **HBM stacks supported** reach about 20 on the 14-reticle package due in 2028 [9].
- **Bonder throughput** depends on the toolmakers. ASMPT alone targets 35 to 40 percent of the thermocompression bonding market [16].
- **Substrate size** limits everything above it. A 14-reticle interposer needs a resin substrate that will not warp under it, and that pushes ABF build-up film, the insulating material in the substrate's layers, to its limits. See Substrates and PCBs.

FOUR THINGS TO REMEMBER

— Many leading accelerators, Nvidia's Blackwell among them, need more logic than one chip can hold, so packaging joins

22 Inside Beijing's Chipmaking Offensive. Center for Security and Emerging Technology (CSET) · 14 July 2025 [A] cset.georgetown.edu/article/inside-beijings-chipmaking-offensive

9 TSMC Debuts A13 Technology at 2026 North America Technology Symposium. TSMC · 23 April 2026 [A] pr.tsmc.com/english/news/3302

8 Nvidia's Blackwell Reworked - Shipment Delays & GB200A Reworked Platforms. SemiAnalysis · 4 August 2024 [B] semianalysis.com/2024/08/04/nvidias-blackwell-reworked-shipment

16 ASMPT Announces 2025 Annual Results AI-Driven Structural Growth Underpins Group Performance. ASMPT · 4 March 2026 [A] asmpt.com/en/investor-relations/news-events/asmpt-announces-2025-annual-results

several chips and their memory.

- TSMC packages nearly every AI accelerator, in Taiwan.
- Packaging was one of the main limits on AI chip output in 2025. In July 2026 TSMC said its packaging capacity was still so tight that it limited customers' growth.
- No US site yet packages Nvidia-class accelerators with HBM on CoWoS, though Intel makes its own Foveros 3D packages in volume in New Mexico.

CHAPTER
12
STAGE 12 OF 15

Substrates and PCBs

Ajinomoto, a Japanese food company, told investors in June 2026 that it holds over 95 percent of the market for the insulating film inside high-end chip packages.

IN PLAIN TERMS

A package substrate is the adapter between a chip and the circuit board under it. The chip's connections are finer than a human hair, but the board is wired in millimeters. The substrate is built in layers, like plywood, around a stiff core of glass cloth and resin. Thin sheets of epoxy film are pressed onto both faces of the core. A laser burns holes through each sheet, and copper plated into the holes carries signals from one layer to the next. Every added layer is another set of holes that can fail. A handful of firms build these substrates. Japan's Ajinomoto dominates the film, and Japan's Nitto Boseki dominates the low-expansion glass cloth in the core.

IN SHORT

Ajinomoto counts 22 layers of its film in a 2026 AI package substrate about 100 mm square, against 18 on a high-performance computing substrate about 70 mm square around 2023 [1]. Its third site for making the film, in Gifu Prefecture, is scheduled to start operating in 2032 to meet demand after 2030 [1]. Nitto Boseki says market demand seems to have risen about two to three times since it set its plan for fiscal 2024 to 2027 [3]. Apple and Qualcomm are among the customers worried about Japanese glass cloth supply [11].

1 Ajinomoto, Growth Strategy for the Electronic Materials Business, 30 June 2026. Ajinomoto · 30 June 2026 [A] ajinomoto.co.jp/company/en/ir/event/business_briefing/main/011117/teaserItems1/01/linkList...

3 Nittobo Group Integrated Report 2025. Nitto Boseki (Nittobo) · 16 September 2025 [A] nittobo.co.jp/eng/ir/pdf/Nittobo_Integrated-Report_2025.pdf

11 Apple and Qualcomm fret over strained supplies of Japan's glass cloth. Nikkei Asia · 14 January 2026 [B] asia.nikkei.com/business/technology/tech-asia/apple-and-qualcomm-fret-over-strained-suppli...

CHOKEPOINT CARD | EXTREME CONCENTRATION | HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
Extreme	Hard Rated on the insulating film inside the substrate. Ajinomoto supplies over 95 percent of it, by its own June 2026 count. Every substrate maker would have to re-approve a replacement layer by layer.	**Ibiden is spending about JPY 500 billion on high-performance substrate capacity from fiscal 2026 to 2028, about JPY 220 billion of it in the first phase**

WHO LEADS

- JP **Ajinomoto** Told investors in June 2026 that its build-up film has held over 95% of its market since launch
- JP **Nitto Boseki** Claims an overwhelming advantage in low thermal expansion glass
- JP **Ibiden** Says it typically takes close to 100% of each new generation of interposer-type substrates at launch, a share that may fall by 20 to 30% after about three to six months
- TW **Unimicron** Build-up substrates were 52% of sales in Q2 2026
- KR **Samsung Electro-Mechanics** Package solution sales KRW 771.6bn in Q2 2026, up 37%

WHERE IT IS MADE

- JP **Japan** Build-up film, low-expansion glass cloth, Ibiden and Shinko substrate plants in Gifu and Nagano, and a plant in Chiba that Shinko agreed to buy in September 2026
- TW **Taiwan** Unimicron, Nan Ya PCB and Kinsus, substrate makers that sell to outside customers
- KR **South Korea** Samsung Electro-Mechanics, LG Innotek, and a planned glass-core joint venture with Sumitomo Chemical's Dongwoo Fine-Chem
- CN **China** Most of TTM's circuit board competitors are based in China and Taiwan. Chinese firms are moving into special glass cloth
- AT **Austria** AT&S, which says it is the only company in Europe able to develop and produce package substrates locally

WHY SUBSTITUTION IS SLOW

A few firms build substrates, and money can expand their plants. The materials inside the substrate are harder to replace. Ajinomoto makes the thin insulating film between the substrate's wiring layers. The film has been the standard since 1999. Ajinomoto told investors in June 2026 that the film has held over 95 percent of its market since launch. Beneath that film is Nitto Boseki's glass cloth, made from a glass that expands little with heat. A newcomer can build a line to make a rival film. But every substrate maker and chip customer would then have to approve it, in each layer of each product.

WHERE CHINA STANDS

TTM, an American circuit board maker, says its circuit board competitors are mostly based in China and Taiwan. Nitto Boseki, a Japanese glass fiber maker, says Chinese and Taiwanese companies are

aggressively catching up with it in the special glass sector. Ajinomoto, a Japanese company, told investors in June 2026 that its build-up film has held over 95 percent of its market since launch.

WHERE THE US STANDS

The White House's 2021 supply chain review listed no American firm among the key substrate makers. A May 2024 report by the Semiconductor Industry Association and Boston Consulting Group counted none among the major package substrate vendors. Since the 2021 review, the Commerce Department has made a final award of up to $75 million to Absolics, a subsidiary of Korea's SKC, for a glass substrate plant in Covington, Georgia. As of October 2026, the award's project page at NIST, a Commerce Department agency, expected the plant's production capacity to come online in 2027. TTM Technologies, an American circuit board maker, calls itself one of the largest technology manufacturers in North America and says it offers package substrates. Data center computing was 24 percent of its sales in 2025, up from 14 percent in 2023. Intel said in 2023 that it planned glass substrates for the latter part of this decade.

The silicon interposer wires chips together inside a package. Beyond the interposer, the wiring runs through a build-up substrate made of epoxy and glass, then a printed circuit board of thirty layers, then a server rack.

How it works

The substrate maker drills the core and plates it with copper before the build-up starts. In each cycle the maker then presses on a film of epoxy, burns holes through it with a laser to connect the layers, and adds copper wiring on top [1]. Each film is about 10 micrometers thick, and a human hair is about 80 micrometers across [2].

Ajinomoto counts 22 layers of its film in a 2026 AI package substrate about 100 mm square, against 18 on a high-performance computing substrate about 70 mm square around 2023 [1]. PC substrates had six layers of the film when PCs first adopted it in 1999 [1]. Silicon and epoxy swell at different rates when heated, so the larger the package, the more it warps. Nitto Boseki says its low-expansion T-glass

1 Ajinomoto, Growth Strategy for the Electronic Materials Business, 30 June 2026. Ajinomoto · 30 June 2026 [A] ajinomoto.co.jp/company/en/ir/event/business_briefing/main/011117/teaserItems1/01/linkList...

2 The Ajinomoto Group's Unexpected Role in Semiconductor Manufacturing: The Insulating Film "ABF" Born from "AminoScience". Ajinomoto · 5 March 2026 [A] ajinomoto.com/stories/the-ajinomoto-groups-unexpected-role-in-semiconductor-manufacturing-...

cloth is highly rated as an effective fix for that warping, and that its use in substrates has increased significantly [3].

Variants and trade-offs

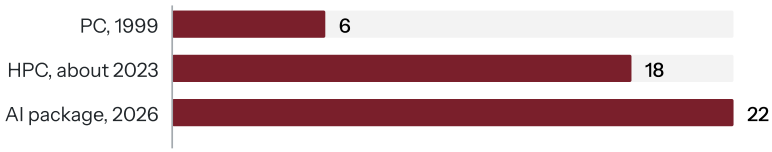
BT substrates

Substrates built on bismaleimide-triazine laminate, or BT, are the other main family. Mitsubishi Gas Chemical, which makes BT materials, says they go mostly into thin packages for smartphone processors, radio modules and memory [4]. The company says BT also goes into flip-chip substrates for computer CPUs and GPUs [4]. In these the chip is mounted face down, and BT laminate forms the core between layers of build-up film [4]. In an October 2024 investor presentation, the company estimated that its BT resins held about 40 percent of the substrate materials market [4].

ABF build-up substrates

Ajinomoto Build-up Film, or ABF, is the insulating film between the layers of package substrates for logic chips, used in PCs, network equipment, servers and AI equipment [1]. Ajinomoto first developed it in 1999 and says major chipmakers have used it as the de facto standard ever since [5]. Asked on Ajinomoto's own site whether its share was nearer 30 or 50 percent, the electronic materials division answered, "More like 95 percent market share" [2]. It said practically all high-performance computers and servers rely on the film [2]. In June 2026 Ajinomoto told investors that the film has held a share of over 95 percent since its launch [1].

Layers of build-up film in one package substrate LAYERS



Source: Ajinomoto, Growth Strategy for the Electronic Materials Business, June 2026

3 Nittobo Group Integrated Report 2025. Nitto Boseki (Nittobo) · 16 September 2025 [A] nittobo.co.jp/eng/ir/pdf/Nittobo_Integrated-Report_2025.pdf

4 Electronics Materials Business (Mitsubishi Gas Chemical business briefing). Mitsubishi Gas Chemical · 2 October 2024 [A] mgc.co.jp/eng/ir/files/2410022e.pdf

1 Ajinomoto, Growth Strategy for the Electronic Materials Business, 30 June 2026. Ajinomoto · 30 June 2026 [A] ajinomoto.co.jp/company/en/ir/event/business_briefing/main/011117/teaserItems1/01/linkList...

5 Ajinomoto, Business Briefing: ABF-Based Growth Strategy in ICT, 12 June 2023. Ajinomoto · 12 June 2023 [A] ajinomoto.co.jp/company/en/ir/event/business_briefing/main/01113/teaserItems1/01/linkList/...

2 The Ajinomoto Group's Unexpected Role in Semiconductor Manufacturing: The Insulating Film "ABF" Born from "AminoScience". Ajinomoto · 5 March 2026 [A] ajinomoto.com/stories/the-ajinomoto-groups-unexpected-role-in-semiconductor-manufacturing-...

Glass core substrates

A glass core expands less with heat than a resin core and stays flatter, so chipmakers and substrate makers are both trying it [6].

- **Absolics**, a subsidiary of Korea's SKC, has a final Commerce Department award of up to $75 million for a 120,000-square-foot glass substrate plant in Covington, Georgia [7]. As of October 2026, the award's project page at NIST, a Commerce Department agency, expected first deliveries to customers in 2025 and production capacity to come online in 2027 [7].
- **Intel** announced glass substrates in September 2023 for the latter part of this decade [8].
- **Samsung Electro-Mechanics** signed a joint venture agreement with Sumitomo Chemical's subsidiary Dongwoo Fine-Chem in July 2026 to make glass core substrates [9]. The partners plan to have a supply system in place by the second half of fiscal 2027 [9].

The materials underneath

Ajinomoto's Healthcare and Others segment, which includes electronic materials such as the film, lifted business profit 45.1 percent to JPY 66.2 billion in the year to March 2026 [10]. Ajinomoto attributes the rise mainly to higher electronic materials revenue and higher profit in Bio-Pharma Services & Ingredients [10]. It has set an investment budget of JPY 25 billion between 2023 and 2030 to expand film production [1]. It says there are no concerns about its current supply chain overall [1].

6 Samsung Electro-Mechanics Signs MOU with Sumitomo Chemical Group to Establish a Joint Venture for Glass Core Used in Package Substrates. Samsung Electro-Mechanics · 5 November 2025 [A] samsungsem.com/global/newsroom/news/view.do

7 Absolics (Georgia). NIST · 5 July 2024 [A] nist.gov/chips/absolics-georgia-covington

8 Intel Unveils Industry-Leading Glass Substrates to Meet Demand for More Powerful Compute. Intel · 18 September 2023 [A] download.intel.com/newsroom/archive/2025/en-us-2023-09-18-intel-unveils-industryleading-gl...

9 Sumitomo Chemical Subsidiary to Establish of a Joint Venture with Samsung Electro-Mechanics for the Glass Core Substrate Business. Sumitomo Chemical · 2 July 2026 [A] sumitomo-chem.co.jp/english/news/detail/20260702e_2.html

10 Consolidated Financial Results for the Fiscal Year Ended March 31, 2026 (Under IFRS). Ajinomoto · 7 May 2026 [A] ajinomoto.co.jp/company/en/ir/library/result/main/014/teaserItems1/0/linkList/06/link/FY25...

1 Ajinomoto, Growth Strategy for the Electronic Materials Business, 30 June 2026. Ajinomoto · 30 June 2026 [A] ajinomoto.co.jp/company/en/ir/event/business_briefing/main/011117/teaserItems1/01/linkList...

In May 2026 it announced the acquisition of land in Gifu Prefecture for a third production site, scheduled to start operating in 2032 to meet demand after 2030 [1].

The glass cloth in the core also has a dominant supplier. Nitto Boseki's T-glass is stiff and expands little when heated, and the company calls it essential for making high-density package substrates [3]. Laminate makers such as Mitsubishi Gas Chemical coat the glass cloth with resin and press it with copper foil into the core laminate that substrate makers build on [4].

Nitto Boseki claims an overwhelming advantage in quality and capacity for low-expansion glass [3]. It planned about JPY 80 billion of capital spending for fiscal 2024 to 2027, and says market demand seems to have risen about two to three times since that plan was set [3]. In August 2025 it decided to expand glass fiber cloth production at its Fukushima Enterprise Center, with production scheduled to begin in the fourth quarter of its fiscal 2026 [3]. Apple and Qualcomm are among the customers worried about Japanese glass cloth supply [11].

Nitto Boseki says Chinese and Taiwanese companies are aggressively catching up with it in the special glass sector [3]. Nikkei reports that Chinese chip materials makers are competing with Japanese rivals for a market it sizes at $73 billion [12].

Who makes it

A few Asian companies build substrates for other firms. The White House's 2021 hundred-day supply chain review listed build-up substrate among the materials that commenters called vulnerable [13]. It named key substrate makers including Ibiden, Nanya, Shinko, Samsung, Unimicron, Shennan Circuits, Zhuhai Yueya and AKM, none of them American [13].

1 Ajinomoto, Growth Strategy for the Electronic Materials Business, 30 June 2026. Ajinomoto · 30 June 2026 [A] ajinomoto.co.jp/company/en/ir/event/business_briefing/main/011117/teaserItems1/01/linkList...

3 Nittobo Group Integrated Report 2025. Nitto Boseki (Nittobo) · 16 September 2025 [A] nittobo.co.jp/eng/ir/pdf/Nittobo_Integrated-Report_2025.pdf

4 Electronics Materials Business (Mitsubishi Gas Chemical business briefing). Mitsubishi Gas Chemical · 2 October 2024 [A] mgc.co.jp/eng/ir/files/2410022e.pdf

11 Apple and Qualcomm fret over strained supplies of Japan's glass cloth. Nikkei Asia · 14 January 2026 [B] asia.nikkei.com/business/technology/tech-asia/apple-and-qualcomm-fret-over-strained-suppli...

12 China chip material makers battle Japan rivals for $73bn market. Nikkei Asia · 1 July 2026 [B] asia.nikkei.com/business/tech/semiconductors/china-chip-material-makers-battle-japan-rival...

13 Building Resilient Supply Chains, Revitalizing American Manufacturing, and Fostering Broad-Based Growth: 100-Day Reviews under Executive Order 14017. The White House (Biden administration archive) · June 2021 [A] bidenwhitehouse.archives.gov/wp-content/uploads/2021/06/100-day-supply-chain-review-report...

A May 2024 report by the Semiconductor Industry Association and Boston Consulting Group counted the major package substrate vendors by headquarters [14]. It found two in Japan, two in Taiwan, one in mainland China and none in the United States [14]. TTM Technologies, an American circuit board maker, says it also offers package substrates [15].

- **Ibiden** will invest about JPY 500 billion in high-performance package substrates from fiscal 2026 to fiscal 2028, starting with about JPY 220 billion, mainly at its Gama plant [16]. It planned mass production from fiscal 2027 [16]. In May 2026 management said Ibiden typically takes close to 100 percent of a new generation of interposer-type substrates at launch [17]. It added that this share may decline by 20 to 30 percent after about three to six months [17]. On silicon-bridge substrates, it said wiring the underside is hard enough that Ibiden expects to keep its technical lead at least through fiscal 2030 [17].
- **Unimicron** reported NT$42.9 billion of sales in the second quarter of 2026, 52 percent of it build-up substrates and 61 percent of it AI data center work [18].
- **Samsung Electro-Mechanics** grew package solution sales 37 percent year on year to KRW 771.6 billion in the second quarter of 2026, on substrates for AI accelerators, server processors and cars [19].

In February 2025 a company owned mainly by funds of JIC Capital, a subsidiary of Japan Investment Corporation, launched a tender offer to take Shinko, a Japanese substrate maker, private [20]. The offer covered the shares Fujitsu did not own, and

14 Emerging Resilience in the Semiconductor Supply Chain. Boston Consulting Group and Semiconductor Industry Association · May 2024 A semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc…

15 TTM TECHNOLOGIES INC, Form 10-K annual report for the period ended 2025-12-29 (10-K). U.S. Securities and Exchange Commission (filing by TTM TECHNOLOGIES INC) · 17 February 2026 A sec.gov/Archives/edgar/data/1116942/000119312526051976/ttmi-20251229.htm

16 Ibiden, Notice Regarding Capital Investment Plan for High-Performance IC Package Substrates. Ibiden · 3 February 2026 A ibiden.com/company/2026/02/notice-regarding-capital-investment-plan-for-high-performance-i…

17 Ibiden, questions and answers from the financial presentation for the year ended 31 March 2026. Ibiden · 12 May 2026 A ibiden.com/ir/items/en_QA_FY25Q4.pdf

18 Unimicron 2026 Q2 Earnings Conference. Unimicron Technology · July 2026 A unimicron.com/files/money/Earnings/en/2026-Q2-consolidated-en.pdf

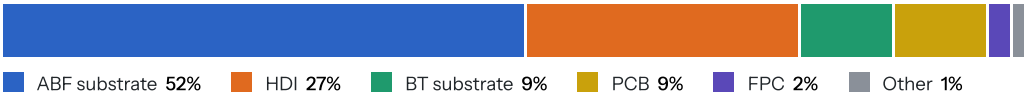
19 Samsung Electro-Mechanics, Q2 2026 business results. Samsung Electro-Mechanics · 30 July 2026 A m.samsungsem.com/global/newsroom/news/view.do

20 Announcement Regarding Commencement of the Tender Offer for the Shares of Shinko Electric Industries Co., Ltd. (Securities Code: 6967) by JICC-04, Ltd.. Japan Investment Corporation · 17 February 2025 A jiccapital.co.jp/en/news/.assets/E_20250217_JIC_JICC_PressRelease.pdf

Fujitsu agreed to sell its 50.02 percent stake back to Shinko [20]. The tender offer succeeded in March 2025 [21].

In September 2026 Shinko agreed to buy a plant in Mobara, Chiba Prefecture, from Japan Display, a display maker [22]. Shinko said it would make the site a new core plant for advanced chip packages and would also expand capacity at its plants in Nagano [22]. Production at the Mobara plant is scheduled to start in fiscal 2028 [22].

Unimicron sales by technology, second quarter 2026 %



ABF substrate 52% HDI 27% BT substrate 9% PCB 9% FPC 2% Other 1%

Source: Unimicron 2026 Q2 earnings conference

The boards under the GPU

TTM Technologies says it regularly builds printed circuit boards of more than 30 layers and can go past 70 [15]. Those layers carry power to the GPUs and route fast signals between them, so how much signal the laminate absorbs is a main design constraint. Nitto Boseki also makes NE-glass, a low-dielectric glass cloth that absorbs less signal, used in switches and routers in data centers [3].

Data center computing was 24 percent of TTM's $2.91 billion of fiscal 2025 revenue, against 14 percent two years earlier [15]. TTM says its circuit board competitors are mostly based in China and Taiwan [15]. In its annual report for 2025, TTM said it was building a plant in Syracuse, New York, to make ultra-high-density circuit boards for national security needs [15]. It expected volume

20 Announcement Regarding Commencement of the Tender Offer for the Shares of Shinko Electric Industries Co., Ltd. (Securities Code: 6967) by JICC-04, Ltd.. Japan Investment Corporation · 17 February 2025 A jiccapital.co.jp/en/news/.assets/E_20250217_JIC_JICC_PressRelease.pdf

21 (Progress Report on Disclosed Matters) Notice Regarding Completion of JICC-04's Tender Offer for Shares of a Consolidated Subsidiary of Fujitsu Limited (SHINKO ELECTRIC INDUSTRIES CO., LTD.). Fujitsu · 19 March 2025 A pr.fujitsu.com/jp/news/2025/03/en/19-1.pdf

22 Notice Concerning the Acquisition of Japan Display Inc.'s Mobara Fab. Shinko Electric Industries · 30 September 2026 A shinko.co.jp/english/news/docs/20260930-01_en.pdf

15 TTM TECHNOLOGIES INC, Form 10-K annual report for the period ended 2025-12-29 (10-K). U.S. Securities and Exchange Commission (filing by TTM TECHNOLOGIES INC) · 17 February 2026 A sec.gov/Archives/edgar/data/1116942/000119312526051976/ttmi-20251229.htm

3 Nittobo Group Integrated Report 2025. Nitto Boseki (Nittobo) · 16 September 2025 A nittobo.co.jp/eng/ir/pdf/Nittobo_Integrated-Report_2025.pdf

production there in the second half of 2026 [15]. Thailand aims to be the biggest circuit board hub outside China as board makers diversify production [23].

The chokepoint

The few firms that build substrates can add capacity by spending money, and they are spending it. The materials inside the substrate are harder to replace. Ajinomoto dominates the film, and Nitto Boseki dominates the low-expansion glass cloth in the core. Both are adding capacity on plans that run for years.

Packages are also getting bigger. TSMC told investors in July 2026 that its CoWoS packaging roadmap goes beyond 14 times the area a lithography machine can print in one shot [24]. A bigger package warps more and needs more layers, so it has more holes that can break the circuit. That lowers yield, the share of substrates that come out working, and each substrate is already expensive before a chip is mounted on it.

Key evaluation criteria

- **Layer count** sets how much current and how many signals the substrate can carry. Each added layer is another chance to lose the part.
- **Body size**, the substrate's length and width, affects every other criterion. Once a body is much larger than the chips it carries, warping and handling replace the wiring as the main limit on what can be built.
- **Coefficient of thermal expansion**, how much a material swells when heated, must be close enough to silicon's that the package survives repeated heating and cooling. T-glass in the core provides that match.
- **Dielectric properties**, how much signal the core laminate absorbs, set the speeds the substrate and the board underneath can run.
- **Qualification time** protects the incumbents. Each substrate maker and chip customer has to approve a second supplier's film or laminate before using it.

FOUR THINGS TO REMEMBER

— A package substrate connects a chip's fine contacts to a circuit board wired in millimeters.

— One Japanese supplier leads the market for the substrate's insulating film, and another leads the market for its glass cloth.

15 TTM TECHNOLOGIES INC, Form 10-K annual report for the period ended 2025-12-29 (10-K). U.S. Securities and Exchange Commission (filing by TTM TECHNOLOGIES INC) · 17 February 2026 [A]
sec.gov/Archives/edgar/data/1116942/000119312526051976/ttmi-20251229.htm

23 Thailand carves out a less-glamorous AI niche: printed circuit boards. Nikkei Asia · 3 September 2025 [B]
asia.nikkei.com/business/technology/tech-asia/thailand-carves-out-a-less-glamorous-ai-nich…

24 Q2 2026 Taiwan Semiconductor Manufacturing Co Ltd Earnings Call (Chinese, English) — edited transcript. LSEG StreetEvents (transcript of a TSMC earnings call), via TSMC · 16 July 2026 [A]
investor.tsmc.com/english/encrypt/files/encrypt_file/reports/2026-08/3e494f0c14dd0890f897a…

- Replacing the film would mean re-approval by every substrate maker and chip customer.
- The White House's 2021 supply chain review named no American firm among the key substrate makers. Absolics is building a glass substrate plant in Georgia, and a Commerce Department agency expects its production capacity to come online in 2027.

CHAPTER
13
STAGE 13 OF 15

Test and Assembly

Testing finds bad chips before they go into an expensive package. Japan's Advantest took about 65 percent of 2025 spending on chip testers, by its own estimate.

IN PLAIN TERMS

Before a chip can be sold, a machine called a tester has to check that it works. While the chips are still on the wafer, thousands of needles press onto small metal pads on each one. The tester sends signals through the needles and reads what comes back, and any chip that fails goes no further. A modern AI accelerator is a dozen expensive pieces bonded into one package, and if one piece is bad, all of them are thrown away. So the industry now tests chips earlier, at higher temperatures and for longer, to make weak ones fail during testing before they reach a finished product. Two firms, Advantest in Japan and Teradyne in the United States, sell most of the testers.

IN SHORT

Advantest argues that switching to another tester maker means rebuilding a customer's whole test setup, so market share moves slowly [36]. Georgetown's CSET finds that between 2019 and 2024 Chinese toolmakers gained 5 percentage points of share in test tools [26]. Across assembly, test and packaging tools, they still lag significantly behind foreign suppliers [26]. Four of the ten largest chip assembly and test contractors by 2024 revenue are headquartered in China [26].

36 **Investors Guide April 25, 2025.** Advantest · 13 May 2025 [A] advantest.com/document/en/investors/ir-library/investors-guide/Investors_Guide_2504E.pdf

26 **Inside Beijing's Chipmaking Offensive.** Center for Security and Emerging Technology (CSET) · 14 July 2025 [A] cset.georgetown.edu/article/inside-beijings-chipmaking-offensive

CHOKEPOINT CARD

HIGH CONCENTRATION

MODERATELY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
High	Moderate No test or assembly tool has only one maker, but switching testers means rebuilding the customer's whole test setup.	**ASE's packaging and test business (about $12.1 billion) and Amkor ($6.71 billion) billed about $18.8 billion between them in 2025** ASE spent $4.1 billion on machinery, buildings, facilities and automation in the first half of 2026

WHO LEADS

- JP **Advantest** About 65% of the tester market in 2025, by its own estimate
- US **Teradyne** Second in testers, with $2.52bn of semiconductor test revenue in 2025, including system-level testers and service
- TW **ASE** The largest assembly and test contractor. About 59 percent of its $20.6bn of 2025 revenue came from packaging and test and 40 percent from electronics manufacturing
- US **Amkor** $6.71bn of 2025 revenue
- JP **Disco** 87% of the $1.9bn market for dicing saws and wafer grinders in 2025, on CSET's analysis of TechInsights data

WHERE IT IS MADE

- TW **Taiwan** The largest assembly and test cluster, home to ASE, Powertech, KYEC and ChipMOS
- JP **Japan** Advantest testers, Accretech and Tokyo Electron probers, Disco and Accretech dicing and grinding, Micronics Japan probe cards
- US **United States** Test equipment from Teradyne, FormFactor, Cohu and Aehr, and Amkor's Arizona packaging plant
- CN **China** Four of the top ten assembly and test firms by 2024 revenue are headquartered in China
- MY **Malaysia** Penang and Kulim. Malaysia said in July 2026 that it ships 13% of the world's packaged chips

WHY SUBSTITUTION IS POSSIBLE

Advantest and Teradyne sell most of the testers between them, but customers have approved machines from both. The smaller kinds of test and assembly tool each have several makers. Advantest tells its investors that customers are highly cautious about switching, because a new tester means rebuilding their whole setup for development, evaluation and volume production. For a new tester maker, that switching cost is a one-time barrier, and getting its machines approved by customers would take two to five years. Four of the top ten outsourced assembly and test firms by 2024 revenue are based in China.

WHERE CHINA STANDS

CSET at Georgetown finds that between

2019 and 2024 Chinese toolmakers gained 5 percentage points of share in test tools and 4 points in fabrication tools built for advanced packaging. But they lost share in bonding and other conventional assembly and packaging tools. They still lag significantly behind foreign suppliers. CSET counts four of the top ten assembly and test firms by 2024 revenue as headquartered in China.	WHERE THE US STANDS American firms, among them Teradyne, FormFactor, Cohu and Aehr, are strong in test machines and probe cards, the circuit boards whose needles touch each chip on the wafer. They are weak in assembly. Amkor is building a $7 billion packaging and test campus in Arizona and has a long-term deal to run packaging and test there for TSMC.

Turning silicon into a finished, packaged chip takes 1,000 to 1,500 steps, depending on how complex the chip is, the Semiconductor Industry Association estimated in 2025 [1]. Advantest's investor guide of July 2026 says test is the only step that puts electrical signals through the chip [2]. The earlier stages of chipmaking make dies, the individual chips on a wafer. The back end of the industry, assembly and test, checks which of them work and puts the good ones into packages.

How it works

Each round of test catches failures the round before it cannot see.

- **Wafer sort.** Thousands of needles land on each die while it is still on the wafer. A tester sends in test patterns and checks the answers, and dies that fail go no further.
- **Burn-in.** The part runs hot and at raised voltage for minutes to days, on the wafer or in the package, so that early-life failures happen in the factory, before the part ships [3].
- **Final test.** The tester checks the packaged part again, because packaging adds failures of its own.
- **System-level test.** The finished part sits in a socket and runs something close to a real workload, because some failures appear only when everything runs at once.

As the parts in a package grow more valuable, finding a bad die before packaging matters more [4]. Aehr's wafer-level burn-in systems look for early-life and hidden failures before dies are cut from the wafer [4]. Catching those failures early lowers

1 Testimony of David Isaacs, Semiconductor Industry Association, hearing on “Trade in Critical Supply Chains”. U.S. Senate Committee on Finance · 14 May 2025 A finance.senate.gov/download/051425-isaacs-testimony&download=1

2 Investors Guide July 29, 2026. Advantest · 30 July 2026 A advantest.com/document/en/investors/ir-library/investors-guide/Investors_Guide_2607E.pdf

3 AEHR TEST SYSTEMS, Form 10-K annual report for the period ended 2026-05-29. U.S. Securities and Exchange Commission (filing by AEHR TEST SYSTEMS) · 27 July 2026 A sec.gov/Archives/edgar/data/1040470/000165495426006919/aehr_10k.htm

4 Aehr Receives $22 Million Follow-On Order for AI Processor Wafer-Level Burn-In Systems. Aehr Test Systems · 12 August 2026 A aehr.com/2026/08/aehr-receives-22-million-follow-on-order-for-ai-processor-wafer-level-bur...

the risk of packaging high-bandwidth memory and an expensive substrate with a processor that may later fail [4]. Aehr's lead AI customer is shifting burn-in for its AI accelerators from the finished system to the wafer [5]. The FOX-XP systems that customer ordered from Aehr in August 2026 each burn in nine 300 mm wafers at a time [4].

Variants and trade-offs

Wafer sort and known-good die

The hard part of wafer sort is the probe card, the circuit board that holds the needles. FormFactor says the high-frequency capability of its probe cards lets it compete favorably in testing high-bandwidth memory, a stack of eight, twelve or sixteen memory dies [6]. For a large logic die every needle must sit in the same plane, survive repeated touchdowns without drifting, and do it hot. Each card is built for one chip design, so every new design needs a new card and a new qualification.

Final test, burn-in and system-level test

The bigger the package, the more one bad die can cost. TSMC has told investors that its CoWoS packaging roadmap runs beyond 14 times the area a lithography machine can print in one shot [7]. Every added die is another chance to lose the whole assembly, along with its memory and substrate [4]. ASE, the largest assembly and test contractor, ran 6,797 testers in the second quarter of 2025 and 8,348 a year later [8].

Automated test equipment: share of the ~$9.0bn tester market, 2025 %



Advantest 65% Teradyne and all others 35%

Source: Advantest FY2025 results briefing, April 2026

4 Aehr Receives $22 Million Follow-On Order for AI Processor Wafer-Level Burn-In Systems. Aehr Test Systems · 12 August 2026 A aehr.com/2026/08/aehr-receives-22-million-follow-on-order-for-ai-processor-wafer-level-bur…

5 Aehr Test Systems Reports Fiscal 2026 Fourth Quarter and Full Year Financial Results with Record Quarterly Bookings and $100 Million Effective Backlog. Aehr Test Systems · 14 July 2026 A aehr.com/2026/07/aehr-test-systems-reports-fiscal-2026-fourth-quarter-and-full-year-financ…

6 FORMFACTOR INC, Form 10-K annual report for the period ended 2025-12-27 (10-K). U.S. Securities and Exchange Commission (filing by FORMFACTOR INC) · 20 February 2026 A sec.gov/Archives/edgar/data/1039399/000103939926000009/form-20251227.htm

7 Q2 2026 Taiwan Semiconductor Manufacturing Co Ltd Earnings Call (Chinese, English) — edited transcript. LSEG StreetEvents (transcript of a TSMC earnings call), via TSMC · 16 July 2026 A investor.tsmc.com/english/encrypt/files/encrypt_file/reports/2026-08/3e494f0c14dd0890f897a…

8 ASE Technology Holding Co., Ltd., Form 6-K report of foreign private issuer for the period ended 2026-07-30. U.S. Securities and Exchange Commission (filing by ASE Technology Holding Co., Ltd.) · 30 July 2026 A sec.gov/Archives/edgar/data/1122411/000095010326011351/dp250868_6k.htm

Who makes it

Advantest puts its own share at about 65 percent of an estimated $9.0 billion tester market in 2025 [9]. Its share of system-on-chip (SoC) testers, the machines that test processors and other non-memory chips, rose about ten points in a year to 66 percent [9]. It held about 60 percent of memory test [9].

Teradyne, Advantest's main rival, took $2.52 billion of semiconductor test revenue in 2025, service included [10]. Of that, $1.89 billion came from SoC testers and $505 million from memory testers [10]. Teradyne's total revenue grew 13 percent to $3.19 billion in 2025, with fourth-quarter revenue up 44 percent on AI demand [11].

Where the tester money went in calendar 2025 $BN



SoC testers 6.9 Memory testers 2.1

Source: Advantest FY2025 results briefing, April 2026

The rest of the test and assembly equipment market splits into niches, each with its own suppliers.

- **Probe cards.** FormFactor took $785.0 million of revenue in fiscal 2025 [12]. Technoprobe's revenue grew 15.7 percent to EUR 628.4 million in 2025, about 38 percent of it from AI, and it aims to double capacity by the end of 2027 [13].

9 Advantest, presentation notes for the FY2025 (year ended 31 March 2026) results briefing, 27 April 2026. Advantest · 27 April 2026 [A] advantest.com/document/en/investors/ir-library/result/JE_BIZ_260427_note.pdf

10 TERADYNE, INC, Form 10-K annual report for the period ended 2025-12-31 (10-K). U.S. Securities and Exchange Commission (filing by TERADYNE, INC) · 19 February 2026 [A] sec.gov/Archives/edgar/data/97210/000119312526059002/ter-20251231.htm

11 Teradyne Reports Fourth Quarter and Full Year 2025 Results. Teradyne · 2 February 2026 [A] investors.teradyne.com/news-events/press-releases/detail/433/teradyne-reports-fourth-quart…

12 Revenue from Contract with Customer, Excluding Assessed Tax (us-gaap:RevenueFromContractWithCustomerExcludingAssessedTax) reported by FormFactor, Inc., XBRL company concept data. U.S. Securities and Exchange Commission (XBRL data for FormFactor, Inc.) [A] data.sec.gov/api/xbrl/companyconcept/CIK0001039399/us-gaap/RevenueFromContractWithCustomer…

13 Technoprobe, press release: the Board of Directors approves the draft statutory and consolidated annual report as at 31 December 2025. Technoprobe · 18 March 2026 [A] technoprobe.com/wp-content/uploads/2026/03/PR-FY-2025_.pdf

- **Probers**, the machines that move each wafer under the probe card and keep it at a set temperature. Tokyo Seimitsu, which sells under the Accretech name, says it has the largest global share in probers [14]. It also makes dicing machines and grinders [14]. Tokyo Electron makes probers too [15].
- **Handlers and burn-in**, the machines that feed parts into a tester and heat them. Cohu took $453.0 million of revenue in fiscal 2025 [16]. Aehr's revenue was $50.0 million in fiscal 2026, but it booked a record $60.7 million of orders in the fourth quarter alone [5]. It expects revenue of $130 million to $150 million in fiscal 2027 [5].
- **Dicing and grinding.** Dicing saws cut wafers into dies, and grinders thin them. Georgetown University's CSET, analyzing TechInsights data, found that in 2025 Disco held 87 percent and Accretech 8.1 percent of the $1.9 billion market for these machines [17]. In December 2025 Disco announced a saw that cuts packages up to 400 by 400 mm [18].
- **Bonding.** Kulicke and Soffa took $654.1 million of revenue in fiscal 2025 [19]. Besi's second-quarter 2026 revenue rose 68.7 percent to EUR 249.9 million on hybrid bonding and data center demand [20].

Thermocompression bonders use heat and force to press memory dies into a stack and to attach chiplets to substrates. ASMPT says it remained the leader in that market in 2025 and aims for a 35 to 40 percent share [21]. In its 2025 results it

14 Tokyo Seimitsu Integrated Report 2025. Accretech · 3 December 2025 [A] accretech.com/en/ir/index/main/00/teaserItems1/00/linkList/02/link/IntegratedReport2025_E....

15 Precio™ Series. Tokyo Electron [A] tel.com/product/precio.html

16 Revenue from Contract with Customer, Including Assessed Tax (us-gaap:RevenueFromContractWithCustomerIncludingAssessedTax) reported by COHU, INC., XBRL company concept data. U.S. Securities and Exchange Commission (XBRL data for COHU, INC.) [A] data.sec.gov/api/xbrl/companyconcept/CIK0000021535/us-gaap/RevenueFromContractWithCustomer...

5 Aehr Test Systems Reports Fiscal 2026 Fourth Quarter and Full Year Financial Results with Record Quarterly Bookings and $100 Million Effective Backlog. Aehr Test Systems · 14 July 2026 [A] aehr.com/2026/07/aehr-test-systems-reports-fiscal-2026-fourth-quarter-and-full-year-financ...

17 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

18 DISCO develops the DFD6080 fully automatic dicing saw for work sizes up to 400 mm square (最大400mm角のパッケージ切断に対応したダイシングソー「DFD6080」を開発). DISCO Corporation · 15 December 2025 [A] disco.co.jp/eg/news/corp/20251215_1.html

19 Revenues (us-gaap:Revenues) reported by KULICKE AND SOFFA INDUSTRIES, INC., XBRL company concept data. U.S. Securities and Exchange Commission (XBRL data for KULICKE AND SOFFA INDUSTRIES, INC.) [A] data.sec.gov/api/xbrl/companyconcept/CIK0000056978/us-gaap/Revenues.json

20 BE Semiconductor Industries N.V. Announces Q2-26 and H1-26 Results. BE Semiconductor Industries (Besi) · 23 July 2026 [A] besi.com/investor-relations/press-releases/2025/details-1/be-semiconductor-industries-nv-a...

21 ASMPT Announces 2025 Annual Results AI-Driven Structural Growth Underpins Group Performance. ASMPT · 4 March 2026 [A] asmpt.com/en/investor-relations/news-events/asmpt-announces-2025-annual-results

said it expects the market to reach $1.6 billion by 2028 [21]. ASMPT took a repeat order for fifteen chip-to-substrate bonders in December 2025 [22]. Hanmi Semiconductor in Korea also sells thermocompression bonders, including dual bonders for high-bandwidth memory [23]. See Memory and HBM.

The back end's geography

Taiwanese firms lead outsourced assembly and test, the OSAT business, and Chinese firms are large challengers. ASE, the biggest, billed $20.6 billion in 2025 [24]. About 59 percent of that came from packaging and test and 40 percent from its electronics manufacturing arm [25]. CSET counts four of the top ten OSATs by 2024 revenue as headquartered in China [26].

Malaysia's investment agency said in July 2026 that the country ships 13 percent of the world's packaged chips and wants 7 percent of advanced packaging shipments by 2035 [27]. A 2024 study by the Semiconductor Industry Association and Boston Consulting Group gave Malaysia 7 percent of global assembly, test and packaging capacity in 2022, the largest share in Southeast Asia [28]. The study measured capacity by factory floor space at contractors and chipmakers' own plants [28].

ASE spent $2.7 billion on machinery and $1.4 billion on buildings, facilities and automation in the first half of 2026 [29]. It says revenue from leading-edge advanced packaging is running ahead of its $3.5 billion forecast for the year [29]. Powertech,

21 ASMPT Announces 2025 Annual Results AI-Driven Structural Growth Underpins Group Performance. ASMPT · 4 March 2026 A asmpt.com/en/investor-relations/news-events/asmpt-announces-2025-annual-results

22 ASMPT Secures Additional Orders for Fifteen Chip-to-Substrate Thermo-Compression Bonding Tools Driven by AI Tailwind. ASMPT · 22 December 2025 A asmpt.com/en/investor-relations/news-events/asmpt-secures-additional-orders-for-fifteen-ch…

23 Hanmi Semiconductor, products page. Hanmi Semiconductor A hanmisemi.com/en/product/index.html

24 Revenue (ifrs-full:Revenue) reported by ASE Technology Holding Co., Ltd., XBRL company concept data. U.S. Securities and Exchange Commission (XBRL data for ASE Technology Holding Co., Ltd.) A data.sec.gov/api/xbrl/companyconcept/CIK0001122411/ifrs-full/Revenue.json

25 ASE Technology Holding Co., Ltd., Form 6-K report of foreign private issuer for the period ended 2026-02-05. U.S. Securities and Exchange Commission (filing by ASE Technology Holding Co., Ltd.) · 5 February 2026 A sec.gov/Archives/edgar/data/1122411/000095010326001661/dp241159_6k.htm

26 Inside Beijing's Chipmaking Offensive. Center for Security and Emerging Technology (CSET) · 14 July 2025 A cset.georgetown.edu/article/inside-beijings-chipmaking-offensive

27 Advanced Packaging: How Malaysia is Packaging the Future of AI. Malaysian Investment Development Authority · 1 July 2026 A mida.gov.my/advanced-packaging-how-malaysia-is-packaging-the-future-of-ai

28 Emerging Resilience in the Semiconductor Supply Chain. Boston Consulting Group and Semiconductor Industry Association · May 2024 A semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc…

29 ASE Technology Holding Co., Ltd., Form 6-K report of foreign private issuer for the period ended 2026-07-30. U.S. Securities and Exchange Commission (filing by ASE Technology Holding Co., Ltd.) · 30 July 2026 A sec.gov/Archives/edgar/data/1122411/000095010326011353/dp250875_6k.htm

another Taiwanese assembly and test firm, aims to run the first panel-level packaging line for AI chips in 2027 [30].

Amkor, an American assembly and test contractor, took $6.71 billion of revenue in 2025 [31]. Its revenue in the second quarter of 2026 alone was $1.90 billion, up 26 percent year on year [32]. It has raised the investment in its Arizona campus to $7 billion [33].

In October 2024 TSMC and Amkor signed a memorandum of understanding for TSMC to buy packaging and test services from Amkor's planned plant in Peoria, Arizona [34]. TSMC would use those services for its customers, especially those using its Phoenix fabs [34]. The two announced a long-term partnership in June 2026 [35].

Testers installed at ASE, the largest assembly and test contractor TESTERS



Source: ASE Technology Holding, second quarter 2026 earnings release (SEC Form 6-K)

The chokepoint

Unlike lithography, no test or assembly tool has a single supplier. Advantest argues that switching to another tester maker means rebuilding the whole test setup used for development, evaluation and production, so market share moves slowly [36].

30 Powertech eyes world's first panel-level packaging for AI chips in 2027. Nikkei Asia · 27 August 2026 [B] asia.nikkei.com/business/tech/semiconductors/powertech-eyes-world-s-first-panel-level-pack...

31 AMKOR TECHNOLOGY, INC., Form 8-K current report for the period ended 2026-02-09 (8-K). U.S. Securities and Exchange Commission (filing by AMKOR TECHNOLOGY, INC.) · 9 February 2026 [A] sec.gov/Archives/edgar/data/1047127/000104712726000007/amkr123125erex-991.htm

32 AMKOR TECHNOLOGY, INC., Form 8-K current report for the period ended 2026-07-27 (8-K). U.S. Securities and Exchange Commission (filing by AMKOR TECHNOLOGY, INC.) · 27 July 2026 [A] sec.gov/Archives/edgar/data/1047127/000104712726000043/amkr6302026erex-991.htm

33 Amkor Technology Breaks Ground on New Semiconductor Advanced Packaging and Test Campus in Arizona; Expands Investment to $7 Billion. Amkor Technology [A] ir.amkor.com/news-releases/news-release-details/amkor-technology-breaks-ground-new-semicon...

34 Amkor and TSMC to Expand Partnership and Collaborate on Advanced Packaging in Arizona. TSMC · 4 October 2024 [A] pr.tsmc.com/english/news/3174

35 TSMC and Amkor Technology Announce Long Term Partnership to Accelerate Advanced Packaging in the United States. Amkor Technology [A] ir.amkor.com/news-releases/news-release-details/tsmc-and-amkor-technology-announce-long-te...

36 Investors Guide April 25, 2025. Advantest · 13 May 2025 [A] advantest.com/document/en/investors/ir-library/investors-guide/Investors_Guide_2504E.pdf

CSET finds that between 2019 and 2024 Chinese toolmakers gained 5 percentage points of share in test tools and 4 points in fabrication tools built for advanced packaging [26]. But they lost share in bonding and other conventional assembly and packaging tools [26]. They still lag significantly behind foreign suppliers [26].

ASE's packaging and test business took about $12.1 billion from outside customers in 2025 [37]. With Amkor's $6.71 billion, the two billed about $18.8 billion between them, compared with global semiconductor sales of $791.7 billion [38]. That is less than 3 percent of chip sales, for a step every chip has to pass through.

Key evaluation criteria

- **Test coverage** is the share of hidden failures caught before they reach a package that cannot be taken apart.
- **Test time per part** sets throughput and cost, and large dies, memory stacks and chiplets all push it up.
- **Probe card flatness and life** set the cost of wafer sort. Every needle must stay in the same plane across a full die and hold that over a long life.
- **Thermal control** gets harder as accelerators draw more power. A kilowatt-class accelerator has to be kept at a set temperature in the test socket while it runs at full speed.
- **Capacity and geography** set how many chips can be tested and packaged, and where. Taiwan, China and Malaysia hold most of the floor space, and new capacity takes years.

THREE THINGS TO REMEMBER

— Chips are tested early and often because one bad chip can ruin a package of expensive ones.

— Advantest and Teradyne account for most spending on chip testers.

— Four of the top ten outsourced assembly and test firms by 2024 revenue are based in China.

26 Inside Beijing's Chipmaking Offensive. Center for Security and Emerging Technology (CSET) · 14 July 2025 [A] cset.georgetown.edu/article/inside-beijings-chipmaking-offensive

37 ASE Technology Holding Co., Ltd., Form 20-F annual report for the period ended 2025-12-31. U.S. Securities and Exchange Commission (filing by ASE Technology Holding Co., Ltd.) · 1 April 2026 [A] sec.gov/Archives/edgar/data/1122411/000119312526135585/d50802d20f.htm

38 Global Annual Semiconductor Sales Increase 25.6% to $791.7 Billion in 2025. Semiconductor Industry Association · 6 February 2026 [A] semiconductors.org/global-annual-semiconductor-sales-increase-25-6-to-791-7-billion-in-2025

CHAPTER
14
STAGE 14 OF 15

Systems and Networking

AI chips are wired together 72 to a rack and run as one computer. Many firms can assemble the racks, but few can make the lasers for the optical links between them.

IN PLAIN TERMS

One AI chip is not enough to train a model. So seventy-two of them are wired together in a rack, a cabinet the size of a wardrobe, and run as one computer. Each chip sits on its own board and draws more than a thousand amps, several times the current a whole house is wired to carry. The chips have to swap results with each other all the time. A copper wire can carry those signals only a few meters before they fade, so the longer links send them as light down glass fiber. Many companies can build the rack, but only a few can make the lasers that send the light.

IN SHORT

In March 2026 Nvidia agreed to invest $2 billion in the laser maker Lumentum, alongside a multibillion-dollar purchase commitment and rights to future capacity for advanced laser components [21]. Indium phosphide is used in high-speed lasers for optical communications, and China produced 70 percent of the world's indium in 2025 [23].

On 4 February 2025 China added indium phosphide substrates, the wafers the lasers are grown on, to its export control list [29]. AXT, a US-listed maker of those wafers, does all its manufacturing in China [29]. In its filing for the quarter to June 2026, AXT said it could not estimate when it would resume shipping the wafers to the United States [29].

21 NVIDIA Announces Strategic Partnership With Lumentum to Develop State-of-the-Art Optics Technology. Nvidia · 2 March 2026 [A] nvidianews.nvidia.com/news/nvidia-announces-strategic-partnership-with-lumentum-to-develop...

23 Mineral Commodity Summaries 2026: Indium. U.S. Geological Survey · 5 February 2026 [A] pubs.usgs.gov/periodicals/mcs2026/mcs2026-indium.pdf

29 AXT INC, Form 10-Q quarterly report for the period ended 2026-06-30. U.S. Securities and Exchange Commission (filing by AXT INC) · 13 August 2026 [A] sec.gov/Archives/edgar/data/1051627/000143774926027677/axti20260630_10q.htm

CHOKEPOINT CARD

HIGH CONCENTRATION

MODERATELY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
High	Moderate Rated on the two most concentrated parts: Nvidia controls NVLink and its NVSwitch chips, which join the accelerators in a rack, and few firms make the lasers for the optical links.	**Nvidia agreed in March 2026 to invest $2B each in Coherent and Lumentum, alongside multibillion-dollar purchase commitments for their laser products**

WHO LEADS

- US **Nvidia** NVLink and NVSwitch for links inside a rack; InfiniBand and Spectrum-X Ethernet between racks
- US **Broadcom** Switch chips sold to any buyer (Tomahawk, Jericho) and custom AI accelerators
- TW **Foxconn** World's largest AI server provider, by its own account; builds in Taiwan, Mexico, Texas
- CN **InnoLight** High-speed optical modules; sites in Suzhou, Taiwan and Thailand
- US **Coherent** Indium phosphide lasers and data center transceivers

WHERE IT IS MADE

- US **United States** Switch and accelerator silicon, system design, indium phosphide lasers (Coherent's Sherman, Texas fab)
- TW **Taiwan** Contract design and rack assembly; high-layer-count circuit boards
- JP **Japan** Laser maker Lumentum's wafer fab in Sagamihara; low-loss board laminate
- CN **China** Optical module assembly and test; most of the world's indium; export permits on indium phosphide substrates
- TH **Thailand** Chinese module makers' offshore transceiver plants
- MX **Mexico** Rack integration for the North American market

WHY SUBSTITUTION IS POSSIBLE

Rack assembly can be moved, though the move would likely take several quarters, an unsourced estimate. The two concentrated parts are inside the rack, and each has an open alternative or new investment. UALink is a published open specification for the same job as NVLink. It runs 200 Gb/s per lane and addresses up to 1,024 endpoints. AMD's Helios rack, launched in July 2026, runs it over Ethernet. Since May 2025 Nvidia's NVLink Fusion program has let other firms' custom chips use NVLink, including CPUs from Fujitsu and Qualcomm that connect to Nvidia GPUs. That brings more chips onto Nvidia's link, so the dependence on Nvidia stays. For the lasers, Nvidia agreed in March 2026 to invest $2 billion each in Coherent and Lumentum, alongside multibillion-dollar purchase commitments and rights to future capacity.

WHERE CHINA STANDS

China is strong in two parts of the optical link, the modules and the indium their lasers need. Chinese firms such as InnoLight assemble high-speed optical modules in volume, at sites including Suzhou and Thailand. In 2025 China produced about 70 percent of the world's indium, the raw material for the lasers inside those modules. In February 2025 it put indium under new export restrictions and began requiring export permits for indium phosphide substrates, the wafers the lasers are grown on.

WHERE THE US STANDS

American firms make the switch chips, the links that join accelerators inside a rack, and the lasers. Some of their wafer fabs are abroad, including Lumentum's in Japan. Nvidia said in August 2026 that Wistron's new plant in Fort Worth, Texas, was producing its GB300 superchips and that Foxconn was building a factory in Houston for its AI systems. In July 2026 the Federal Communications Commission barred new authorizations for devices containing logic-bearing components made by firms on its Covered List.

Before a chip package can compute, it has to be built into a machine. Building that machine is a separate industry from making the package.

How it works

An accelerator (the AI chip) ships as a module, the chip on a small board with its power supply and memory stacks. Nvidia calls its version SXM, and the open-standards equivalent is the Open Compute Project's accelerator module. Eight modules bolt onto a baseboard that carries their power and the switch chips, which pass data between any two modules at full speed.

Rack-scale systems drop the baseboard and put the chips on trays stacked in the rack. An Nvidia GB300 NVL72 holds 18 compute trays plus nine switch trays, 72 GPUs in all, and draws up to 142 kW [1]. The newer Vera Rubin NVL72 uses the same rack footprint [2].

Inside the rack the trays reach each other over copper. Between neighboring racks, the signal can travel up to 7 meters over active electrical cables, a type of copper cable sold by firms such as Credo [3]. Longer links run on glass fiber, with an optical transceiver at each end. A laser in the sending transceiver turns the electrical signal into pulses of light, and a detector in the receiving one turns the pulses back into current.

Every AI cluster runs two networks that are not interchangeable. Scale-up connects anywhere from a few dozen accelerators (72 in an Nvidia NVL72 rack) to thousands (9,216 in a Google Ironwood superpod) closely enough that they can read each other's memory. It runs at terabytes per second. Scale-out joins those groups to one another at hundreds of gigabits per second. A byte is eight bits, so scale-up is tens of times faster than scale-out.

1 System Hardware & Components. Nvidia [A] docs.nvidia.com/enterprise-reference-architectures/nvl72-ai-factory/latest/components.html

2 NVIDIA, Partners Drive Next-Gen Efficient Gigawatt AI Factories in Buildup for Vera Rubin. Nvidia · 13 October 2025 [A] blogs.nvidia.com/blog/gigawatt-ai-factories-ocp-vera-rubin

3 Credo Technology Group Holding Ltd, Form 10-K annual report for the period ended 2026-05-02 (10-K). U.S. Securities and Exchange Commission (filing by Credo Technology Group Holding Ltd) · 15 June 2026 [A] sec.gov/Archives/edgar/data/1807794/000162828026043303/crdo-20260502.htm

A model too big for one accelerator is split across many. The parts that exchange data constantly stay inside the scale-up group, while whole blocks of layers can sit on different groups across the slower scale-out network [4]. So scale-up bandwidth sets how efficiently a large model trains.

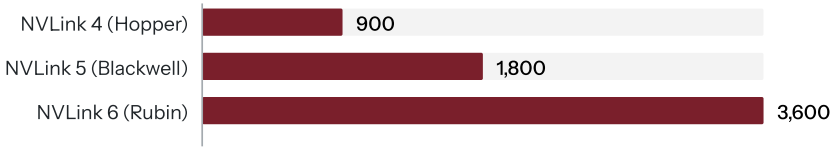
Variants and trade-offs

Scale-up: NVLink, UALink and Google's torus

The designs differ in how many accelerators can share one pool of memory.

- **NVLink and NVSwitch.** Nvidia's per-GPU bandwidth went from 900 GB/s in NVLink 4 to 1,800 in NVLink 5, and the rack's own network moves 130 TB/s on Blackwell [5]. Nvidia's Rubin launch release of January 2026 gives NVLink 6 as 3,600 GB/s per GPU and 260 TB/s for a Vera Rubin NVL72 rack [6]. As of October 2026, Nvidia's NVLink product page still lists preliminary figures of 3,000 GB/s and 216 TB/s [5].
- **UALink.** The open specification runs 200 Gb/s per lane, four lanes to a station for 800 Gb/s, and can address up to 1,024 endpoints in one group [7]. AMD's Helios rack, launched in July 2026, joins 72 GPUs and runs UALink over Ethernet for its scale-up network [8].
- **Google's 3D torus.** Google wires 64 of its Ironwood TPU accelerators into each rack, called a cube [9]. Inside the cube each chip links directly to six others, in a three-dimensional grid that Google calls a 3D torus [9]. Optical circuit switches join cubes into superpods of 9,216 chips that can bypass a failed cube or link [9].

NVLink bandwidth per GPU by generation GB/S



Source: NVIDIA NVLink product page (NVLink 4 and 5); NVIDIA Rubin launch release, January 2026 (NVLink 6)

4 Parallelisms Guide — Megatron Bridge. Nvidia [A] docs.nvidia.com/nemo-framework/user-guide/latest/nemotoolkit/features/parallelisms.html

5 NVLink & NVLink Switch for Advanced Multi-GPU Communication. Nvidia · 20 April 2026 [A] nvidia.com/en-us/data-center/nvlink

6 NVIDIA Kicks Off the Next Generation of AI With Rubin. Nvidia [A] nvidianews.nvidia.com/news/rubin-platform-ai-supercomputer

7 UALink™ 200G 1.0 Specification Overview. UALink Consortium [A] ualinkconsortium.org/blog/ualink-200g-1-0-specification-overview-802

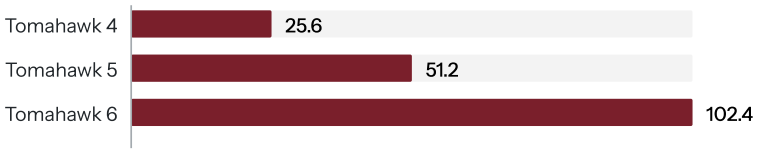
8 AAI 2026: AMD Launches AMD Helios Rackscale Solution for Frontier AI. AMD Newsroom · 23 July 2026 [A] newsroom.amd.com/news/aai-2026-helios-update

9 Inside the Ironwood TPU codesigned AI stack. Google Cloud · 6 November 2025 [A] cloud.google.com/blog/products/compute/inside-the-ironwood-tpu-codesigned-ai-stack

Scale-out: InfiniBand versus Ethernet

InfiniBand, the specialized network built for supercomputers, was the main scale-out network when AI clusters were small and uniform. Ethernet switch capacity grew quickly. Broadcom, which sells its switch chips to any buyer, doubled their bandwidth in each generation, from 25.6 Tbps on one chip in Tomahawk 4 [10] to 51.2 in Tomahawk 5 [11] and 102.4 in Tomahawk 6 [12]. In June 2025 the Ultra Ethernet Consortium published its 1.0 specification, a set of Ethernet-based communication protocols for AI and supercomputing that covers transport, congestion control and direct memory access [13].

Broadcom Tomahawk switching capacity on one chip TBPS



Source: Broadcom product releases for Tomahawk 4, 5 and 6

Nvidia sells both networks. Its data center revenue reached $89.0 billion in the quarter to July 2026, up 117 percent from a year earlier [14]. In the quarter before, it credited its data center growth to Blackwell 300 products and to demand for its "InfiniBand, Spectrum-X Ethernet, and NVLink" solutions [15].

Optics: pluggable, then co-packaged

Almost every link longer than 7 meters is optical. A cluster needs several transceivers, the modules that turn electricity into light and back, for each accelerator. Their power cost is one reason Broadcom and Nvidia are moving the optics out of pluggable modules and into the switch chip's own package, a design called co-packaged optics. Broadcom's Tomahawk 6 Davisson cuts optical interconnect power by 70 percent, more than 3.5 times lower than traditional pluggable modules [16]. Nvidia's Spectrum-X Photonics switches, due in 2026, reach 512 ports of 800 gigabits and use four times fewer lasers than traditional designs [17].

10 Broadcom Ships Tomahawk 4, Industry's Highest Bandwidth Ethernet Switch Chip at 25.6 Terabits per Second. Broadcom [A] investors.broadcom.com/news-releases/news-release-details/broadcom-ships-tomahawk-4-indust...

11 Broadcom Ships Tomahawk 5, Industry's Highest Bandwidth Switch Chip to Accelerate AI/ML Workloads. Broadcom [A] investors.broadcom.com/news-releases/news-release-details/broadcom-ships-tomahawk-5-indust...

12 Broadcom Ships Tomahawk 6: World's First 102.4 Tbps Switch. Broadcom [A] investors.broadcom.com/news-releases/news-release-details/broadcom-ships-tomahawk-6-worlds...

13 Ultra Ethernet Consortium (UEC) Launches Specification 1.0 Transforming Ethernet for AI and HPC at Scale. Ultra Ethernet Consortium · 11 June 2025 [A] ultraethernet.org/ultra-ethernet-consortium-uec-launches-specification-1-0-transforming-et...

14 NVIDIA CORP, Form 10-Q quarterly report for the period ended 2026-07-26 (10-Q). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 26 August 2026 [A] sec.gov/Archives/edgar/data/1045810/000104581026000075/nvda-20260726.htm

15 NVIDIA CORP, Form 10-Q quarterly report for the period ended 2026-04-26 (10-Q). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 20 May 2026 [A] sec.gov/Archives/edgar/data/1045810/000104581026000052/nvda-20260426.htm

16 Broadcom Announces Tomahawk 6 – Davisson, the Industry's First 102.4-Tbps Ethernet Switch with Co-Packaged Optics. Broadcom [A] investors.broadcom.com/news-releases/news-release-details/broadcom-announces-tomahawkr-6-d...

17 NVIDIA Announces Spectrum-X Photonics, Co-Packaged Optics Networking Switches to Scale AI Factories to Millions of GPUs. Nvidia · 18 March 2025 [A] nvidianews.nvidia.com/news/nvidia-spectrum-x-co-packaged-optics-networking-switches-ai-fac...

Who makes it

Optical modules are assembled where labor and land are cheap, and the Chinese makers went offshore early. InnoLight has run a plant in Thailand since 2019, alongside its sites in Suzhou and Taiwan [18]. The Thai subsidiary of Eoptolink, another Chinese module maker, lists sites in Chonburi and Rayong [19].

The Federal Communications Commission publishes a Covered List of equipment and services judged to pose an unacceptable risk to US national security or to the security and safety of US persons [20]. Its July 2026 order bars new authorizations for devices that contain logic-bearing hardware components produced by firms on that list [20]. Commenters in the record named optical transceivers as one such component [20]. The same order notes that the key players in the optical transceiver market, Coherent and InnoLight among them, do not appear on the Covered List [20].

The lasers inside those modules come from far fewer firms. In March 2026 Nvidia agreed to invest $2 billion in one of them, Lumentum, alongside a multibillion-dollar purchase commitment and future capacity access rights for advanced laser components [21]. It made the same $2 billion investment in Coherent, another laser maker, with a multibillion-dollar purchase commitment for advanced laser and optical networking products [22].

Indium phosphide is used in the high-speed laser diodes and photodetectors for optical communications [23]. China accounted for 70 percent of world indium production in 2025 [23]. In February 2025 China's Ministry of Commerce placed indium under new export restrictions [23].

18 InnoLight Technology, company overview page. InnoLight Technology [A] innolight.com/about

19 Eoptolink Technology, contacts page. Eoptolink Technology [A] eoptolink.com/about-us/contacts

20 Protecting Against National Security Threats to the Communications Supply Chain through the Equipment Authorization Program: Third Report and Order and Third Further Notice of Proposed Rulemaking (FCC-26-50). U.S. Federal Communications Commission · 23 July 2026 [A] docs.fcc.gov/public/attachments/FCC-26-50A1.pdf

21 NVIDIA Announces Strategic Partnership With Lumentum to Develop State-of-the-Art Optics Technology. Nvidia · 2 March 2026 [A] nvidianews.nvidia.com/news/nvidia-announces-strategic-partnership-with-lumentum-to-develop…

22 NVIDIA and Coherent Announce Strategic Partnership to Develop Optics Technology to Scale Next-Generation Data Center Architecture. Nvidia · 2 March 2026 [A] nvidianews.nvidia.com/news/nvidia-and-coherent-announce-strategic-partnership-to-develop-o…

23 Mineral Commodity Summaries 2026: Indium. U.S. Geological Survey · 5 February 2026 [A] pubs.usgs.gov/periodicals/mcs2026/mcs2026-indium.pdf

Two US firms, Broadcom and Marvell, design both custom accelerators and switch chips. Broadcom sells the Ethernet switch chips that other firms build their switches around, along with custom AI accelerators. Its AI semiconductor revenue reached $16.7 billion in the third quarter of fiscal 2026, up 221 percent from a year earlier [24]. Marvell's five-year agreement with Amazon Web Services, announced in December 2024, covers custom AI chips alongside the optical and networking parts that go with them [25].

Taiwanese firms do much of the physical assembly. Foxconn booked NT$8.1 trillion of revenue in 2025, up 18 percent, with cloud and networking among the drivers [26].

Some assembly is moving to the United States. In April 2025 Nvidia said it was building supercomputer manufacturing plants in Texas, with Foxconn in Houston and Wistron in Dallas [27]. It expected mass production to ramp up in the following 12 to 15 months [27]. In August 2026 Nvidia said Wistron's new plant in Fort Worth, Texas, near Dallas, was producing its GB300 superchips and Foxconn was building its factory in Houston [28].

The chokepoint

Several contract manufacturers on three continents assemble racks. The power shelves, copper busbars, cables and connectors have many suppliers. In its filing for the quarter to April 2026, Nvidia said the complexity of rack-scale systems had caused production delays [15]. Moving Foxconn's share of rack assembly to other assemblers such as Quanta and Wistron would likely take several quarters, though no source puts a figure on it.

Four components that go into the rack take longer to replace.

24 Broadcom Inc., Form 8-K current report for the period ended 2026-09-02 (8-K). U.S. Securities and Exchange Commission (filing by Broadcom Inc.) · 2 September 2026 A sec.gov/Archives/edgar/data/1730168/000173016826000076/avgo-08022026x8kxex99.htm

25 Marvell Technology, Inc., Form 8-K current report for the period ended 2024-12-02 (8-K). U.S. Securities and Exchange Commission (filing by Marvell Technology, Inc.) · 2 December 2024 A sec.gov/Archives/edgar/data/1835632/000183563224000193/final2024_12x02xmarvell-.htm

26 Hon Hai Technology Group (Foxconn) Announces FY2025 & 4Q25 Financial Results. Hon Hai Precision Industry (Foxconn) · 16 March 2026 A honhai.com/en-us/press-center/press-releases/latest-news/1978

27 NVIDIA to Manufacture American-Made AI Supercomputers in US for First Time. Nvidia · 14 April 2025 A blogs.nvidia.com/blog/nvidia-manufacture-american-made-ai-supercomputers-us

28 NVIDIA and Partners Build in America, for America. Nvidia · 5 August 2026 A blogs.nvidia.com/blog/nvidia-and-partners-build-in-america-for-america

15 NVIDIA CORP, Form 10-Q quarterly report for the period ended 2026-04-26 (10-Q). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 20 May 2026 A sec.gov/Archives/edgar/data/1045810/000104581026000052/nvda-20260426.htm

- **Indium phosphide laser capacity.** In March 2026 Nvidia agreed to invest $2 billion in Coherent, alongside a multibillion-dollar purchase commitment and future access and capacity rights for Coherent's advanced laser and optical networking products [22]. In 2025 China produced 70 percent of the world's indium, which those lasers need [23]. It put indium under new export restrictions in February 2025 [23].
- **Indium phosphide substrates.** China added indium phosphide substrates, the wafers the lasers are grown on, to its export control list on 4 February 2025 [29]. AXT, a US-listed maker of these wafers, does all its manufacturing in China [29]. In its filing for the quarter to June 2026, AXT called the export permits the most significant challenge it faces [29]. It said it could not estimate when it would resume shipping the wafers to the United States [29].
- **NVSwitch.** NVLink and its switch chip, NVSwitch, are Nvidia's own designs. UALink is an open standard designed for accelerator communications within one scale-up group [7]. Since May 2025 Nvidia's NVLink Fusion program has let other firms' custom chips use NVLink, including CPUs from Fujitsu and Qualcomm that connect to Nvidia GPUs [30].
- **High-layer-count circuit boards.** The boards, and the low-loss laminate they are built from, are concentrated in Taiwan and Japan. See Substrates and PCBs.

Optical module assembly is concentrated, but firms choose where to do it based on cost and policy. It has moved before, as the Chinese makers' Thai plants show [18].

Key evaluation criteria

- **Scale-up domain size** is how many accelerators share one pool of memory. The UALink spec allows up to 1,024 endpoints [7]. An Ironwood superpod has 9,216 chips [9].
- **Bandwidth per accelerator** limits how efficiently one model can be split across chips. NVLink 6 carries 3,600 GB/s per GPU [6].
- **Picojoules per bit** is the energy it takes to move one bit of data. Broadcom says traditional pluggable optics consume more power as AI traffic grows, and its

22 NVIDIA and Coherent Announce Strategic Partnership to Develop Optics Technology to Scale Next-Generation Data Center Architecture. Nvidia · 2 March 2026 [A] nvidianews.nvidia.com/news/nvidia-and-coherent-announce-strategic-partnership-to-develop-o…

23 Mineral Commodity Summaries 2026: Indium. U.S. Geological Survey · 5 February 2026 [A] pubs.usgs.gov/periodicals/mcs2026/mcs2026-indium.pdf

29 AXT INC, Form 10-Q quarterly report for the period ended 2026-06-30. U.S. Securities and Exchange Commission (filing by AXT INC) · 13 August 2026 [A] sec.gov/Archives/edgar/data/1051627/000143774926027677/axti20260630_10q.htm

7 UALink™ 200G 1.0 Specification Overview. UALink Consortium [A] ualinkconsortium.org/blog/ualink-200g-1-0-specification-overview-802

30 NVIDIA Unveils NVLink Fusion for Industry to Build Semi-Custom AI Infrastructure With NVIDIA Partner Ecosystem. Nvidia [A] nvidianews.nvidia.com/news/nvidia-nvlink-fusion-semi-custom-ai-infrastructure-partner-ecos…

18 InnoLight Technology, company overview page. InnoLight Technology [A] innolight.com/about

9 Inside the Ironwood TPU codesigned AI stack. Google Cloud · 6 November 2025 [A] cloud.google.com/blog/products/compute/inside-the-ironwood-tpu-codesigned-ai-stack

6 NVIDIA Kicks Off the Next Generation of AI With Rubin. Nvidia [A] nvidianews.nvidia.com/news/rubin-platform-ai-supercomputer

co-packaged switch cuts optical interconnect power by 70 percent [16].

- **Serviceability** is how easily a failed part can be swapped. It changes with co-packaged optics, which move the optics off pluggable modules and onto the switch package. Broadcom's Davisson switch keeps its lasers in field-replaceable modules [16].
- **Supply concentration** is in the components. Many firms assemble the modules, but far fewer can make the lasers inside them.

FOUR THINGS TO REMEMBER

- An AI rack wires 72 chips together to run as one computer.
- Almost every link longer than 7 meters carries its signal as light, because copper fades over longer runs.
- Nvidia's NVLink switch chips and the lasers are the hardest parts to replace. Rack assembly can move to another firm, though the move would take an estimated several quarters.
- China produced about 70 percent of the world's indium in 2025 and has required export permits for indium phosphide substrates since February 2025.

16 Broadcom Announces Tomahawk 6 – Davisson, the Industry's First 102.4-Tbps Ethernet Switch with Co-Packaged Optics. Broadcom [A] investors.broadcom.com/news-releases/news-release-details/broadcom-announces-tomahawkr-6-d...

CHAPTER
15
STAGE 15 OF 15

Data Centers and Power

In 2024 a large transformer commonly took three years to arrive. In June 2026, big power users, mostly data centers, sought about 474 GW of grid connections in Texas.

IN PLAIN TERMS

A data center feeds electricity to the machines inside it and carries away the heat they give off. One rack of AI chips, a cabinet taller than a door, draws up to 142 kilowatts [2]. Every watt of it comes back out as heat, and liquid piped through the rack carries the heat away [2]. A large training site already holds more than a hundred thousand chips [3]. The newest sites are planned in gigawatts [4]. One gigawatt is the whole output of a large power station, and OpenAI's planned sites add up to more than nine of them by 2029 [4]. Getting that much electricity to the site is harder than putting up the building. Transformers and heavy switches connect a site to the grid, and turbines make the power. All of them come from a few suppliers and take years to arrive.

IN SHORT

Measured by capacity, US factories supplied 29.2 percent of the large power transformers bought in the US in 2023, and imports the rest [21]. GE Vernova's gas equipment backlog and reserved factory slots stood at 116 GW in mid-2026, against planned output of 20 GW a year from the third quarter [13].

ERCOT, the grid operator for most of Texas, counted 9,042 MW of big power users approved to energize by March 13, 2026 [18]. Judging by peak monthly use, it believed 3,883 MW of that was operating [18]. On August 3, 2026, Governor Greg Abbott ordered an audit of every data center seeking a connection before any data center project moves forward [19].

2 System Hardware & Components. Nvidia [A] docs.nvidia.com/enterprise-reference-architectures/nvl72-ai-factory/latest/components.html

3 Colossus 1, AI data centers directory. Epoch AI [A] epoch.ai/data/ai-data-centers/directory/colossus-1

4 OpenAI Stargate: where the US sites stand. Epoch AI · 17 April 2026 [A] epoch.ai/publications/openai-stargate-where-the-us-sites-stand

21 Large Power Transformers from South Korea, Investigation No. 731-TA-1189 (Second Review), Publication 5531. U.S. International Trade Commission · August 2024 [A] usitc.gov/publications/701_731/pub5531.pdf

13 GE Vernova Inc., Form 8-K current report for the period ended 2026-07-22 (8-K). U.S. Securities and Exchange Commission (filing by GE Vernova Inc.) · 22 July 2026 [A] sec.gov/Archives/edgar/data/1996810/000199681026000147/gevpressrelease2q26.htm

18 Electric Reliability Council of Texas, Large Load Interconnection Status Update, 13 March 2026. Electric Reliability Council of Texas · 13 March 2026 [A] ercot.com/files/docs/2026/03/12/March-TAC-Report.pdf

19 Governor Abbott Directs Comprehensive Data Center Audit. Office of the Texas Governor [A] gov.texas.gov/news/post/governor-abbott-directs-comprehensive-data-center-audit

CHOKEPOINT CARD

HIGH CONCENTRATION

MODERATELY HARD TO SUBSTITUTE

CONCENTRATION	SUBSTITUTABILITY	PRICE OR MARKET SIZE
High	Moderate Nvidia names nine data center power system providers for its gigawatt designs. But a large transformer commonly took three years to arrive in 2024, and GE Vernova's gas equipment backlog and reserved factory slots stood at 116 GW in mid-2026.	**Large power transformer lead times of 36 months were commonly quoted in 2024, against under a year before the pandemic** GE Vernova's gas equipment backlog and reserved factory slots grew from 100 GW to 116 GW in the second quarter of 2026

WHO LEADS

- US **GE Vernova** 116 GW of gas equipment backlog and reserved factory slots at mid-2026. Also makes transformers and switchgear, including through Prolec GE
- CH **Hitachi Energy** Large power transformers and high-voltage switchgear
- DE **Siemens Energy** Grid technologies, transformers, turbines
- JP **Mitsubishi Power** Power generation equipment, including gas turbines. GE Vernova names it a key competitor, alongside Siemens Energy
- US **Vertiv** Data center power and cooling systems. A partner in Nvidia's gigawatt AI factory designs
- FR **Schneider Electric** Electrical distribution and cooling. A partner in Nvidia's 800-volt DC designs

WHERE IT IS MADE

- US **United States** Most new AI capacity. The ERCOT grid in Texas and the PJM grid in the Mid-Atlantic and Midwest carry the growth in demand
- CH **Switzerland** Hitachi Energy and ABB transformer and switchgear engineering
- DE **Germany** Siemens Energy grid technologies and turbines
- TW **Taiwan** Delta Electronics power shelves, busbars and thermal modules
- JP **Japan** Mitsubishi Heavy Industries power equipment, including gas turbines. Hitachi group transformers

WHY SUBSTITUTION IS POSSIBLE

Nvidia names nine data center power system providers for its gigawatt designs, so a buyer can choose among suppliers. But money cannot shorten the wait. In 2024, large power transformer lead times of 36 months were commonly quoted, and the longest reached 60 months. Before the pandemic, an order took under a year. An International Energy Agency survey published in 2025 found waits of up to four years worldwide. GE Vernova's gas equipment backlog and reserved factory slots reached 116 GW in mid-2026. It expects its gas turbine output to reach 20 GW a year in the third quarter of 2026.

WHERE CHINA STANDS	WHERE THE US STANDS
China has spare generating capacity and a large domestic industry that makes transformers and switchgear, and it can connect new sites quickly. The supply of accelerators limits its AI buildout.	The United States has the demand and the money but cannot shorten the wait for equipment. Transformers and turbines both take years to arrive. Measured by capacity, US factories supplied 29.2 percent of the large power transformers bought in the US in 2023, and imports supplied the rest. No restarted nuclear reactor is serving data centers yet.

A wafer passes through a leading-edge fab in months, but in 2024 a large power transformer commonly took three years to arrive [1]. Before the pandemic, the same order took under a year [1].

How it works

The rack is the basic unit of AI compute. Nvidia's GB300 NVL72 rack draws up to 142 kW through eight 33 kW power shelves, the rack's own power supplies [2]. Colossus 1, an AI data center in Memphis, Tennessee, held an estimated 230,000 Nvidia GPUs by July 2025 [3].

Between the grid and the rack, a substation steps the high voltage down with transformers. Switchgear, the building's heavy circuit breakers, can cut any part of the site off in a fault. Battery-backed supplies keep the racks running for the seconds it takes backup generators to start. Distribution boards split the feed among the rows of racks.

At the back of each rack, a busbar carries the current down to the server trays. It is a solid copper bar, because a cable would melt at that current. Every piece of equipment from the substation to the busbar is heavy and has its own lead time, the wait between order and delivery.

Heat leaves in water. Cold plates, metal blocks with coolant running through them, sit on the GPU packages and pass the heat to the building's water loop, which releases it outdoors [5]. Nvidia's Rubin racks run coolant at up to 45 °C, warm enough that in many climates the loop can release its heat without turning on chillers [5].

Nvidia is also moving the power that feeds its racks to 800 volts of direct current. Raising the voltage cuts the current needed for the same power, and the heat in

1 U.S. Department of Energy Large Power Transformer Resilience Report to Congress, July 2024. U.S. Department of Energy · 10 July 2024 [A] energy.gov/sites/default/files/2024-10/EXEC-2022-001242%20-%20Large%20Power%20Transformer%...

2 System Hardware & Components. Nvidia [A] docs.nvidia.com/enterprise-reference-architectures/nvl72-ai-factory/latest/components.html

3 Colossus 1, AI data centers directory. Epoch AI [A] epoch.ai/data/ai-data-centers/directory/colossus-1

5 Hotter Than a Hot Tub: The 45°C Breakthrough to Cool AI's Biggest Machines. Nvidia · 22 June 2026 [A] blogs.nvidia.com/blog/liquid-cooling-ai-factories

copper comes from current. So the same conductor carries over 150 percent more power, and the 200 kg busbars that feed a single rack are no longer needed [6].

Variants and trade-offs

Cooling

- **Air.** Heat exchangers in the rack door work when density, the power each rack draws, is moderate, but Nvidia's GB300 NVL72 rack is liquid cooled [2].
- **Direct-to-chip liquid.** Nvidia's Rubin platform is 100 percent liquid cooled, and every cloud provider and data center operator building for it is moving to liquid [5].
- **Immersion.** The servers sit in a bath of fluid. It allows higher density still, but the hardware is hard to service.

Nvidia says the cooling supply chain is keeping pace with the move to liquid [5]. Motivair, Schneider Electric's cooling division, has worked alongside Nvidia's product roadmap for nearly a decade [5].

Rack power

Nvidia's 800-volt DC architecture is designed for racks from 100 kW to more than a megawatt, with megawatt racks from 2027 [7]. Nvidia is making the change because the 54-volt distribution in existing racks begins to hit physical limits once racks exceed 200 kW [7].

Nvidia names nine data center power system providers for its 800-volt gigawatt designs, among them ABB, Eaton, Hitachi Energy, Mitsubishi Electric, Schneider Electric, Siemens and Vertiv [6].

A megawatt rack changes the busbar, the floor loading, the coolant flow, the leak detection and the stored energy that keeps the rack running through a power dip. Nvidia's Vera Rubin NVL72 rack design holds 20 times more energy storage to keep power steady [6].

Where the electrons come from

- **Grid interconnection.** The slow path, with long queues to connect [4].

6 NVIDIA, Partners Drive Next-Gen Efficient Gigawatt AI Factories in Buildup for Vera Rubin. Nvidia · 13 October 2025 A blogs.nvidia.com/blog/gigawatt-ai-factories-ocp-vera-rubin

2 System Hardware & Components. Nvidia A docs.nvidia.com/enterprise-reference-architectures/nvl72-ai-factory/latest/components.html

5 Hotter Than a Hot Tub: The 45°C Breakthrough to Cool AI's Biggest Machines. Nvidia · 22 June 2026 A blogs.nvidia.com/blog/liquid-cooling-ai-factories

7 NVIDIA 800 VDC Architecture Will Power the Next Generation of AI Factories. Nvidia · 20 May 2025 A developer.nvidia.com/blog/nvidia-800-v-hvdc-architecture-will-power-the-next-generation-of...

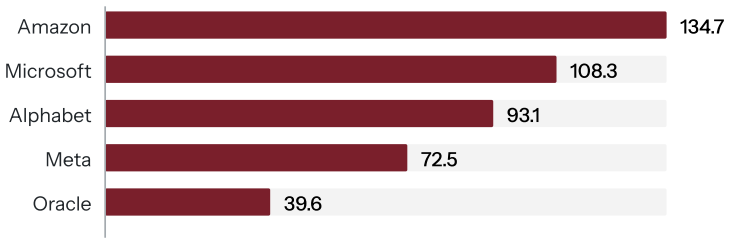
4 OpenAI Stargate: where the US sites stand. Epoch AI · 17 April 2026 A epoch.ai/publications/openai-stargate-where-the-us-sites-stand

- **Behind-the-meter gas**, generation on the site itself. The fast path, which at least three of OpenAI's seven Stargate sites will use to avoid grid queues [4].
- **Nuclear.** Existing plants already sell power to data centers [8]. Talen Energy supplies Amazon's data center campus next to the Susquehanna nuclear plant in Pennsylvania under a June 2025 contract expected to reach 1,920 MW no later than 2032 [8]. Restarts and new reactors are still to come. Constellation signed a power purchase agreement with Microsoft to restart Three Mile Island Unit 1 as the 835 MW Crane Clean Energy Center [9]. The plant shut in 2019 [10]. A $1 billion federal loan that closed in November 2025 partly finances its restart, which is subject to Nuclear Regulatory Commission licensing approvals [10]. No restarted reactor is serving data centers yet.

Who makes it

The largest data center operators pay for the buildout. Capital spending by Alphabet, Amazon, Meta, Microsoft and Oracle came to $448.2 billion in 2025 [11]. If the trend continues, it will reach $770 billion in 2026 [11].

Capital spending by five large US data center operators, 2025 $B



Source: Epoch AI, hyperscaler capex trend, February 2026, from company filings

Epoch AI, a research group, projects that OpenAI's seven US Stargate sites will exceed 9 GW by 2029 [4]. It expects the site in Abilene, Texas, to grow from an estimated 0.3 GW in April to 1.2 GW in the fourth quarter of 2026 [4].

Equipment makers limit how fast the buildout can go. GE Vernova builds large gas turbines, and its 2025 annual report names Siemens Energy and Mitsubishi Power among the key competitors of its power generation business [12]. Hitachi Energy,

4 OpenAI Stargate: where the US sites stand. Epoch AI · 17 April 2026 [A] epoch.ai/publications/openai-stargate-where-the-us-sites-stand

8 Talen Energy Corp, Form 8-K current report for the period ended 2025-06-11 (8-K). U.S. Securities and Exchange Commission (filing by Talen Energy Corp) · 11 June 2025 [A] sec.gov/Archives/edgar/data/1622536/000162828025030559/a20250611pressreleasebusin.htm

9 Constellation Ahead of Schedule for Launch of Crane Clean Energy Center. Constellation Energy · 19 February 2025 [A] constellationenergy.com/news/2025/constellation-ahead-of-schedule-for-launch-of-crane-clea…

10 Energy Department Closes Loan to Restart Nuclear Power Plant in Pennsylvania. U.S. Department of Energy · 17 November 2025 [A] energy.gov/articles/energy-department-closes-loan-restart-nuclear-power-plant-pennsylvania

11 Hyperscaler capex has quadrupled since GPT-4's release. Epoch AI · 26 February 2026 [A] epoch.ai/data-insights/hyperscaler-capex-trend

12 GE Vernova Inc., Form 10-K annual report for the period ended 2025-12-31 (10-K). U.S. Securities and Exchange Commission (filing by GE Vernova Inc.) · 29 January 2026 [A] sec.gov/Archives/edgar/data/1996810/000199681026000015/gev-20251231.htm

Siemens Energy, GE Vernova, ABB, Eaton and Schneider Electric supply transformers and switchgear.

GE Vernova bought the remaining half of Prolec GE, a transformer maker, in the first half of 2026 [13]. GE Vernova's gas equipment backlog and reserved factory slots grew from 100 GW to 116 GW in the second quarter of 2026 [13]. In the first half of the year, data center orders for its electrification business, which makes grid equipment, passed $5 billion, more than double its total for 2025 [13].

GE Vernova's planned annual gas turbine output GW/YR



Source: GE Vernova second-quarter 2026 results, July 2026

Capacity and demand

US data centers used 192 TWh in 2024, 4.7 percent of national electricity, according to a June 2026 estimate by Lawrence Berkeley National Laboratory that leaves out cryptocurrency mining [14]. The lab's reference case, its main scenario, projects 464 TWh in 2028 and 649 TWh in 2030, which would be 11.8 percent of forecast US electricity use [14]. Its scenarios for 2030 range from 521 to 843 TWh, or 9.5 percent to 15.3 percent [14].

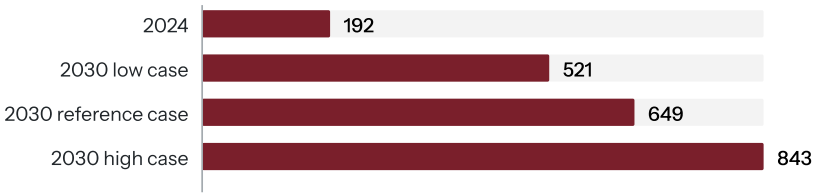
In its September 2026 Short-Term Energy Outlook, the Energy Information Administration expects record US electricity sales of 4,135 TWh in 2026 and 4,211 TWh in 2027 [15]. Sales would grow by almost 2 percent in each year, driven by data centers and manufacturing [15]. The same outlook lowers its 2027 electricity sales forecast for the region that includes Texas by 4.7 percent from its August forecast [15]. It notes that Texas has paused connecting new data centers to the grid while it audits proposed projects [15].

13 GE Vernova Inc., Form 8-K current report for the period ended 2026-07-22 (8-K). U.S. Securities and Exchange Commission (filing by GE Vernova Inc.) · 22 July 2026 [A] sec.gov/Archives/edgar/data/1996810/000199681026000147/gevpressrelease2q26.htm

14 United States Data Center Energy Usage Report: 2025 Update. Lawrence Berkeley National Laboratory · 18 June 2026 [A] energy.gov/documents/united-states-data-center-energy-usage-report-2025-update

15 Short-Term Energy Outlook, September 2026. U.S. Energy Information Administration · 8 September 2026 [A] eia.gov/outlooks/steo/archives/sep26.pdf

US data center electricity use, 2024 and 2030 projection range TWH



Source: Lawrence Berkeley National Laboratory, June 2026

Chip efficiency is still improving. Google's Ironwood TPU, an AI chip, delivers twice the performance per watt of Trillium, the TPU before it [16]. A token is the chunk of text a model reads or writes. Energy per token can fall while total energy use still rises, if the number of tokens rises faster.

The chokepoint

Far more capacity has been announced than is connected and running. ERCOT, the grid operator for most of Texas, was tracking about 474 GW of large loads (big power users) seeking a connection in June 2026, roughly 90 percent of it data centers [17]. ERCOT's large load update of March 13, 2026, counted 9,042 MW approved to energize, but judging by peak monthly use, ERCOT believed only 3,883 MW of it was operating [18]. In a preliminary review in July 2026, ERCOT found about 205 GW of the queue eligible for Batch Zero, its first study of connection requests as a group [17]. On August 3, 2026, Texas Governor Greg Abbott ordered the state's Public Utility Commission and ERCOT to audit every data center in the interconnection process before any data center project moves forward [19].

Projects in that queue also wait on equipment.

- **Large power transformers.** In 2024, lead times of 36 months were commonly quoted and the longest reached 60 months, against under a year before the pandemic [1]. By a Commerce Department estimate, 82 percent of the units put into US service in 2019 were imported [1]. In an industry survey published in 2025, the International Energy Agency found that large power transformers

16 Inside the Ironwood TPU codesigned AI stack. Google Cloud · 6 November 2025 A cloud.google.com/blog/products/compute/inside-the-ironwood-tpu-codesigned-ai-stack

17 Electric Reliability Council of Texas, ERCOT update to the Senate Committee on Business and Commerce, 29 July 2026. Electric Reliability Council of Texas · 29 July 2026 A ercot.com/files/docs/2026/07/29/ERCOT-Senate-July-29-Panel-1-Assessing-The-Grid.pdf

18 Electric Reliability Council of Texas, Large Load Interconnection Status Update, 13 March 2026. Electric Reliability Council of Texas · 13 March 2026 A ercot.com/files/docs/2026/03/12/March-TAC-Report.pdf

19 Governor Abbott Directs Comprehensive Data Center Audit. Office of the Texas Governor A gov.texas.gov/news/post/governor-abbott-directs-comprehensive-data-center-audit

1 U.S. Department of Energy Large Power Transformer Resilience Report to Congress, July 2024. U.S. Department of Energy · 10 July 2024 A energy.gov/sites/default/files/2024-10/EXEC-2022-001242%20-%20Large%20Power%20Transformer%...

took up to four years to secure worldwide [20]. It also found that their average lead times had almost doubled since 2021 [20]. Measured by capacity, US factories supplied 29.2 percent of the transformers of 60 MVA and above bought in the US in 2023, and imports supplied the rest [21].

- **Gas turbines.** GE Vernova's gas equipment backlog and reserved factory slots stood at 116 GW in mid-2026 [13]. It expects its annual gas turbine output to reach 20 GW in the third quarter of 2026, 24 GW in 2028 and 30 GW in 2030 [13].
- **Firm new generation**, the kind that runs whatever the weather. The 835 MW Crane plant shut in 2019, and its restart depends on Nuclear Regulatory Commission licensing approvals [10]. The NRC has approved a fuel license amendment for the plant, and Constellation expects to restart it in 2027 [22].

Memory chips are also in short supply. Micron's quarterly report for the period ended May 28, 2026, says AI-driven growth in demand for memory and storage is outpacing industry supply [23]. The memory shortage and the wait for power equipment both limit how much of the 2026 capital spending turns into working compute.

Key evaluation criteria

- **Time to power** counts the months from a signed lease until the power is switched on. Developers compete on this number.
- **Transformer and turbine lead time** covers the two queues that set the schedule. Large power transformer lead times of 36 months were commonly quoted in 2024 [1]. GE Vernova's gas equipment backlog and reserved factory slots stood at 116 GW in mid-2026, against planned output of 20 GW a year from the third quarter [13].
- **Rack density supported** sets which racks an existing hall can take. A GB300 NVL72 rack draws up to 142 kW and is liquid cooled [2].
- **Firmness of supply** tells whether a site's power is backed by delivered equipment, such as a transformer, or exists only on paper, as a signed interconnection agreement.

20 Executive summary – Building the Future Transmission Grid. IEA [A] iea.org/reports/building-the-future-transmission-grid/executive-summary

21 Large Power Transformers from South Korea, Investigation No. 731-TA-1189 (Second Review), Publication 5531. U.S. International Trade Commission · August 2024 [A] usitc.gov/publications/701_731/pub5531.pdf

13 GE Vernova Inc., Form 8-K current report for the period ended 2026-07-22 (8-K). U.S. Securities and Exchange Commission (filing by GE Vernova Inc.) · 22 July 2026 [A] sec.gov/Archives/edgar/data/1996810/000199681026000147/gevpressrelease2q26.htm

10 Energy Department Closes Loan to Restart Nuclear Power Plant in Pennsylvania. U.S. Department of Energy · 17 November 2025 [A] energy.gov/articles/energy-department-closes-loan-restart-nuclear-power-plant-pennsylvania

22 Constellation Reports Second Quarter 2026 Results. Constellation Energy · 6 August 2026 [A] constellationenergy.com/news/2026/08/constellation-reports-second-quarter-2026-results.html

23 MICRON TECHNOLOGY INC, Form 10-Q quarterly report for the period ended 2026-05-28 (10-Q). U.S. Securities and Exchange Commission (filing by MICRON TECHNOLOGY INC) · 25 June 2026 [A] sec.gov/Archives/edgar/data/723125/000072312526000015/mu-20260528.htm

1 U.S. Department of Energy Large Power Transformer Resilience Report to Congress, July 2024. U.S. Department of Energy · 10 July 2024 [A] energy.gov/sites/default/files/2024-10/EXEC-2022-001242%20-%20Large%20Power%20Transformer%...

2 System Hardware & Components. Nvidia [A] docs.nvidia.com/enterprise-reference-architectures/nvl72-ai-factory/latest/components.html

- **Energy per token** is the electricity a model uses for each token it reads or writes. Chips with higher performance per watt, such as Google's Ironwood TPU, lower it [16].

FOUR THINGS TO REMEMBER

- One AI rack draws up to 142 kilowatts, too much heat for air, so racks are cooled with liquid.
- The newest AI sites are planned in gigawatts, and each gigawatt is the output of a large power station.
- Power equipment is one of the slow parts of building a data center. In 2024 lead times of 36 months were commonly quoted for a large transformer.
- China has power to spare, but the supply of AI chips limits its buildout.

16 **Inside the Ironwood TPU codesigned AI stack.** Google Cloud · 6 November 2025 [A] cloud.google.com/blog/products/compute/inside-the-ironwood-tpu-codesigned-ai-stack

CHAPTER
16
THE ECONOMICS

The Economics of AI Chips

Epoch AI estimated in December 2025 that Nvidia's B200 AI chip cost about $6,400 to build, mostly for memory and packaging. Its processor silicon cost about $900.

IN PLAIN TERMS

A bill of materials is the list of parts inside a product and what each one costs to build. Epoch AI estimated in December 2025 that Nvidia's B200, an AI chip from its Blackwell line, cost about $6,400 to build [1]. The two logic dies, the pieces of silicon that do the calculating, cost about $900 of that [1]. The stacked memory beside them accounted for about half, and the packaging that wires the dies and memory together cost more than the dies [1]. Three firms, SK hynix, Samsung and Micron, dominate the world market for that memory [3]. One Taiwanese firm, TSMC, does that packaging [1].

IN SHORT

Reports in March 2024 put the B200's selling price at $30,000 to $40,000 per chip [1]. Epoch AI notes that most Blackwell revenue comes from servers and rack-scale systems, which may carry lower margins [1]. It says Nvidia's Blackwell margin may therefore be lower than the gap between cost and price suggests [1].

The estimate's memory price of $14 to $17 per gigabyte predates the 2026 rise in memory makers' sales and margins [5]. Micron's gross margin, the share of sales left after production costs, rose from 39.8 percent for the year to August 2025 to 86.8 percent in the quarter to September 2026 [11]. Nvidia's was 75.0 percent in the quarter to July 2026 [7].

1 NVIDIA's B200 costs around $6,400 to produce. Epoch AI · 10 December 2025 [A] epoch.ai/data-insights/b200-cost-breakdown

3 Understanding the Biden Administration's Updated Export Controls. Center for Strategic and International Studies · 11 December 2024 [A] csis.org/analysis/understanding-biden-administrations-updated-export-controls

5 SK hynix Announces 2Q26 Financial Results. SK hynix [A] news.skhynix.com/en/q2-2026-business-results

11 Micron Technology, Inc. Reports Record Fiscal Fourth-Quarter and Full-Year 2026 Results. Micron Technology · 30 September 2026 [A] investors.micron.com/news/press-release/2026/Micron-Technology-Inc--Reports-Record-Fiscal-...

7 NVIDIA Announces Financial Results for Second Quarter Fiscal 2027. Nvidia · 26 August 2026 [A] nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2027

In December 2025 Epoch AI estimated that building one of Nvidia's Blackwell B200 chips cost about $6,400 [1]. Reports in March 2024 put its selling price at $30,000 to $40,000 per chip [1]. Epoch AI's chip sales model, with data through October 2025, puts Nvidia's compute revenue per GPU in its GB200 rack systems at $33,000 to $42,000 [2]. That figure leaves out the networking gear in those racks [2]. Some stages behind that $6,400 cost a lot because they are hard to make, some because capacity is short, and some because one firm sets the price.

How it works

Three numbers set the cost of a die, the finished rectangle of circuitry cut out of a wafer.

- **Wafer price.** TSMC does not publish it. Nvidia makes Blackwell chips on a custom-built TSMC 4NP process [4]. Epoch AI estimates $17,000 for a 12-inch (300 mm) wafer on that process [1].
- **Dies per wafer.** A wafer is a circle and dies are rectangles, so the edge of the wafer goes to waste. The bigger the die, the fewer fit. Each of the two logic dies in a B200 is about 800 mm2 [1]. Nvidia calls them reticle-limited, meaning they are near the largest size a scanner can print in one shot [4].
- **Yield.** Yield is the share of dies that work. Defects land at random, so a bigger die is more likely to catch one. Epoch AI models Blackwell die yield at 40 to 70 percent, centered on 60 [1].

Together, those three numbers put the two dies inside a B200 at about $900 of the roughly $6,400 build cost [1]. The estimate counts only variable costs and leaves out fixed costs such as research and development [1].

Variants and trade-offs

Memory is the largest line

The largest line on the bill is the 192 GB of high-bandwidth memory beside those dies, a generation called HBM3E [1]. At $14 to $17 per gigabyte it comes to about $2,900, which Epoch AI counts as half of the roughly $6,400 build cost [1]. Epoch AI says HBM prices are negotiated deal by deal and can vary significantly with

1 NVIDIA's B200 costs around $6,400 to produce. Epoch AI · 10 December 2025 A epoch.ai/data-insights/b200-cost-breakdown

2 AI chip sales documentation. Epoch AI · 8 January 2026 A epoch.ai/data/ai-chip-sales-documentation/methodology#nvidia

4 NVIDIA Blackwell Architecture. Nvidia · 18 December 2025 A nvidia.com/en-us/data-center/technologies/blackwell-architecture

volume and contract length [1]. The $14 to $17 range dates from December 2025, before the 2026 rise in memory makers' sales and margins described below [5].

Three firms, SK hynix, Samsung and Micron, dominate the world market for that memory [3]. SK hynix says demand for its products has exceeded its available supply in recent quarters [6].

Packaging, substrate and test

CoWoS-L, the TSMC process that mounts dies and memory on one carrier and wires them together, adds about $1,100, more than the logic dies it carries [1]. Epoch AI models packaging yield at 65 to 95 percent [1]. A package that fails testing loses its logic dies and memory stacks [1]. Spreading that loss over the good units adds about $1,000 to each B200 [1]. Power delivery hardware, the circuit board, and final assembly and testing add $480 [1]. See Advanced packaging for why packaging capacity runs out first.

The B200 bill of materials

LINE	ESTIMATE	RANGE
HBM3E, 192 GB	$2,900	$2,800-3,100
CoWoS-L packaging	$1,100	$1,000-1,200
Packaging yield loss	$1,000	$430-1,700
Logic die, two on TSMC 4NP	$900	$720-1,200
Power delivery, board, assembly and test	$480	$370-600
Manufacturing cost	~$6,400	$5,700-7,300
Reported selling price per chip, March 2024	$30,000-40,000	
Nvidia compute revenue per GPU in GB200 racks, data to October 2025	$33,000-42,000	

Epoch AI models each line as a range, deriving the packaging line from TSMC's advanced packaging revenue and Nvidia's share of CoWoS capacity [1]. The estimate covers the accelerator module only and leaves out servers and rack-scale systems [1]. Epoch AI notes that most Blackwell revenue comes from those systems, which

1 NVIDIA's B200 costs around $6,400 to produce. Epoch AI · 10 December 2025 A epoch.ai/data-insights/b200-cost-breakdown

5 SK hynix Announces 2Q26 Financial Results. SK hynix A news.skhynix.com/en/q2-2026-business-results

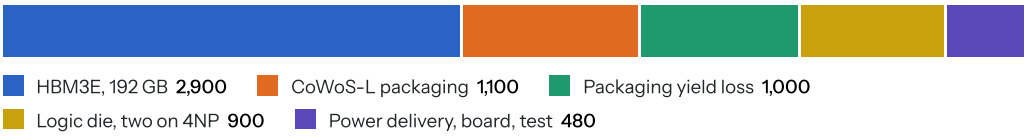
3 Understanding the Biden Administration's Updated Export Controls. Center for Strategic and International Studies · 11 December 2024 A csis.org/analysis/understanding-biden-administrations-updated-export-controls

6 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 A sec.gov/Archives/edgar/data/2120882/000119312526299963/d32785d424b4.htm

may carry lower margins [1]. So, Epoch AI says, Nvidia's actual margin on Blackwell may be lower than the gap between this cost and the reported chip price suggests [1].

The B200 is used here because Epoch AI's line-by-line estimate covers it. In its results for the quarter to July 2026, Nvidia said its Vera Rubin platform was ramping into full production [7]. Each Rubin GPU carries up to 288 GB of HBM4, a newer generation of the stacked memory [8].

Where the ~$6,400 of B200 manufacturing cost goes, Epoch AI estimate, December 2025 $



Source: Epoch AI B200 cost model, from a $17,000 4NP wafer, 40–70% die yield, $14–17 per GB of memory, and a packaging line derived from TSMC's advanced packaging revenue

From package to rack to cluster

At rack scale, cost is measured per GPU-hour, one GPU running for one hour. For an operator that owns the machine, SemiAnalysis puts a Vera Rubin NVL72 at $3.57 per GPU-hour, against $1.84 for a GB200 and $2.36 for a GB300 [9]. Each generation costs its owner more per hour than the last [9].

But SemiAnalysis finds Rubin cheaper per token than GB200 and GB300 at every interactivity level it compared [9]. A token is a short chunk of text, and interactivity is how many tokens per second each user receives. That finding rests on early results from the cloud provider CoreWeave running the DeepSeek R1 model, and SemiAnalysis has not verified them [9].

One level up, a cluster of 100,000 of Nvidia's earlier H100 GPUs costs over $4 billion in servers [10]. It needs more than 150 MW of data center capacity and uses 1.59 TWh a year, which costs $123.9 million at $0.078 per kWh [10]. In a rail-optimized network design, 98,304 optical transceivers link its GPUs to the first layer of switches, called leaf switches [10].

1 NVIDIA's B200 costs around $6,400 to produce. Epoch AI · 10 December 2025 [A] epoch.ai/data-insights/b200-cost-breakdown

7 NVIDIA Announces Financial Results for Second Quarter Fiscal 2027. Nvidia · 26 August 2026 [A] nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2027

8 Inside NVIDIA Rubin GPU Architecture: Powering the Era of Agentic AI. Nvidia · 21 July 2026 [A] developer.nvidia.com/blog/inside-nvidia-rubin-gpu-architecture-powering-the-era-of-agentic…

9 Vera Rubin NVL72 vs GB200 NVL72? Inference TCO & Architecture Analysis. SemiAnalysis · 23 July 2026 [B] inferencex.semianalysis.com/blog/vera-rubin-nvl72-vs-gb200-nvl72-inference

10 100,000 H100 Clusters: Power, Network Topology, Ethernet vs InfiniBand, Reliability, Failures, Checkpointing. SemiAnalysis · 17 June 2024 [B] newsletter.semianalysis.com/p/100000-h100-clusters-power-network

Who captures the value

FIRM	STAGE	QUARTER ENDED	Gross margin
Micron	HBM and DRAM	September 2026	86.8%
Nvidia	Accelerator design	July 2026	75.0%
TSMC	Wafers and CoWoS	June 2026	67.7%
KLA	Inspection	June 2026	61.4%
ASML	Lithography	June 2026	54.0%
Lam Research	Etch	June 2026	51.7%
Applied Materials	Deposition	July 2026	50.3%
Ibiden	Substrates	June 2026	36.4%
ASE Technology	Packaging and test segment	June 2026	27.3%
Amkor	OSAT	June 2026	16.8%

Each figure comes from the firm's own earnings release [11] [7] [12] [13] [14] [15] [16] [17] [18] [19]. The KLA and Ibiden releases give no percentage, so their figures are gross profit divided by revenue [13] [17]. For ASE the table shows its packaging and test segment [18]. The company as a whole had a gross margin of 21.0 percent, because about a third of its revenue came from electronics manufacturing services, which earned 8.9 percent [18].

SK hynix's earnings release gives no gross margin. Its operating margin was 76 percent in the quarter to June 2026, and its revenue was up 257 percent from a

11 Micron Technology, Inc. Reports Record Fiscal Fourth-Quarter and Full-Year 2026 Results. Micron Technology · 30 September 2026 [A] investors.micron.com/news/press-release/2026/Micron-Technology-Inc--Reports-Record-Fiscal-...

7 NVIDIA Announces Financial Results for Second Quarter Fiscal 2027. Nvidia · 26 August 2026 [A] nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2027

12 TSMC Reports Second Quarter EPS of NT$27.25. TSMC · 16 July 2026 [A] pr.tsmc.com/english/news/3326

13 KLA Corporation Reports Fiscal 2026 Fourth Quarter and Full Year Results. KLA Corporation · 28 July 2026 [A] ir.kla.com/news-events/press-releases/detail/518/kla-corporation-reports-fiscal-2026-fourt...

14 ASML reports €9.3 billion total net sales and €2.9 billion net income in Q2 2026. ASML · 15 July 2026 [A] asml.com/en/news/press-releases/2026/q2-2026-financial-results

15 Lam Research Corporation Reports Financial Results for the Quarter Ended June 28, 2026. Lam Research · 29 July 2026 [A] newsroom.lamresearch.com/2026-07-29-Lam-Research-Corporation-Reports-Financial-Results-for...

16 Applied Materials Announces Third Quarter 2026 Results. Applied Materials · 13 August 2026 [A] ir.appliedmaterials.com/news-releases/news-release-details/applied-materials-announces-thi...

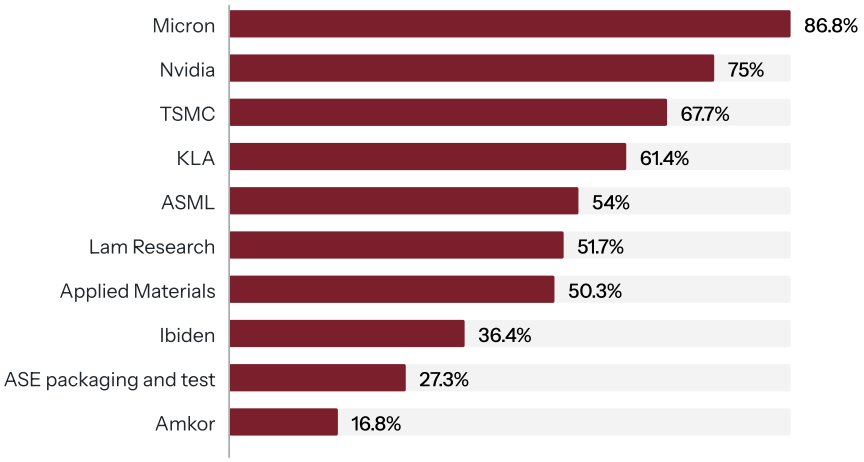
17 Consolidated Financial Results for the First Quarter Ended June 30, 2026. Ibiden · 4 August 2026 [A] ibiden.com/ir/items/tannshinn2026Q1.pdf

18 ASE Technology Holding Co., Ltd. Reports Its Unaudited Consolidated Financial Results for the Second Quarter of 2026. PR Newswire · 30 July 2026 [A] prnewswire.com/news-releases/ase-technology-holding-co-ltd-reports-its-unaudited-consolida...

19 AMKOR TECHNOLOGY, INC., Form 8-K current report for the period ended 2026-07-27 (8-K). U.S. Securities and Exchange Commission (filing by AMKOR TECHNOLOGY, INC.) · 27 July 2026 [A] sec.gov/Archives/edgar/data/0001047127/000104712726000043/amkr6302026erex-991.htm

year earlier [5].

Gross margin, each firm's most recent reported quarter %



Source: Each firm's own quarterly earnings release; quarters end between June and September 2026; ASE's bar is its packaging and test segment (TSMC shown)

In the table, design, equipment and leading-edge wafers earn gross margins of 50.3 to 75.0 percent, substrates 36.4 percent, and assembly and test 16.8 to 27.3 percent. In 2026 Micron passed all of them. Its gross margin rose from 39.8 percent for the year to August 2025 to 86.8 percent in the quarter to September 2026 [11]. Nvidia's was 75.0 percent in the quarter to July 2026 [7].

The Semiconductor Industry Association, using monthly data compiled by World Semiconductor Trade Statistics, counted $791.7 billion of global chip sales in 2025 [20]. In its June 2026 forecast, World Semiconductor Trade Statistics put 2025 at $795.6 billion and 2026 at $1.51 trillion, with memory growing about 250 percent and logic 37 percent [21].

The industry group SEMI counted $135.1 billion of semiconductor equipment billings in 2025 [22]. Capital spending at Alphabet, Amazon, Meta, Microsoft and Oracle came to $448.2 billion in 2025, the sum of Epoch AI's quarterly figures [23]. Epoch AI says that spending would reach $770 billion in 2026 if the trend continues [23]. These totals overlap, so they cannot be added together.

5 SK hynix Announces 2Q26 Financial Results. SK hynix [A] news.skhynix.com/en/q2-2026-business-results

11 Micron Technology, Inc. Reports Record Fiscal Fourth-Quarter and Full-Year 2026 Results. Micron Technology · 30 September 2026 [A] investors.micron.com/news/press-release/2026/Micron-Technology-Inc--Reports-Record-Fiscal-...

7 NVIDIA Announces Financial Results for Second Quarter Fiscal 2027. Nvidia · 26 August 2026 [A] nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2027

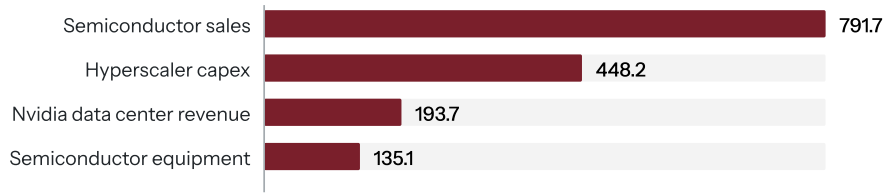
20 Global Annual Semiconductor Sales Increase 25.6% to $791.7 Billion in 2025. Semiconductor Industry Association · 6 February 2026 [A] semiconductors.org/global-annual-semiconductor-sales-increase-25-6-to-791-7-billion-in-2025

21 Global Semiconductor Market Surges Beyond USD 1.5 Trillion in 2026 Driven by Extraordinary Memory Expansion. World Semiconductor Trade Statistics · 2 June 2026 [A] wsts.org/esraCMS/extension/media/f/WST/7618/WSTS_FC-Release-2026-May.pdf

22 SEMI Reports Global Semiconductor Equipment Billings Reached $135 Billion in 2025, Up 15% Year-on-Year. SEMI · 7 April 2026 [A] semi.org/en/SEMI-Reports-Global-Semiconductor-Equipment-Billings-Reached-135-Billion-in-2025

23 Hyperscaler capex has quadrupled since GPT-4's release. Epoch AI · 26 February 2026 [A] epoch.ai/data-insights/hyperscaler-capex-trend

Selected market totals, 2025 $B



Source: Semiconductor Industry Association on chip sales; Epoch AI quarterly capital spending at Alphabet, Amazon, Meta, Microsoft and Oracle, summed for 2025; Nvidia data center revenue for its fiscal year to January 2026; SEMI on equipment billings

Where the money goes

In Epoch AI's estimate for the B200, packaging ($1,100) and packaging yield loss ($1,000) together cost more than twice as much as the logic dies ($900) [1].

The table below sorts each stage by whether its cost falls as volume grows. The classes are a judgment drawn from the sources in this chapter and the stage chapters.

STAGE	CLASS	WHY
EUV lithography	Structurally expensive	One supplier, ASML. See Lithography.
Masks and pellicles	Structurally expensive	Every design needs its own mask set, and every EUV layer a flawless mask. See Photomasks and pellicles.
Metrology and inspection	Structurally expensive	KLA held about five times its nearest rival's share of these tools in 2025, by CSET's count. See Metrology and inspection.
HBM	Structurally expensive	SK hynix says HBM needs more wafer input than ordinary DRAM. See Memory and HBM.
Advanced packaging, CoWoS-L	Structurally expensive	Modeled packaging yield is 65 to 95 percent, and a failed package loses its dies and memory.
ABF substrates	Structurally expensive	Investment can add substrate lines, but Ajinomoto said in June 2026 that the insulating film it makes for these substrates has held over 95 percent of its market since launch [24]. See Substrates and PCBs.
Power and cooling plant	Structurally expensive	Grid connections face long queues, and GE Vernova's gas equipment backlog and slot reservations reached 116 GW in mid-2026. See Data centers and power.
Leading-edge logic wafers	Compressible with scale	A very large die yields well only once a process matures. See Transistors and the front end.
Design and EDA	Compressible with scale	A design is paid for once, and its cost is spread across every chip sold. See Chip design, EDA and IP.
Deposition and etch tools	Compressible with scale	Most deposition and etch steps have three or four established tool makers. See Deposition and etch.

1 NVIDIA's B200 costs around $6,400 to produce. Epoch AI · 10 December 2025 [A] epoch.ai/data-insights/b200-cost-breakdown

24 Ajinomoto, Growth Strategy for the Electronic Materials Business, 30 June 2026. Ajinomoto · 30 June 2026 [A] ajinomoto.co.jp/company/en/ir/event/business_briefing/main/011117/teaserItems1/01/linkList...

Optical transceivers	Compressible with scale	Nvidia agreed in March 2026 to invest $2 billion each in laser makers Coherent and Lumentum, with rights to future capacity. See Systems and networking.
Test and assembly (OSAT)	Already commoditized	Gross margins of 16.8 to 27.3 percent at Amkor and ASE's packaging and test segment, in the table above.
Wafers and bulk gases	Already commoditized	Several firms make approved 300 mm wafers, and the top five served about 75 percent of the market in 2025. See Silicon and wafers.
Rack integration (ODM)	Already commoditized	Several contract manufacturers on three continents assemble racks. See Systems and networking.

Key evaluation criteria

- **Cost per good die.** The wafer price divided by the dies on that wafer that work. A larger die or a lower yield raises it.
- **Memory dollars per gigabyte, and gigabytes per part.** Memory is the largest line on the bill, and three firms make most of it.
- **Packaging yield.** Every point lost throws away finished dies and the memory attached to them [1].
- **Share of the bill the vendor does not control.** Nvidia sets its own margin but buys memory, packaging and substrates at another firm's price.
- **Rack watts.** A GB300 NVL72 rack can draw up to 142 kW, and the data center has to supply and cool all of it [25].
- **Compressibility class.** Whether a line's cost falls when volume doubles. If it does not, capacity limits supply.

THREE THINGS TO REMEMBER

— Epoch AI estimated in December 2025 that a B200 cost about $6,400 to build, against a selling price of $30,000 to $40,000 per chip reported in March 2024.

— The stacked memory is the most expensive part, about half of the cost.

— Packaging costs more than the logic chips, because a failed package wastes everything inside it.

1 NVIDIA's B200 costs around $6,400 to produce. Epoch AI · 10 December 2025 [A] epoch.ai/data-insights/b200-cost-breakdown

25 System Hardware & Components. Nvidia [A] docs.nvidia.com/enterprise-reference-architectures/nvl72-ai-factory/latest/components.html

CHAPTER
17
THE CONTEST

Geopolitics

Since 2018 the United States has used export controls to limit China's access to advanced chips and chipmaking tools, and China has been building its own.

IN PLAIN TERMS

An export control is a rule that requires the government's permission for a sale. The government names an item and a buyer, and from then on its own companies must hold a license before they ship that item to that buyer. The American government extends the rule to goods made abroad with American technology. Few governments can use such a rule to much effect, because it works only when their own firms sell something the buyer cannot get elsewhere. The machines, software and chemicals used to make an advanced chip come from a short list of firms in the United States, the Netherlands, Japan and South Korea. China can copy the simpler machines within a few years. It cannot yet copy the machine that prints the finest circuits, or the chemicals that machine needs. It has also yet to show it can make high-bandwidth memory, the stacked memory beside an AI chip, in volume.

IN SHORT

Since 2022 the United States has restricted sales to China of advanced AI chips and the equipment used to make them, and Japan and the Netherlands have aligned their own controls [1]. China still cannot get extreme ultraviolet (EUV) scanners, the machines that print the finest circuits on the most advanced chips [75]. In September 2025 the Center for Strategic and International Studies (CSIS) judged that making its own EUV machines remained firmly out of China's reach despite massive state investment [75]. But Chinese toolmakers supplied 35 percent of China's own chipmaking equipment market in 2025, up from 10 to 15 percent before 2022 [43].

The most advanced chips come mainly from Taiwan. In 2022 Taiwan held 69 percent of the world's capacity for making logic chips on processes below 10 nanometers, by the count of the Semiconductor Industry Association and Boston Consulting Group [7].

The United States is still changing its export controls, and China is still building its response to them.

1 Where the Chips Fall: U.S. Export Controls Under the Biden Administration from 2022 to 2024. Center for Strategic and International Studies · 12 December 2024 A csis.org/analysis/where-chips-fall-us-export-controls-under-biden-administration-2022-2024

75 Breakthroughs or Boasts? Assessing Recent Chinese Lithography Advancements. Center for Strategic and International Studies A csis.org/blogs/strategic-technologies-blog/breakthroughs-or-boasts-assessing-recent-chines...

43 China's Localization Drive in Semiconductors Gains Impetus from Allied Chip Export Controls. Center for Strategic and International Studies · 24 March 2026 A csis.org/analysis/chinas-localization-drive-semiconductors-gains-impetus-allied-chip-expor...

7 Emerging Resilience in the Semiconductor Supply Chain. Boston Consulting Group and Semiconductor Industry Association · May 2024 A semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc...

The shape of the contest

In September 2022 National Security Advisor Jake Sullivan called securing "as large a lead as possible" in technologies such as AI a national security imperative [1]. The October rules that followed were the first to apply his doctrine of a small yard with a high fence, meaning strict controls on a narrow set of technologies [1]. Those rules went beyond the old aim of staying two technology nodes ahead of China and tried to push China's capability back below the level it had already reached [2].

Whether a sale needed a license once depended on who was buying and why. The 2022 rules also made it depend on where the item was going [3]. They required a license for any advanced AI chip bound for China, and for tools bound for Chinese fabs beyond set technology thresholds [3]. Most applications were presumed denied [3]. But equipment for the China fabs of firms headquartered in the United States or in allied and partner countries was reviewed case by case [3].

Chokepoints and who holds them

SEGMENT	LEADER	SHARE	BASIS
EUV lithography	Netherlands	100.0%	revenue, 2025
Immersion DUV lithography	Netherlands	99.3%	revenue, 2025
Metrology and inspection	United States	71.7%	revenue, 2025
Deposition	United States	60.3%	revenue, 2025
Etch and clean	United States	52.6%	revenue, 2025
Photoresist	Japan	78%	revenue, 2023
Silicon wafers	Japan	53%	revenue, 2023
Leading-edge logic, below 10 nm	Taiwan	69%	capacity, 2022
Foundry, all nodes	Taiwan	78%	market share, 2025

1 Where the Chips Fall: U.S. Export Controls Under the Biden Administration from 2022 to 2024. Center for Strategic and International Studies · 12 December 2024 [A] csis.org/analysis/where-chips-fall-us-export-controls-under-biden-administration-2022-2024

2 Choking off China's Access to the Future of AI. Center for Strategic and International Studies · 11 October 2022 [A] csis.org/analysis/choking-chinas-access-future-ai

3 Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items; Supercomputer and Semiconductor End Use; Entity List Modification. Federal Register (Commerce Department; Industry and Security Bureau) · 13 October 2022 [A] federalregister.gov/documents/2022/10/13/2022-21658/implementation-of-additional-export-co…

HBM	South Korea	82.5%	revenue, 2025: SK hynix 63.2, Samsung 19.3, Micron 17.4
Logic chip design	United States	65%	sales, 2022
Mature logic, 28 nm and above	China	33%	capacity, 2022

- **Equipment.** CSET's ETO Supply Chain Explorer, TechInsights 2025 revenue by parent-company headquarters [4].
- **Materials.** Country shares of photoresist and silicon wafers as Japan's Ministry of Economy, Trade and Industry published them in December 2025, from Fuji Keizai's materials survey, by parent-company headquarters [5]. Japan held 77 percent of photoresist sales in 2021 and 78 percent in 2023, and 52 and 53 percent of silicon wafer sales in those years, so both shares barely changed [5].
- **Taiwan.** The 78 percent is the Taiwan Semiconductor Industry Association's 2025 count, as the Stimson Center gives it [6]. The 69 percent is the Semiconductor Industry Association and Boston Consulting Group's 2022 measurement of logic capacity below 10 nm, the most recent the report gives [7]. The 92 percent quoted more often is the same measure for 2019, from the two groups' April 2021 report [8].
- **Memory.** SK hynix's US listing prospectus splits 2025 high-bandwidth memory revenue between SK hynix at 63.2 percent, Samsung at 19.3 and Micron at 17.4, so Korean firms held 82.5 percent [9]. By the first quarter of 2026 the Korean share was down to 76.9 percent [9]. The underlying figures come from the market research firm IDC, relayed by the prospectus.
- **Design.** CSET's ETO Supply Chain Explorer, from WSTS and SIA figures on 2022 sales of logic chips and microprocessors, still the latest year it gives for this segment [10]. On a wider measure that counts every type of chip, memory and

4 Documentation: Semiconductor Supply Chain Dataset. Emerging Technology Observatory (ETO) [A] eto.tech/dataset-docs/chipexplorer

5 METI, Semiconductor and digital industry strategy: future direction (半導体・デジタル産業戦略の今後の方向性), 23 December 2025. Ministry of Economy, Trade and Industry (Japan) · 23 December 2025 [A] meti.go.jp/policy/mono_info_service/joho/conference/semicon_digital/0014/handeji14-4.pdf

6 Why Taiwan Fears 'America First' Risks Eroding Its 'Silicon Shield'. Stimson Center · 10 October 2025 [A] stimson.org/2025/why-taiwan-fears-america-first-risks-eroding-its-silicon-shield

7 Emerging Resilience in the Semiconductor Supply Chain. Boston Consulting Group and Semiconductor Industry Association · May 2024 [A] semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc…

8 Strengthening the Global Semiconductor Supply Chain in an Uncertain Era. Boston Consulting Group and Semiconductor Industry Association · April 2021 [A] semiconductors.org/wp-content/uploads/2021/05/BCG-x-SIA-Strengthening-the-Global-Semicondu…

9 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 [A] sec.gov/Archives/edgar/data/0002120882/000119312526299963/d32785d424b4.htm

10 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

analog included, US-headquartered firms took 53.4 percent of global chip sales in 2025 [11].

- **China's mature nodes.** The 33 percent is the 2022 industry measurement of logic capacity at 28 nm and above, and the same report projects 37 percent by 2032 [7]. In a February 2025 filing the Semiconductor Industry Association counted capacity for every type of chip made at 28 nm and above, memory and analog included [12]. On that count China's share was 33 percent in 2023, up from 19 percent in 2015 [12].

Advanced packaging is missing from the table because nobody publishes a split of its capacity by country and the Taiwanese chipmaker TSMC will not give one. No country is self-sufficient, so the United States and China can each use export controls to hurt the other.

The US export-control campaign, 2018 to today

The Commerce Department built the controls one rule at a time, each closing a gap the last one left.

- **17 August 2020, the Foreign Direct Product Rule.** In May 2019 Commerce cut Huawei and 68 affiliates off from US technology by adding them to the Entity List, its roster of buyers that need a license [13]. Huawei bought foreign chips instead. Commerce then extended the license requirement to any foreign item made with US technology and bound for Huawei [14]. That reach makes American controls extraterritorial.
- **7 October 2022, the first country-wide rules.** Two new export classifications, 3A090 and 3B090, capped the performance of exportable AI chips and set thresholds for advanced fabs [3]. Those thresholds were FinFET or 16/14 nm logic, 18 nm half-pitch DRAM and 128-layer NAND [3]. The rules also required a

11 2026 State of the U.S. Semiconductor Industry. Semiconductor Industry Association · 2 July 2026 A semiconductors.org/wp-content/uploads/2026/07/SIA_2026_State-of-the-Semiconductor-Industry...

7 Emerging Resilience in the Semiconductor Supply Chain. Boston Consulting Group and Semiconductor Industry Association · May 2024 A semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc...

12 Section 301 Legacy Investigation SIA Comments. Semiconductor Industry Association · 6 February 2025 A semiconductors.org/wp-content/uploads/2025/02/USTR-2024-0024-00109674-CAT-5016-Public-Docu...

13 Addition of Entities to the Entity List. Federal Register (Commerce Department; Industry and Security Bureau) · 21 May 2019 A federalregister.gov/documents/2019/05/21/2019-10616/addition-of-entities-to-the-entity-list

14 Addition of Huawei Non-U.S. Affiliates to the Entity List, the Removal of Temporary General License, and Amendments to General Prohibition Three (Foreign-Produced Direct Product Rule). Federal Register (Commerce Department; Industry and Security Bureau) · 20 August 2020 A federalregister.gov/documents/2020/08/20/2020-18213/addition-of-huawei-non-us-affiliates-t...

3 Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items; Supercomputer and Semiconductor End Use; Entity List Modification. Federal Register (Commerce Department; Industry and Security Bureau) · 13 October 2022 A federalregister.gov/documents/2022/10/13/2022-21658/implementation-of-additional-export-co...

license for US persons to support chip development or production at those advanced Chinese fabs [3].

- **17 October 2023, the density fix.** Nvidia had designed its China-specific chips to stay under the 2022 limit on interconnect speed, how fast a chip passes data to other chips [15]. So Commerce's Bureau of Industry and Security (BIS) dropped that limit [15]. Since then BIS has identified controlled chips by total processing performance (TPP), a measure of computing power, by performance density, and by whether they are designed or marketed for data centers [15]. Under the 2023 rule, a data-center chip needs a license at a TPP of 4,800, or at 1,600 with a performance density of 5.92 [15].
- **2 December 2024, the widest rule.** The rules issued that day covered 24 equipment types and three software tools, added the first controls on high-bandwidth memory, and put 140 entities on the Entity List [16].
- **15 January 2025, AI Diffusion.** The rule required a license to ship the most powerful AI chips and computers, 3A090.a and 4A090.a, to any destination worldwide, then reopened access through exceptions [17]. License Exception AIA covered the nineteen destinations named in supplement no. 5 to part 740 [17]. License Exception LPP allowed capped volumes elsewhere, and no exception covered China [17]. On 13 May 2025, two days before the compliance date, BIS said it would not enforce the rule [18]. It withdrew none of the text. The worldwide license requirement and the 790 million TPP per-country allocation still stand in the Code of Federal Regulations [19]. So do License Exceptions AIA, ACM and

3 Implementation of Additional Export Controls: Certain Advanced Computing and Semiconductor Manufacturing Items; Supercomputer and Semiconductor End Use; Entity List Modification. Federal Register (Commerce Department; Industry and Security Bureau) · 13 October 2022 A federalregister.gov/documents/2022/10/13/2022-21658/implementation-of-additional-export-co…

15 Explainer: The Commerce Department's October 2023 Export Control Update. Center for Security and Emerging Technology (CSET) · 4 December 2023 A cset.georgetown.edu/article/bis-2023-update-explainer

16 Commerce Strengthens Export Controls to Restrict China's Capability to Produce Advanced Semiconductors for Military Applications. U.S. Bureau of Industry and Security · 2 December 2024 A bis.gov/press-release/commerce-strengthens-export-controls-restrict-chinas-capability-prod…

17 Framework for Artificial Intelligence Diffusion. Federal Register (Commerce Department; Industry and Security Bureau) · 15 January 2025 A federalregister.gov/documents/2025/01/15/2025-00636/framework-for-artificial-intelligence-…

18 Department of Commerce Announces Rescission of Biden-Era Artificial Intelligence Diffusion Rule, Strengthens Chip-Related Export Controls. U.S. Bureau of Industry and Security · 13 May 2025 A bis.gov/press-release/department-commerce-announces-rescission-biden-era-artificial-intell…

19 15 CFR 742.6, Regional stability. Electronic Code of Federal Regulations A ecfr.gov/current/title-15/subtitle-B/chapter-VII/subchapter-C/part-742/section-742.6

LPP and supplement no. 5 [20]. BIS says it enforces the worldwide requirement only against Country Groups D:1, D:4 and D:5 and against firms headquartered in D:5 or Macau[21]. No replacement rule that sorts countries into tiers has appeared. Bilateral deals took the place of its caps.

On 9 April 2025 the government told Nvidia it needed a license to ship its H20 chip to China and to the D:5 group of arms-embargoed countries [22]. The requirement also covered any company headquartered or ultimately owned in those places [22]. That condition meant no Chinese buyer could route an order through a foreign subsidiary [22]. Nvidia took a $4.5 billion charge for excess H20 inventory and purchase obligations [23].

The government granted H20 licenses in August 2025 [23]. Nvidia's filing for the quarter ended 27 July 2025 said US officials expected the government to receive 15 percent of the revenue from licensed H20 sales, though no regulation codifies that requirement [23]. Nvidia's annual report for the year to 25 January 2026 put its H20 revenue under those licenses at about $60 million [24]. China restricted those sales, and Nvidia said in its filing for the quarter ended 26 July 2026 that it had been unable to sell its H20 inventory [25].

Three more steps followed in early 2026.

- **13 January 2026, effective 15 January.** BIS announced on 13 January that it would move Nvidia's H200 and AMD's MI325X to case-by-case review, on conditions that included independent third-party testing in the United States

20 A ecfr.gov/current/title-15/subtitle-B/chapter-VII/subchapter-C/part-740

21 Enhanced Favorable Treatment for the United Arab Emirates Under the Export Administration Regulations. Federal Register (Commerce Department; Industry and Security Bureau) · 14 July 2026 A federalregister.gov/documents/2026/07/14/2026-14132/enhanced-favorable-treatment-for-the-u…

22 NVIDIA CORP, Form 8-K current report for the period ended 2025-04-09 (8-K). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 15 April 2025 A sec.gov/Archives/edgar/data/1045810/000104581025000082/nvda-20250409.htm

23 NVIDIA CORP, Form 10-Q quarterly report for the period ended 2025-07-27 (10-Q). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 27 August 2025 A sec.gov/Archives/edgar/data/1045810/000104581025000209/nvda-20250727.htm

24 NVIDIA CORP, Form 10-K annual report for the period ended 2026-01-25 (10-K). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 25 February 2026 A sec.gov/Archives/edgar/data/1045810/000104581026000021/nvda-20260125.htm

25 NVIDIA CORP, Form 10-Q quarterly report for the period ended 2026-07-26 (10-Q). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 26 August 2026 A sec.gov/Archives/edgar/data/1045810/000104581026000075/nvda-20260726.htm

[26]. The rule, at 91 FR 1684, took effect on publication two days later [27]. It caps exports to China or Macau at 50 percent of the total processing performance of the same chips shipped to US customers [27].

- **14 January 2026, effective 15 January.** A national-security proclamation under Section 232 put a 25 percent tariff on the same class of chips and exempted almost every American use, from data centers to startups [28]. The H200 licenses require each China-bound chip to pass an inspection in the United States before shipment, so every licensed H200 is imported first and pays the tariff [25]. Nvidia said in its filing for the quarter ended 26 July 2026 that it had been unable to pass any of that tariff on to its customers [25].
- **12 February 2026.** Applied Materials agreed to pay about $252 million, the statutory maximum, over about $126 million of ion implanters routed through Korea to an Entity-Listed Chinese customer [29].

Eight rule changes that shaped US chip export controls

- 2019-05-16
  Huawei and 68 affiliates added to the Entity List
- 2022-10-07
  First country-wide controls on AI chips and fab tools
- 2023-10-17
  Performance density added; A800 and H800 captured
- 2024-12-02
  24 tool types, HBM controls, 140 entity listings
- 2025-01-15
  AI Diffusion rule published, three country tiers
- 2025-05-13
  BIS stops enforcing the AI Diffusion rule
- 2025-08-11
  US officials seek 15% of licensed H20 sales
- 2026-01-13
  H200 and MI325X case-by-case review announced, effective 15 January

Source: BIS press releases, Federal Register rules, a CSET explainer and Nvidia filings, cited in the text

26 Department of Commerce Revises License Review Policy for Semiconductors Exported to China. U.S. Bureau of Industry and Security · 13 January 2026 [A] bis.gov/press-release/department-commerce-revises-license-review-policy-semiconductors-exp…

27 Revision to License Review Policy for Advanced Computing Commodities. Federal Register (Commerce Department; Industry and Security Bureau) · 15 January 2026 [A] federalregister.gov/documents/2026/01/15/2026-00789/revision-to-license-review-policy-for-…

28 Adjusting Imports of Semiconductors, Semiconductor Manufacturing Equipment, and Their Derivative Products Into the United States. Federal Register (Executive Office of the President) · 20 January 2026 [A] federalregister.gov/documents/2026/01/20/2026-01052/adjusting-imports-of-semiconductors-se…

25 NVIDIA CORP, Form 10-Q quarterly report for the period ended 2026-07-26 (10-Q). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 26 August 2026 [A] sec.gov/Archives/edgar/data/1045810/000104581026000075/nvda-20260726.htm

29 Applied Materials to Pay $252 Million Penalty to BIS for Illegally Exporting Semiconductor Manufacturing Equipment. U.S. Bureau of Industry and Security · 11 February 2026 [A] bis.gov/press-release/applied-materials-pay-252-million-penalty-bis-illegally-exporting-se…

Cloud and remote access

Renting out time on a chip does not count as exporting it. Under 15 CFR 734.13, an export is a shipment or transmission out of the United States, or a release of technology or source code to a foreign person in the United States [30]. So a Singapore data center selling time on an H200 already installed there moves nothing across a border [30].

The American rules reach that trade through the firms holding the chips. A January 2026 license for China or Macau requires the applicant to obtain the consignee's know-your-customer procedures [31]. The applicant must also obtain the consignee's list of Infrastructure-as-a-Service customers, the users renting its computing power, in Belarus, China, Cuba, Iran, Macau, North Korea, Russia and Venezuela [31]. The rule calls that information "necessary to prevent unauthorized remote access from end users described in paragraph (dd)(1)(iv)" [31]. The license also depends on a promise that no algorithm trained on the chips is served back to those users [31].

The AI Diffusion rule also addressed the chips' output, trained AI models. It classified the weights of frontier AI models, the numbers a model learns in training, as ECCN 4E091 [17]. It warned American cloud providers that training such a model for the US subsidiary of a foreign-headquartered customer raises a red flag that the weights will leave without a license [17]. It also barred data-center operators with validated end-user status from training such a model outside the nineteen listed destinations without BIS authorization [17]. But it said that API access to a model, and rented capacity for inference, "are not prohibited" [17].

So the controls cover the chip, frontier model weights and the promises a licensed consignee signs. They do not cover computing time sold by the hour on chips lawfully installed abroad.

30 **15 CFR 734.13 -- Export..** Electronic Code of Federal Regulations A ecfr.gov/current/title-15/subtitle-B/chapter-VII/subchapter-C/part-734/section-734.13

31 **Revision to License Review Policy for Advanced Computing Commodities.** U.S. Government Publishing Office (Federal Register) · 15 January 2026 A govinfo.gov/content/pkg/FR-2026-01-15/html/2026-00789.htm

17 **Framework for Artificial Intelligence Diffusion.** Federal Register (Commerce Department; Industry and Security Bureau) · 15 January 2025 A federalregister.gov/documents/2025/01/15/2025-00636/framework-for-artificial-intelligence-...

Allies: the Netherlands, Japan, Korea and Europe

US tool controls alone would only shift the sales to Dutch and Japanese suppliers, so the campaign depends on those two governments.

- **The Netherlands.** The Dutch government began licensing advanced deep-ultraviolet tools on 1 September 2023 and widened the controls in September 2024 [32]. On 1 April 2025 it extended them to measurement and inspection tools [33].
- **Japan.** Japan's controls on advanced manufacturing equipment took effect on 23 July 2023, apply to all destinations and name no country [34].
- **South Korea.** BIS revoked the Validated End-User status that let Samsung's and SK hynix's China subsidiaries and Intel Semiconductor (Dalian) receive US equipment without individual licenses, effective 31 December 2025 [35]. BIS said it would license those fabs to keep running, but not to expand capacity or upgrade technology [36]. SK hynix, which bought the Dalian fab from Intel, received an annual BIS license in December 2025 to ship approved US equipment to its China fabs during 2026 [37].

A BIS rule effective 29 September 2025 extended Entity List restrictions to any company at least 50 percent owned by one or more listed firms [38]. On 30 September the Dutch economy minister invoked the Goods Availability Act against Nexperia, a Nijmegen chipmaker that supplies Europe's car industry [39]. He cited governance failures that risked moving technology out of Europe [39].

32 The Netherlands expands export control measure for advanced semiconductor manufacturing equipment. Government of the Netherlands · 6 September 2024 A government.nl/latest/news/2024/09/06/the-netherlands-expands-export-control-measure-advanc…

33 Klever: export controls on advanced semiconductor manufacturing equipment to be tightened. Government of the Netherlands · 15 January 2025 A government.nl/latest/news/2025/01/15/klever-export-controls-on-advanced-semiconductor-manu…

34 CSIS Translation: Updated Japanese Export Controls on High-Performance Semiconductor Manufacturing Equipment. Center for Strategic and International Studies · 18 July 2023 A csis.org/analysis/csis-translation-updated-japanese-export-controls-high-performance-semic…

35 Revocation of Validated End-User Authorizations in the People's Republic of China. Federal Register (Commerce Department; Industry and Security Bureau) · 2 September 2025 A federalregister.gov/documents/2025/09/02/2025-16735/revocation-of-validated-end-user-autho…

36 Department of Commerce Closes Export Controls Loophole for Foreign-Owned Semiconductor Fabs in China. U.S. Bureau of Industry and Security · 29 August 2025 A bis.gov/press-release/department-commerce-closes-export-controls-loophole-foreign-owned-se…

37 SK hynix Inc., Form 424B4 prospectus for the period ended 2026-07-10 (424(B)(4)). U.S. Securities and Exchange Commission (filing by SK hynix Inc.) · 10 July 2026 A sec.gov/Archives/edgar/data/2120882/000119312526299963/d32785d424b4.htm

38 Expansion of End-User Controls To Cover Affiliates of Certain Listed Entities. Federal Register (Commerce Department; Industry and Security Bureau) · 30 September 2025 A federalregister.gov/documents/2025/09/30/2025-19001/expansion-of-end-user-controls-to-cove…

39 Minister of Economic Affairs invokes Goods Availability Act. Government of the Netherlands · 12 October 2025 A government.nl/latest/news/2025/10/12/minister-of-economic-affairs-invokes-goods-availabili…

Both were suspended in November. BIS suspended its rule for a year on 10 November [40]. The minister suspended his order on 19 November [41]. Europe holds a chokepoint in leading-edge lithography, but none in the mature-node chips its car industry uses.

The MATCH Act, a bill introduced on 8 April 2026, would ban sales of immersion deep-ultraviolet scanners to China [42]. It would also name five Chinese companies, CXMT, Hua Hong, Huawei, SMIC and YMTC, in statute [42]. It would give allies 150 days to match those limits before the United States extends the Foreign Direct Product Rule on its own [42].

China's response

Five government bodies run China's response. The Ministry of Industry and Information Technology (MIIT) sets substitution targets, and in 2024 reportedly told Chinese telecom operators to strip foreign chips out of their networks by 2027 [43]. The National Development and Reform Commission (NDRC) plans China's build-out of computing capacity and co-issued the national computing-power network plan [44]. The Ministry of Commerce (MOFCOM) writes and administers the export controls, including the mineral rules below [45]. The State Administration for Market Regulation (SAMR) is the antitrust regulator and opened an investigation into Nvidia in December 2024 [46]. The Cyberspace Administration of China (CAC) runs the security reviews, and in September 2025 told firms including ByteDance and Alibaba to cancel Nvidia orders [47].

China has retaliated against American suppliers and restricted its exports of minerals. Micron, the American memory maker, was first, failing a CAC review on 21 May 2023 and losing its critical-infrastructure buyers in China [48]. Mineral controls followed.

- **3 December 2024.** MOFCOM's Notice 46 banned exports of gallium, germanium, antimony and superhard materials to the United States, one day after the

40 One Year Suspension of Expansion of End-User Controls for Affiliates of Certain Listed Entities. Federal Register (Commerce Department; Industry and Security Bureau) · 12 November 2025 A federalregister.gov/documents/2025/11/12/2025-19846/one-year-suspension-of-expansion-of-en…

41 Update on invoking Goods Availability Act. Government of the Netherlands · 19 November 2025 A government.nl/documents/2025/11/19/update-on-invoking-goods-availability-act

42 Risch, Ricketts, Kim Introduce MATCH Act; Level the Global Playing Field for U.S. Tech. U.S. Senate Committee on Foreign Relations · 8 April 2026 A foreign.senate.gov/press/rep/release/risch-ricketts-kim-introduce-match-act-level-the-glob…

43 China's Localization Drive in Semiconductors Gains Impetus from Allied Chip Export Controls. Center for Strategic and International Studies · 24 March 2026 A csis.org/analysis/chinas-localization-drive-semiconductors-gains-impetus-allied-chip-expor…

44 China accelerates building of national computing power network. The State Council of the People's Republic of China · 27 December 2023 A english.www.gov.cn/news/202312/27/content_WS658b72afc6d0868f4e8e28ba.html

45 Announcement No.18 of 2025 of The Ministry of Commerce and The General Administration of Customs of The People's Republic of China Announcing the decision to implement export control on some medium and heavy rare earth related items. Ministry of Commerce of the People's Republic of China A english.mofcom.gov.cn/Policies/AnnouncementsOrders/art/2025/art_0dd87cbee7b045bf93fabe6ab2…

46 Cyber Week in Review: December 13, 2024. Council on Foreign Relations · 13 December 2024 A cfr.org/articles/cyber-week-review-december-13-2024

47 Beijing blocks Nvidia sales in vote of confidence for Chinese chipmakers, the White House strikes a deal for 10 percent of Intel, and Commerce voids a $7.4B CHIPS Act grant. Center for Security and Emerging Technology (CSET) · 18 September 2025 A cset.georgetown.edu/newsletter/september-18-2025

48 Micron Aggression: The Right Response to Beijing's Ban on the U.S. Chipmaker. Center for Strategic and International Studies · 22 June 2023 A csis.org/analysis/micron-aggression-right-response-beijings-ban-us-chipmaker

American chip rules [49].

- **4 April 2025.** Announcement 18 required export licenses for seven medium and heavy rare earths, from samarium to yttrium, as metal, oxide, compound or magnet material [45].
- **9 October 2025.** A further announcement extended licensing to foreign-made goods containing Chinese rare earths worth 0.1 percent or more of their value [50]. A foreign firm would then need a Chinese export license to ship its own product [50]. MOFCOM suspended the announcement on 7 November 2025 until 10 November 2026 [51].
- **November 2025.** After Donald Trump and Xi Jinping met at the Asia-Pacific summit, MOFCOM suspended the gallium ban for a year [52].

China repealed none of these controls. The gallium ban and the October rules are paused, and the April licensing rules still apply.

The Chinese government also discouraged customers from buying, importing or using Nvidia's data-center products, and endorsed performance-per-watt and bandwidth standards for accelerators in Chinese data centers [25]. When the United States licensed H200 sales in February 2026, China restricted them [25]. In the quarter ended 26 July 2026, Nvidia's H200 shipments under those licenses came to less than 1 percent of its data-center revenue [25]. Nvidia's data-center revenue that quarter was $89.0 billion, and its guidance for the next quarter assumed no China compute revenue [53].

Chinese toolmakers held 10 to 15 percent of their home market before 2022, 25 percent in 2024 and 35 percent in 2025 [43]. They hold about 40 percent of the Chinese market for etch and deposition tools [43]. At ASML, the Dutch maker of

49 Ministry of Commerce Notice 2024 No. 46: Notice Concerning Strengthening Controls on Exports of Relevant Dual-Use Items to the United States. Center for Security and Emerging Technology (CSET) · 3 December 2024 [A] cset.georgetown.edu/publication/china-rare-earth-export-ban

45 Announcement No.18 of 2025 of The Ministry of Commerce and The General Administration of Customs of The People's Republic of China Announcing the decision to implement export control on some medium and heavy rare earth related items. Ministry of Commerce of the People's Republic of China [A] english.mofcom.gov.cn/Policies/AnnouncementsOrders/art/2025/art_0dd87cbee7b045bf93fabe6ab2...

50 Ministry of Commerce Notice 2025 No. 61: Announcement of the Decision to Implement Controls on Exports of Rare Earth-Related Items to Foreign Countries. Center for Security and Emerging Technology (CSET) · 9 October 2025 [A] cset.georgetown.edu/publication/mofcom-notice-2025-61

51 MOFCOM and GAC Announcement 2025 No. 70, suspending the October 2025 export-control measures, Announcements Nos. 55 to 58, 61 and 62 (商务部、海关总署公告2025年第70号). Ministry of Commerce of the People's Republic of China · 7 November 2025 [A] mofcom.gov.cn/zwgk/zcfb/art/2025/art_b1ec77dd3f0d4762952904df7cdaadec.html

52 The U.S.-China Trade Truce Has Not Solved the Gallium Problem. Center for Strategic and International Studies · 11 May 2026 [A] csis.org/analysis/us-china-trade-truce-has-not-solved-gallium-problem

25 NVIDIA CORP, Form 10-Q quarterly report for the period ended 2026-07-26 (10-Q). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 26 August 2026 [A] sec.gov/Archives/edgar/data/1045810/000104581026000075/nvda-20260726.htm

53 NVIDIA Announces Financial Results for Second Quarter Fiscal 2027. Nvidia · 26 August 2026 [A] nvidianews.nvidia.com/news/nvidia-announces-financial-results-for-second-quarter-fiscal-2027

43 China's Localization Drive in Semiconductors Gains Impetus from Allied Chip Export Controls. Center for Strategic and International Studies · 24 March 2026 [A] csis.org/analysis/chinas-localization-drive-semiconductors-gains-impetus-allied-chip-expor...

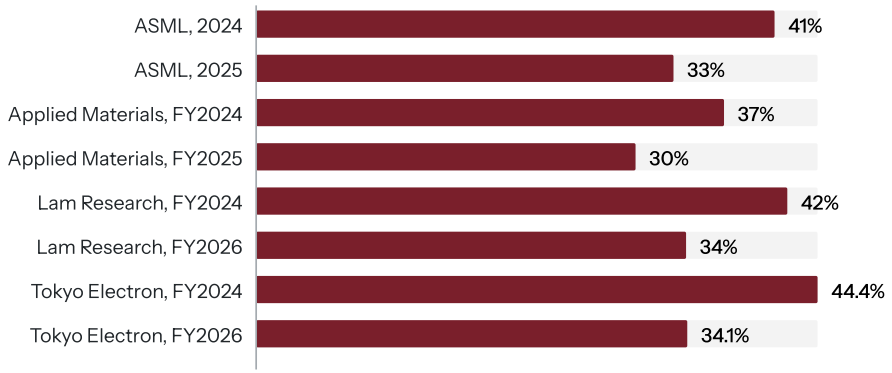
lithography machines, China's share of net system sales fell from 41 percent in 2024 to 33 percent in 2025 [54].

Chinese chipmakers remain behind at the leading edge. SMIC's most advanced process, N+3, reaches the transistor density of TSMC's 6 nm-class process without extreme-ultraviolet scanners [55]. It prints its tightest layers in several passes on deep-ultraviolet tools and relies on redesigned standard cells, the small circuit blocks a chip is built from [55]. In September 2025 the research firm SemiAnalysis expected Huawei to ship about 805,000 of its Ascend AI chips in 2025, with dies made by both TSMC and SMIC [56].

SemiAnalysis judged that memory would limit Huawei's output first. By its count, a stockpile of roughly 13 million memory stacks, most bought before the December 2024 rules took effect, was enough for about 1.6 million Ascend 910C packages [56]. It expected CXMT, China's own memory maker, to make only about 2 million stacks in 2026, enough for 250,000 to 300,000 more [56].

Huawei said in September 2025 that its Ascend 950 chips would carry high-bandwidth memory of its own design [57]. HiBL 1.0 would go on the 950PR, due in the first quarter of 2026, and HiZQ 2.0 on the 950DT, due in the fourth quarter [57]. Huawei gave no figure for how much of that memory it would make [57].

China's share of revenue at the four largest chip-tool makers, peak fiscal year and latest %



Source: ASML Q4 2025 investor presentation, Applied Materials and Lam Research 10-K filings, Tokyo Electron Q4 FY2026 results

54 ASML 2025 fourth-quarter and full-year results (investor presentation). ASML · 28 January 2026 [A] ourbrand.asml.com/m/3136300aa4999bc1/original/2026_01_28_Presentation-Investor-Relations-Q...

55 Is SMIC N+3's Metal Pitch Smaller than Intel 18A's?. SemiAnalysis · 14 June 2026 [B] newsletter.semianalysis.com/p/steel-smic-n3-teardown

56 Huawei Ascend Production Ramp: Die Banks, TSMC Continued Production, HBM is The Bottleneck. SemiAnalysis · 8 September 2025 [B] newsletter.semianalysis.com/p/huawei-ascend-production-ramp

57 Groundbreaking SuperPoD Interconnect: Leading a New Paradigm for AI Infrastructure. Huawei · 18 September 2025 [A] huawei.com/en/news/2025/9/hc-xu-keynote-speech

Industrial policy: CHIPS, Rapidus, Big Fund and the rest

The CHIPS and Science Act authorized about $52.7 billion [43]. By July 2025 Commerce had awarded $30.9 billion of direct funding to 19 companies for 40 projects, covering a median 14.2 percent of each project's capital cost [58]. Commerce counted a federal tax credit, the Advanced Manufacturing Investment Credit, when it sized those grants [58]. Since the One Big Beautiful Bill Act of 2025, the credit covers 35 percent of qualified investment in an advanced chip plant placed in service after 2025, up from 25 percent [58].

On 22 August 2025 the government converted $5.7 billion of unpaid Intel grants and $3.2 billion of Secure Enclave defense money into $8.9 billion of Intel stock, a 9.9 percent stake [59].

The United States has also used tariffs to draw chip investment. The 2026 US-Taiwan trade agreement cuts the reciprocal tariff on Taiwanese goods to 15 percent in exchange for at least $250 billion of Taiwanese investment in US chip production [60]. On 16 July 2026 TSMC announced another $100 billion for four more Arizona fabs at 2 nm or more advanced, bringing its total Arizona investment to $265 billion [61].

The European Commission says the EU still depends on other countries for advanced chip manufacturing [62]. Its Chips Act 2.0 proposal, adopted on 3 June 2026, adds measures to support advanced chip production in the EU [62].

Japan instead backs a homegrown chipmaker, Rapidus. It began prototyping 2 nm gate-all-around transistors at its IIM-1 plant in July 2025 and measured their first electrical characteristics [63]. In February 2026 it raised another 267.6 billion yen, about $1.7 billion, from a government agency and 32 companies, aiming for 2 nm mass production by 2027 [64]. By June 2026 it had raised 424.95 billion yen in equity, the latest 150 billion yen from a state agency [65].

43 China's Localization Drive in Semiconductors Gains Impetus from Allied Chip Export Controls. Center for Strategic and International Studies · 24 March 2026 A csis.org/analysis/chinas-localization-drive-semiconductors-gains-impetus-allied-chip-expor…

58 GAO-26-107882, SEMICONDUCTORS: Information on Projects Funded to Strengthen U.S. Supply Chain. U.S. Government Accountability Office A files.gao.gov/reports/GAO-26-107882/index.html

59 Intel and Trump Administration Reach Historic Agreement to Accelerate American Technology and Manufacturing Leadership. Intel · 22 August 2025 A intc.com/news-events/press-releases/detail/1748/intel-and-trump-administration-reach-histo…

60 U.S.-Taiwan Trade Agreement Leaves Major Questions Open. Council on Foreign Relations · 13 February 2026 A cfr.org/articles/u-s-taiwan-trade-agreement-leaves-major-questions-open

61 TSMC Announces Additional $100 Billion Investment In Arizona. City of Phoenix · 16 July 2026 A phoenix.gov/newsroom/ced-news/tsmc-announces-additional--100-billion-investment-in-arizona…

62 Proposal for the Chips Act 2.0. European Commission A digital-strategy.ec.europa.eu/en/library/proposal-chips-act-20

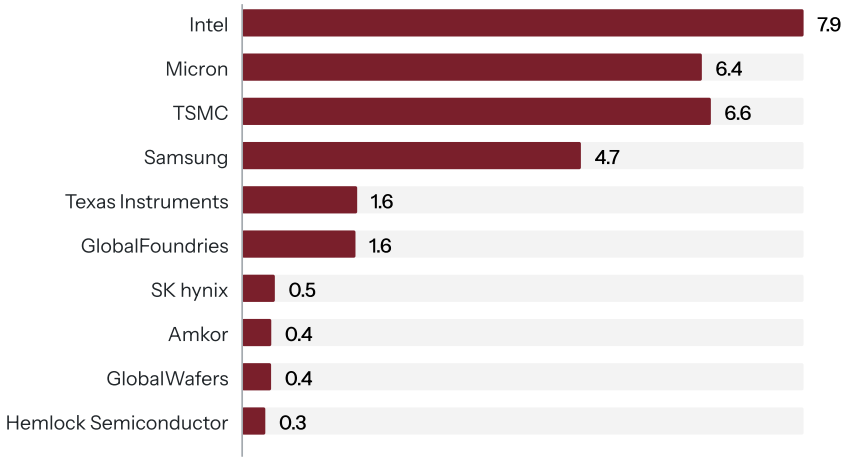
63 Rapidus Achieves Significant Milestone at its State-of-the-Art Foundry with Prototyping of Leading-Edge 2nm GAA Transistors. Rapidus · 18 July 2025 A rapidus.inc/en/news_topics/information/rapidus-achieves-significant-milestone-at-its-state…

64 Rapidus Secures 267.6 Billion Yen in Funding from Japan Government and Private Sector Companies. Rapidus · 27 February 2026 A rapidus.inc/en/news_topics/information/rapidus-secures-267-6-billion-yen-in-funding-from-j…

65 Rapidus Completes 150 Billion Yen Funding Round from Japan Government. Rapidus · 5 June 2026 A rapidus.inc/en/news_topics/information/20260605-2

China spends more than the United States. The latest round of its National Integrated Circuit Industry Investment Fund, the Big Fund, reached $47 billion [66]. CSIS puts cumulative state funding since 2014 at an estimated $150 billion and calls the CHIPS Act substantially smaller in scale [43]. The money has produced a domestic tool industry but no leading-edge node.

CHIPS Act direct funding awarded, by firm, as of July 2025 $B



Source: GAO-26-107882, Semiconductors: Information on Projects Funded to Strengthen U.S. Supply Chain

Taiwan

No policy has reduced the world's reliance on Taiwan. Taiwan's foundries hold 78 percent of the world's contract chipmaking on the Taiwan Semiconductor Industry Association's 2025 count, and TSMC makes its 2 nm chips only in Taiwan [6]. Taiwan held 69 percent of the world's logic capacity below 10 nm in 2022, by a 2024 count from the Semiconductor Industry Association and Boston Consulting Group [7].

Advanced packaging is in Taiwan too. Asked in October 2025 how much CoWoS packaging capacity he would add, TSMC's chairman said only that the company

66 China's Evolving Industrial Policy for AI. RAND Corporation · 26 June 2025 A rand.org/pubs/perspectives/PEA4012-1.html

43 China's Localization Drive in Semiconductors Gains Impetus from Allied Chip Export Controls. Center for Strategic and International Studies · 24 March 2026 A csis.org/analysis/chinas-localization-drive-semiconductors-gains-impetus-allied-chip-expor...

6 Why Taiwan Fears 'America First' Risks Eroding Its 'Silicon Shield'. Stimson Center · 10 October 2025 A stimson.org/2025/why-taiwan-fears-america-first-risks-eroding-its-silicon-shield

7 Emerging Resilience in the Semiconductor Supply Chain. Boston Consulting Group and Semiconductor Industry Association · May 2024 A semiconductors.org/wp-content/uploads/2024/05/Report_Emerging-Resilience-in-the-Semiconduc...

was "working very hard to narrow the gap between the demand and supply" [67]. He called wafer and packaging capacity both "very tight" [67].

The silicon shield argument holds that this concentration deters an invasion, because destroying TSMC would cost an aggressor more than Taiwan is worth. The counterargument is that staying irreplaceable gives Taiwan's government reason to slow the US push to spread chipmaking beyond Taiwan.

Taiwan's rules keep TSMC's newest process at home first, so TSMC's Arizona fabs have run a generation behind its fabs in Taiwan [6]. The Stimson Center expected 2 nm production in Arizona in 2028 at the earliest [6]. TSMC said in October 2025 that it was "preparing to upgrade our technologies faster to N2 and more advanced process technologies in Arizona" [67]. In July 2026 it announced four more fabs there at 2 nm or more advanced [61].

The Arizona fabs narrow the gap in wafer production, and TSMC has two advanced packaging fabs planned there [67]. But a fab without Taiwan's mask shops, chemical suppliers and trained crews is not a substitute. So the silicon shield still holds, though it is weakening.

Where the chips go: the Gulf, Southeast Asia, and smuggling

The United States has approved large AI chip sales to the United Arab Emirates and Saudi Arabia. In November 2025 Commerce authorized two AI companies, G42 of the Emirates and HUMAIN of Saudi Arabia, to buy up to 35,000 Nvidia GB300s each under security and reporting conditions [68].

On 10 July 2026 Commerce moved the Emirates from Country Groups D:3 and D:4 into A:5, the group that holds close US allies [21]. It named the Emirati government, G42 and US-headquartered AI companies as consignees that can receive

67 Q3 2025 Taiwan Semiconductor Manufacturing Co Ltd Earnings Call (Chinese, English) — edited transcript. Refinitiv StreetEvents (transcript of a TSMC earnings call), via TSMC · 16 October 2025 A investor.tsmc.com/english/encrypt/files/encrypt_file/reports/2025-10/6860312f04fd291d0f26b...

6 Why Taiwan Fears 'America First' Risks Eroding Its 'Silicon Shield'. Stimson Center · 10 October 2025 A stimson.org/2025/why-taiwan-fears-america-first-risks-eroding-its-silicon-shield

61 TSMC Announces Additional $100 Billion Investment In Arizona. City of Phoenix · 16 July 2026 A phoenix.gov/newsroom/ced-news/tsmc-announces-additional--100-billion-investment-in-arizona...

68 Statement on UAE and Saudi Chip Exports. U.S. Department of Commerce · November 2025 A commerce.gov/news/press-releases/2025/11/statement-uae-and-saudi-chip-exports

21 Enhanced Favorable Treatment for the United Arab Emirates Under the Export Administration Regulations. Federal Register (Commerce Department; Industry and Security Bureau) · 14 July 2026 A federalregister.gov/documents/2026/07/14/2026-14132/enhanced-favorable-treatment-for-the-u...

controlled AI chips and systems without a license [21]. Other buyers in the Emirates still need licenses for the most powerful chips and computers, 3A090.a and 4A090.a [21]. Saudi buyers still need individual licenses.

Smugglers route chips past the controls through Southeast Asia. Since 14 July 2025 Malaysia has required a permit to export, transship or transit any high-performance US-origin AI chip, plus 30 days' notice under its Strategic Trade Act [69]. On 25 March 2026 the Justice Department announced charges against three people accused of using Thai companies' names to order servers with export-controlled chips for China [70]. The orders included one for $170 million that was never completed [70].

The House select committee on China cites a median estimate of 140,000 chips smuggled to China in 2024 [71]. A bill in Congress, the Chip Security Act, would require location verification on controlled chips within 180 days of enactment [72]. The House Foreign Affairs Committee passed it 42-0 on 26 March 2026 [73]. It has not become law.

Who gains and who pays

- **Taiwan and Korea.** None of these policies will move leading-edge logic, memory stacking or advanced packaging out of Taiwan and Korea before 2030. Reshoring adds capacity elsewhere without taking any away from them.
- **The Netherlands and Japan.** Their toolmakers build the controlled machines, so the two governments gain a say in who gets them. China accounts for a falling share of the firms' sales. At Tokyo Electron, China's share fell from 44.4 percent in the year to March 2024 to 34.1 percent in the year to March 2026 [74]. The company said the share dropped to 26.8 percent in the quarter to March 2026 because spending on leading-edge chips grew faster than spending on older ones [74].

21 Enhanced Favorable Treatment for the United Arab Emirates Under the Export Administration Regulations. Federal Register (Commerce Department; Industry and Security Bureau) · 14 July 2026 [A] federalregister.gov/documents/2026/07/14/2026-14132/enhanced-favorable-treatment-for-the-u...

69 Export, Transshipment and Transit of High-Performance AI Chips of US Origin. Ministry of Investment, Trade and Industry (Malaysia) · 14 July 2025 [A] miti.gov.my/miti/resources/Media%20Release/%5BFINAL%5D_MITI_Press_Stmt_Malaysia_Regulates_...

70 Chinese National and Two U.S. Citizens Charged with Conspiring to Smuggle Artificial Intelligence Technology to China. U.S. Department of Justice · 25 March 2026 [A] justice.gov/opa/pr/chinese-national-and-two-us-citizens-charged-conspiring-smuggle-artific...

71 Protecting U.S. Tech: China Committee and Bipartisan, Bicameral Leaders Unite to Stop CCP AI Chip Smuggling. U.S. House Select Committee on the Chinese Communist Party · 7 July 2025 [A] chinaselectcommittee.house.gov/media/press-releases/protecting-us-tech-china-committee-and...

72 S.1705 – Chip Security Act, 119th Congress (2025-2026), bill text. U.S. Congress (Congress.gov, Library of Congress) [A] congress.gov/bill/119th-congress/senate-bill/1705/text

73 UNANIMOUS: Bipartisan Huizenga Legislation to Curb AI Chip Smuggling Passes Committee. U.S. House of Representatives (Office of Rep. Bill Huizenga) · 26 March 2026 [A] huizenga.house.gov/news/documentsingle.aspx

74 FY2026 (April 2025-March 2026) Financial Announcement, transcript. Tokyo Electron · 30 April 2026 [A] tel.com/ir/library/report/pjuomj00000000tf-att/fy26q4transcript-e.pdf

- **The United States.** Its firms took 65 percent of the world's logic chip sales in 2022 [10]. Counting every type of chip, they took 53.4 percent of global sales in 2025, their highest share since 1984 [11]. Commerce expects the CHIPS-funded projects to lift the US share of global leading-edge logic chip manufacturing from 0 percent in 2022 to 20 percent by 2030 [58].
- **China.** Mature-node chips, tools below the leading edge, DRAM and NAND are moving to Chinese suppliers. Extreme-ultraviolet machines, the photoresist they need and high-bandwidth memory in volume are not. In September 2025 SemiAnalysis expected China, without smuggled memory, to make fewer Ascend chips in 2026 than in 2025 [56].

The Gulf gains without making chips, because the July 2026 rule lets named Emirati buyers receive advanced computing items without a license [21].

The US government also uses tariffs and an equity stake alongside licenses. Every licensed H200 bound for China pays a 25 percent tariff on its way through US inspection [25]. The government agreed in August 2025 to take a 9.9 percent stake in Intel [59]. H200 sales to China go through case-by-case review [27].

What to watch

- **November 2026.** Whether China renews the one-year suspension of its gallium ban or lets it expire [52].
- **The MATCH Act's 150-day deadline.** If it passes, allies will have 150 days to match its limits on deep-ultraviolet tools [42].
- **CXMT's memory ramp.** In September 2025 SemiAnalysis judged that domestic memory would limit China's output of AI chips before logic wafers did [56].
- **SMIC's advanced-node ramp.** In September 2025 SemiAnalysis conservatively estimated SMIC's capacity at 7 nm and below at 45,000 wafers a month by the end of 2025, 60,000 in 2026 and 80,000 in 2027 [56]. SMIC makes those wafers

10 Supply Chain Explorer: Advanced Chips. Emerging Technology Observatory (ETO) [A] chipexplorer.eto.tech

11 2026 State of the U.S. Semiconductor Industry. Semiconductor Industry Association · 2 July 2026 [A] semiconductors.org/wp-content/uploads/2026/07/SIA_2026_State-of-the-Semiconductor-Industry...

58 GAO-26-107882, SEMICONDUCTORS: Information on Projects Funded to Strengthen U.S. Supply Chain. U.S. Government Accountability Office [A] files.gao.gov/reports/GAO-26-107882/index.html

56 Huawei Ascend Production Ramp: Die Banks, TSMC Continued Production, HBM is The Bottleneck. SemiAnalysis · 8 September 2025 [B] newsletter.semianalysis.com/p/huawei-ascend-production-ramp

21 Enhanced Favorable Treatment for the United Arab Emirates Under the Export Administration Regulations. Federal Register (Commerce Department; Industry and Security Bureau) · 14 July 2026 [A] federalregister.gov/documents/2026/07/14/2026-14132/enhanced-favorable-treatment-for-the-u...

25 NVIDIA CORP, Form 10-Q quarterly report for the period ended 2026-07-26 (10-Q). U.S. Securities and Exchange Commission (filing by NVIDIA CORP) · 26 August 2026 [A] sec.gov/Archives/edgar/data/1045810/000104581026000075/nvda-20260726.htm

59 Intel and Trump Administration Reach Historic Agreement to Accelerate American Technology and Manufacturing Leadership. Intel · 22 August 2025 [A] intc.com/news-events/press-releases/detail/1748/intel-and-trump-administration-reach-histo...

27 Revision to License Review Policy for Advanced Computing Commodities. Federal Register (Commerce Department; Industry and Security Bureau) · 15 January 2026 [A] federalregister.gov/documents/2026/01/15/2026-00789/revision-to-license-review-policy-for-...

52 The U.S.-China Trade Truce Has Not Solved the Gallium Problem. Center for Strategic and International Studies · 11 May 2026 [A] csis.org/analysis/us-china-trade-truce-has-not-solved-gallium-problem

42 Risch, Ricketts, Kim Introduce MATCH Act; Level the Global Playing Field for U.S. Tech. U.S. Senate Committee on Foreign Relations · 8 April 2026 [A] foreign.senate.gov/press/rep/release/risch-ricketts-kim-introduce-match-act-level-the-glob...

without extreme-ultraviolet machines, so its tightest layers need several passes [55]. Yield, the share of chips on a wafer that work, is the open question.

- **Arizona's lag.** How far Arizona's newest process stays behind Taiwan's as TSMC's four new fabs there, at 2 nm or more advanced, start production [61].

FIVE THINGS TO REMEMBER

- An export control works only when a country's firms sell something the buyer cannot get elsewhere.
- US controls also cover foreign goods made with American technology.
- The controls have kept China below the leading edge. But Chinese toolmakers' share of China's own chipmaking equipment market rose from 10 to 15 percent before 2022 to 35 percent in 2025.
- China's share of ASML's and Tokyo Electron's sales has fallen since 2024.
- Taiwan held 69 percent of the world's logic capacity below 10 nm in 2022, and TSMC makes its 2 nm chips only there.

55 Is SMIC N+3's Metal Pitch Smaller than Intel 18A's?. SemiAnalysis · 14 June 2026 [B] newsletter.semianalysis.com/p/steel-smic-n3-teardown

61 TSMC Announces Additional $100 Billion Investment In Arizona. City of Phoenix · 16 July 2026 [A] phoenix.gov/newsroom/ced-news/tsmc-announces-additional--100-billion-investment-in-arizona...

Made by **Yashvardhan Sharma** with Claude. Corrections and questions: contact@chipsupplychain.org.

UPDATED 2 OCTOBER 2026

A company filings, governments, standards bodies, universities and research institutions; B SemiAnalysis, Reuters, the Financial Times and Nikkei Asia.

HOW TO CITE

```
Sharma, Y. Chip Supply Chain: How AI Chips Are Made. September 2026.
https://chipsupplychain.org/ (accessed 2026-10-02).
```

```
@misc{chipsupplychain2026,
  title  = {Chip Supply Chain: How AI Chips Are Made},
  author = {Sharma, Yashvardhan},
  year   = {2026},
  month  = {September},
  url    = {https://chipsupplychain.org/},
  note   = {Accessed 2026-10-02}
}
```